

Computer Architecture and Organization

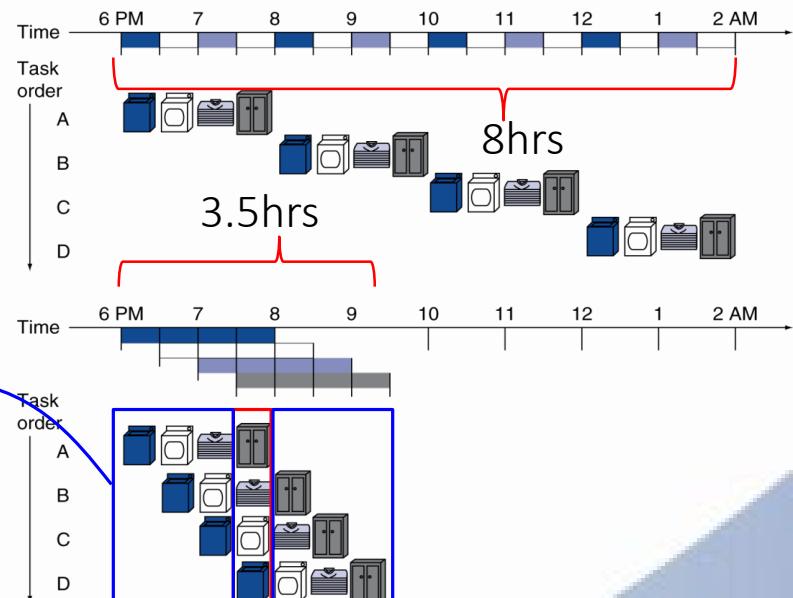
INSTRUCTOR: YAN-TSUNG PENG

DEPT. OF COMPUTER SCIENCE, NCCU

CPU Design: Pipelining

Pipelining Analogy

- Pipelined laundry: overlapping execution
 - Parallelism improves performance (throughput)
 - Not shorten the time needed for completing one load
- Assume each stage takes the same amount of time
- Let each stage take 0.5hr, and 4 stages in total for a task
- If we do four loads separately
 - $4 \text{ loads} * 4 \text{ stages} * 0.5\text{hr} = 8\text{hr}$
- Using pipelining ($k \geq 3$ loads)
 - $(4 \text{ loads}-3)*0.5\text{hr}+(3 \text{ stages} + 3 \text{ stages})*0.5=3.5\text{hr}$
- Using pipelining for infinite loads:
 - $(n \text{ loads}-3)*0.5\text{hr}+(3 \text{ stages} + 3 \text{ stages})*0.5=0.5n+1.5\text{hr}$



$$\text{Max speed-up: } \lim_{n \rightarrow \infty} \frac{2n}{0.5n+1.5} = 4$$

MIPS Pipeline

- Five stages, one step per stage
 1. IF: Instruction fetch from memory
 2. ID: Instruction decode & register read
 3. EX: Execute operation or calculate address
 4. MEM: Access memory operand
 5. WB: Write result back to register

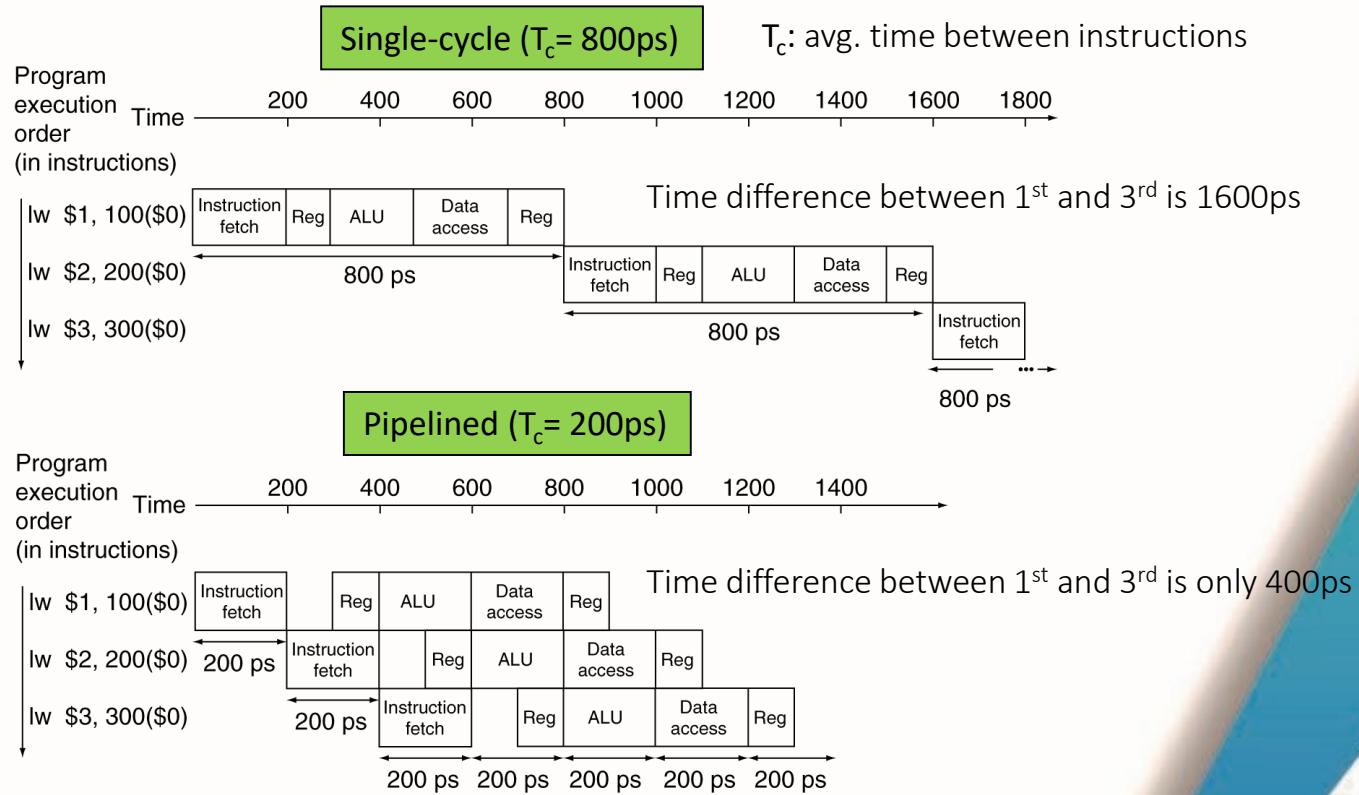
Pipeline Performance

- Assume time for stages is
 - 100ps for register read or write
 - 200ps for other stages
- Compare pipelined datapath with single-cycle datapath (next page)

Instr	IF	ID/REG	EX	MEM	WB	Total time
lw	200ps	100 ps	200ps	200ps	100 ps	800ps
sw	200ps	100 ps	200ps	200ps		700ps
R-format	200ps	100 ps	200ps		100 ps	600ps
beq	200ps	100 ps	200ps			500ps

$$1 \text{ ps} = 10^{-12} \text{ s}$$

Pipeline Performance



Pipeline Speedup

- If all stages are balanced
 - i.e., all take the same time
 - Time between instructions_{pipelined}

$$= \frac{\text{Time between instructions}_{\text{nonpipelined}}}{\text{Number of stages}}$$
- Speedup due to increased throughput
 - Latency (time for each instruction) does not decrease

Instr	IF	ID/REG	EX	MEM	WB	Total time
Iw	200ps	100 ps	200ps	200ps	100 ps	800ps
sw	200ps	100 ps	200ps	200ps		700ps
R-format	200ps	100 ps	200ps		100 ps	600ps
beq	200ps	100 ps	200ps			500ps

When pipelining, each stage takes
200ps, so in total it's 1000ps.

$$\frac{1000 \text{ ps}}{5 \text{ stages}} = 200 \text{ ps}$$

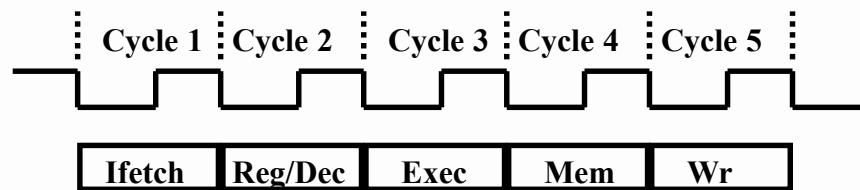
Pipelining and ISA Design

■ MIPS ISA designed for pipelining

- All instructions are 32-bits
 - Easier to fetch and decode in one cycle
 - c.f. x86: 1- to 17-byte instructions
- Few and regular instruction formats
 - Can decode and read registers in one step (in the ID stage)
 - Load and Store addressing (in two stages)
 - Calculate address in 3rd stage (EX)
 - Access memory in 4th stage (MEM)
 - Alignment of memory operands
 - Memory access takes only one cycle

The fact that the instruction formats of MIPS do not vary much and the length of the instructions is the same makes pipelining easy.

Pipeline Stages



IF: Instruction Fetch

Fetch the instruction from the Instruction
Memory

ID: Instruction Decode

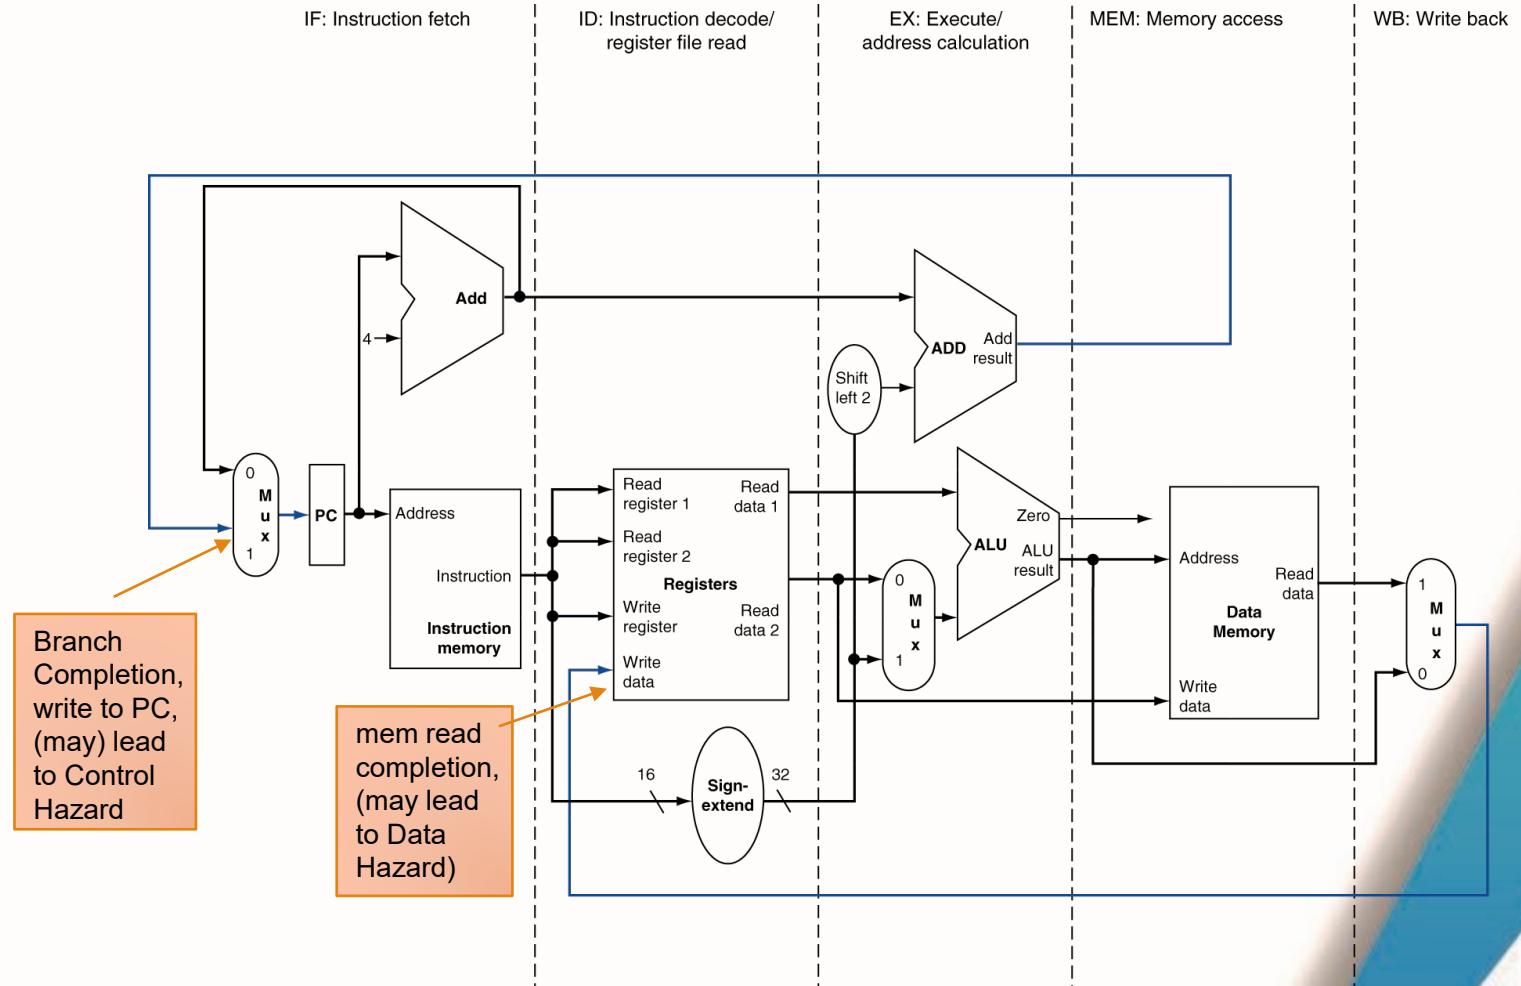
Registers fetch and instruction decode

EX: Calculate the memory address

MEM: Read the data from the Data Memory

WB: Write the data back to the register file

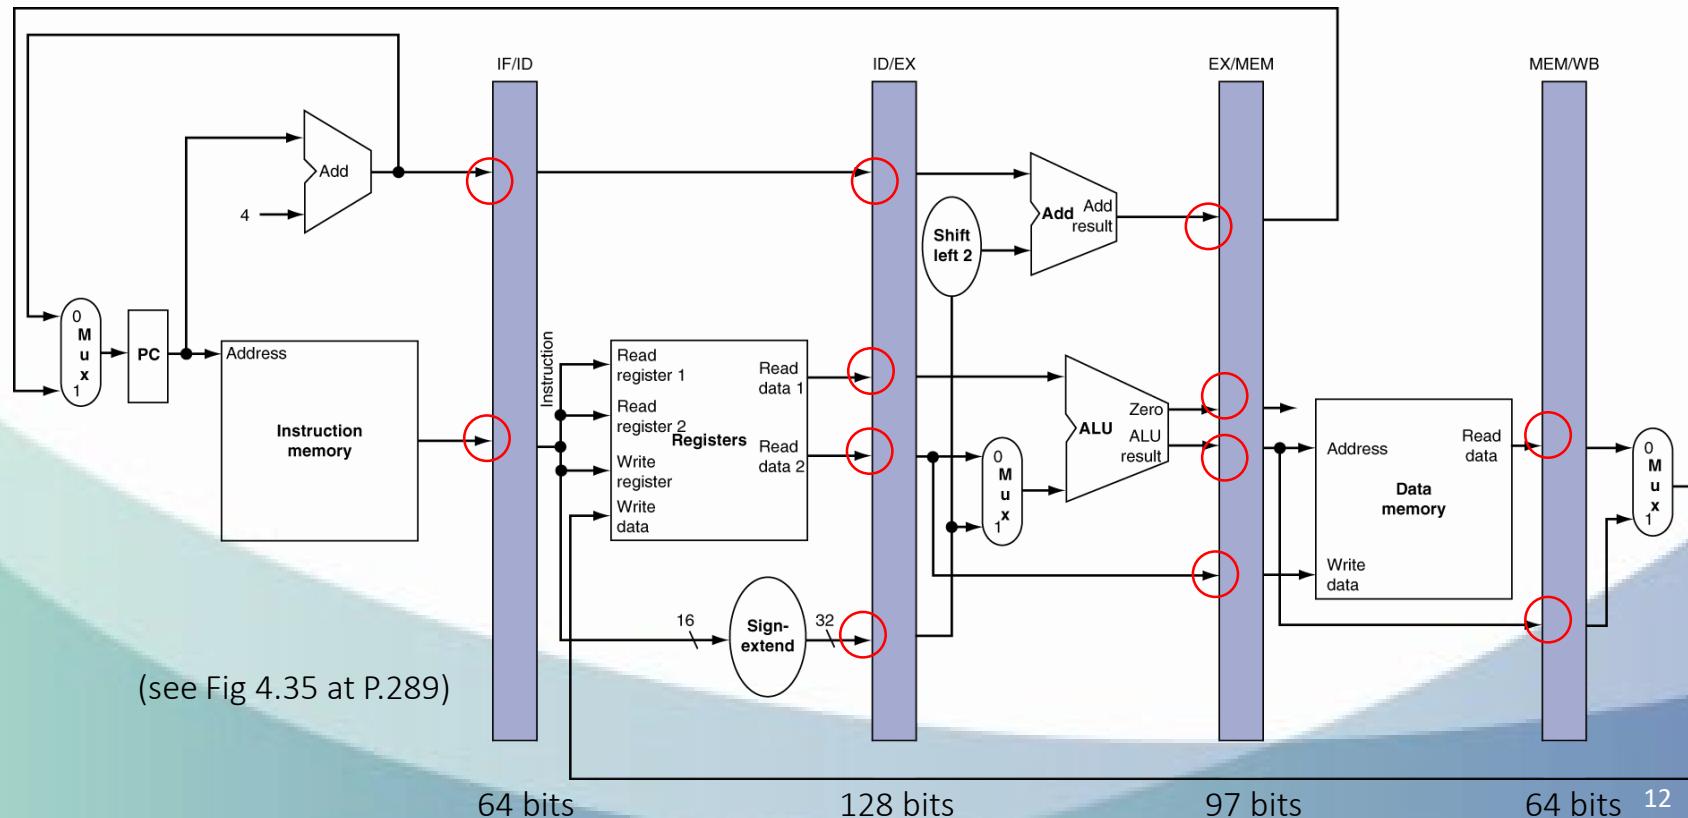
MIPS Pipelined Datapath - What do we need to split an instruction into stages?



Pipelined Version of Datapath

- Need registers between stages
 - To hold information produced in previous cycle (including data and control)

Pipelined registers:
IF/ID, ID/EX, EX/MEM, MEM/WB



Hazards

- Situations that prevent starting the next instruction in the next cycle
- Three Hazards: Structure, data, and control hazards
 - Structure hazards
 - A required resource is busy
 - Competing the same hardware resource in one cycle
 - Data hazard
 - Need to wait for data outputted from previous instruction
 - Control hazard
 - Taken branch decision from the previous instruction needs to void the current instruction

Example - lw

■ lw \$t0, 1200(\$t1)

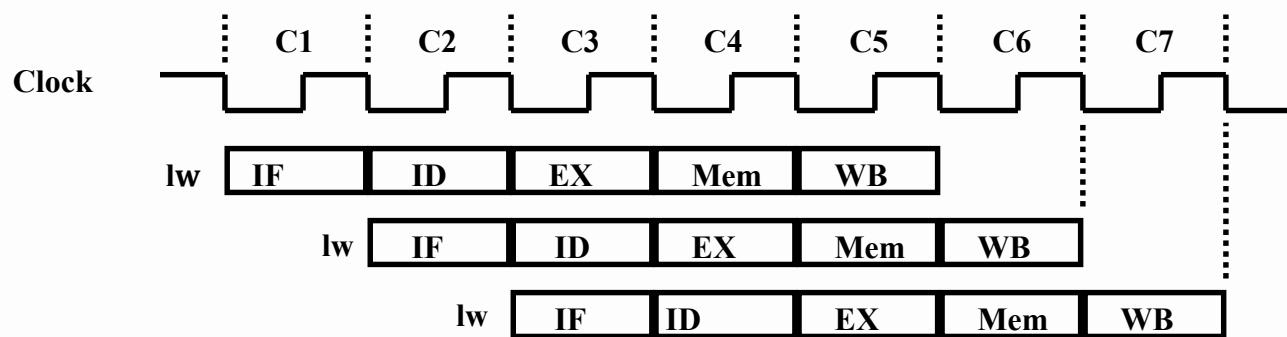
R8=01000 R9=01001
rt rs

$$1200_{ten} = 0000010010110000_{two}$$

100011	01001	01000	0000 0100 0110 0000
6 bits	5 bits	5 bits	16 bits

op rs rt address

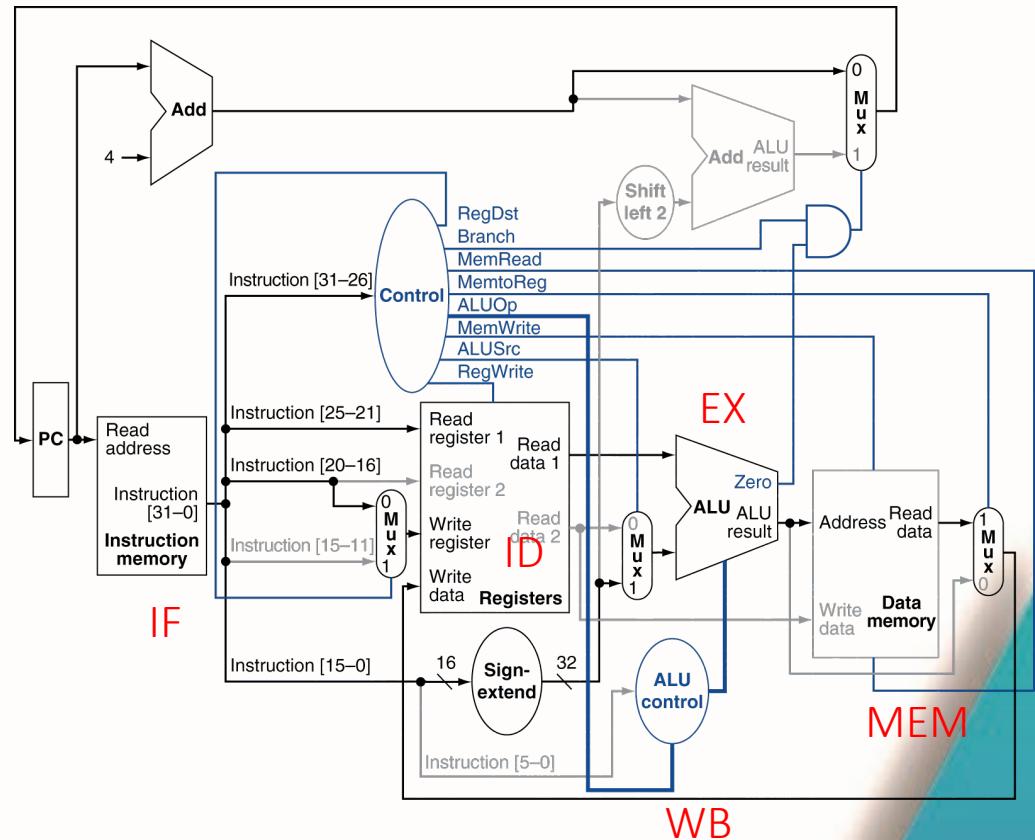
Example - lw



There are 5 functional units in the pipeline datapath:

- IF stage: Instruction Memory
- ID stage: Register File's Read ports
- Ex stage: ALU execution
- MEM stage: Data Memory
- WB stage: Register File's Write port

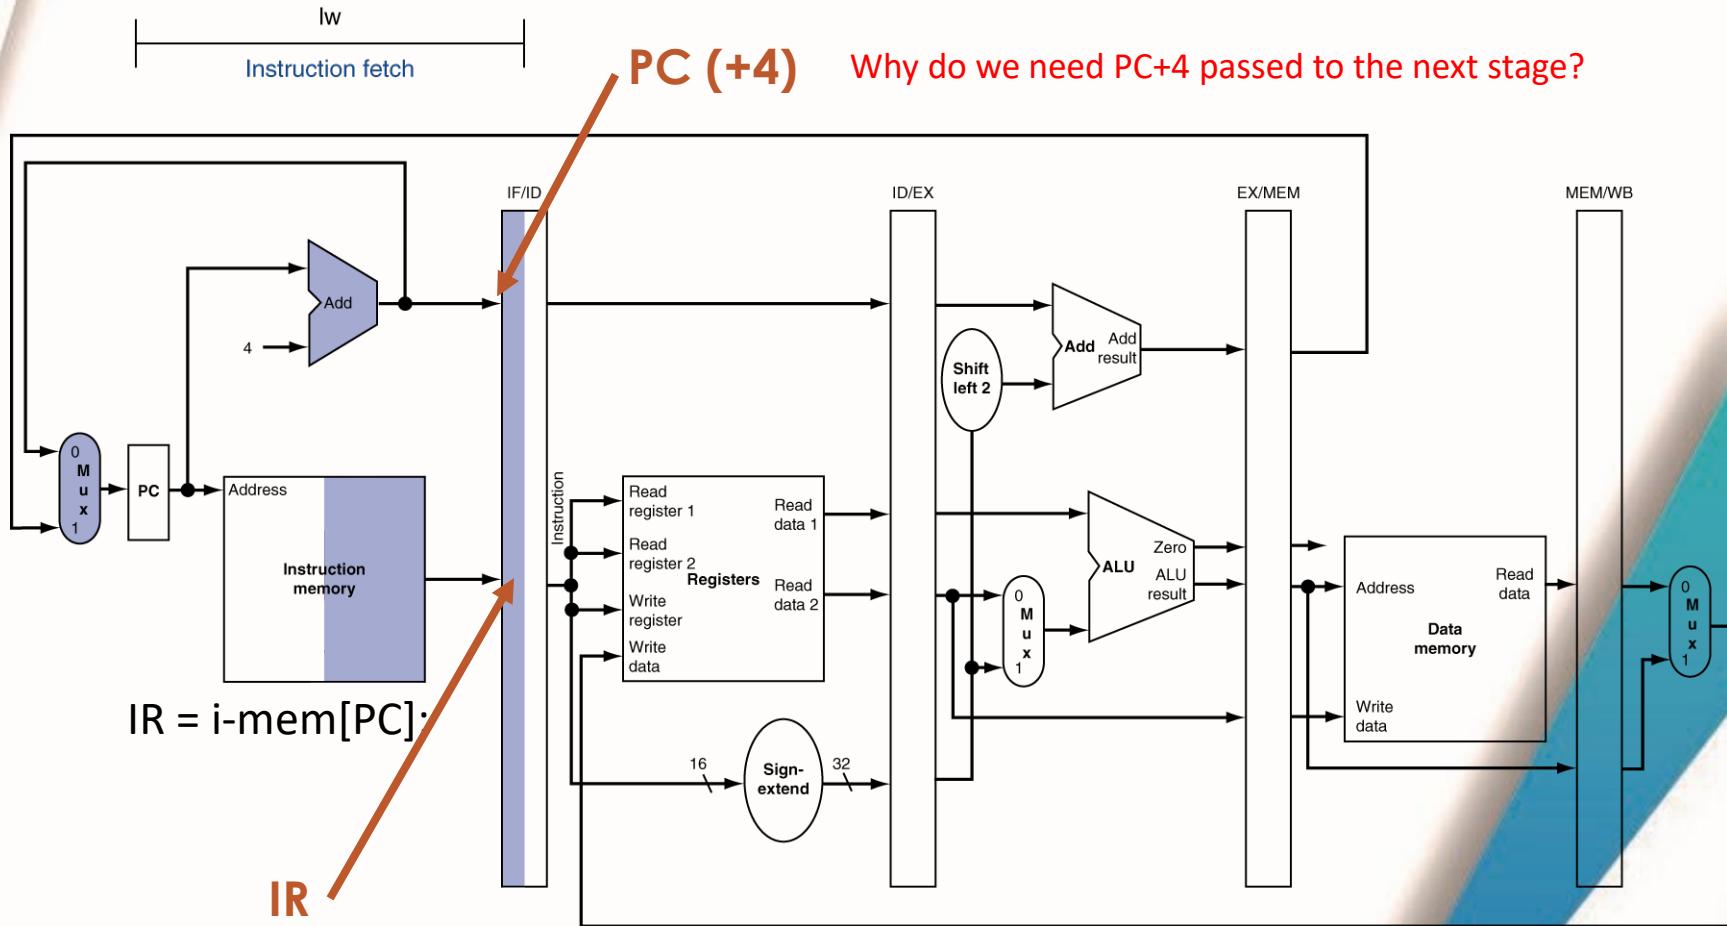
lw \$t0, 1200(\$t1)



- IF: Fetch the instruction from the Instruction Memory
- ID: Registers fetch and instruction decode
- EX: Calculate the memory address
- MEM: Read the data from the Data Memory
- WB: Write the data back to the register file

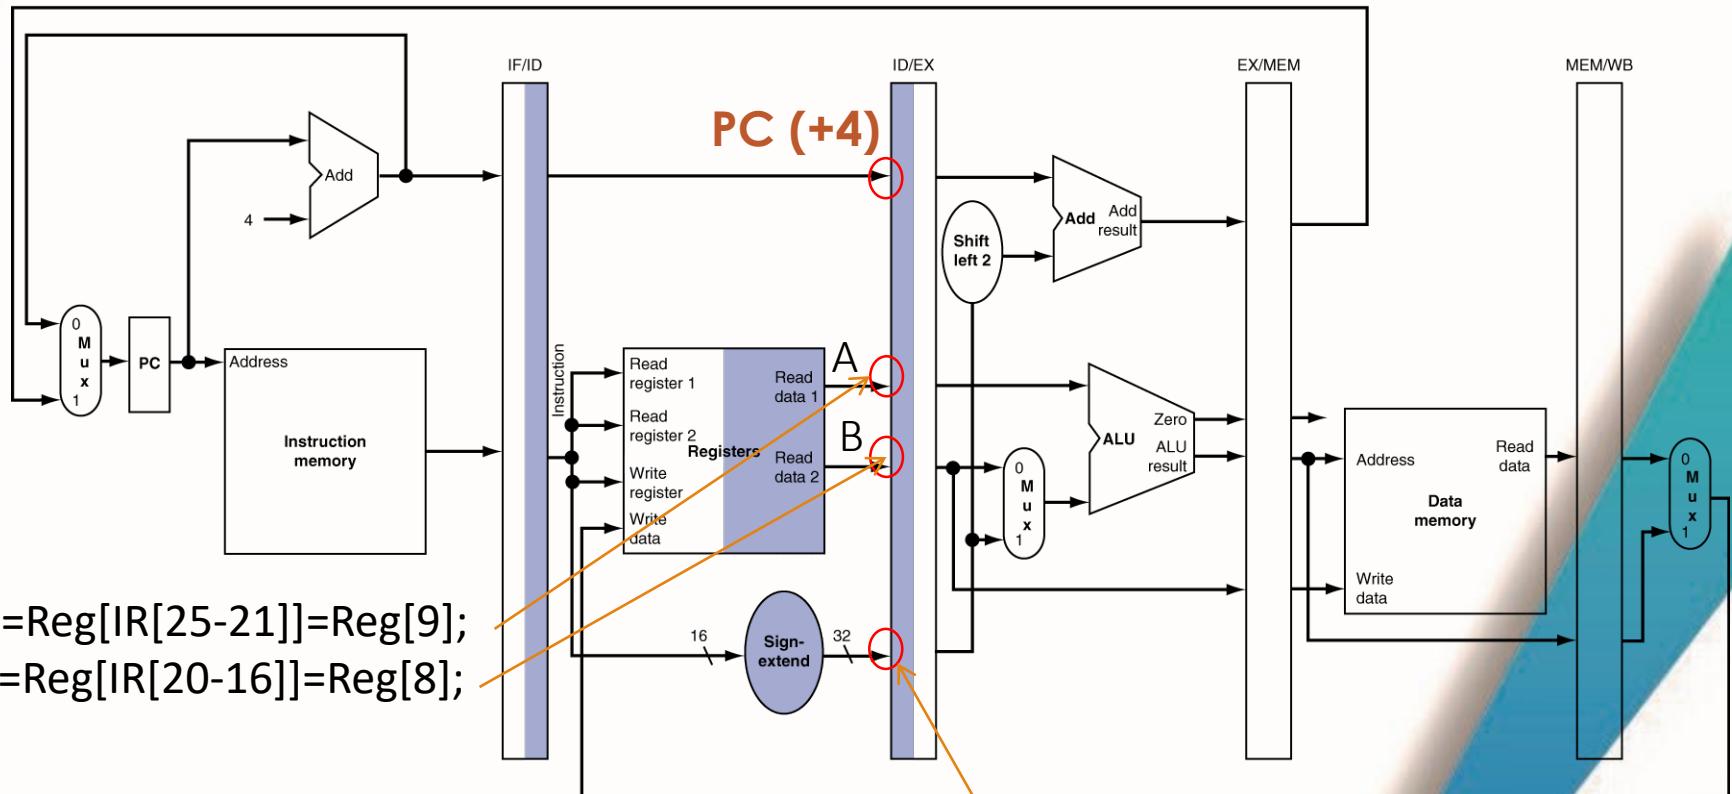
IF Stage - lw

Fetch the instruction from the Instruction Memory

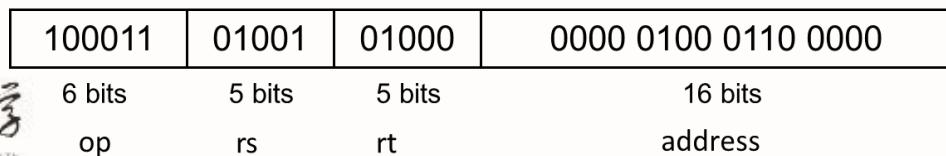


ID Stage - lw

Fetch registers and decode the instruction



$$A = rs = \text{Reg}[\text{IR}[25-21]] = \text{Reg}[9]; \\ B = rt = \text{Reg}[\text{IR}[20-16]] = \text{Reg}[8];$$

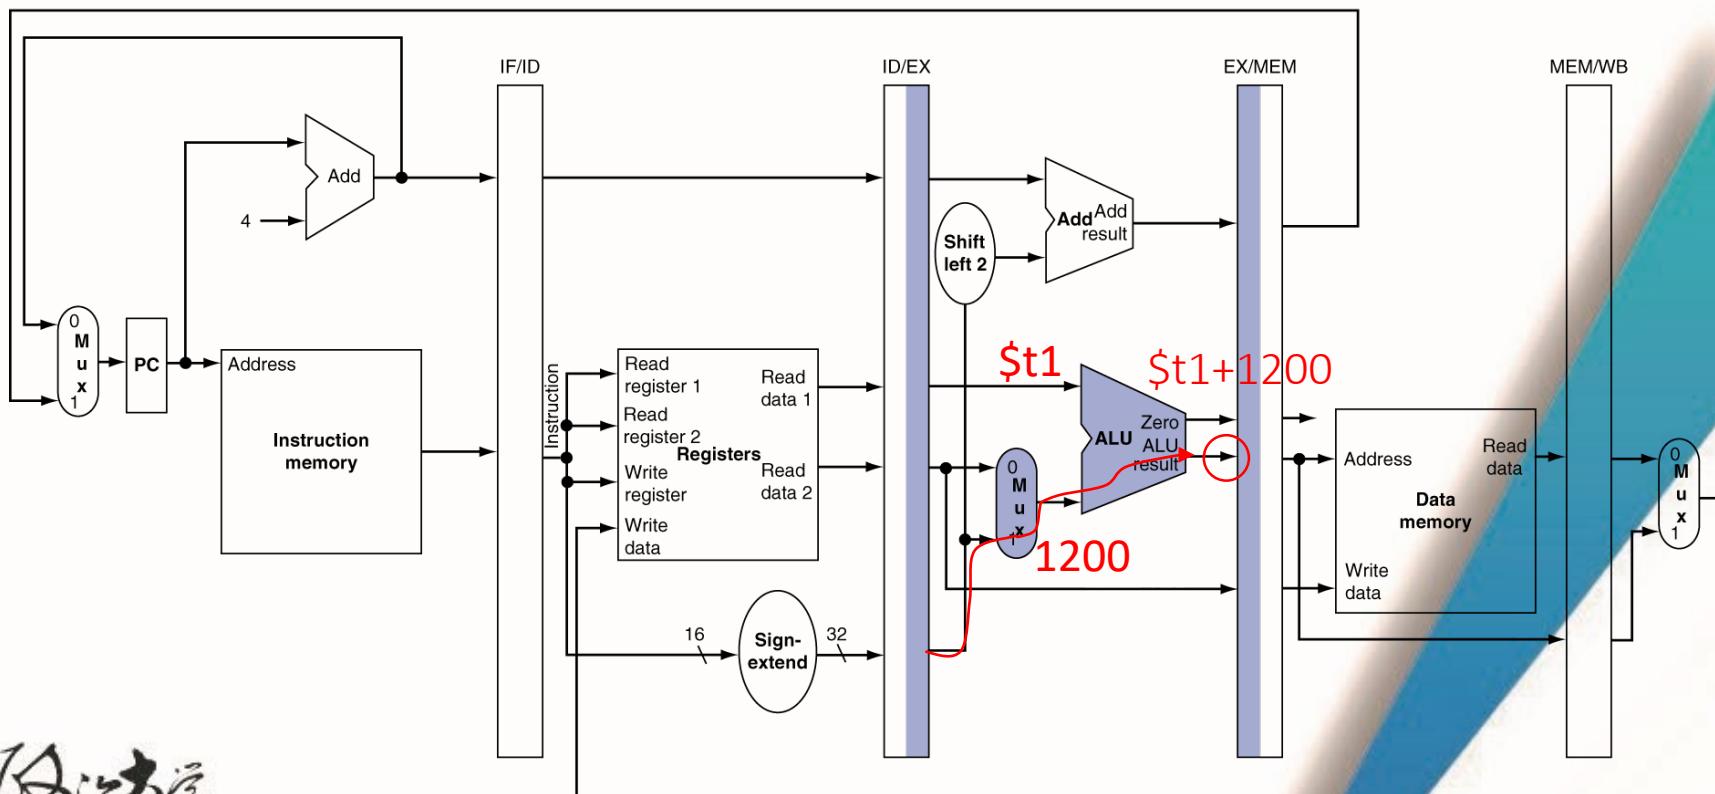


EX Stage - lw

Compute the memory address ($$t1+1200$)

$$\begin{array}{c} \$t1 \quad 1200 \\ \text{ALUout} = A + \text{sign-ext}(IR[15-0]) \end{array}$$

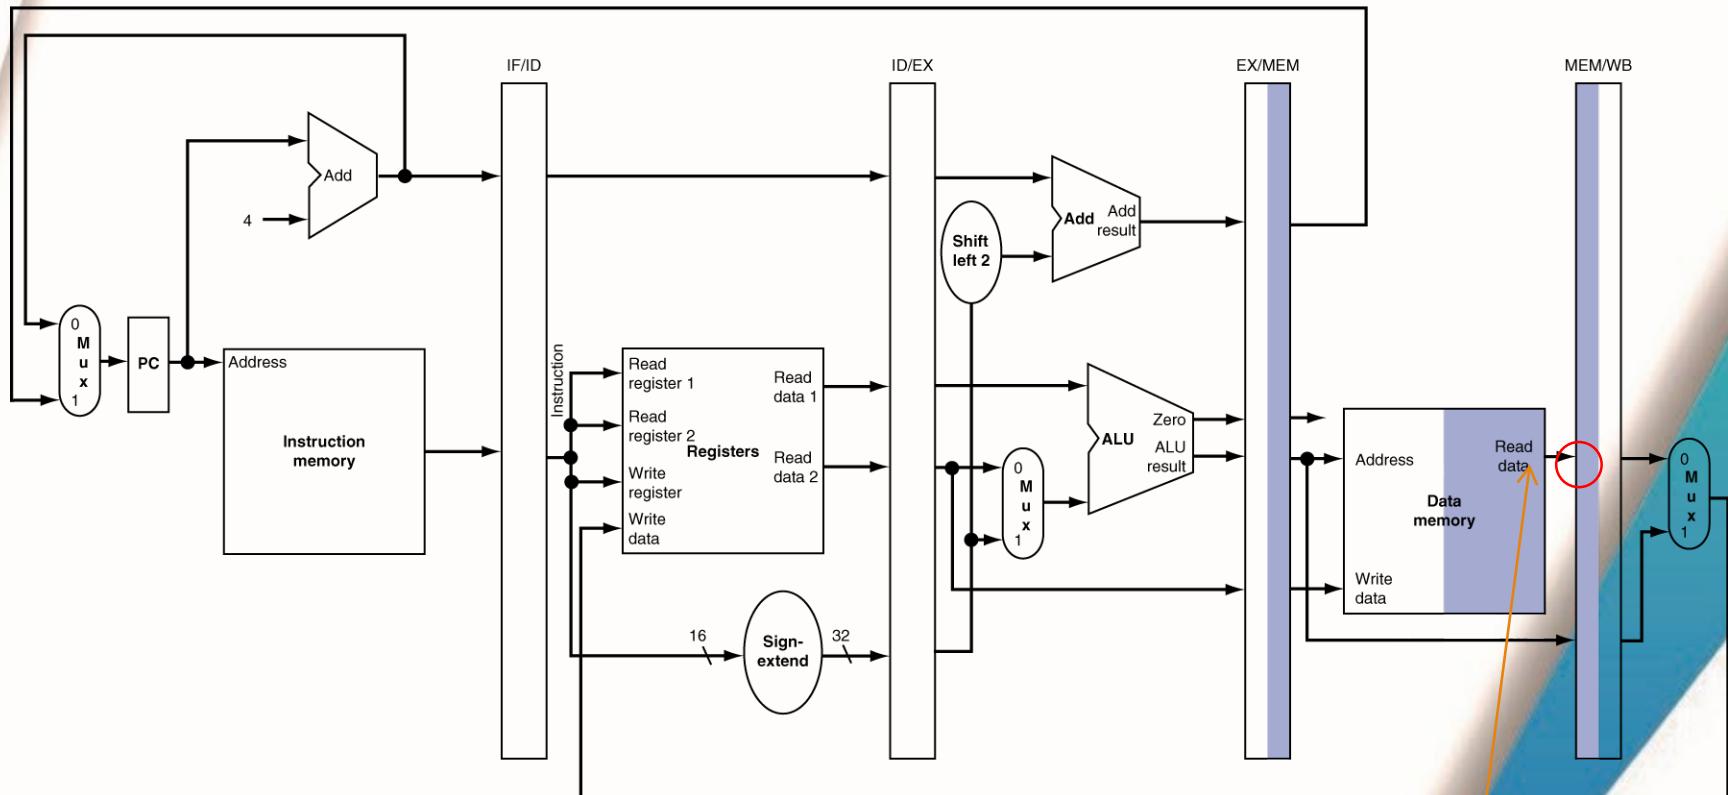
Calculate Base + Offset



MEM Stage - lw

Read data from the Data Memory

lw \$t0, 1200(\$t1)

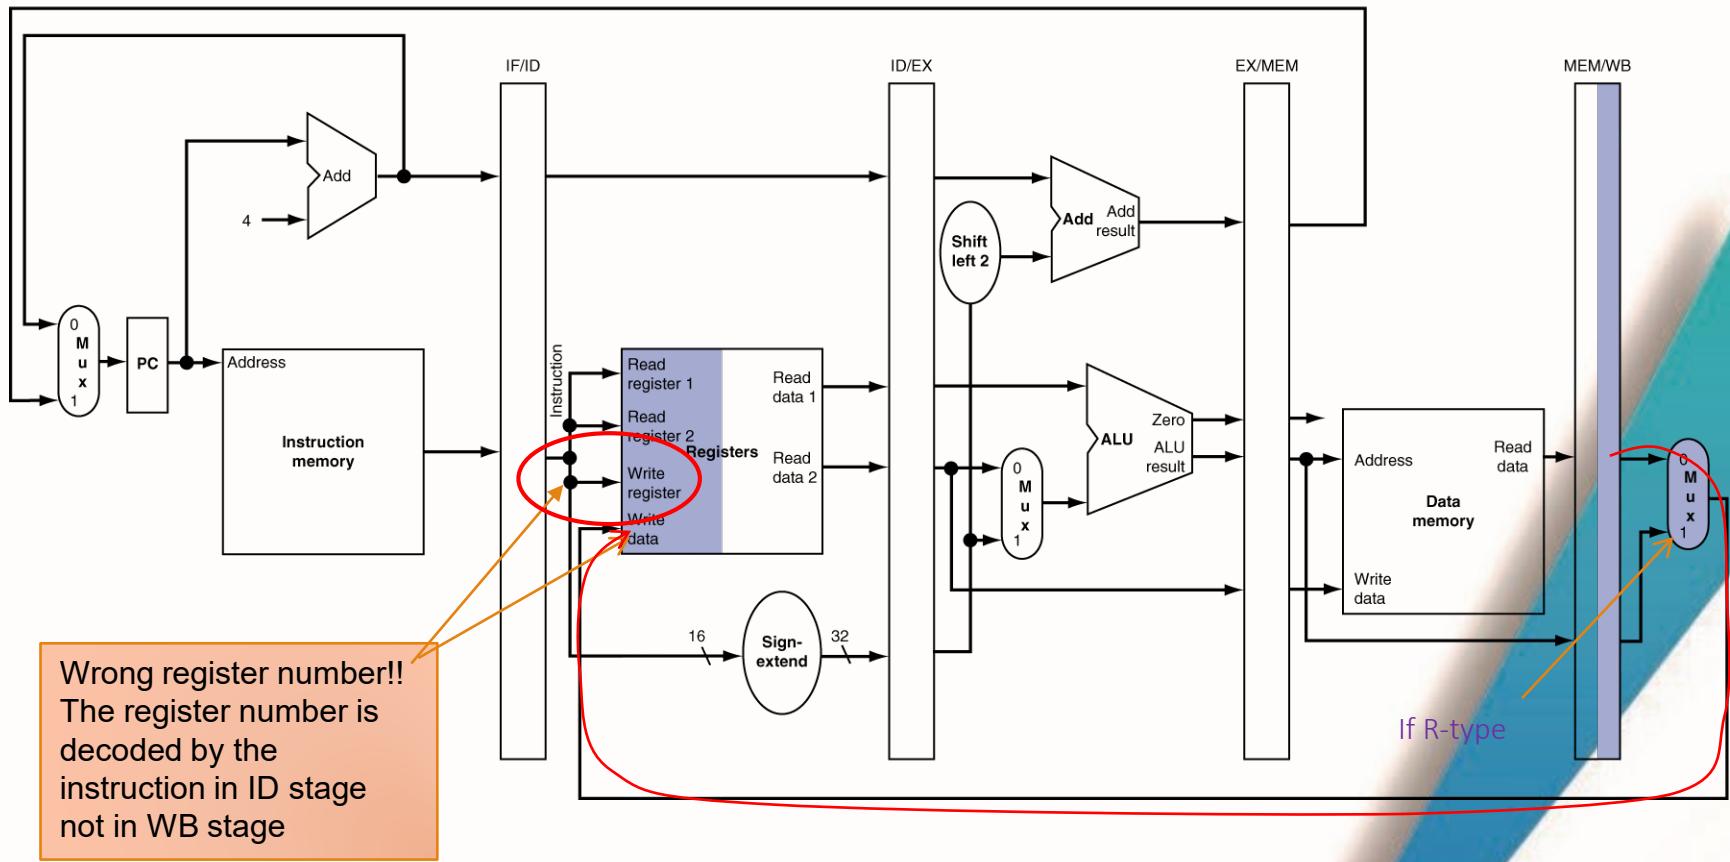


Read MEM[\$t1+1200]

WB Stage - lw

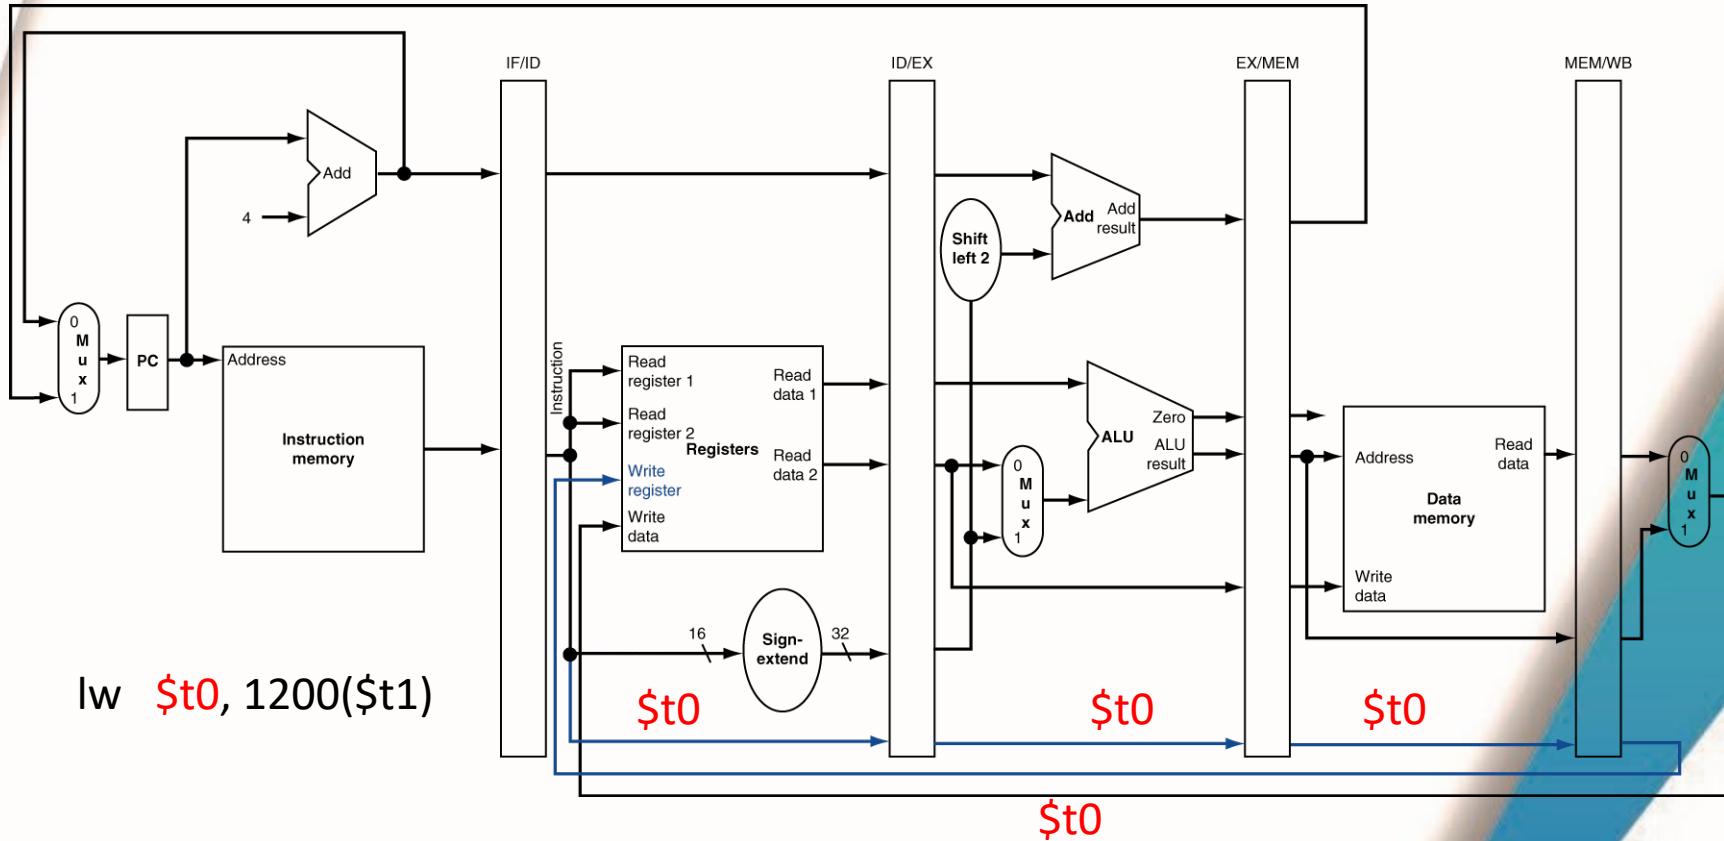
Write the data back to the register

lw
Write back



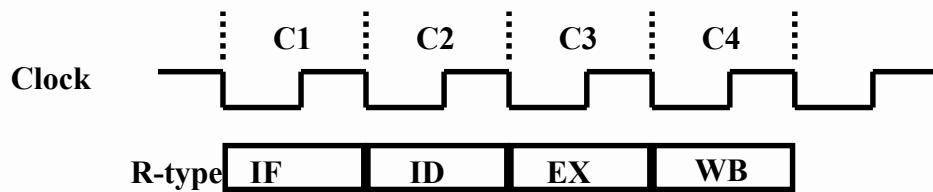
WB Stage - lw

Corrected Datapath for Load



Write register number shall be decided by the instruction in WB stage.

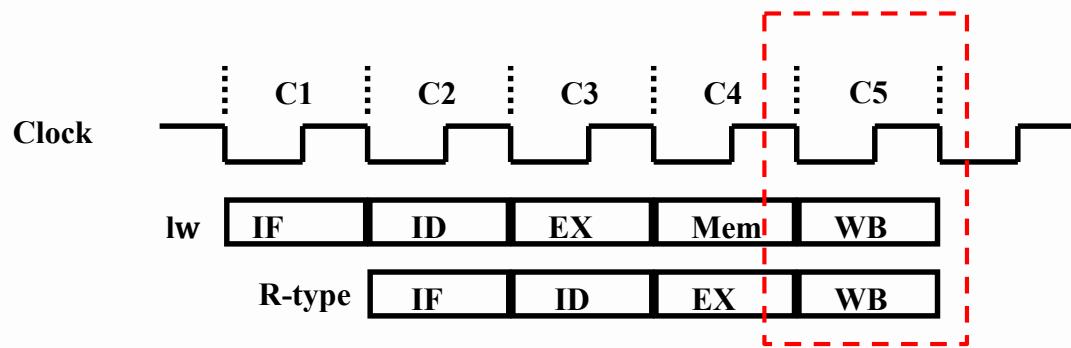
R-type Instructions



There are four stages for a R-type instruction

- IF: fetch the instruction from the Instruction Memory
- ID: registers fetch and instruction decode
- EX: ALU operates on the two register operands
- WB: write ALU output back to the register file

Pipelining R-type and load



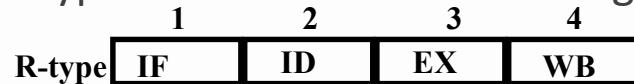
- If R-type takes four cycles and `lw` takes five cycles, we then have a *structural hazard*:
 - `lw` and R-type instruction both write results into the register file on the same stage but the register only has one write port.

Observation

- Structural hazard – two or more instructions try to access the same hardware resource
 - Load writes results into the register on its **5th** stage



- R-type writes results into the register on its **4th** stage

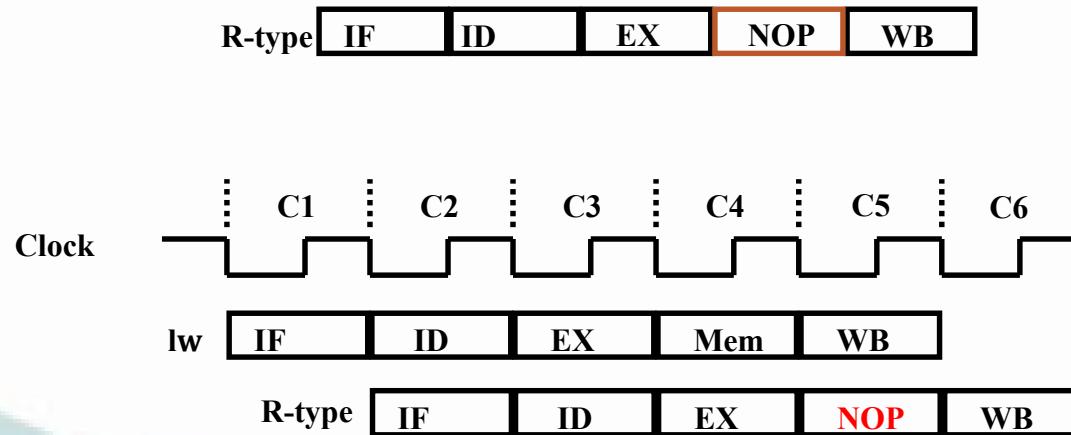


This can be resolved by inserting bubbles or making all the instructions having the same length

Make Instructions having the Same Length

Make R-type instructions to run five stages

- R-type WB at Stage 5
- MEM stage == NOP



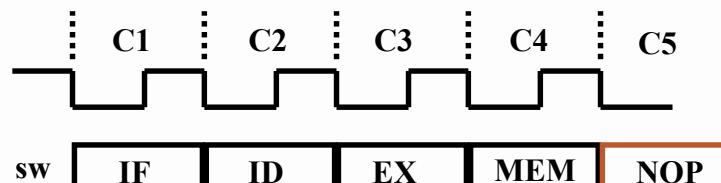
Example – sw

There are four stages for sw

- IF: fetch the instruction from the Instruction Memory
- ID: fetch registers and decode the instruction
- EX: compute the target memory address
- MEM: write the reg data into the Data Memory

Make it having five stages

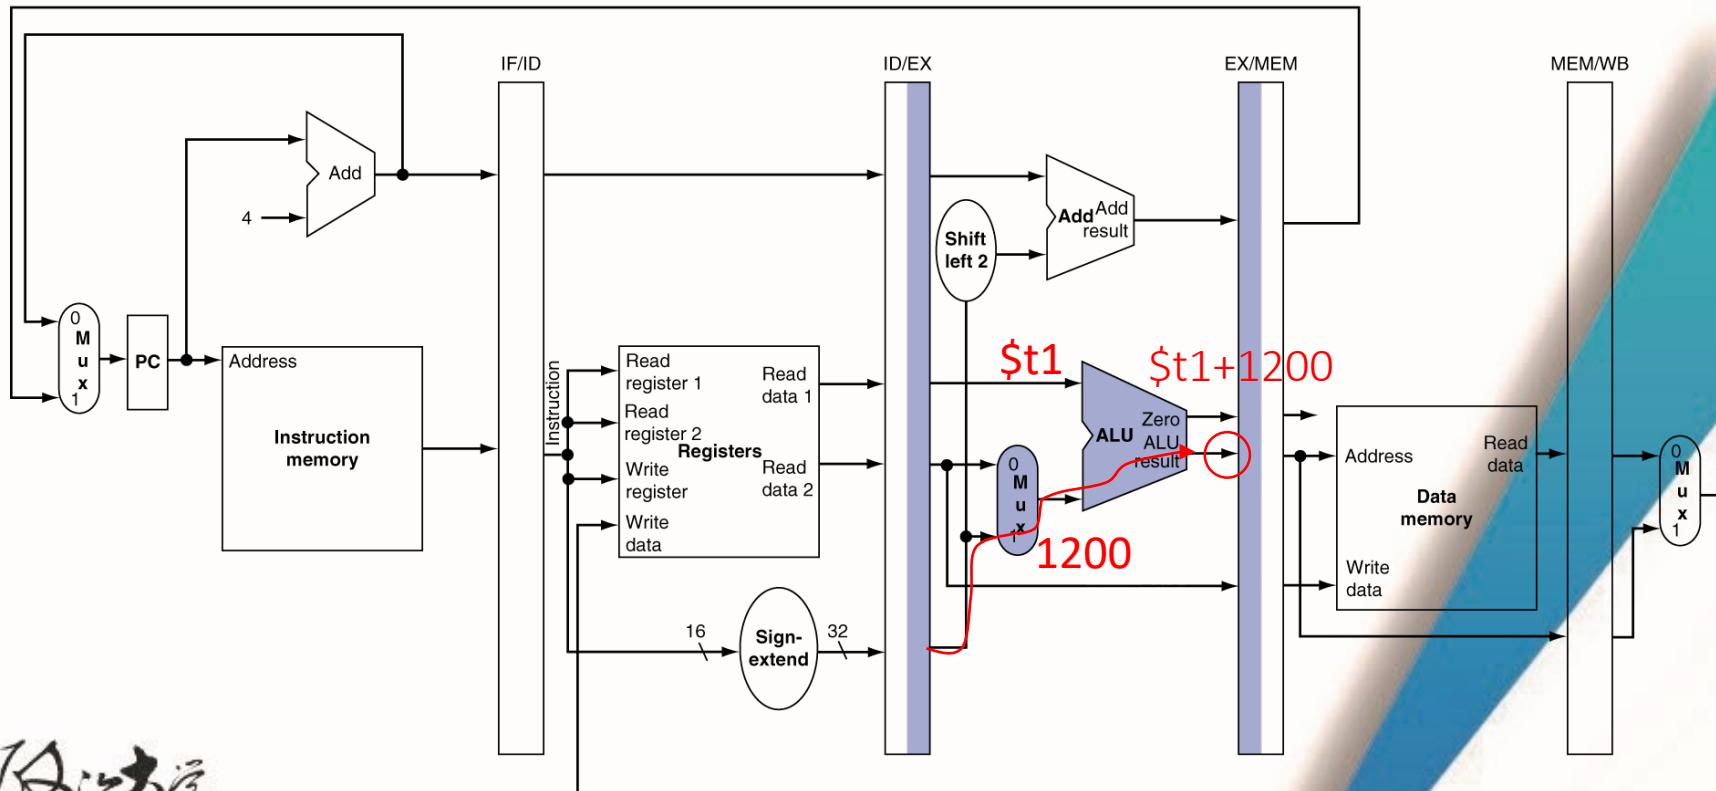
- WB == NOP



EX Stage - SW

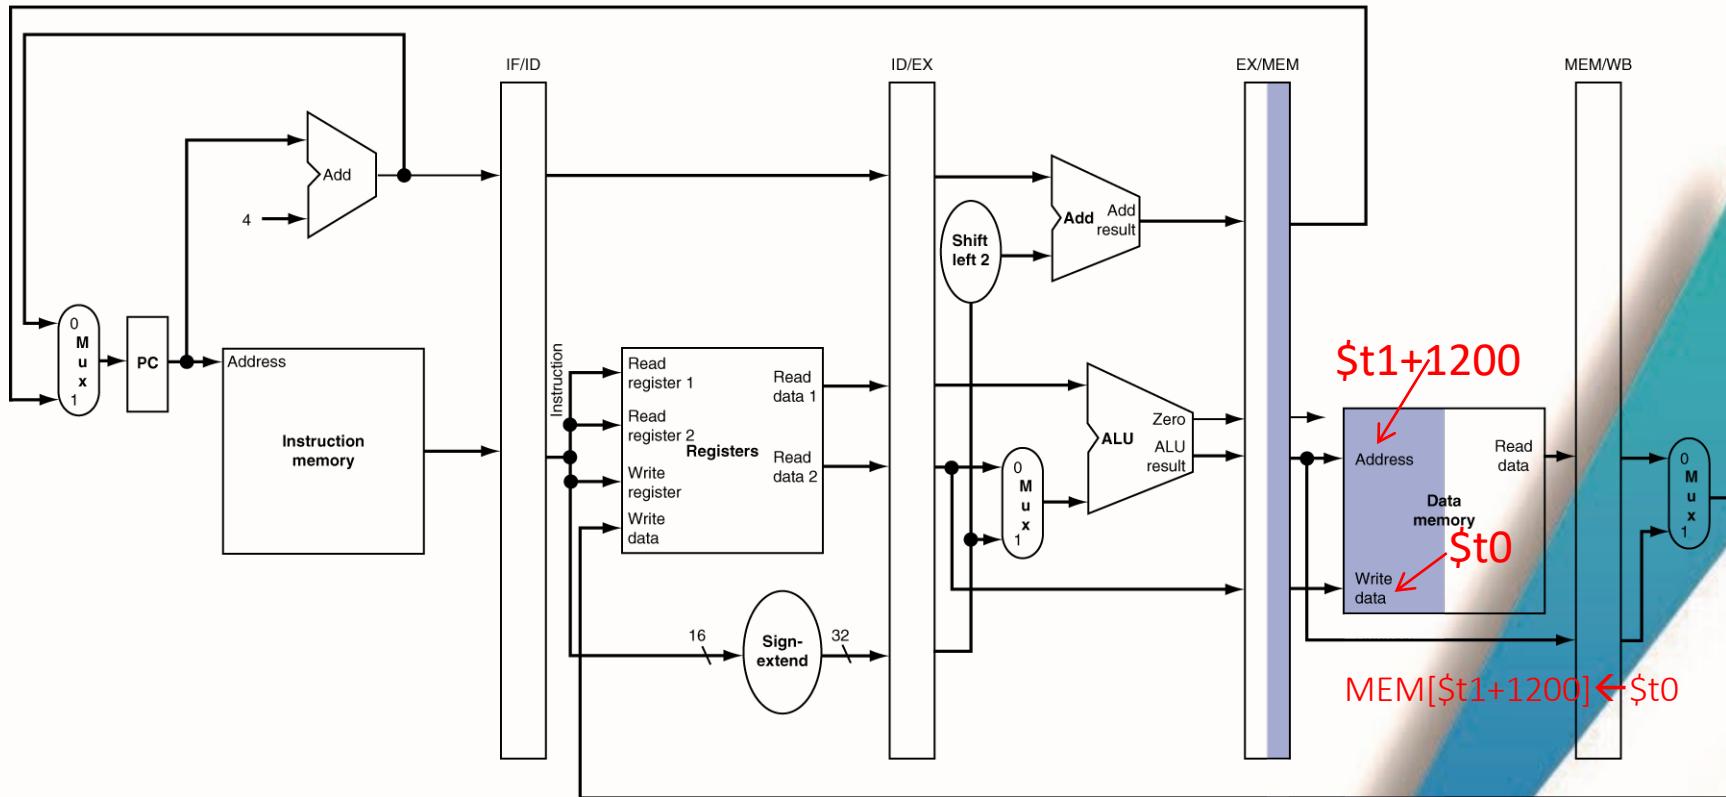
Compute the memory address ($\$t1+1200$)

sw \$t0, 1200(\$t1) |
Execution



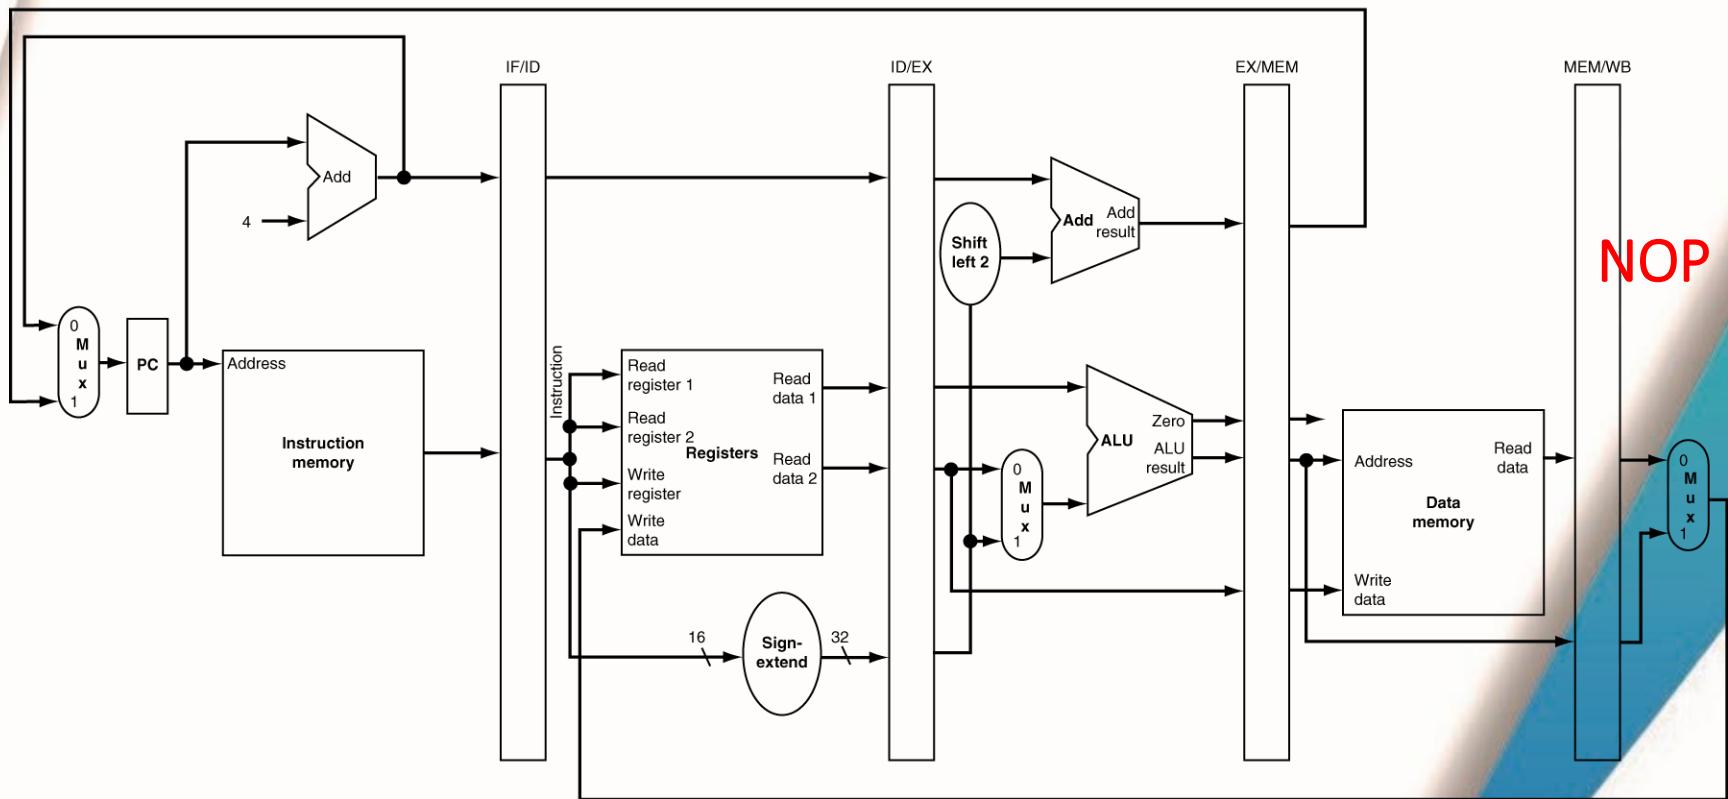
MEM Stage - SW

sw \$t0, 1200(\$t1)



WB Stage - sw

SW
Write-back

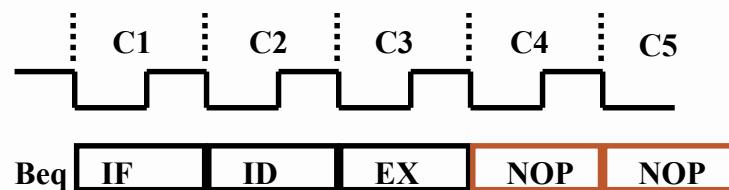


Example - beq

- IF: fetch the instruction from the Instruction Memory
- ID: fetch registers and decode the instruction
- EX:
 - Compares the two register operands with ALU
 - Compute the target address
 - If branch is taken, write the target address into PC; otherwise, write PC+4

Two extra stages:

- MEM: NOP
- WB: NOP



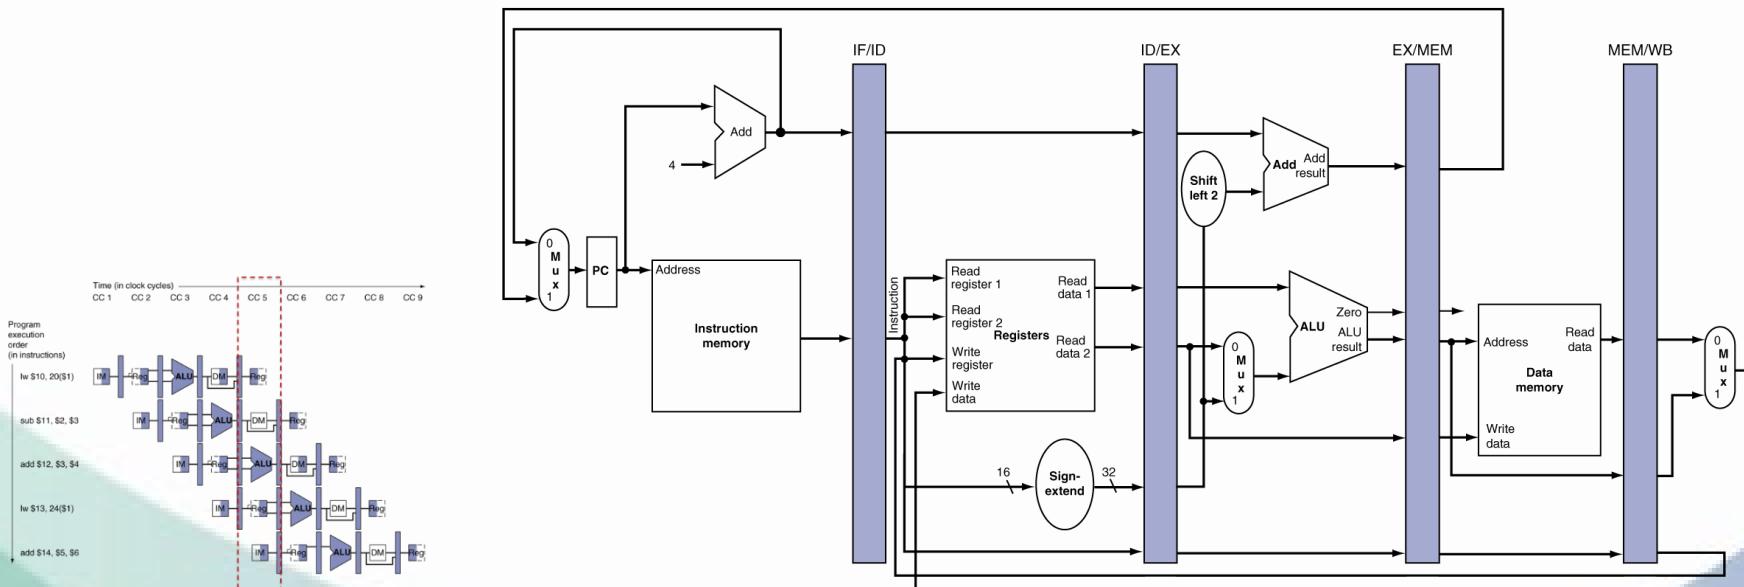
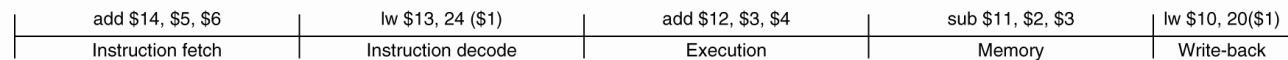
It's possible to calculate the branch address, and update the PC during ID stage with extra hardware added to test registers.

Pipeline Diagrams

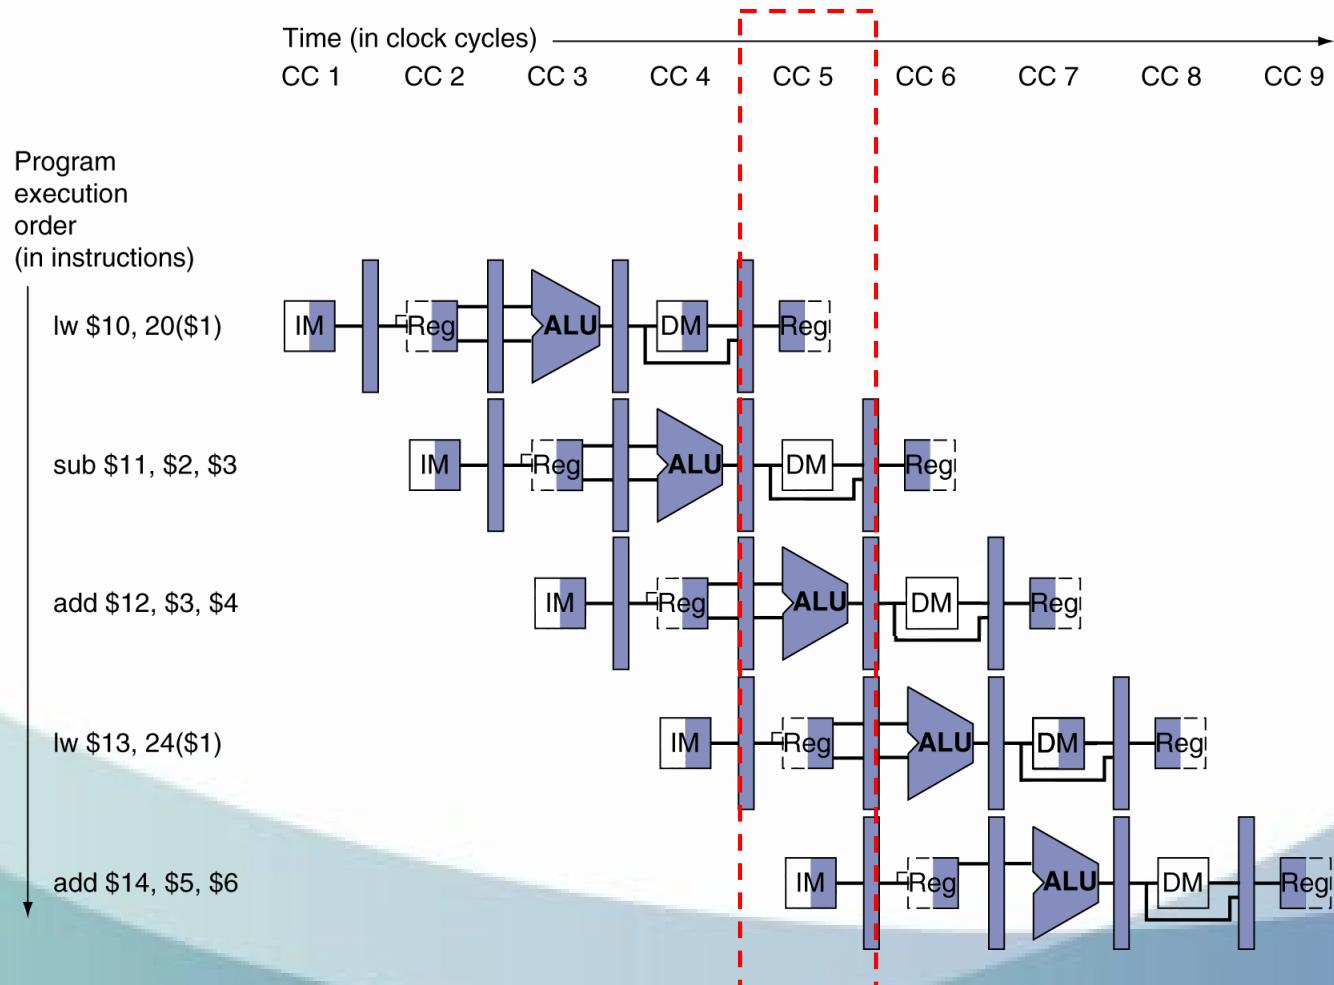
- Cycle-by-cycle flow of instructions through the pipelined datapath
 - “Single-clock-cycle” pipeline diagram
 - Shows pipeline usage in a single cycle
 - Highlight resources used
 - “multi-clock-cycle” diagram
 - Graph of operation over time

Single-Cycle Pipeline Diagram

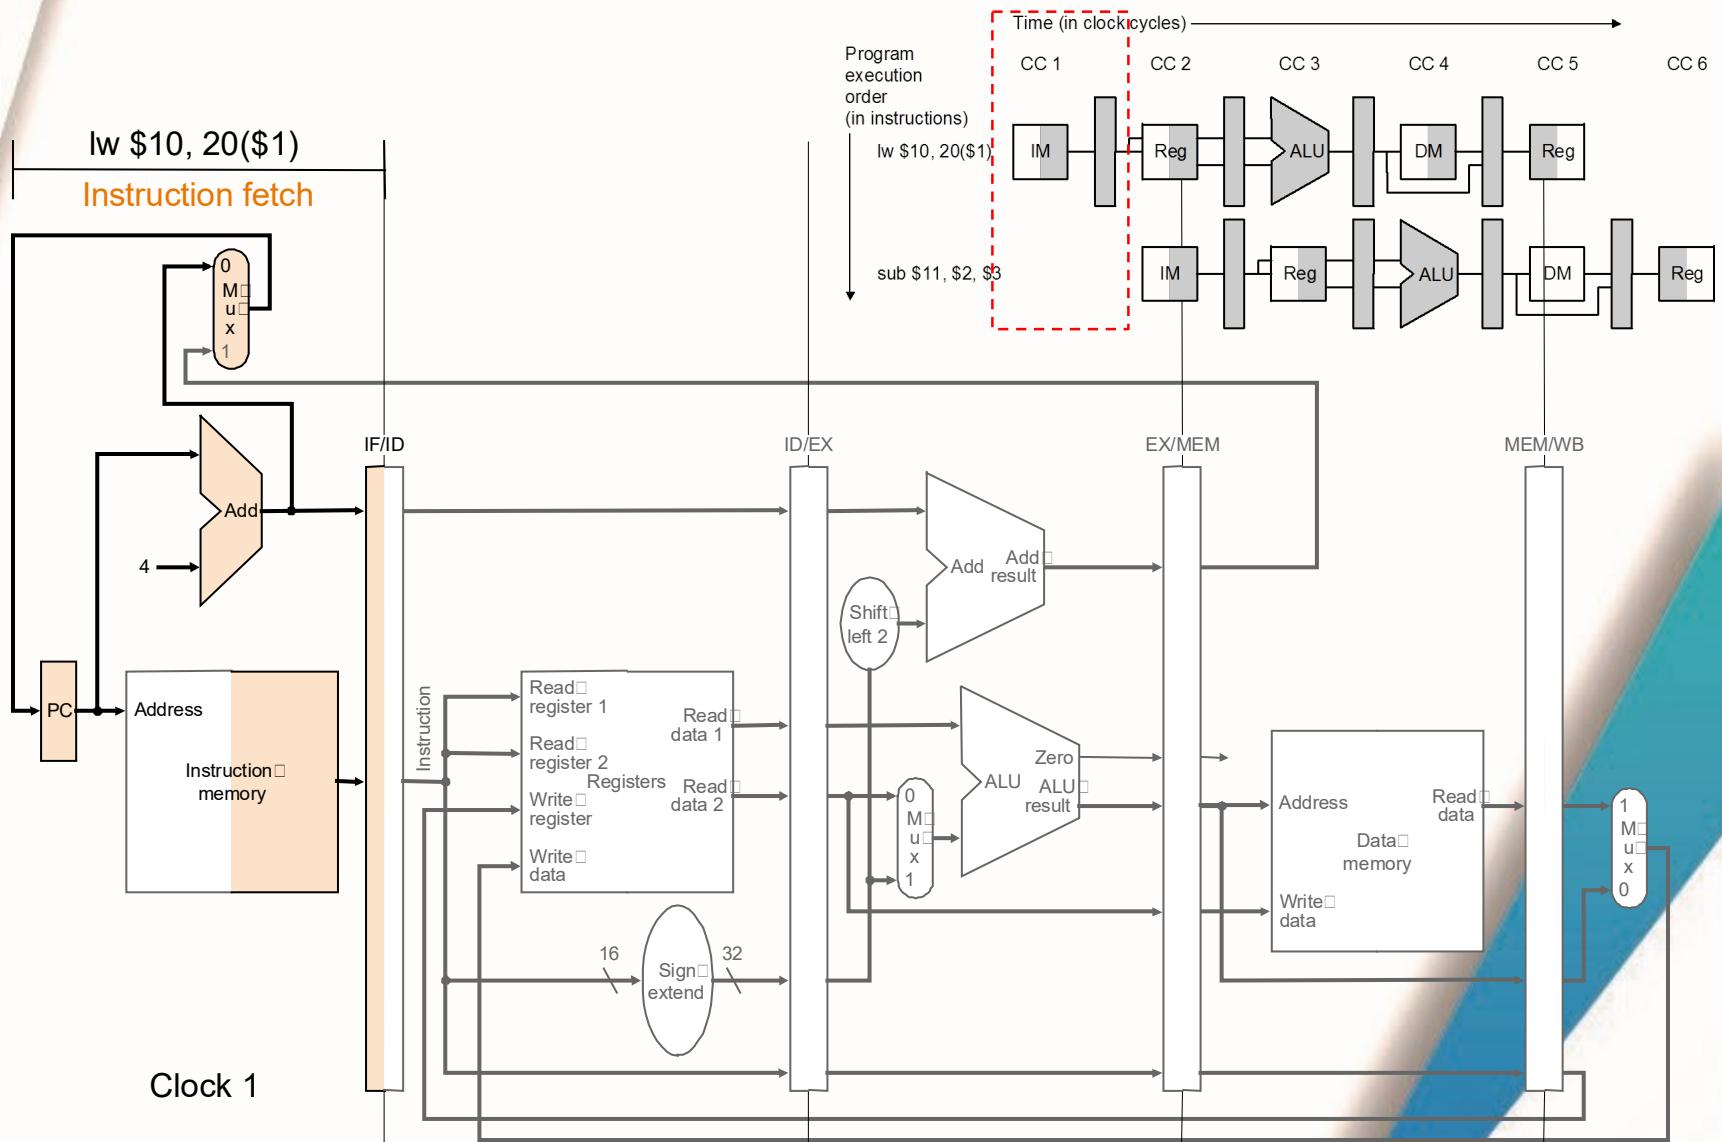
State of pipeline in a given cycle



Multi-Cycle Pipeline Diagram

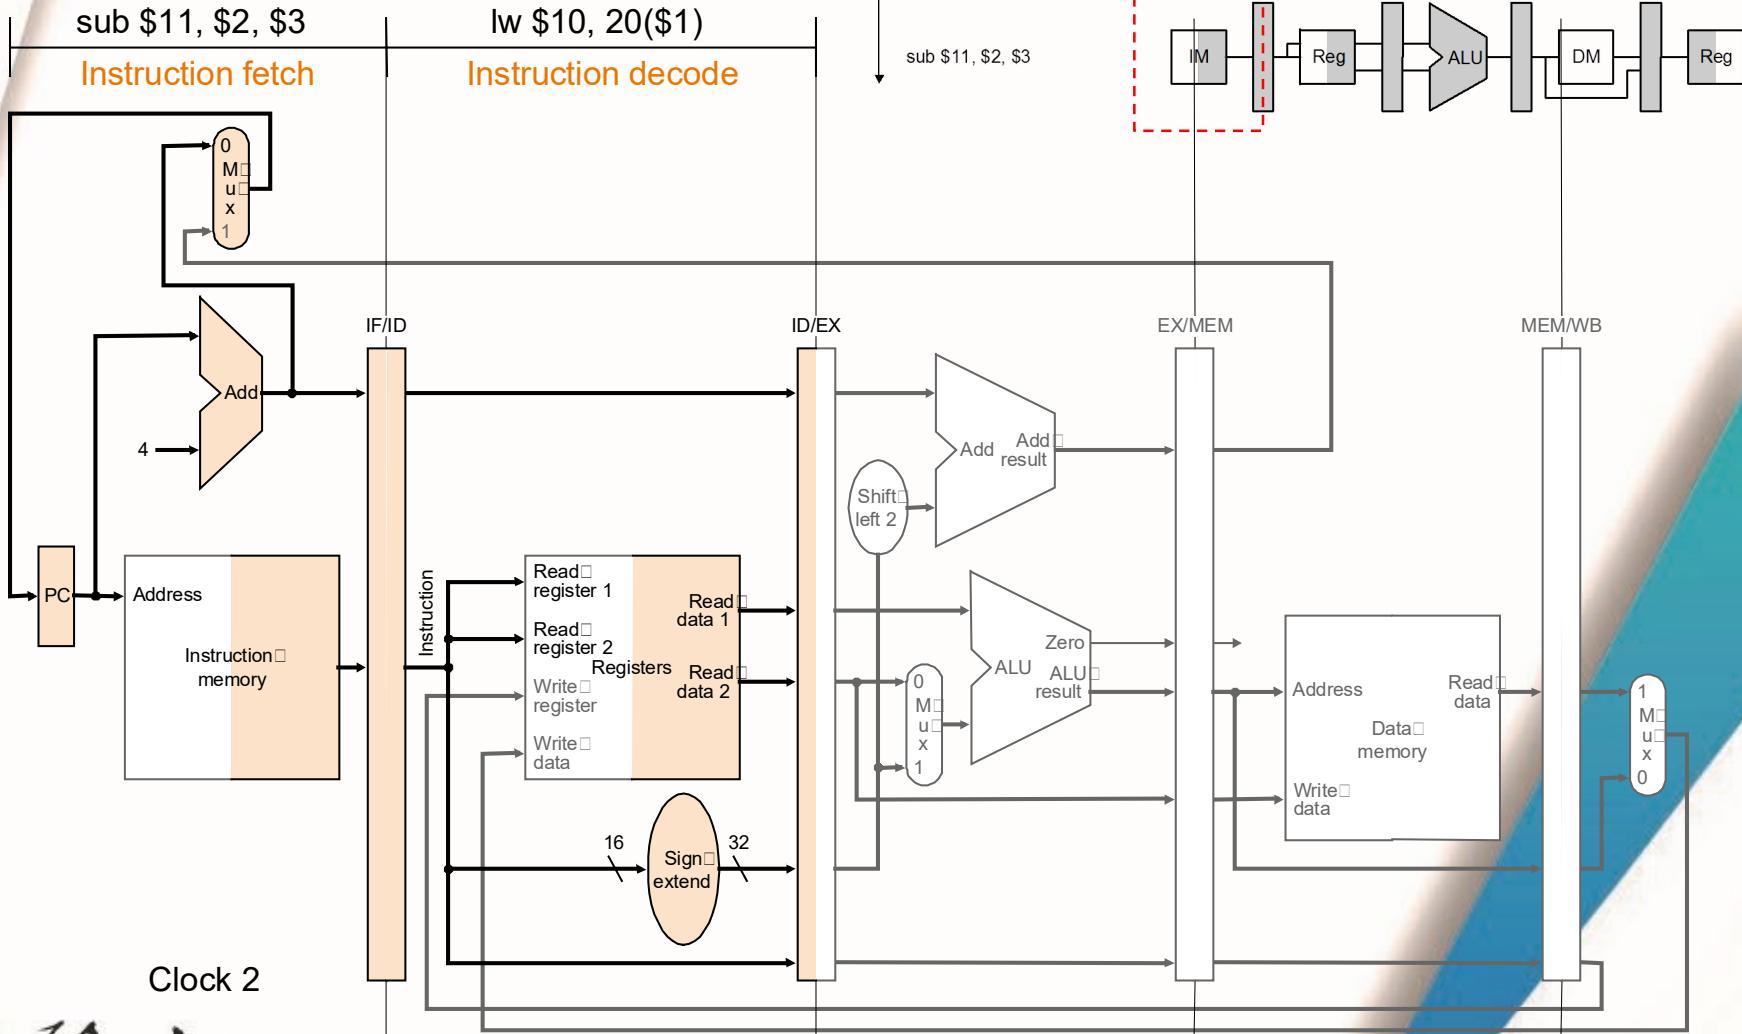


Example – Cycle 1

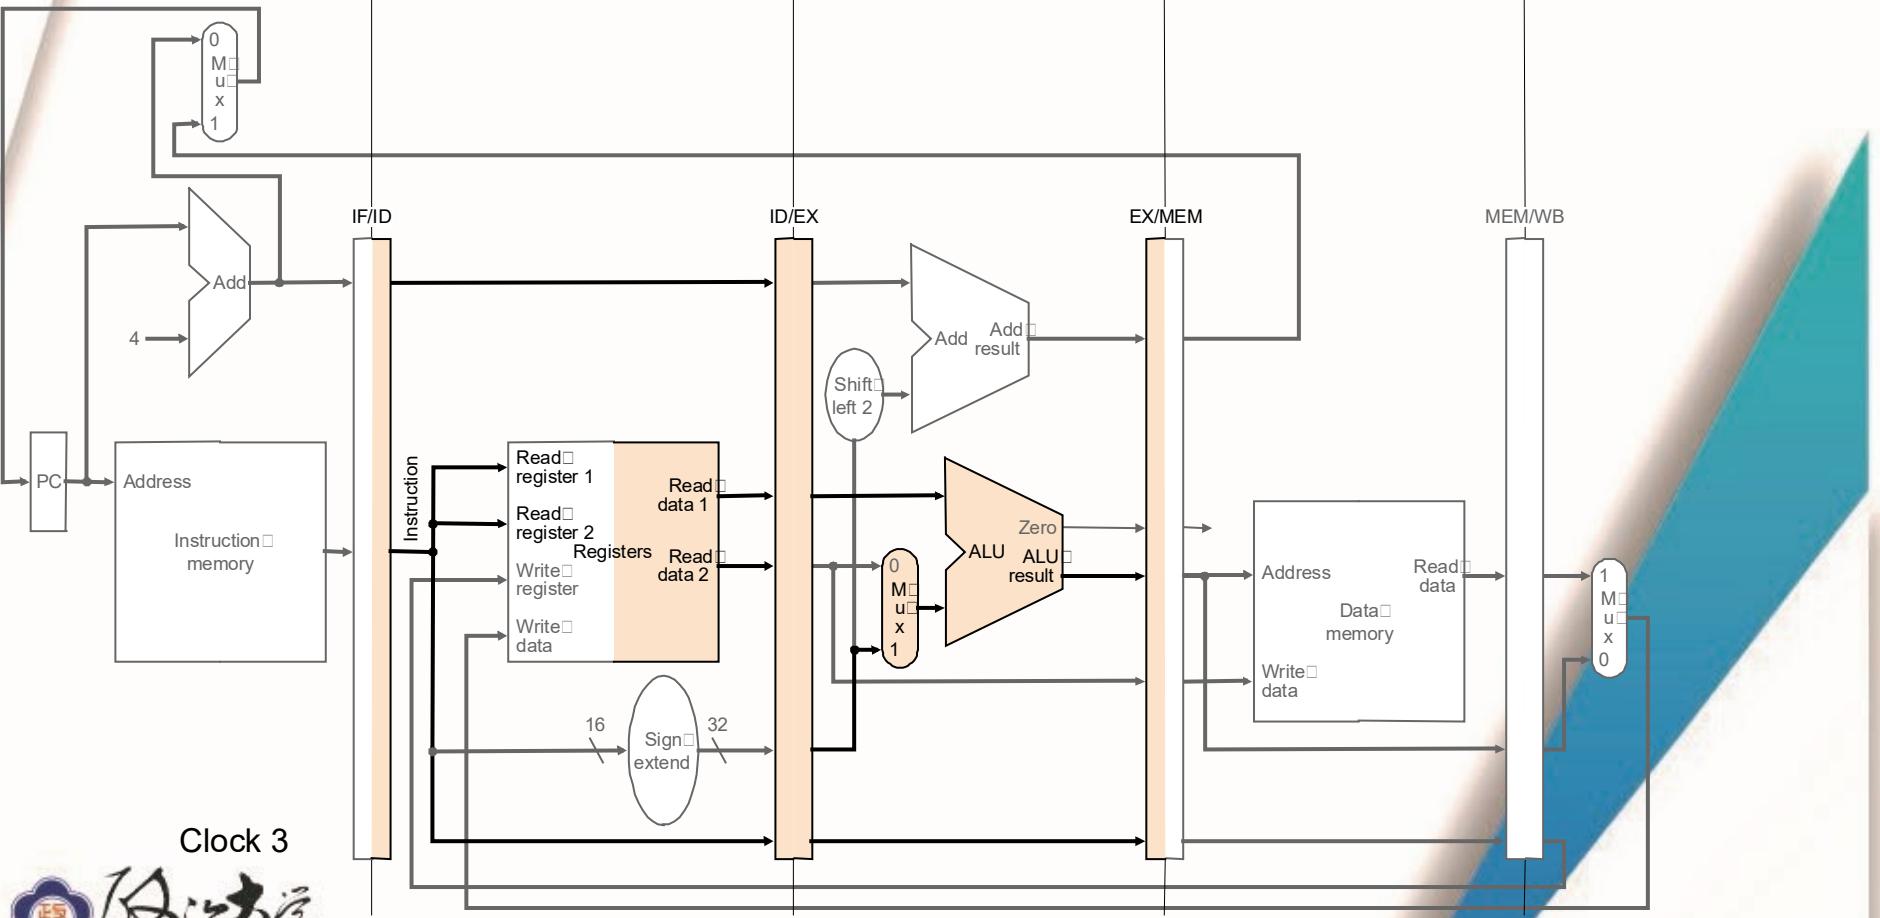


Clock 1

Cycle 2

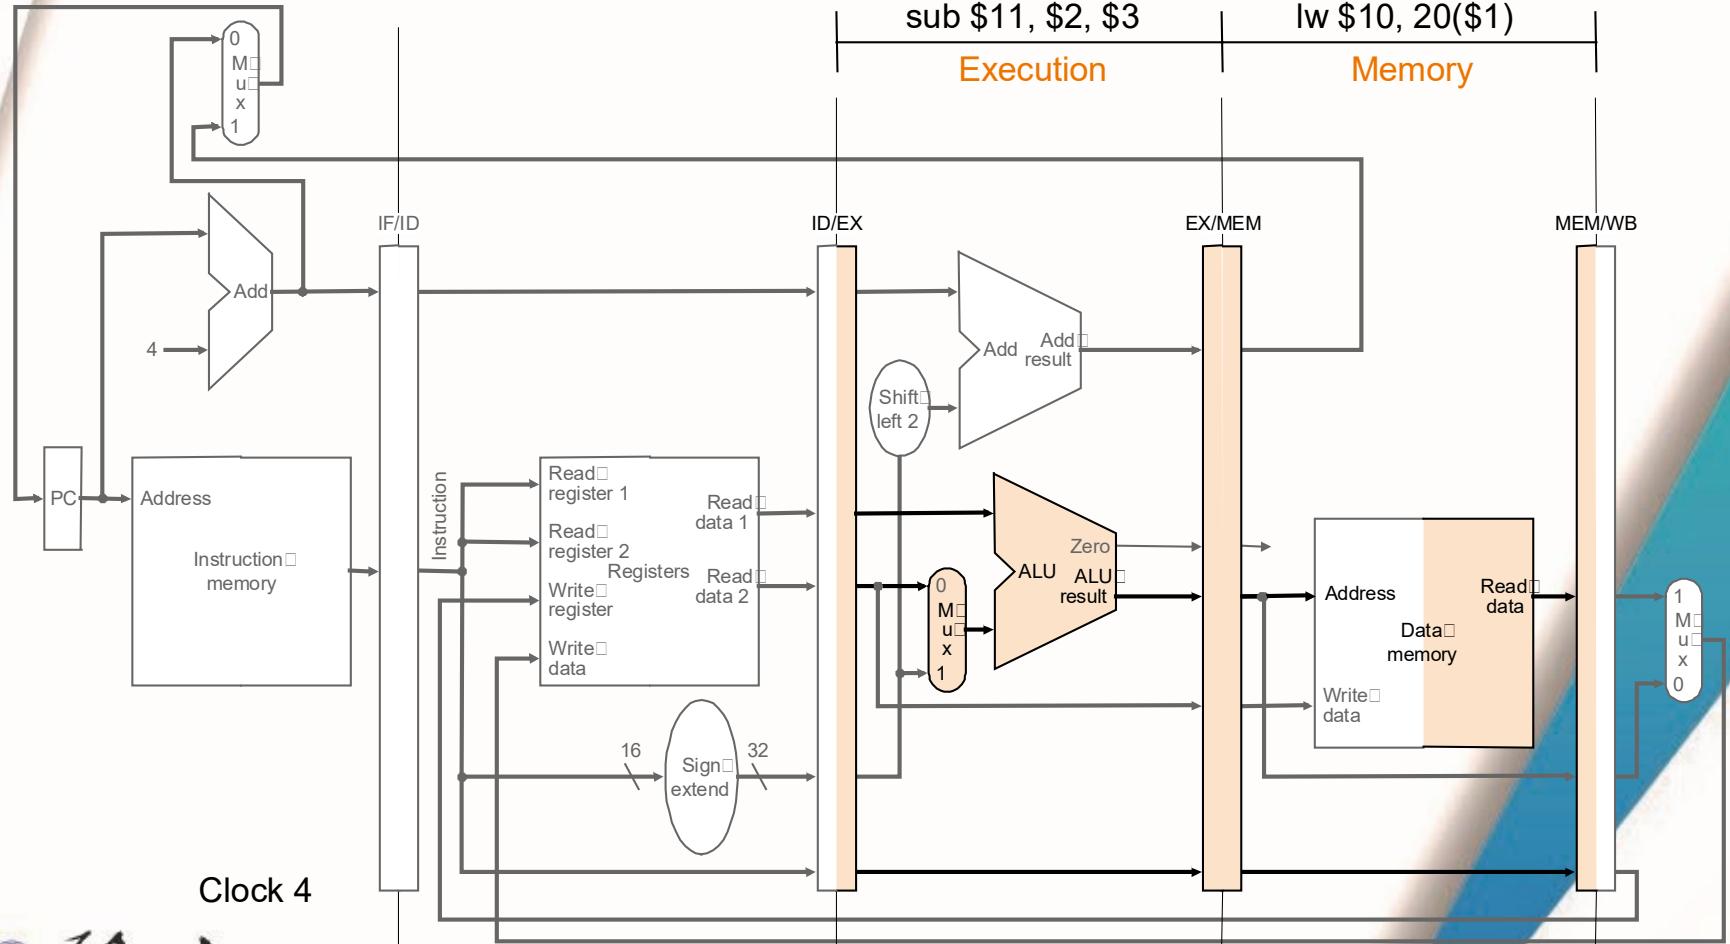
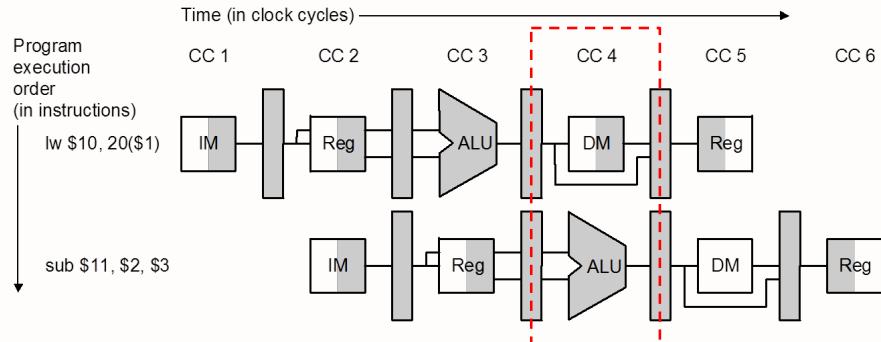


Cycle 3

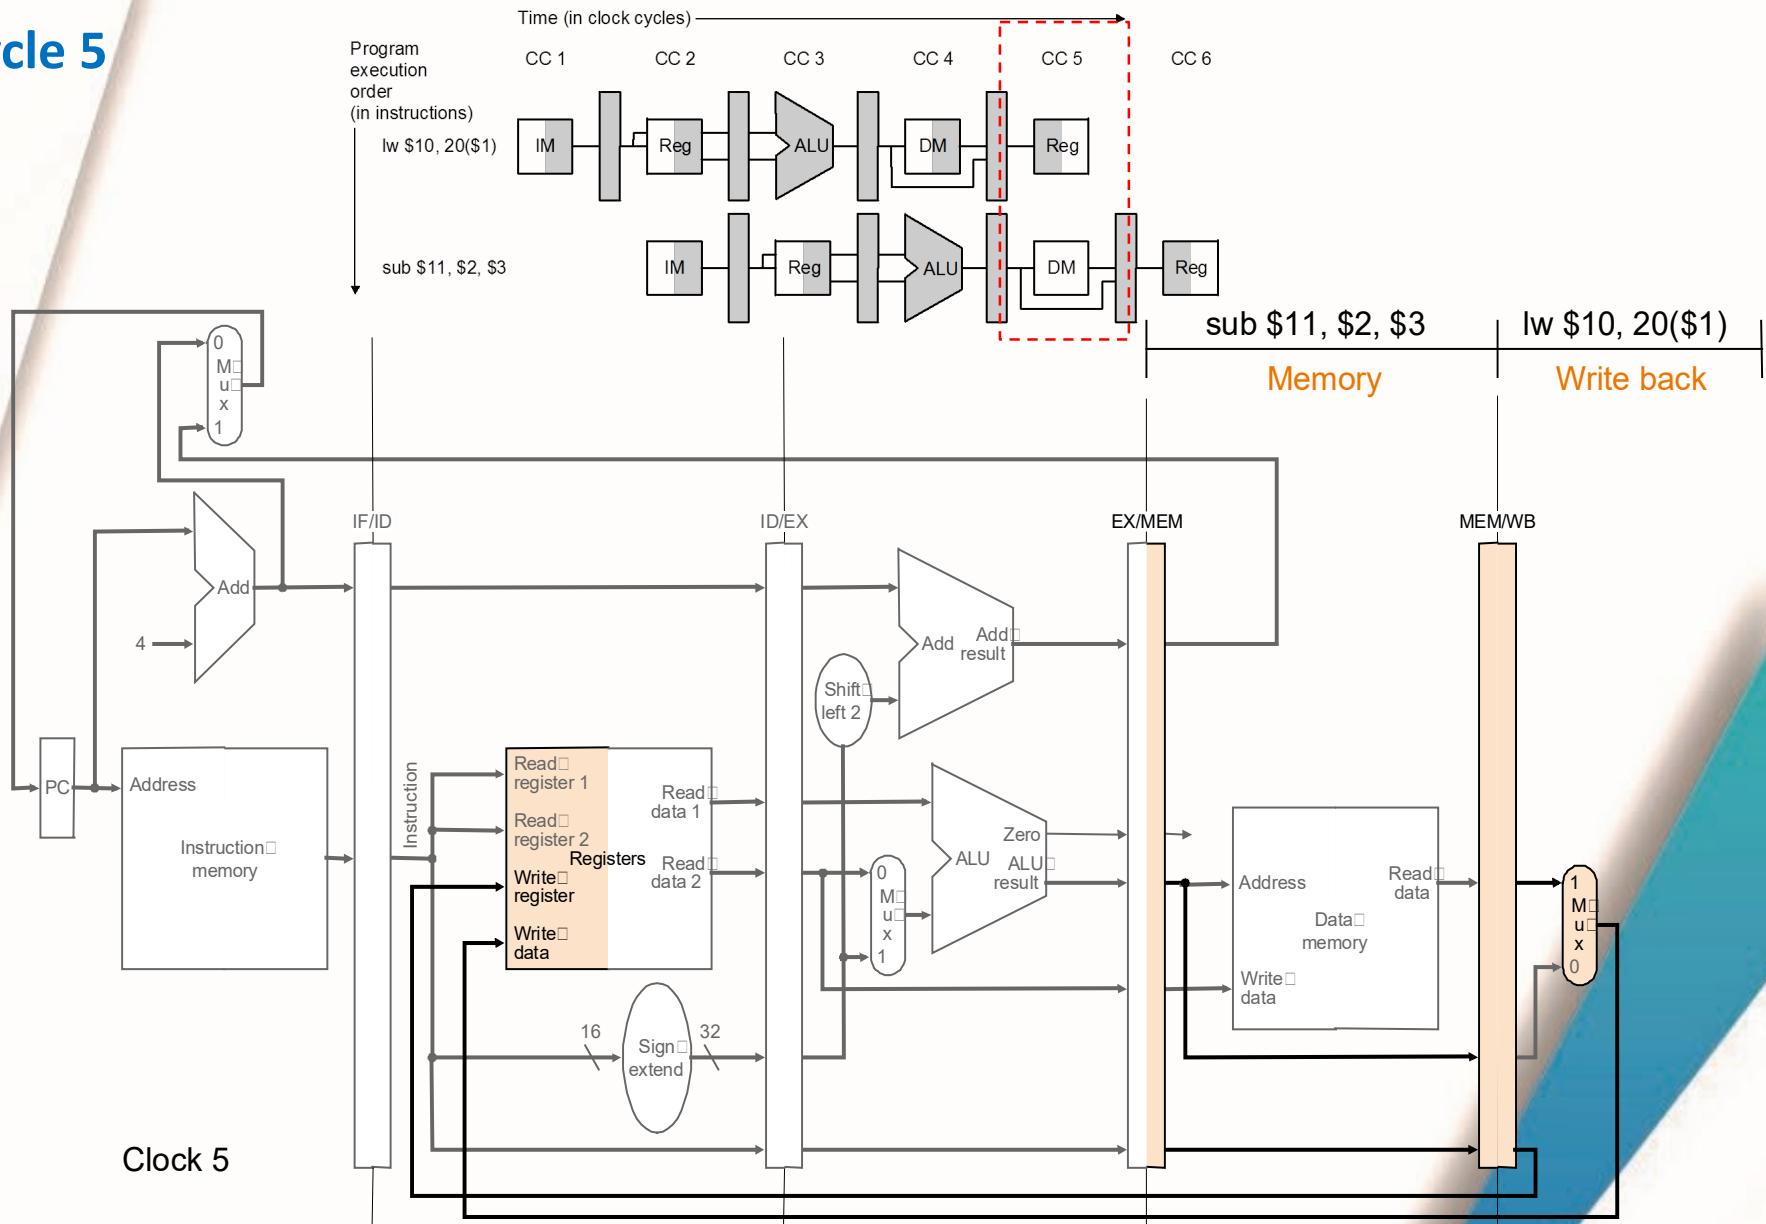


Clock 3

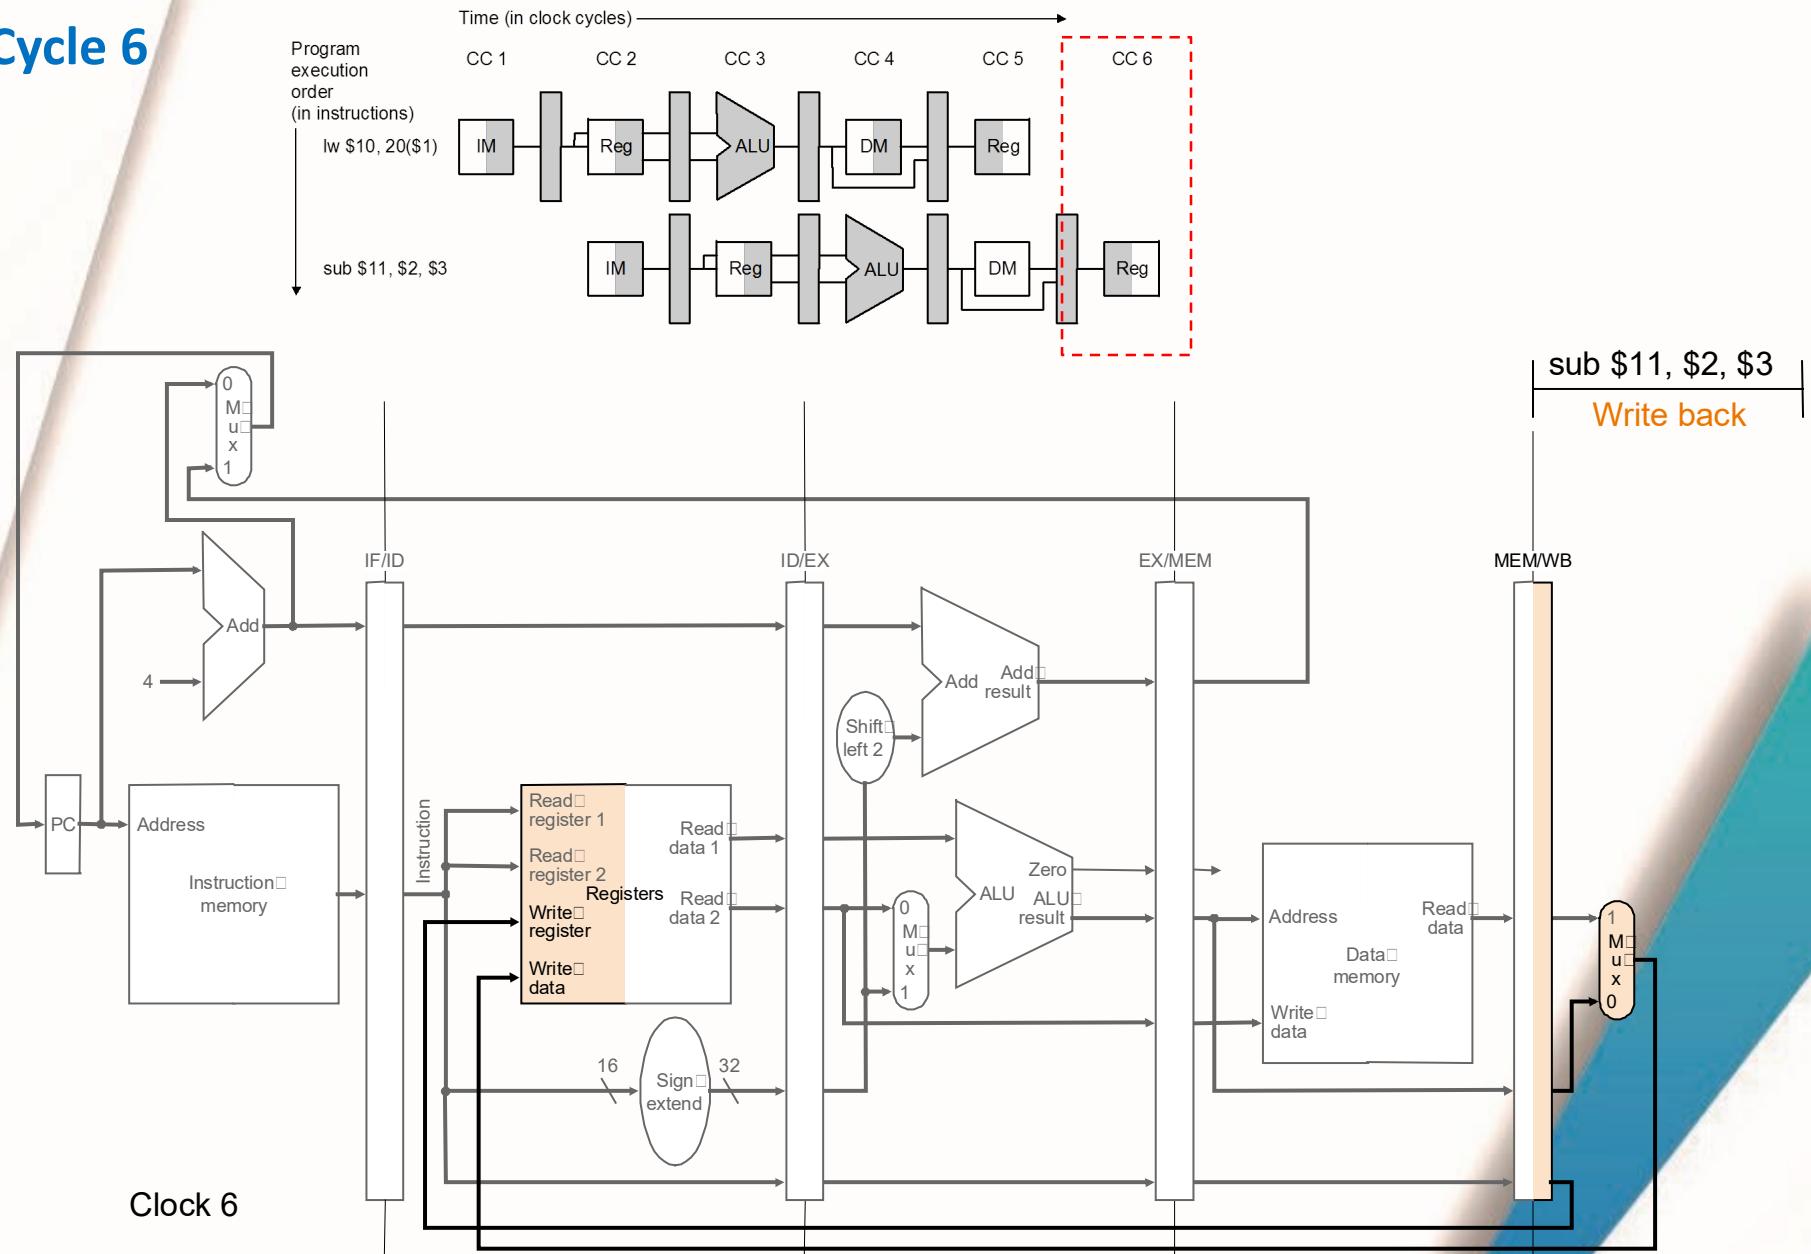
Cycle 4



Cycle 5

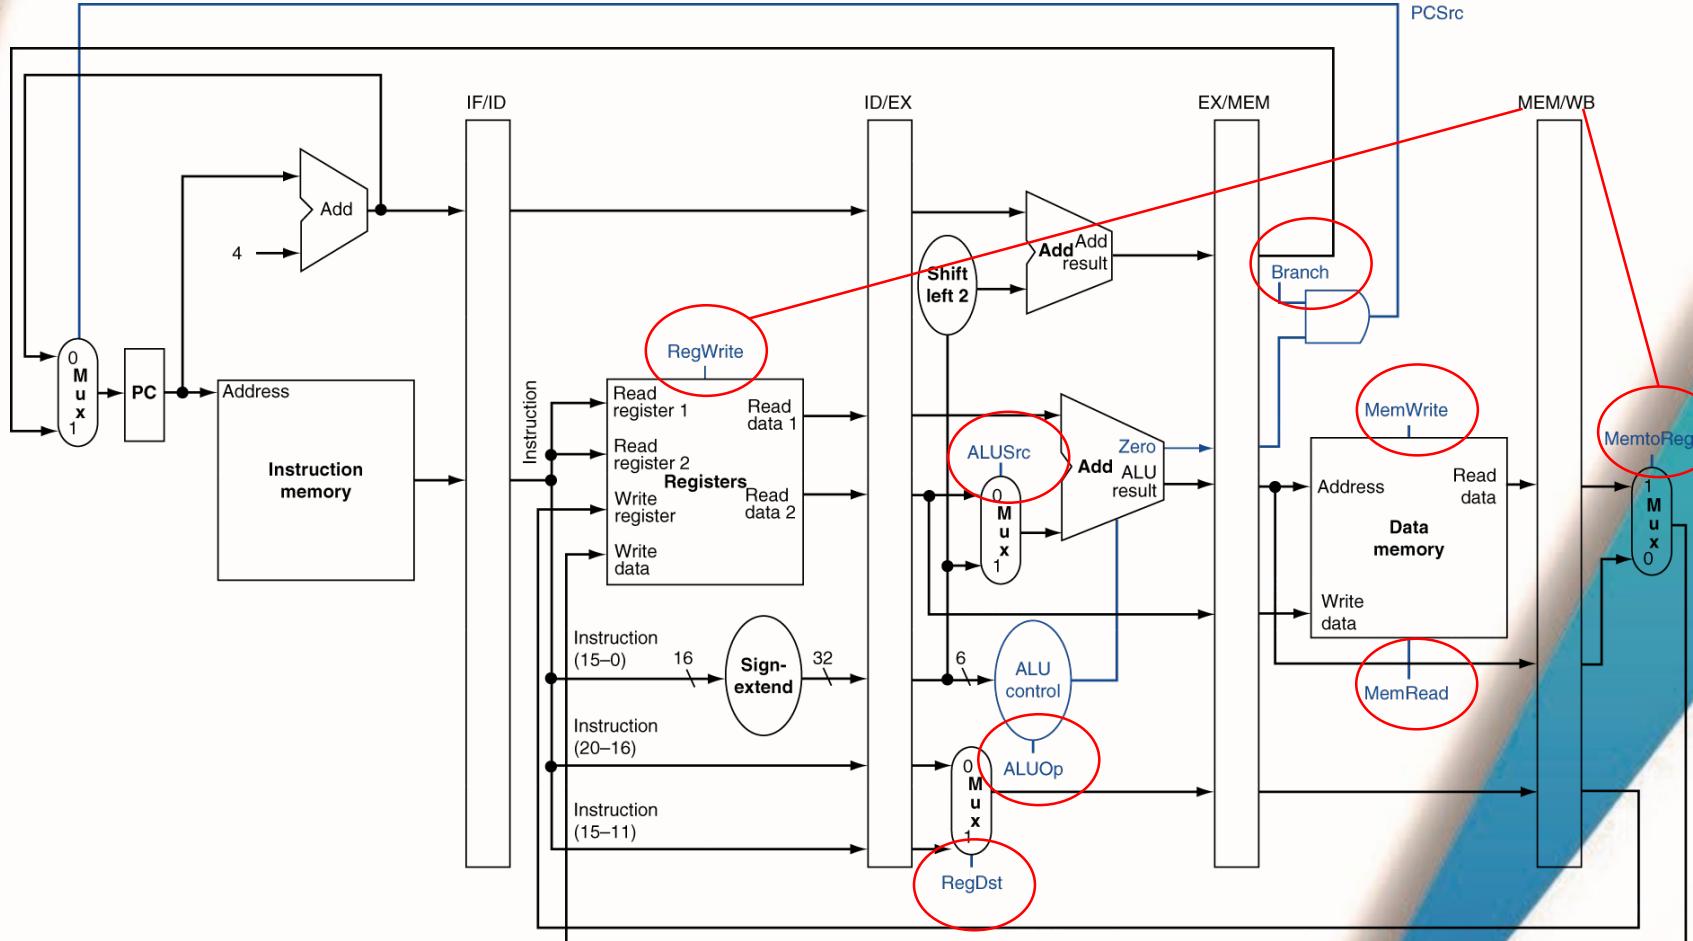


Cycle 6



Pipelined Control (Simplified)

Ignoring branch instructions here for now



EX, MEM, and WB use control signals since all the signals are decoded in ID stage

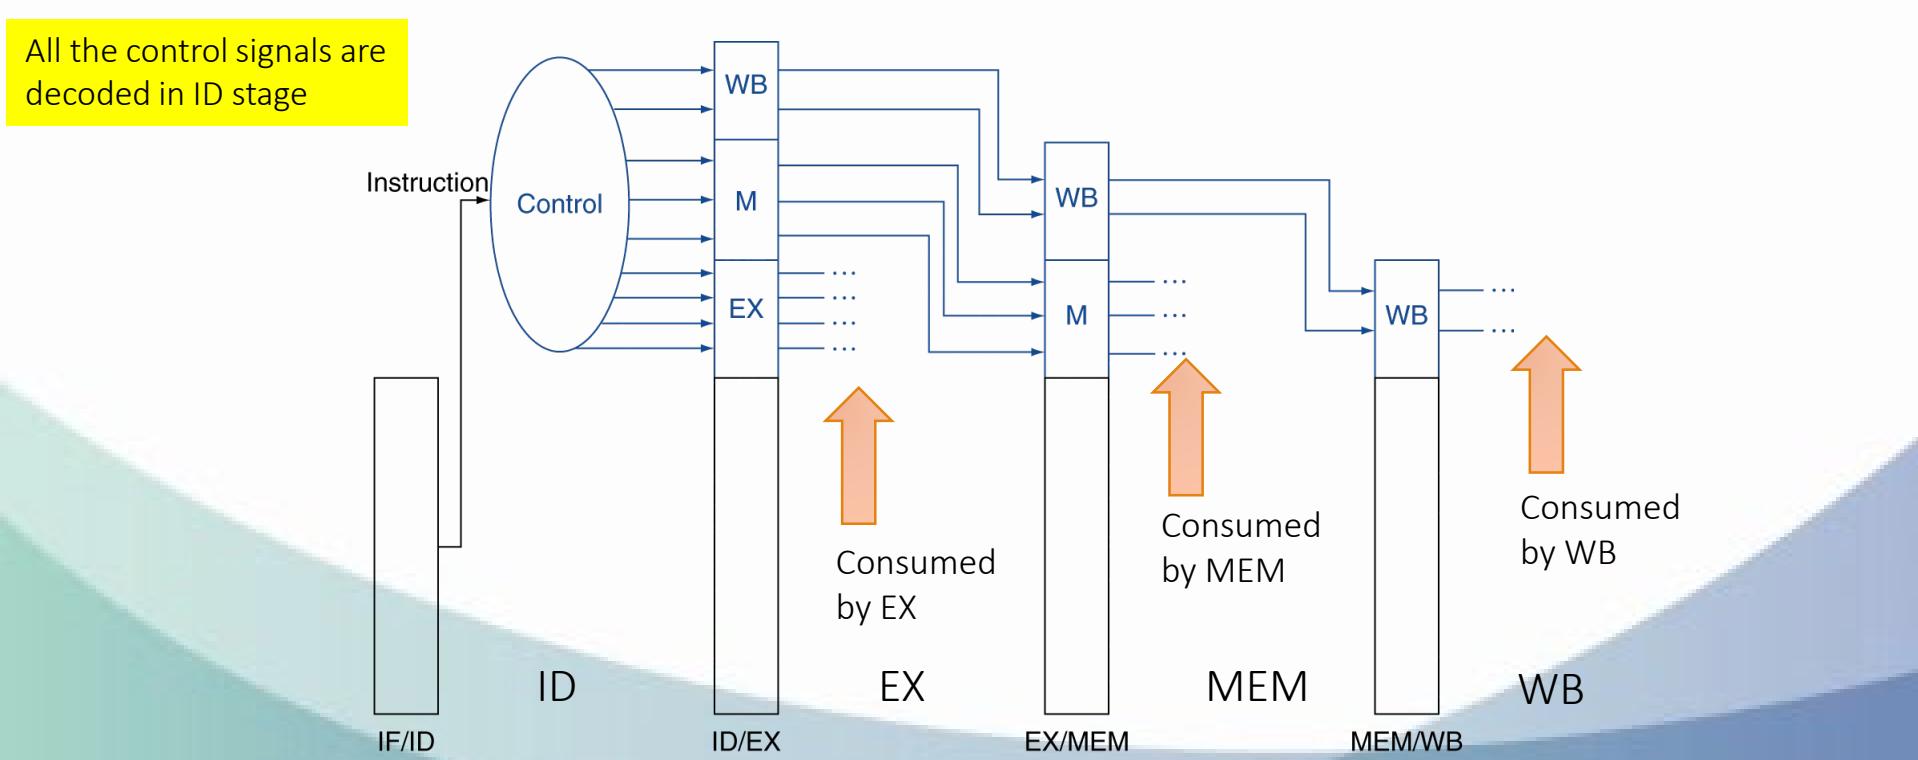
Group Signals According to Stages

- Use the same control signals as those of a single-cycle CPU
- Group them in three categories

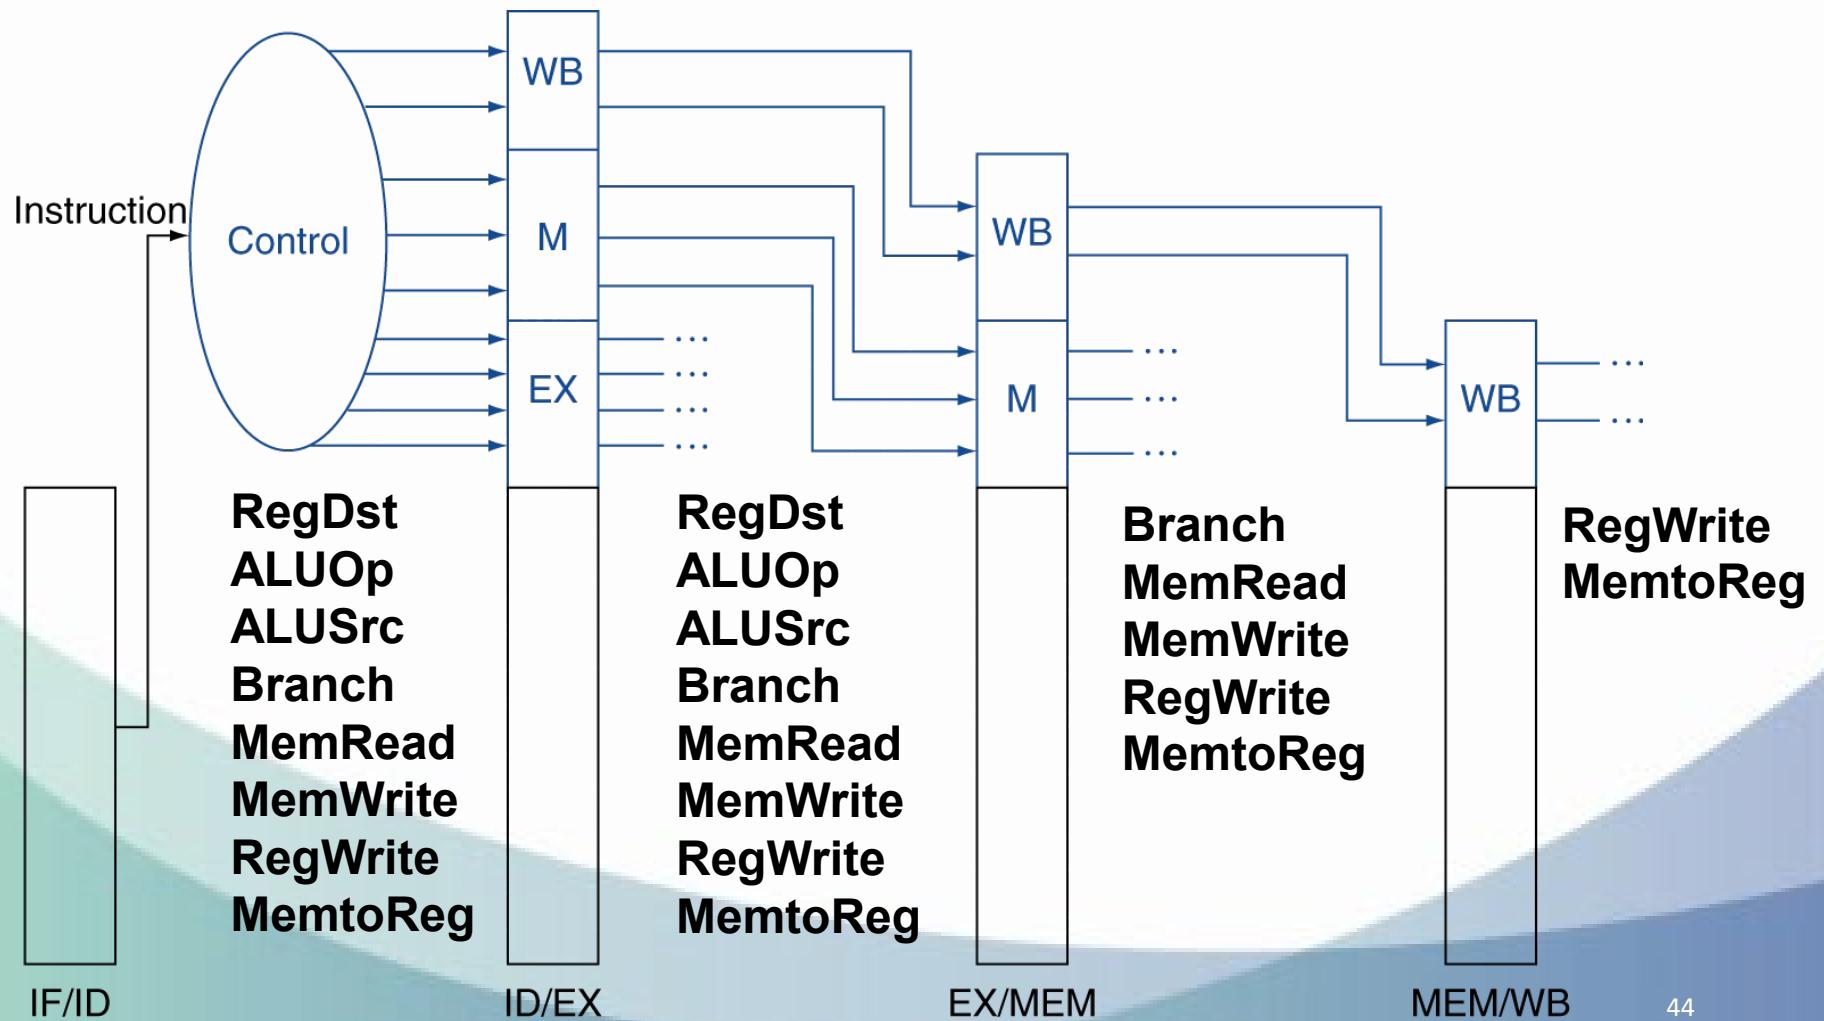
	Execution/Address Calculation stage control lines				Memory access stage control lines			Write-back stage control lines	
	Reg Dst	ALU Op1	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg write	Mem to Reg
R-type	1	1	0	0	0	0	0	1	0
	0	0	0	1	0	1	0	1	1
	X	0	0	1	0	0	1	0	X
	X	0	1	0	1	0	0	0	X

Data Stationary Control

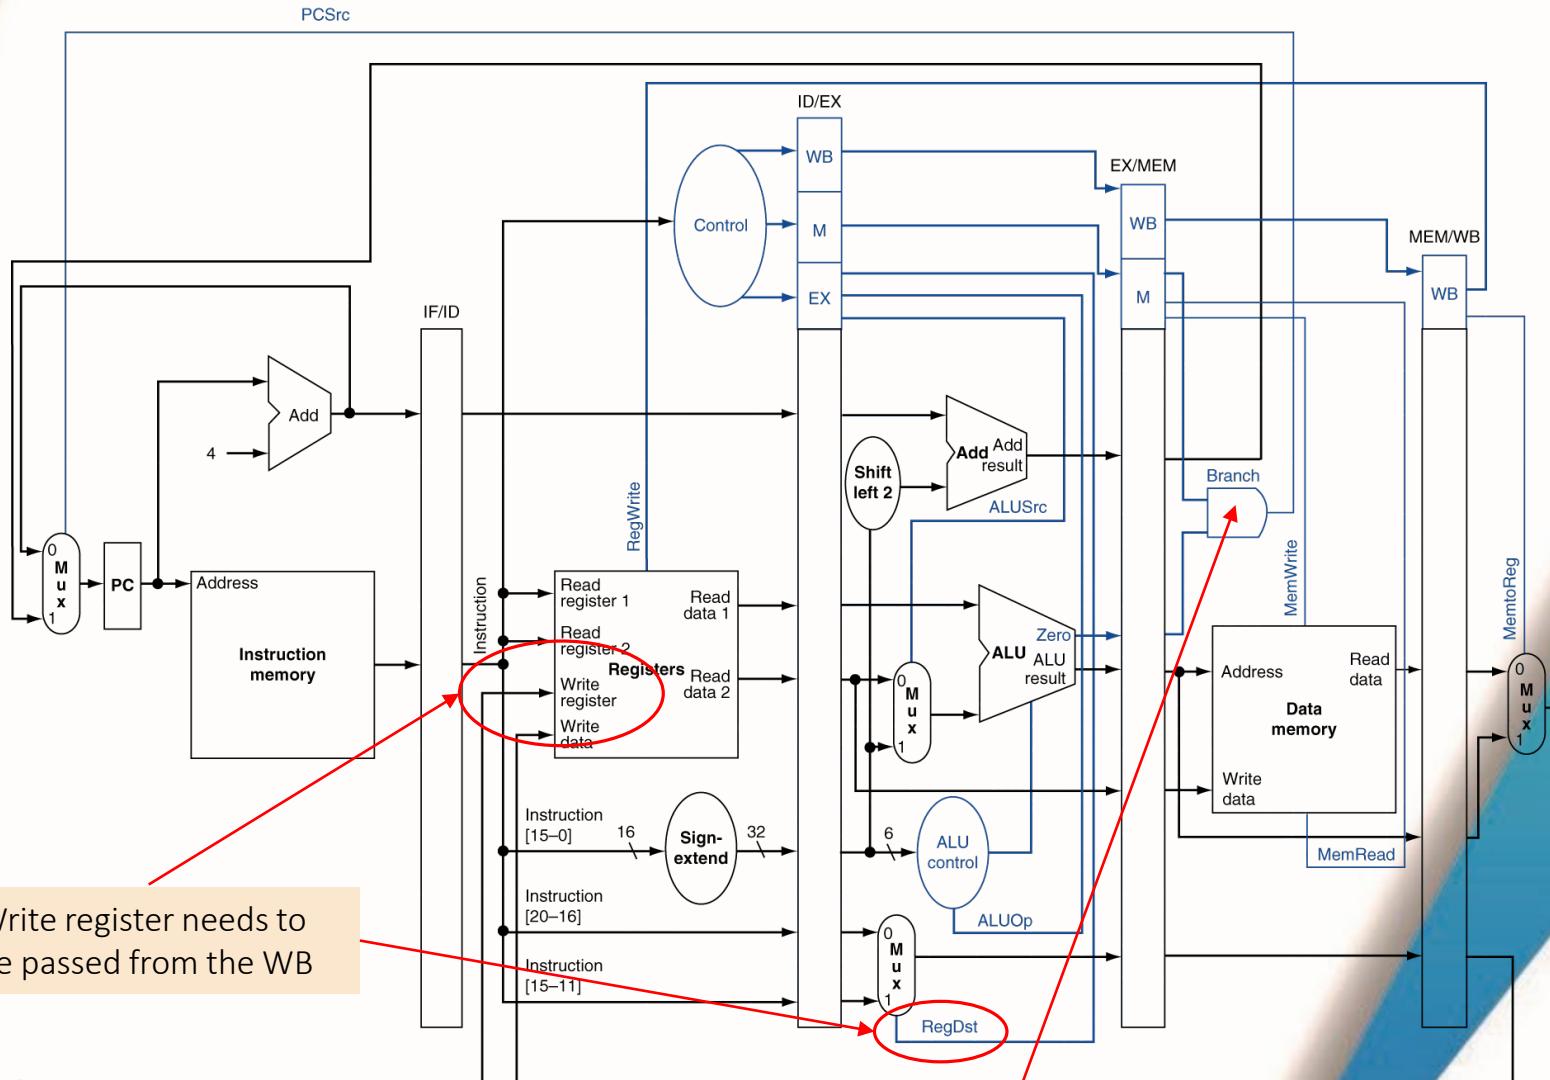
- Control signals derived from instruction
 - As in single-cycle implementation
 - These signals need to be stored into pipeline register



Data Stationary Control



Pipelined Control



Write register needs to be passed from the WB

Branch could be determined in ID stage

Example

lw \$10, 20(\$1)

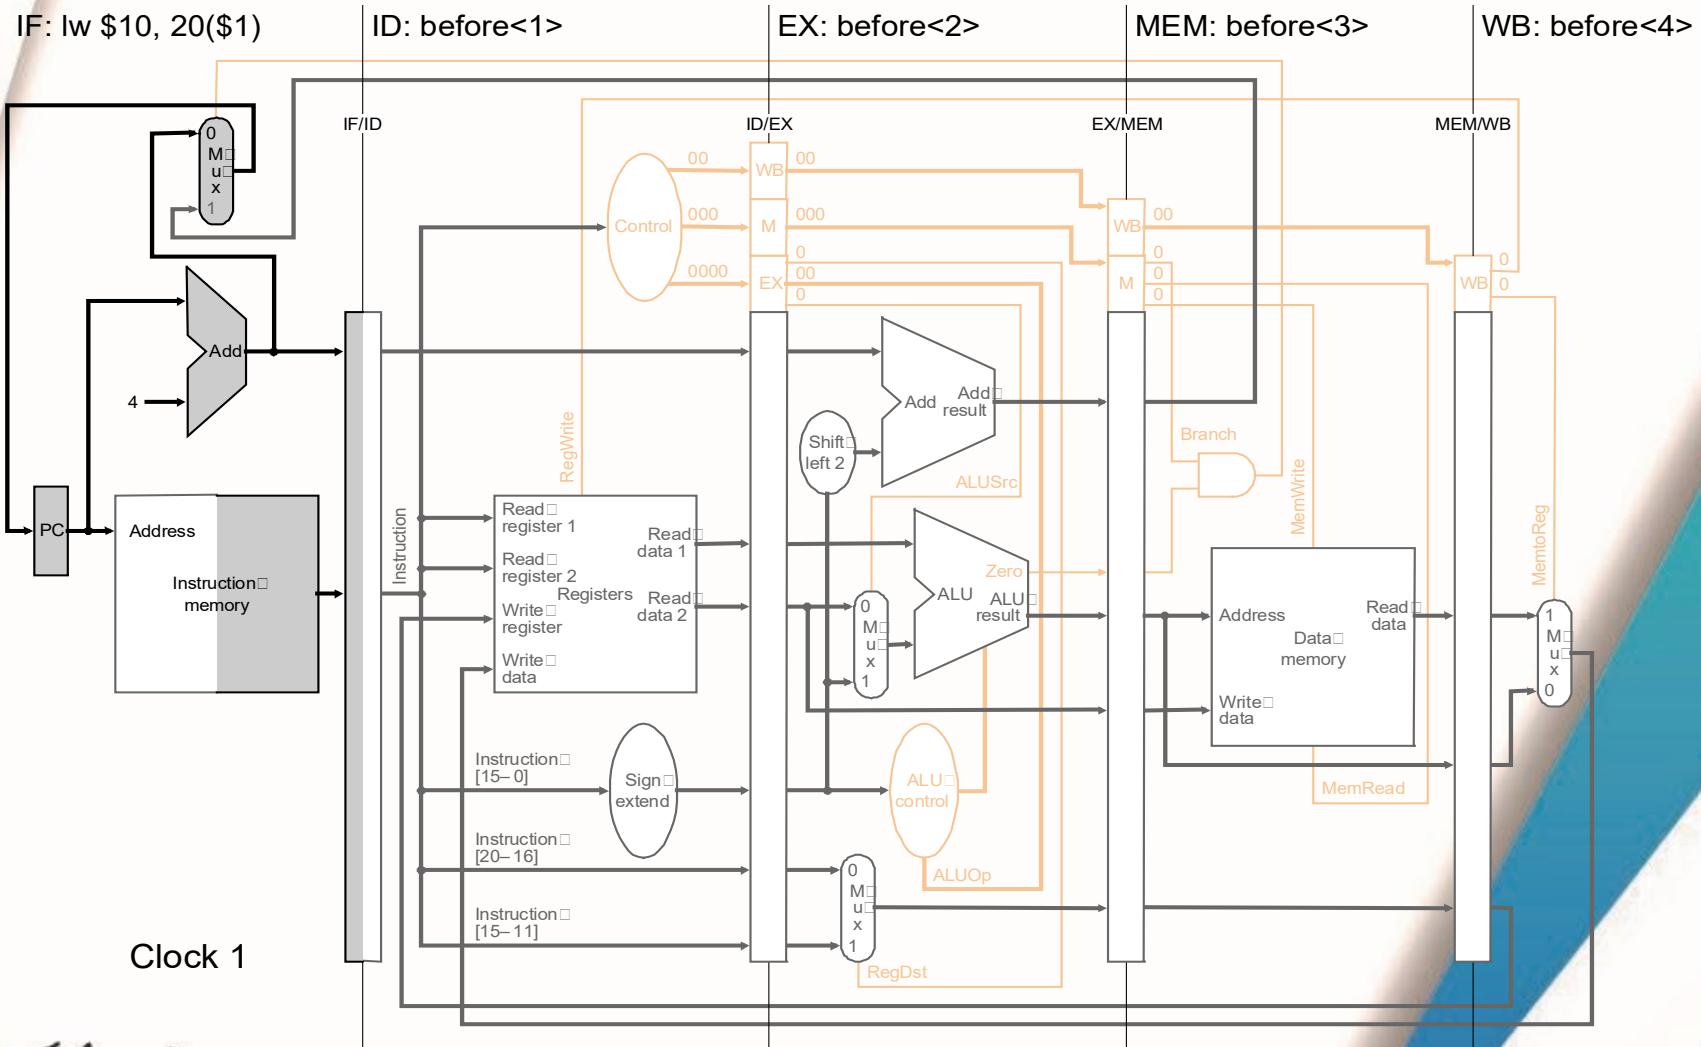
sub \$11, \$2, \$3

and \$12, \$4, \$5

or \$13, \$6, \$7

add \$14, \$8, \$9

Cycle 1

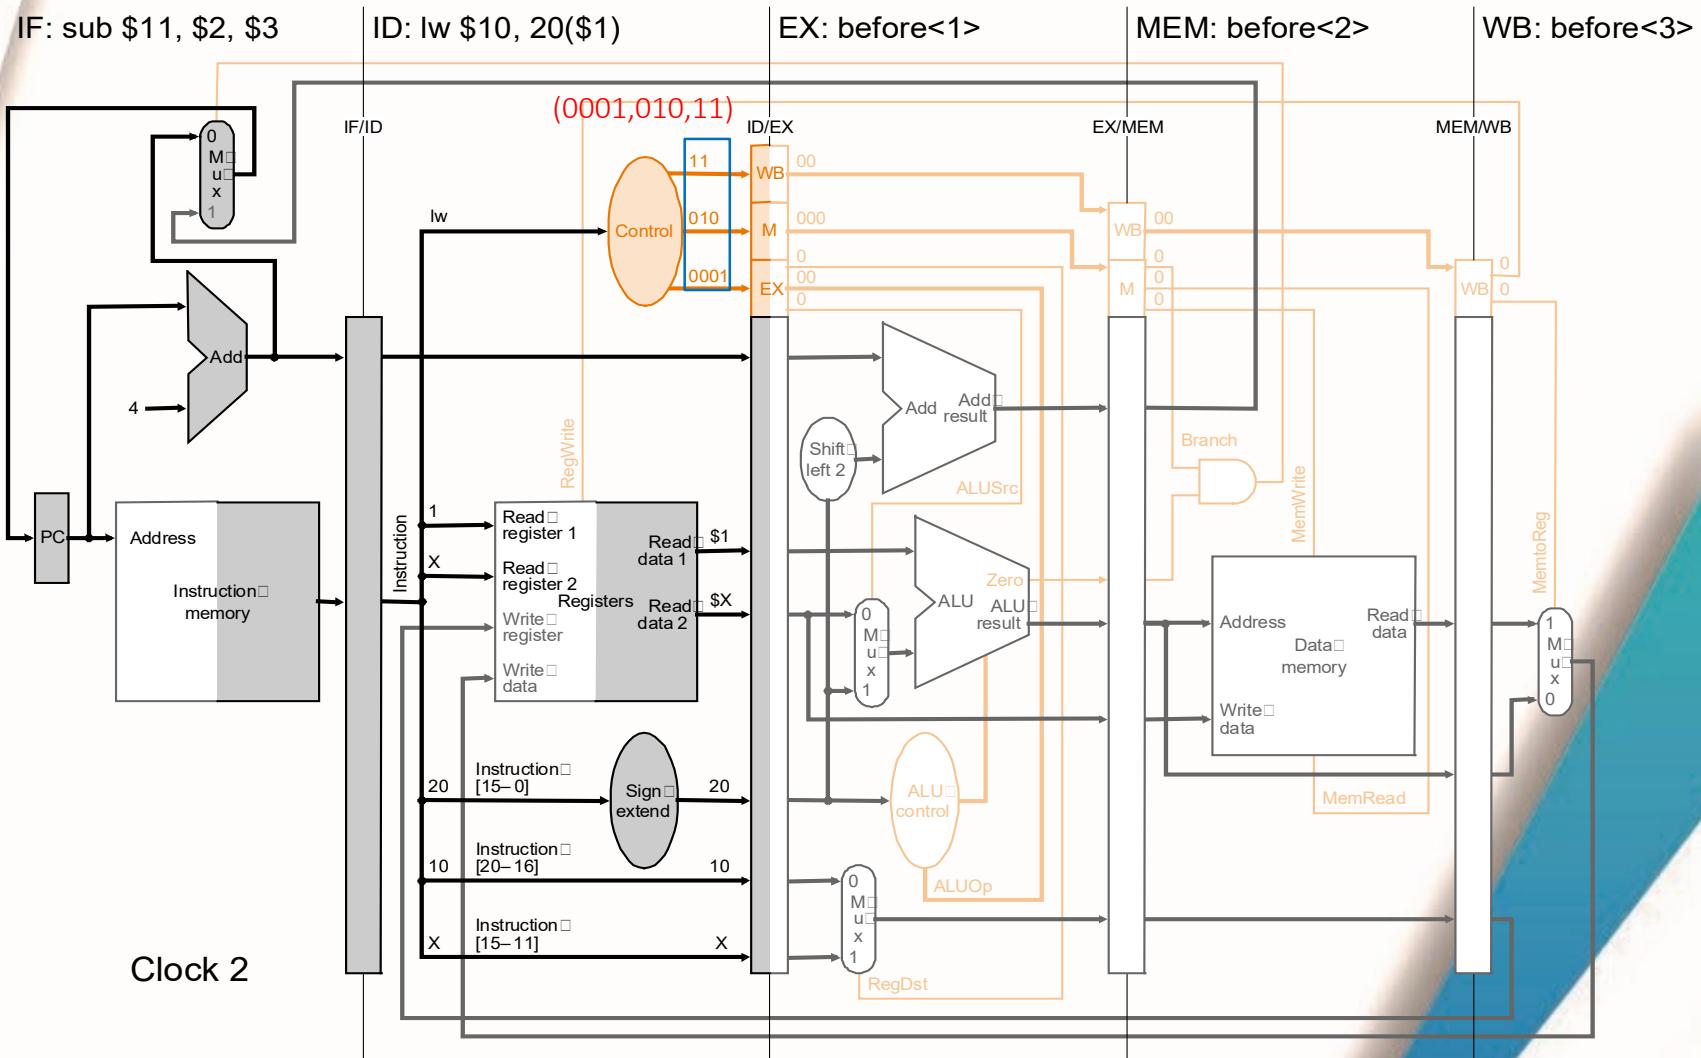


Group Signals According to Stages

Iw \$10, 20(\$1)

		Execution/Address Calculation stage control lines				Memory access stage control lines			Write-back stage control lines	
		Reg Dst	ALU Op1	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg write	Mem to Reg
R-type lw	1	1	0	0	0	0	0	0	1	0
	0	0	0	1	0	1	0	1	1	1
	X	0	0	1	0	0	1	0	0	X
	X	0	1	0	1	0	0	0	0	X

Cycle 2

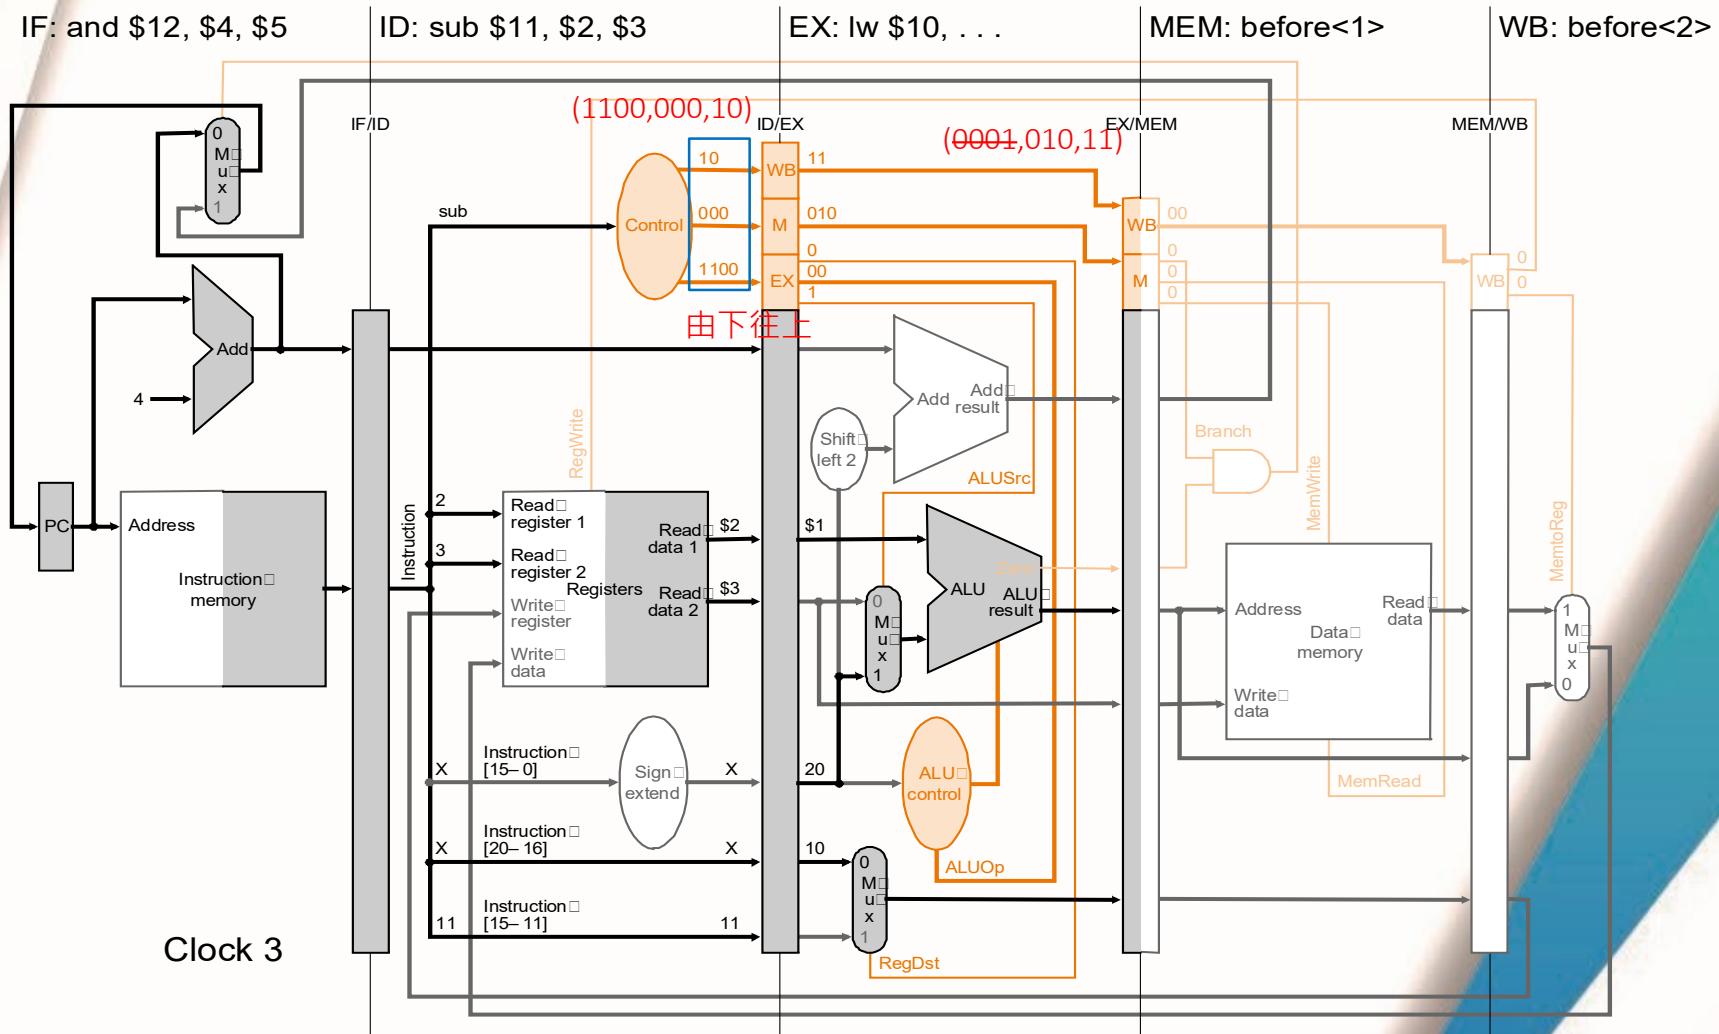


Group Signals According to Stages

sub \$11, \$2, \$3

Execution/Address Calculation stage control lines				Memory access stage control lines			Write-back stage control lines		
	Reg Dst	ALU Op1	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg write	Mem to Reg
R-type	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
sw	X	0	0	1	0	0	1	0	X
beq	X	0	1	0	1	0	0	0	X

Cycle 3

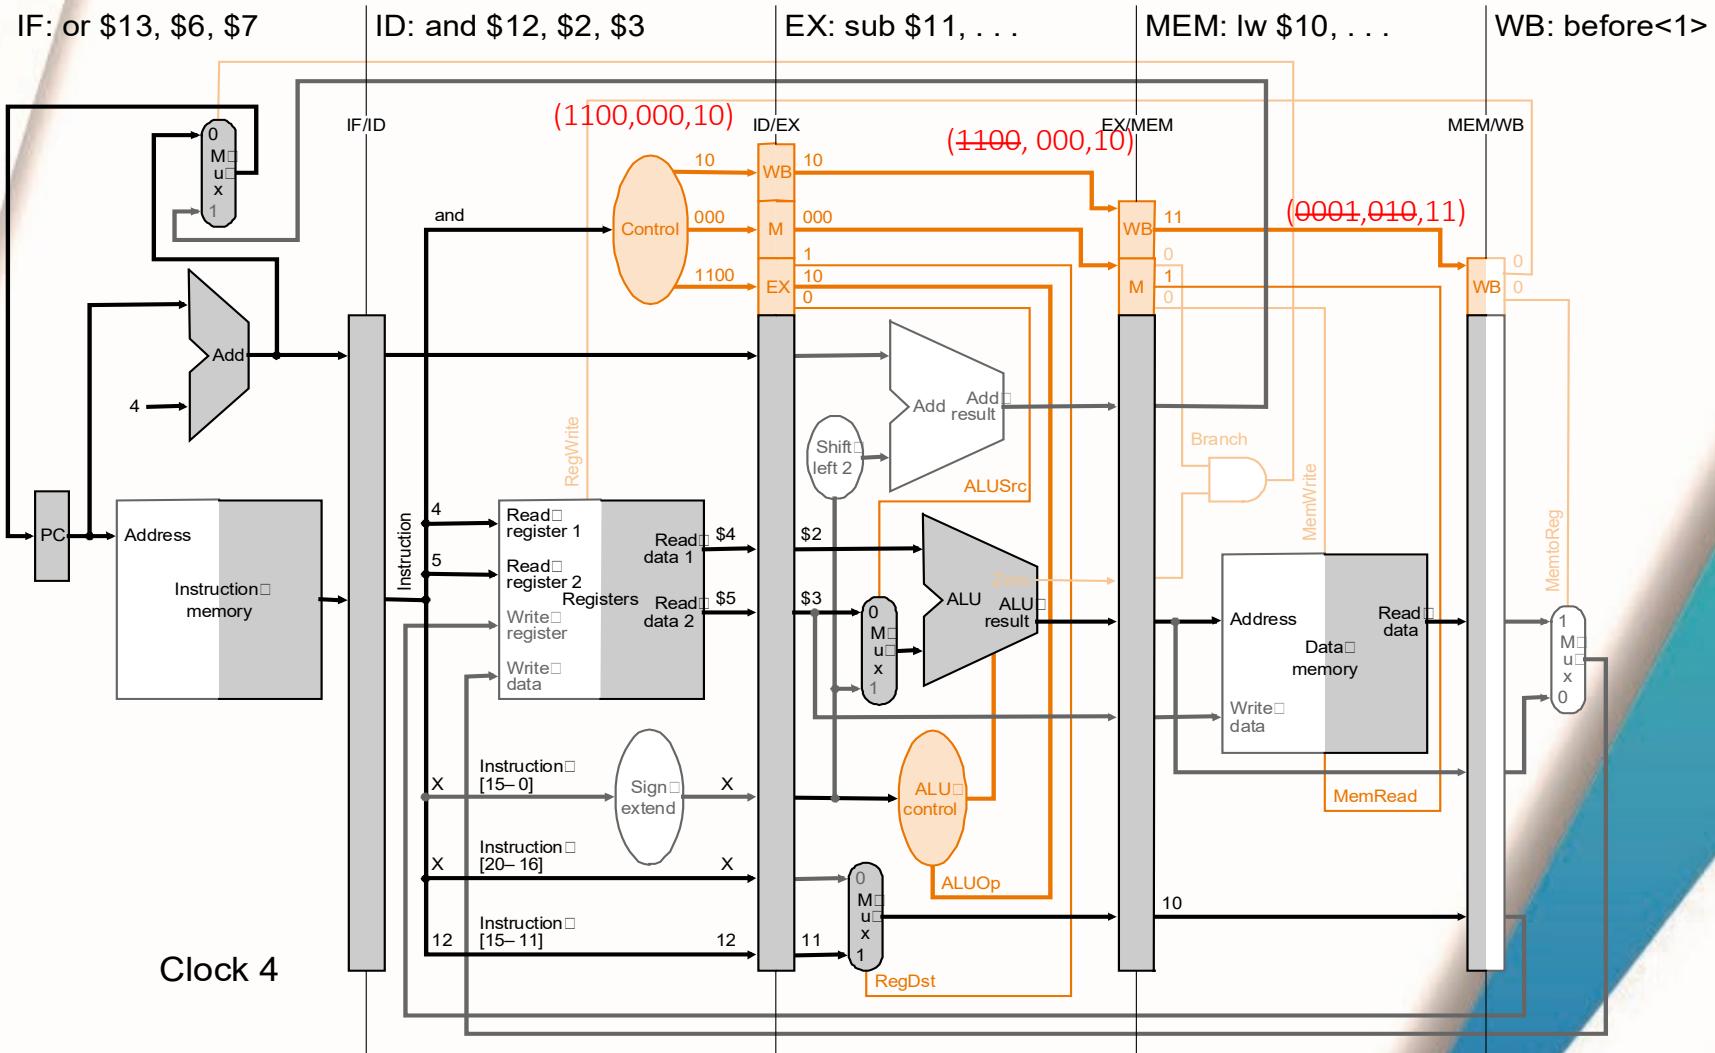


Group Signals According to Stages

and \$12, \$4, \$5

Execution/Address Calculation stage control lines				Memory access stage control lines			Write-back stage control lines	
Reg Dst	ALU Op1	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg write	Mem to Reg
R-type	1	1	0	0	0	0	1	0
Iw	0	0	0	1	0	1	1	1
sw	X	0	0	1	0	0	1	0
beq	X	0	1	0	1	0	0	X

Cycle 4

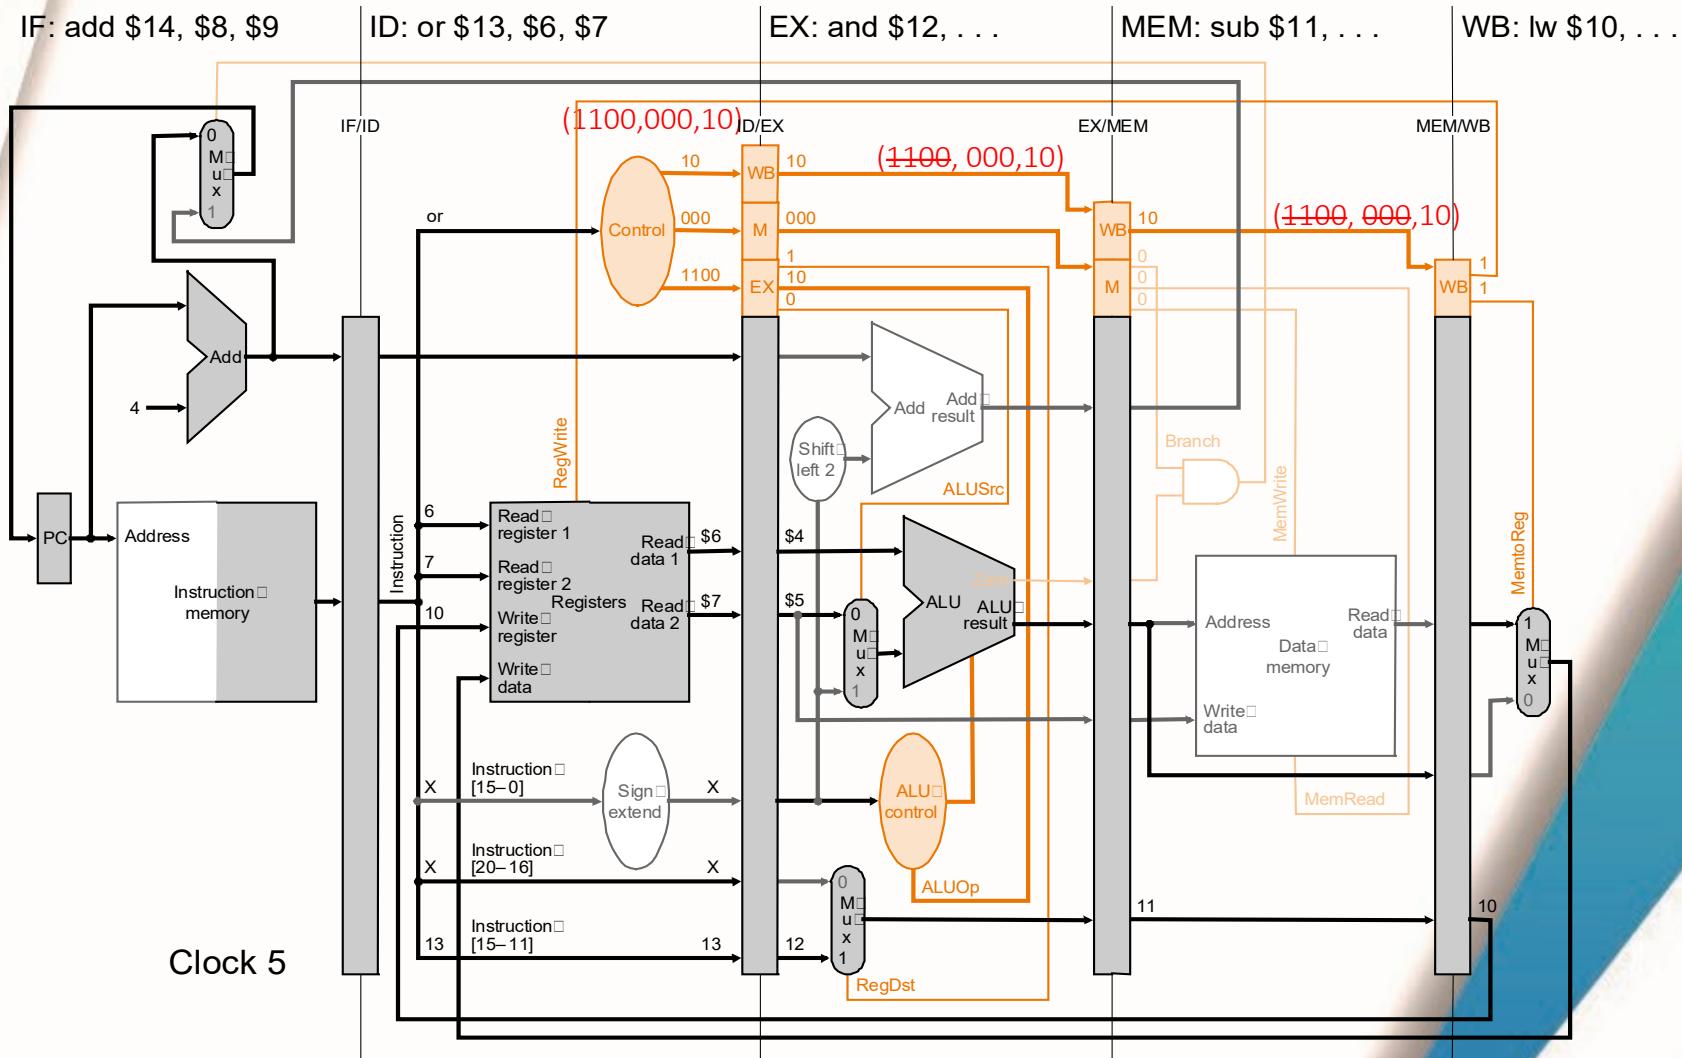


Group Signals According to Stages

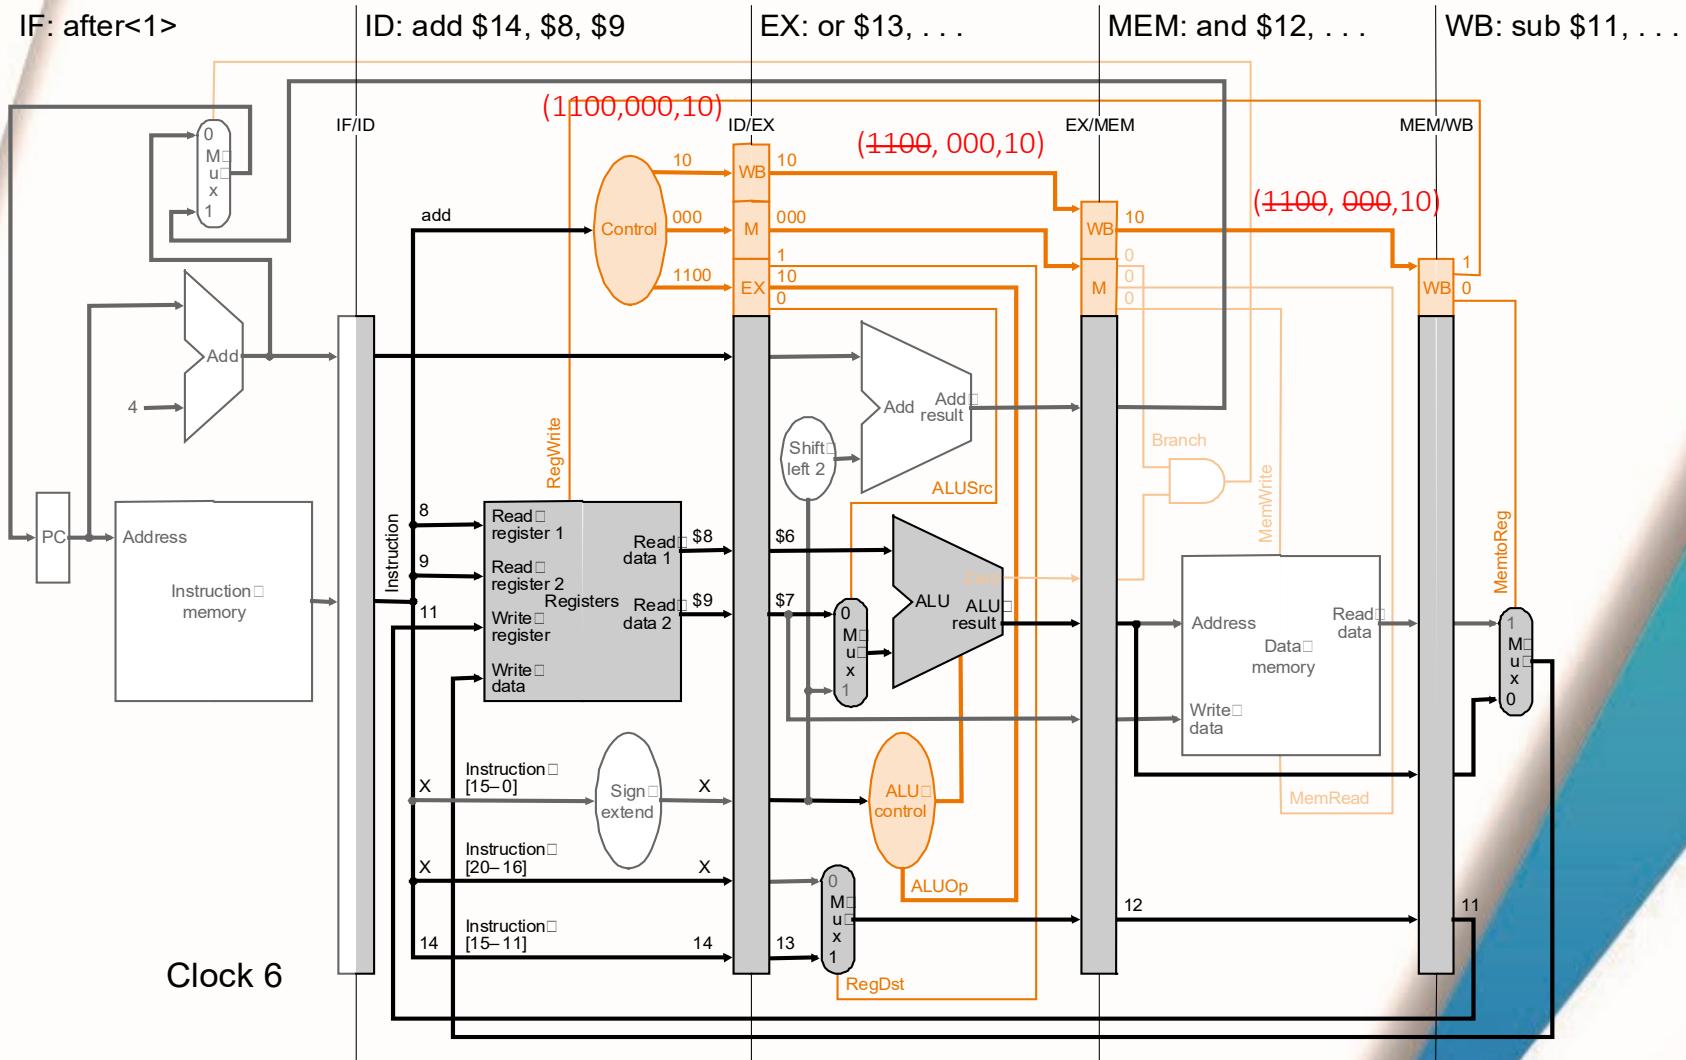
or \$13, \$6, \$7

Execution/Address Calculation stage control lines				Memory access stage control lines			Write-back stage control lines	
Reg Dst	ALU Op1	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg write	Mem to Reg
R-type lw sw beq	1	1	0	0	0	0	1	0
	0	0	0	1	0	1	1	1
	X	0	0	1	0	0	1	0
	X	0	1	0	1	0	0	X

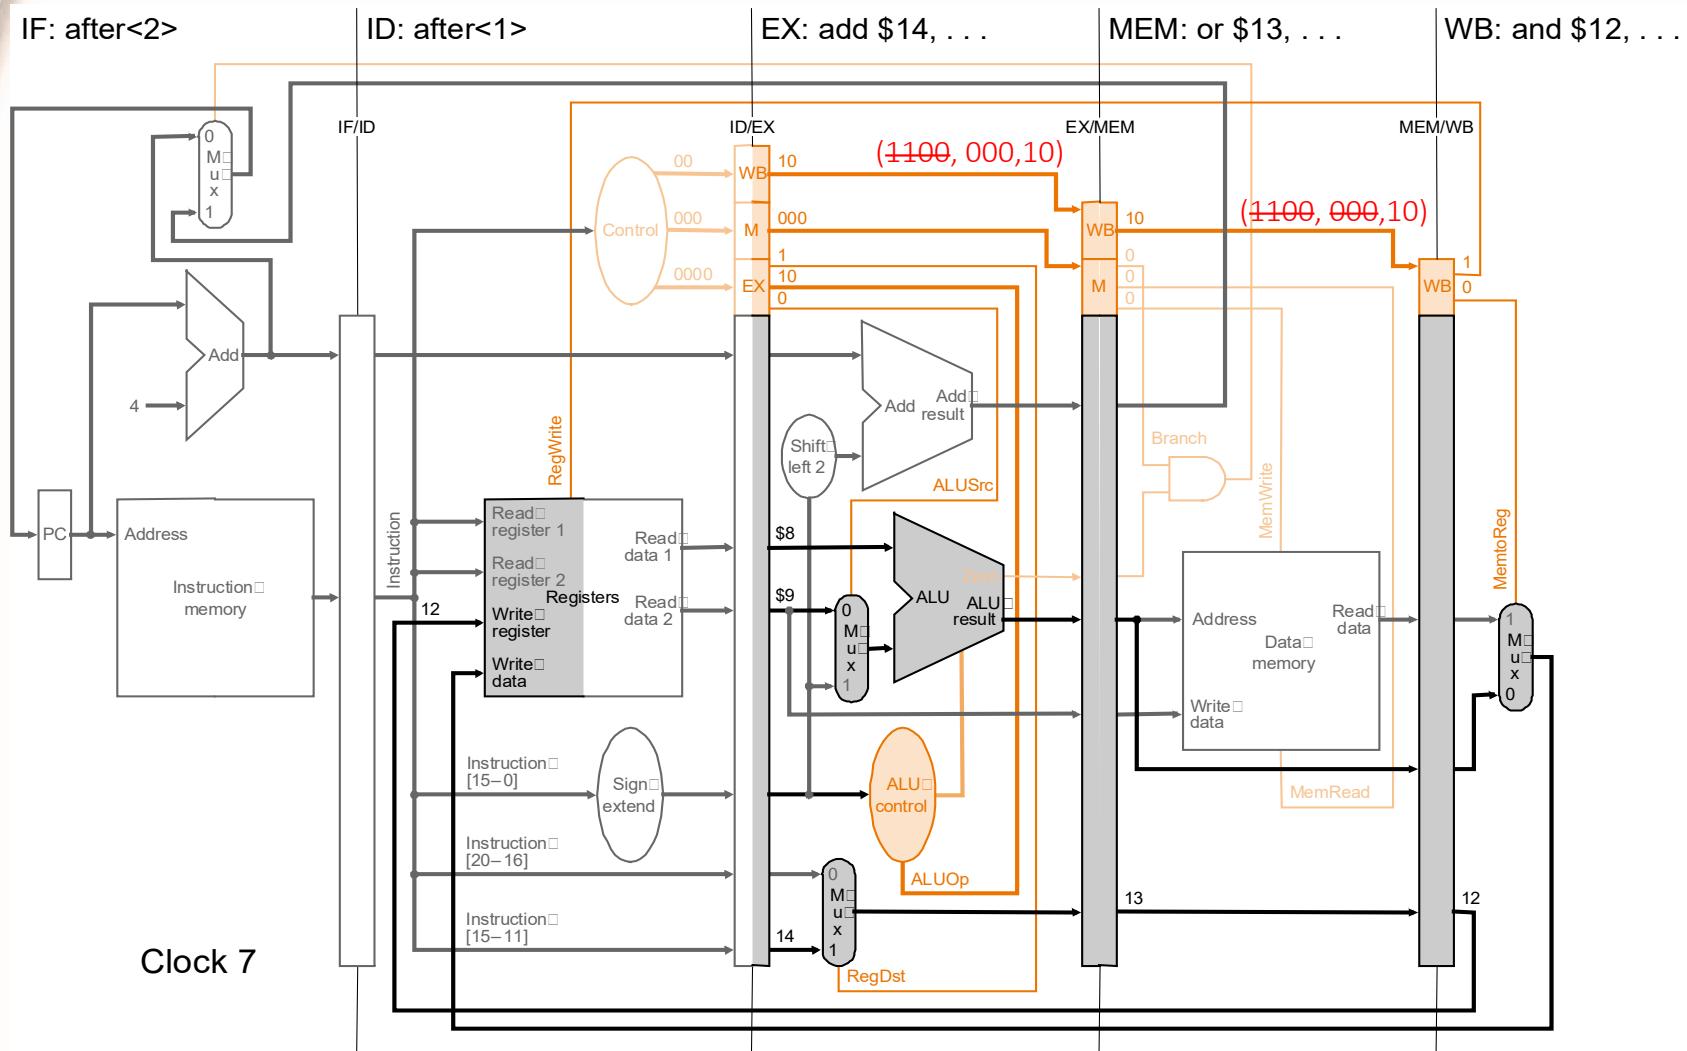
Cycle 5



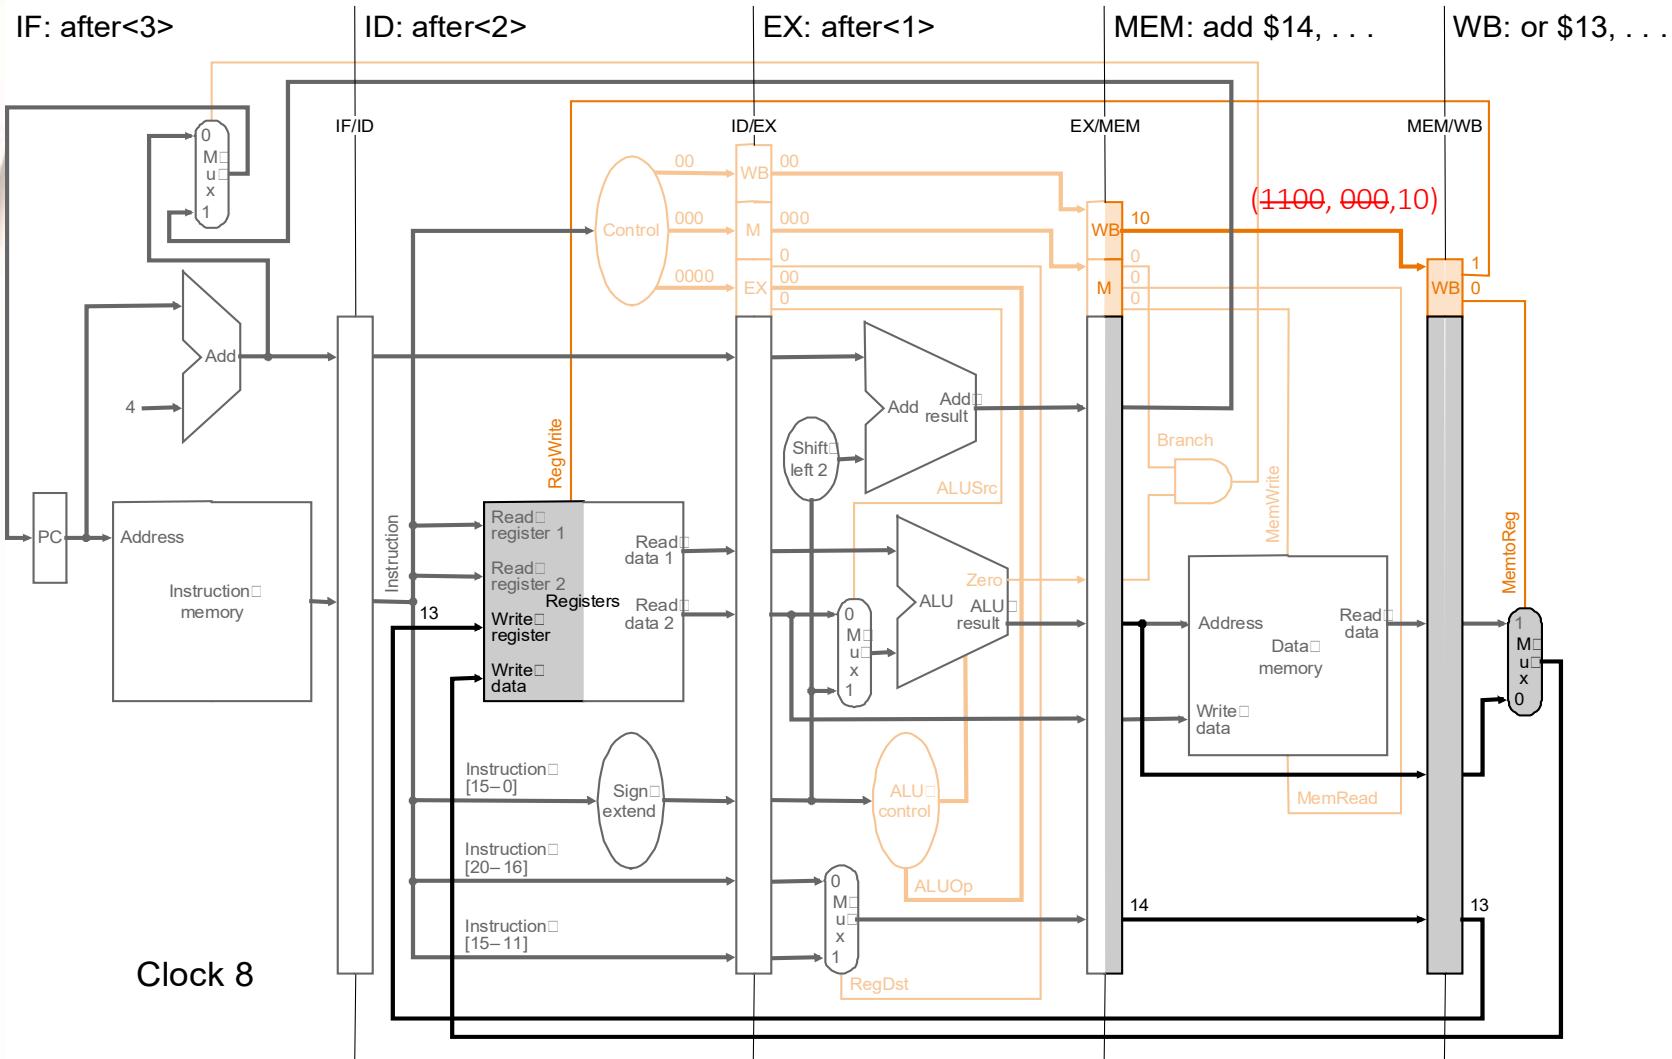
Cycle 6



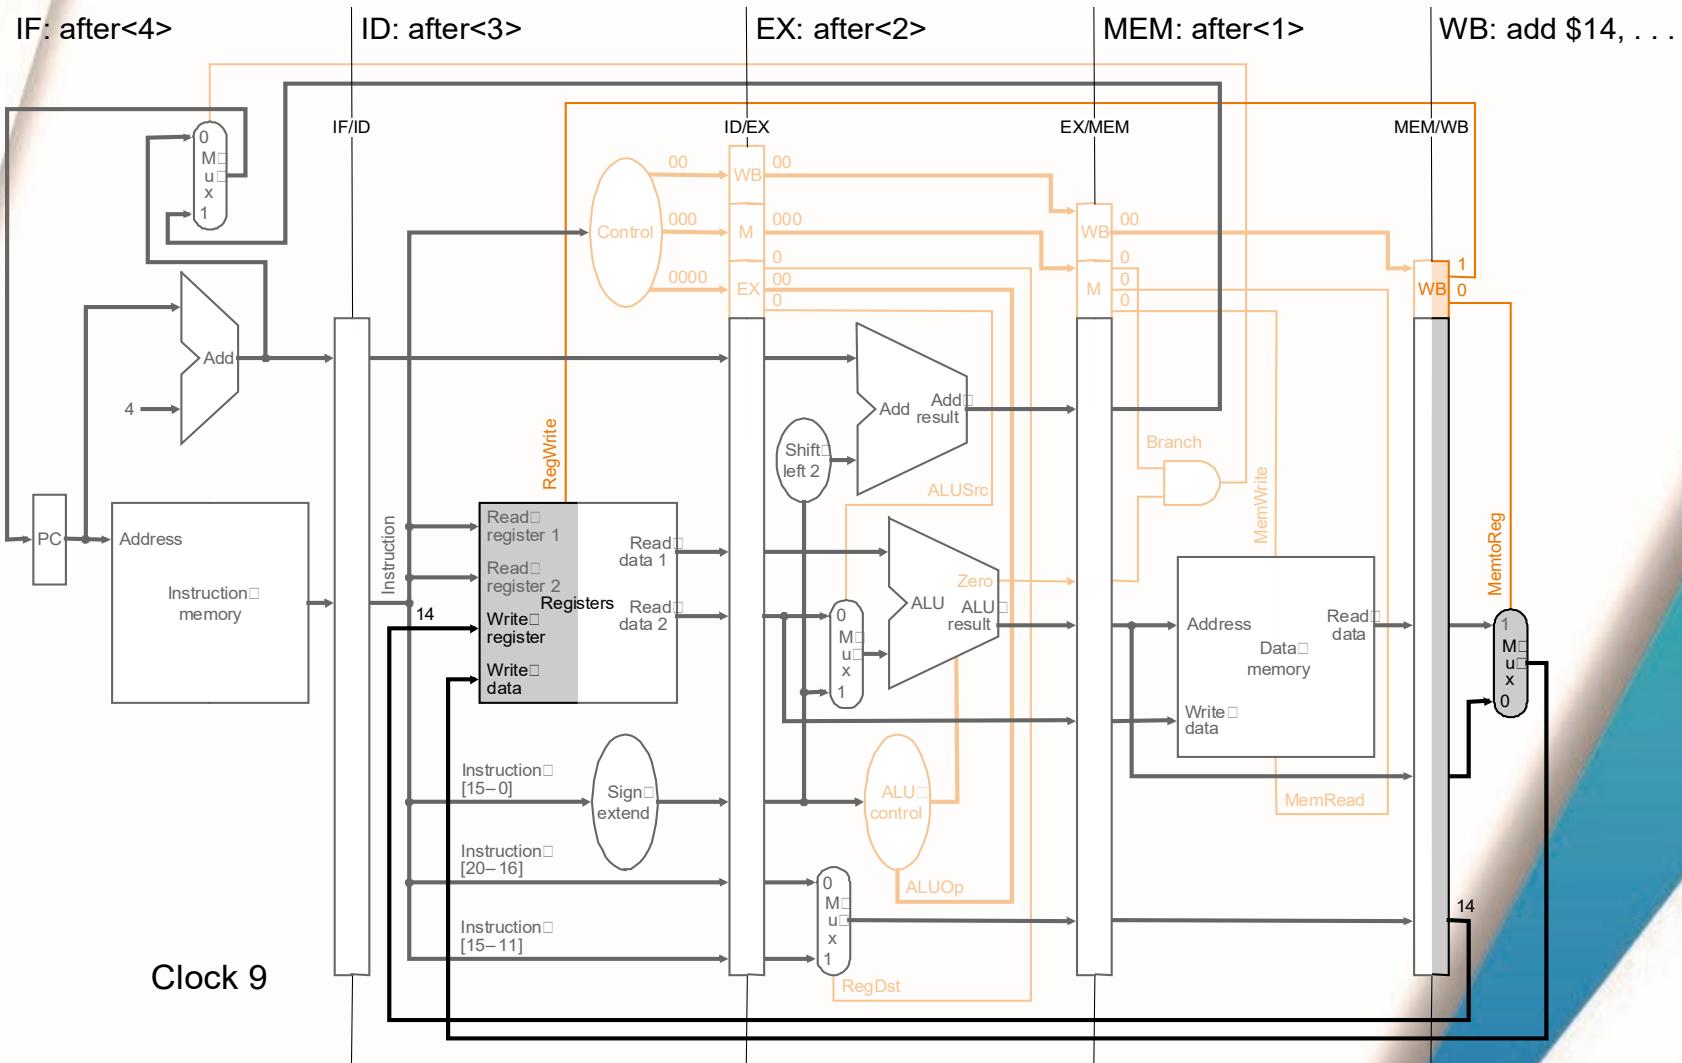
Cycle 7



Cycle 8



Cycle 9



Short Summary

- Pipelining features

- Each stage uses separate resources

- Pipelining

- Executes multiple instructions at the same time
 - Set the clock cycle to the longest stage
 - How to detect hazards and resolve them

- MIPS facilitates pipelining due to

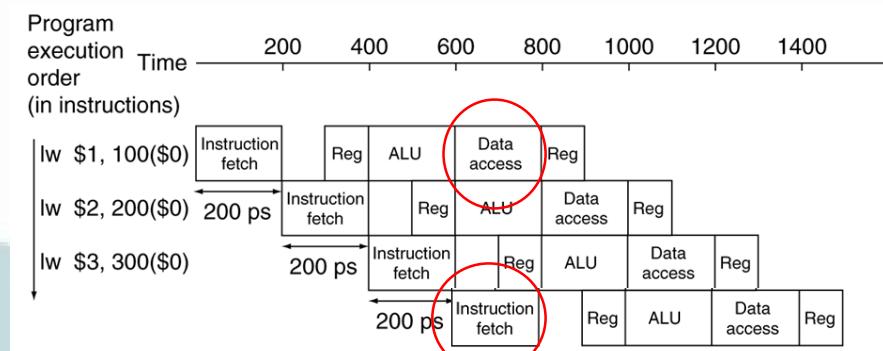
- Only three instruction types and the instructions have the same length
 - Operating memory with only loads and stores

Hazards

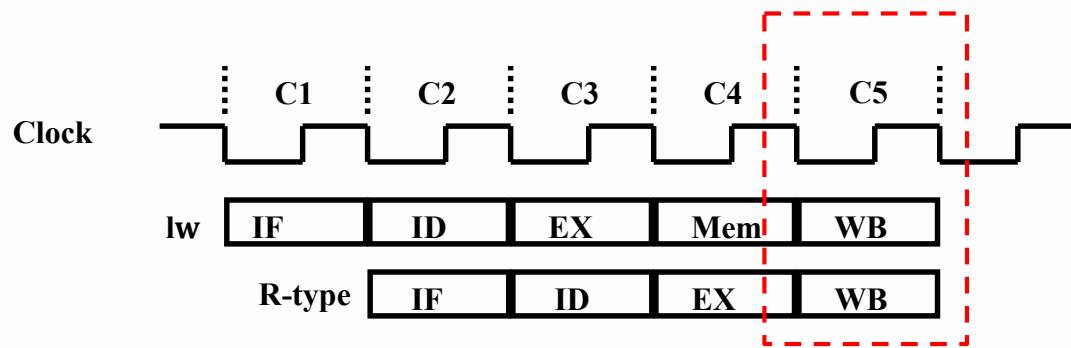
- Situations that prevent starting the next instruction in the next cycle
- Three Hazards: Structure, data, and control hazards
- Structure hazards
 - A required resource is busy
 - Competing the same hardware resource in one cycle
- Data hazard
 - Need to wait for data outputted from previous instruction
 - EX(r-type 、 lw) 、 ID (beq) stage
- Control hazard
 - Taken branch decision from the previous instruction needs to void the current instruction
- Hazards can be resolved by **waiting (stall)**
 1. Hazard detection
 2. Resolving hazards by waiting or other mechanisms.

Structure Hazard

- Resource conflict
- In MIPS pipeline with a single memory
 - Load/store requires data access
 - Instruction fetch would have to *stall* for that cycle
 - Would cause a pipeline “bubble”
- Hence, pipelined datapaths require separate instruction/data memories
 - Or separate instruction/data caches



Pipelining R-type and load

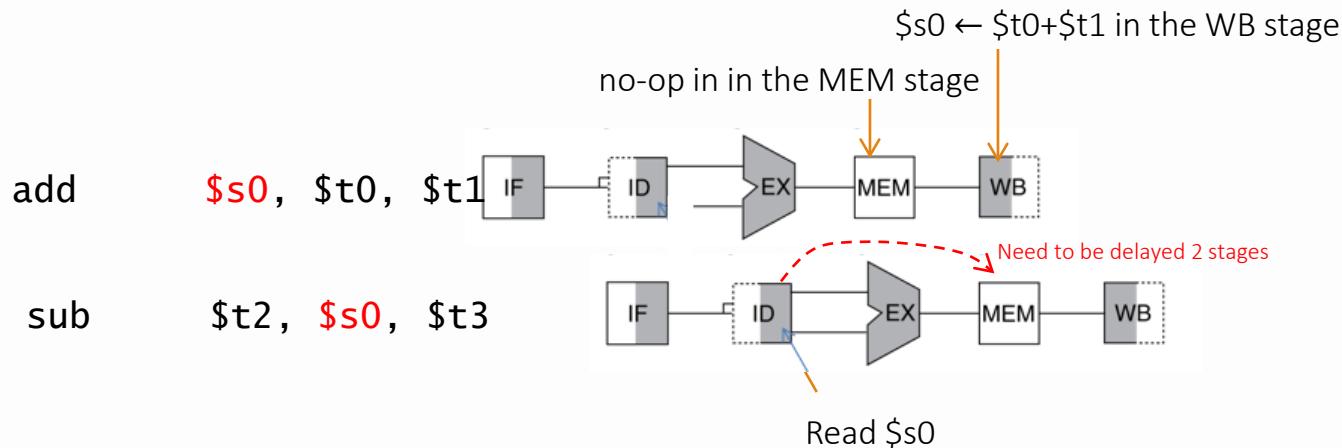


- If R-type takes four cycles and `lw` takes five cycles, we then have a *structural hazard*:
 - `lw` and R-type instruction both write results into the register file on the same stage but the register only has one write port.

Data Hazards

■ Data Hazard

- Attempt to use data before it is ready
- The pipeline must be stalled for one step, waiting for another to complete



Different Data Hazards

- **RAW (read after write):**

The latter instruction needs to read the register after the former instruction writes it

- **WAR (write after read):**

The latter instruction needs to write the register after the former instruction reads it

- Won't happen in MIPS (5 stages)

- **WAW (write after write):**

The latter instruction needs to write the register after the former instruction writes it

- Won't happen in MIPS (5 stages)

Data Hazards

RAW

add

$\$s0, \$t0, \$t1$

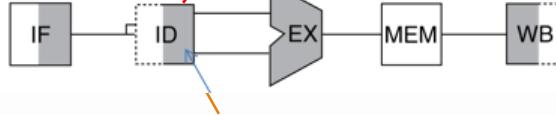


$\$s0 \leftarrow \$t0 + \$t1$ in the WB stage

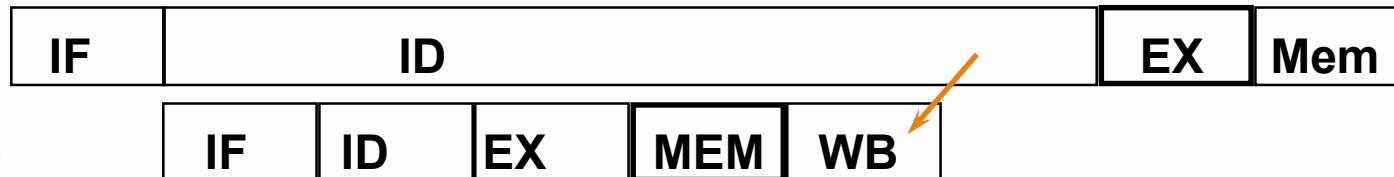
no-op in in the MEM stage

sub

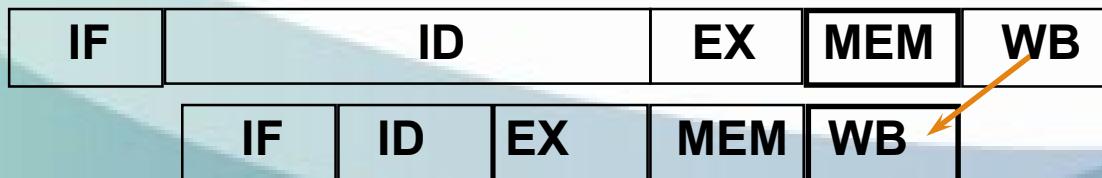
$\$t2, \$s0, \$t3$



WAR



WAW



Read $\$s0$

Resolve Data Hazard

■ Solutions

- Compiler inserts NOP (software solution)
- Forward (hardware solution)
- Stall (hardware solution)

■ Inserting NOPs

- Compiler inserts no operation.

```
add    $s0, $t0, $t1
sub    $t2, $s0, $t3
```



add	\$1,	\$2,	\$3
Nop			
Nop			
sub	\$4,	\$1,	\$5

■ Forward

- Internal forwarding from pipeline registers

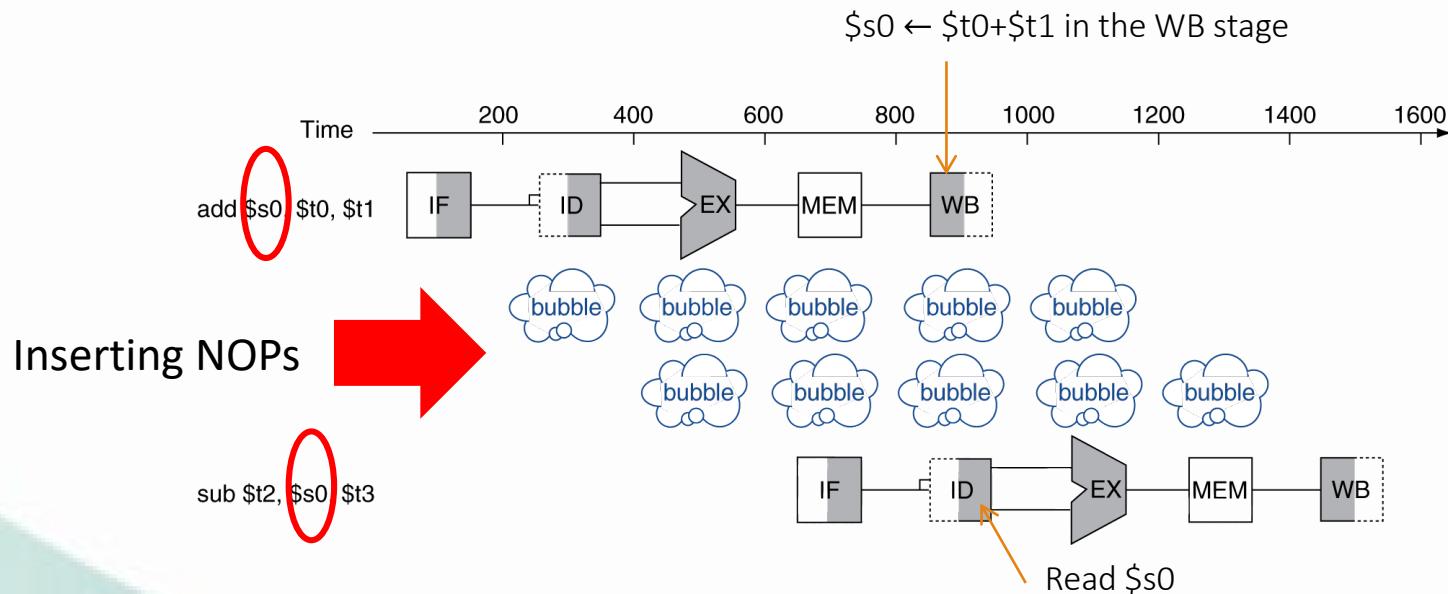
Drawback: low efficiency

■ Stall

- Stall instruction in IF and ID: PC and IF/ID remain changing
=> the stages re-execute the instructions
- What to move into EX: changing EX, MEM, WB control fields of ID/EX pipeline register to 0
 - as control signals propagate, all control signals to EX, MEM, WB are deasserted and no registers or memories are written

Data Hazards

Inserting NOPs - Compiler inserts no operation.



Forwarding

`add $s0, $t0, $t1
sub $t2, $s0, $t3`

(sub) Need to access \$s0 in ID stage

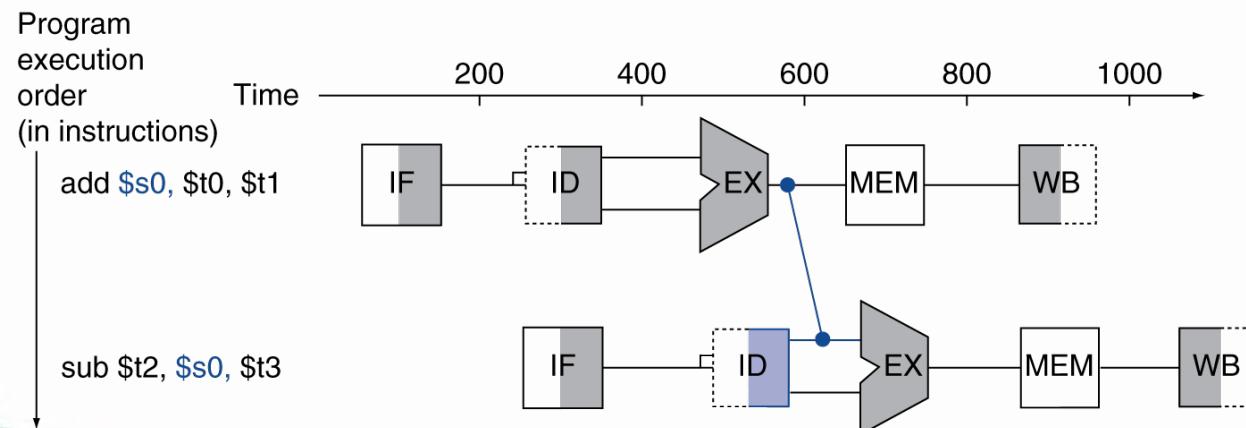
	Step name	Action for R-type instructions	Action for memory-reference instructions	Action for branches	Action for jumps
(IF)	Instruction fetch		IR = Memory[PC] PC = PC + 4		
(ID)	Instruction decode/register fetch		A = Reg [IR[25-21]] B = Reg [IR[20-16]]		
(EXE)	Execution, address computation, branch/jump completion	ALUOut = A op B	ALUOut = A + sign-extend (IR[15-0])	if (A == B) then PC = PC + sign-ext(IR[15-0]) << 2	PC = PC [31-28] II
(MEM)	Memory access or R-type completion	Reg [IR[15-11]] = ALUOut (WB)	Load: MDR = Memory[ALUOut] or Store: Memory [ALUOut] = B		
(WB)	Memory read completion		Load: Reg[IR[20-16]] = MDR		

(add) \$s0 is calculated in EX stage

(add) save the result in \$s0 in WB stage

Forwarding (Bypassing)

- Use the result immediately when computed
 - Don't wait for it to be stored in a register
 - Requires extra connections in the datapath



Hardware Implementation for Resolving Data Hazards

- Consider this sequence:

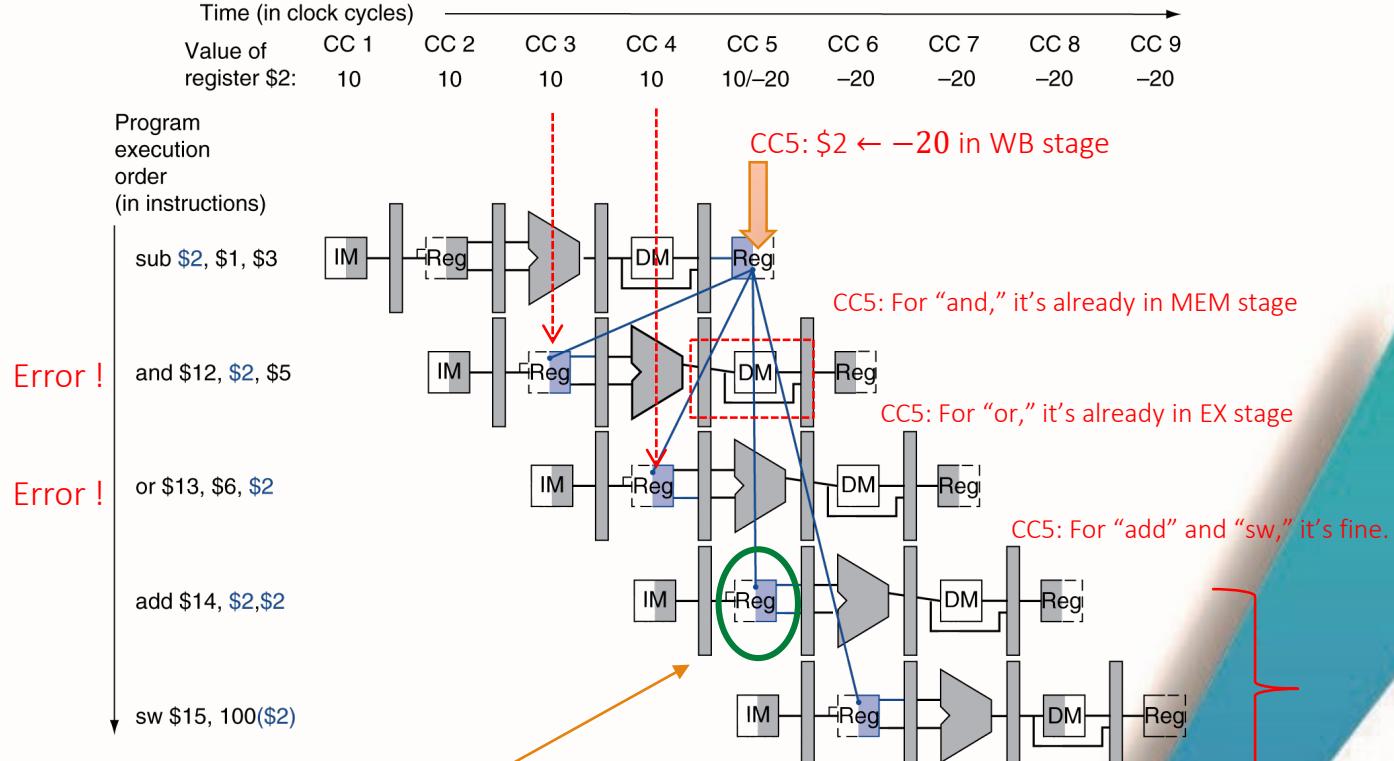
```
sub $2, $1,$3  
and $12,$2,$5  
or $13,$6,$2  
add $14,$2,$2  
sw $15,100($2)
```

Data dependency: \$2

Example

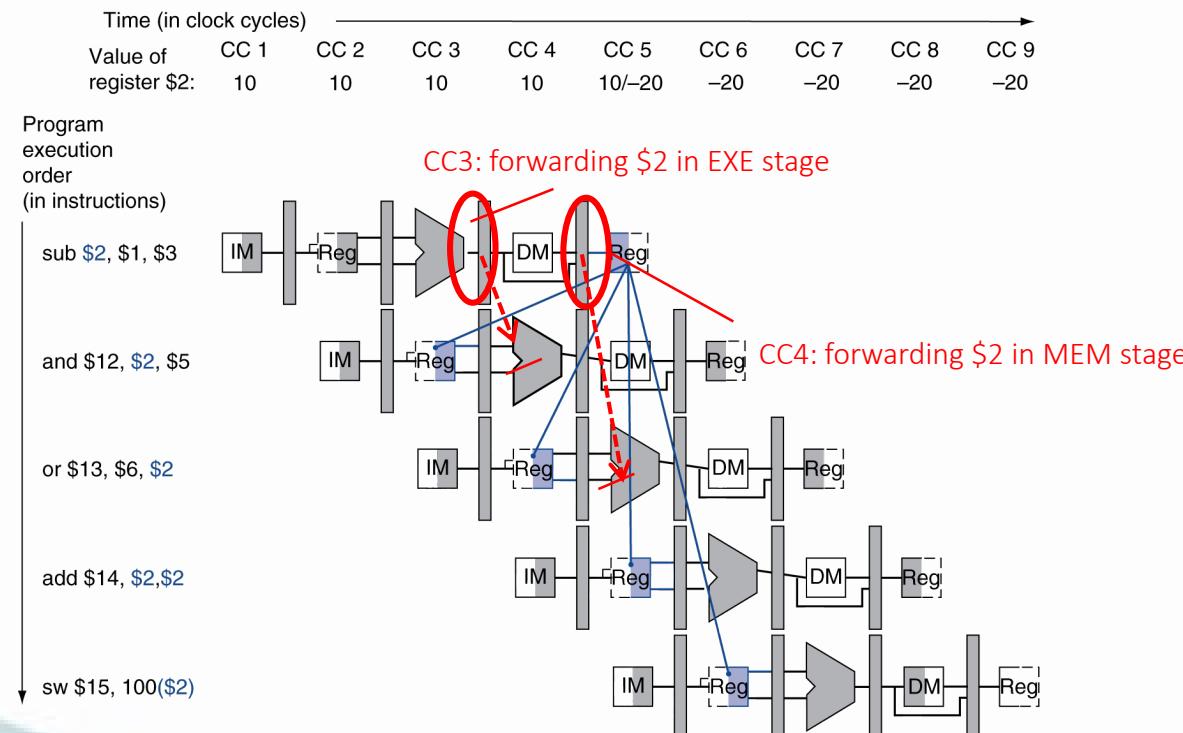
Consider this sequence:

```
sub $2, $1, $3
and $12, $2, $5
or $13, $6, $2
add $14, $2, $2
sw $15, 100($2)
```

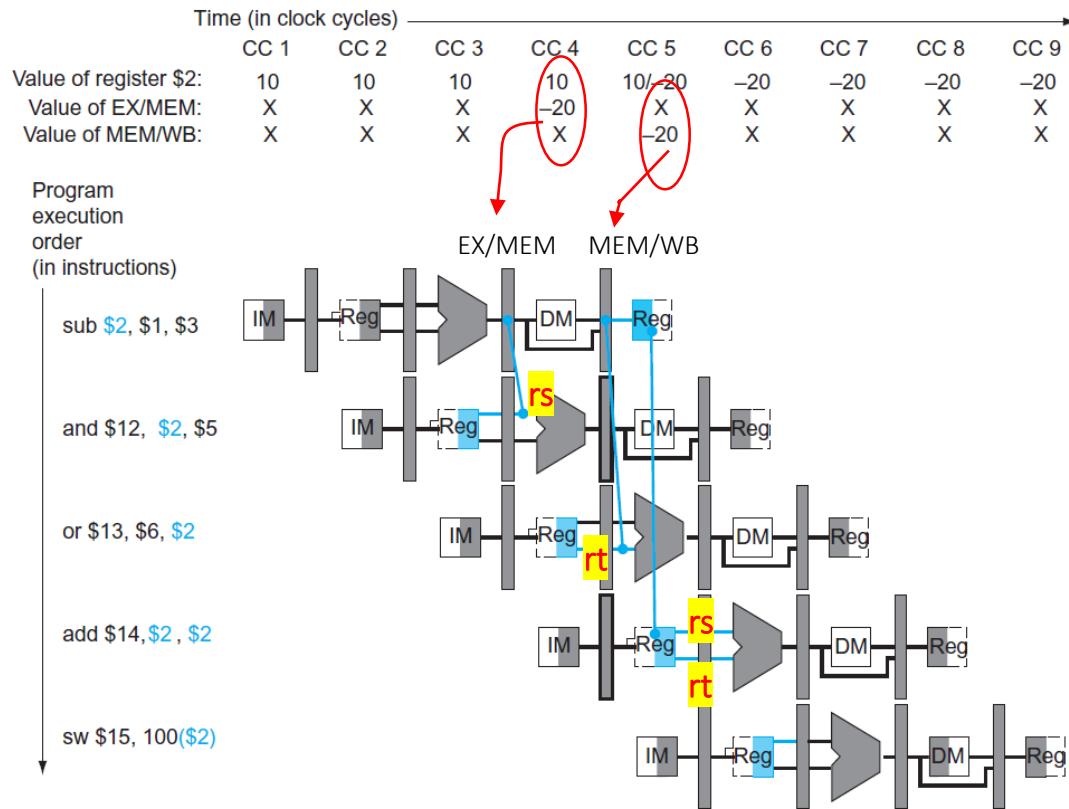


What the clock edge triggers, the data will be written into \$s2 in WB stage of "sub". Then \$s2 is read in ID stage of "add".

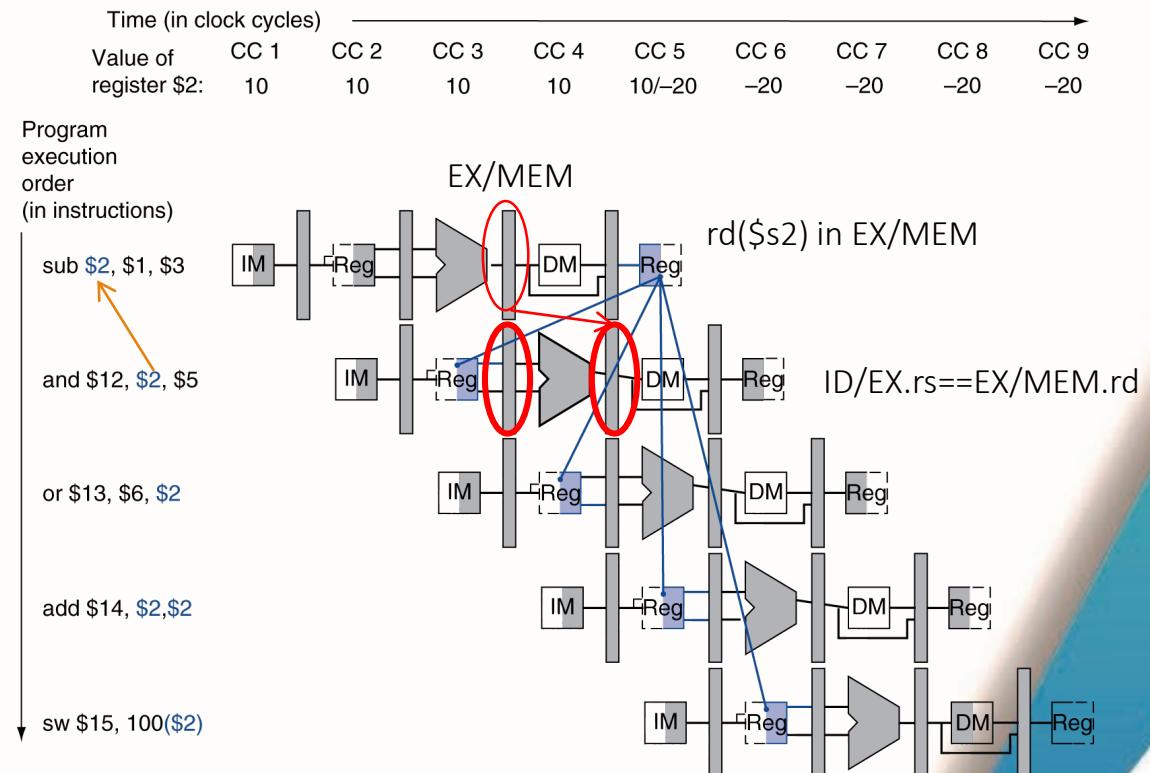
Resolving Data Hazard by Forwarding



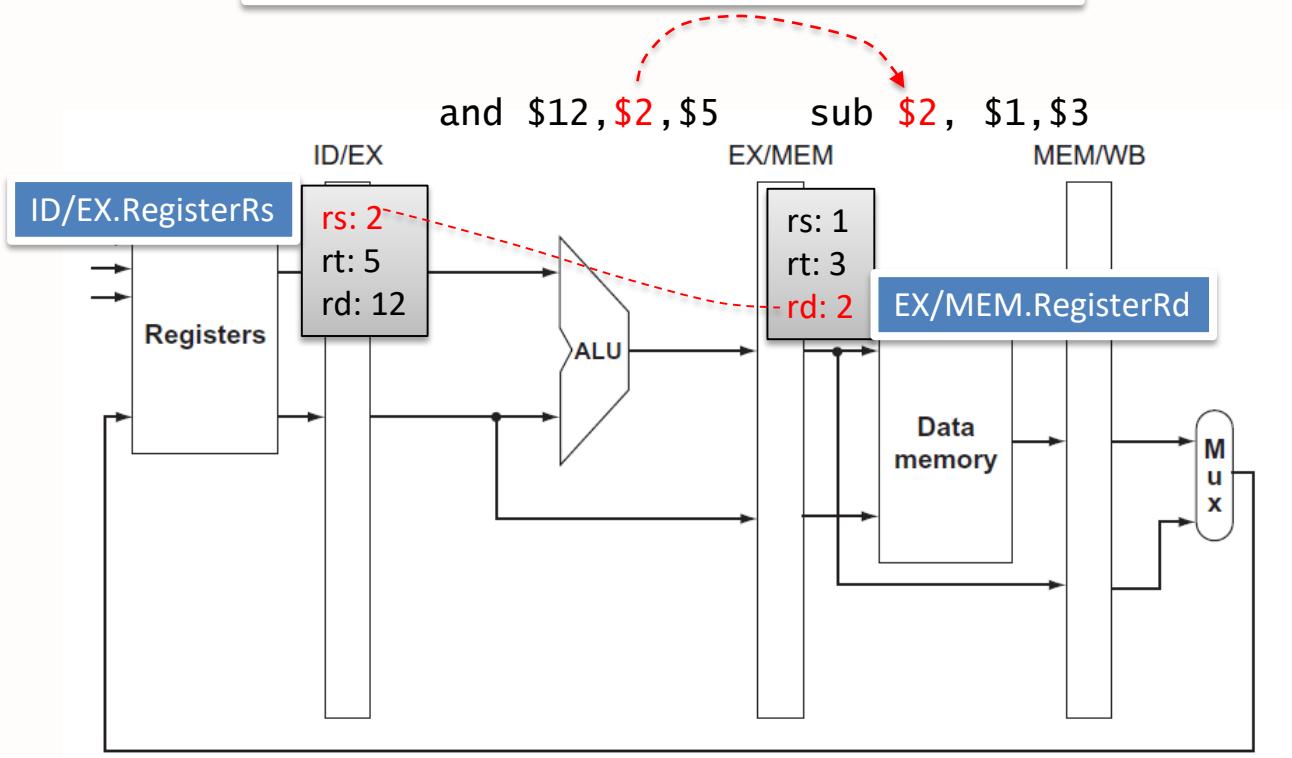
For sub



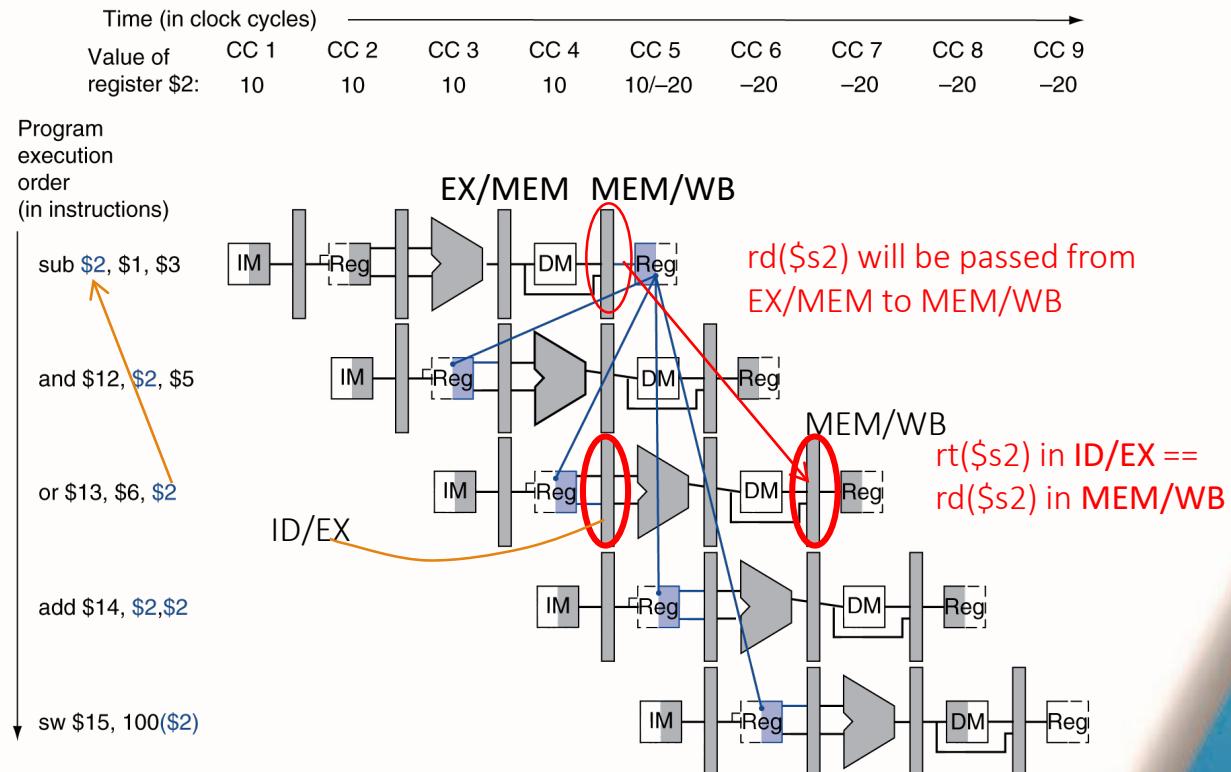
When Data Hazard Occurs - 1



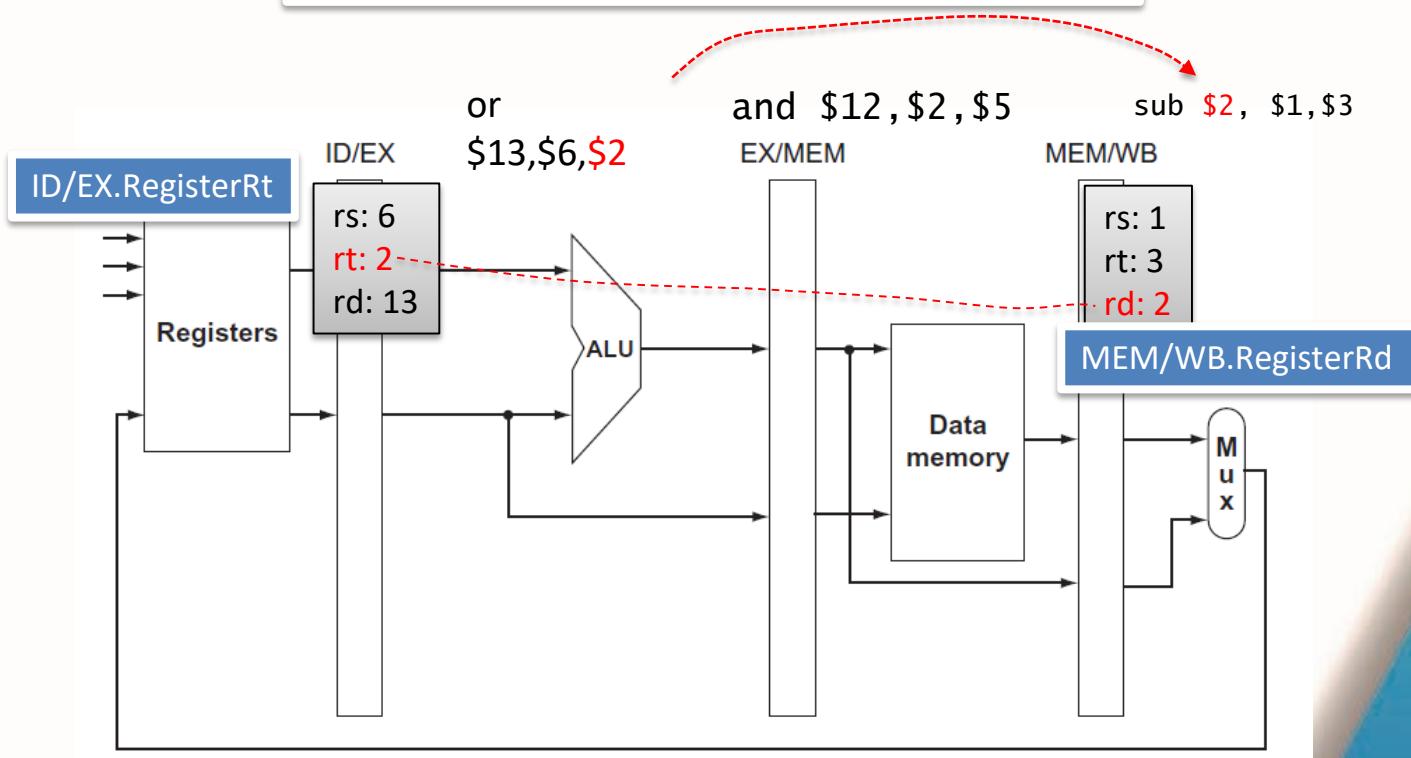
EX/MEM.RegisterRd == ID/EX.RegisterRs



When Data Hazard Occurs - 2



MEM/WB.RegisterRd == ID/EX.RegisterRt



Detecting the Need to Forward

- Data hazards when

It's Reg Num,
not value in Reg

1a. EX/MEM.RegisterRd == ID/EX.RegisterRs

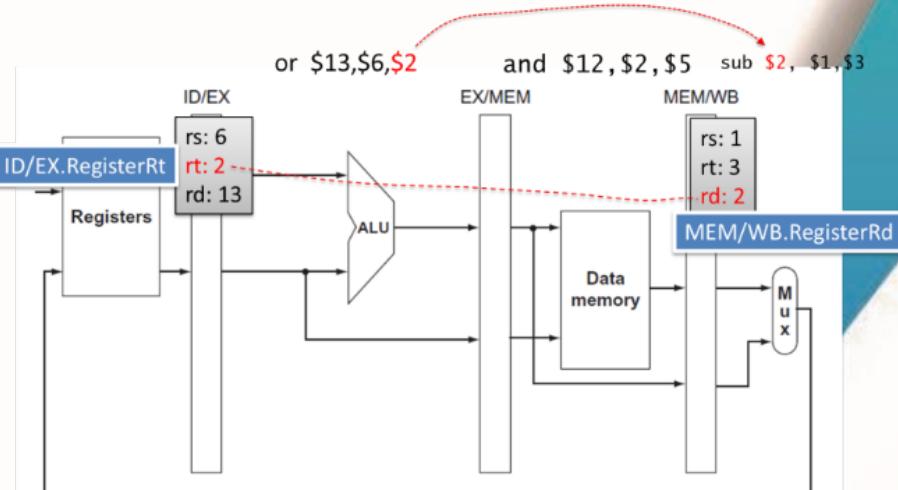
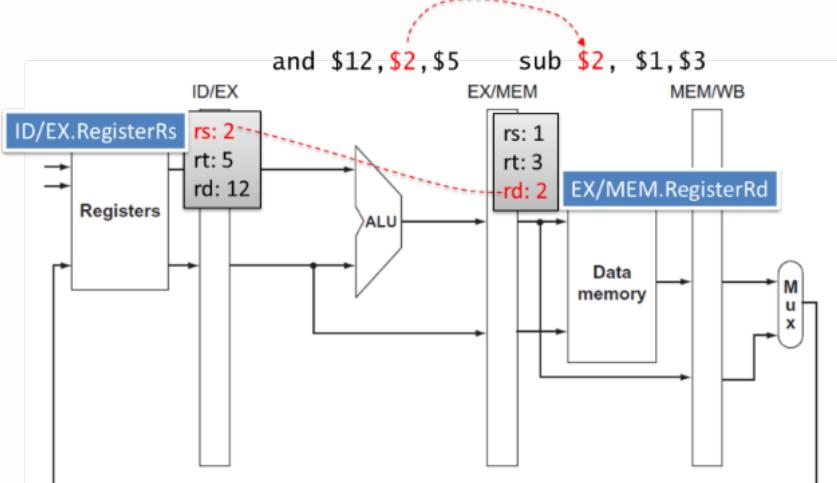
1b. EX/MEM.RegisterRd == ID/EX.RegisterRt

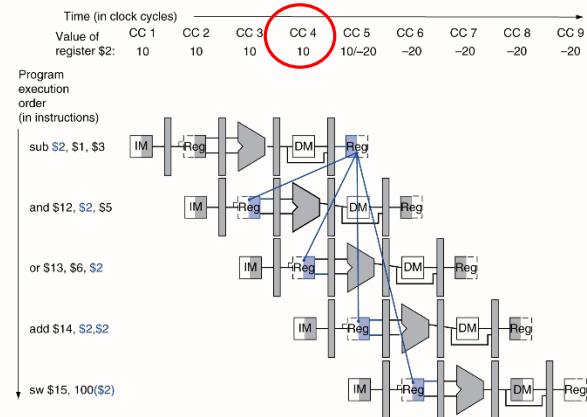
2a. MEM/WB.RegisterRd == ID/EX.RegisterRs

2b. MEM/WB.RegisterRd == ID/EX.RegisterRt

Fwd from
EX/MEM
pipeline reg

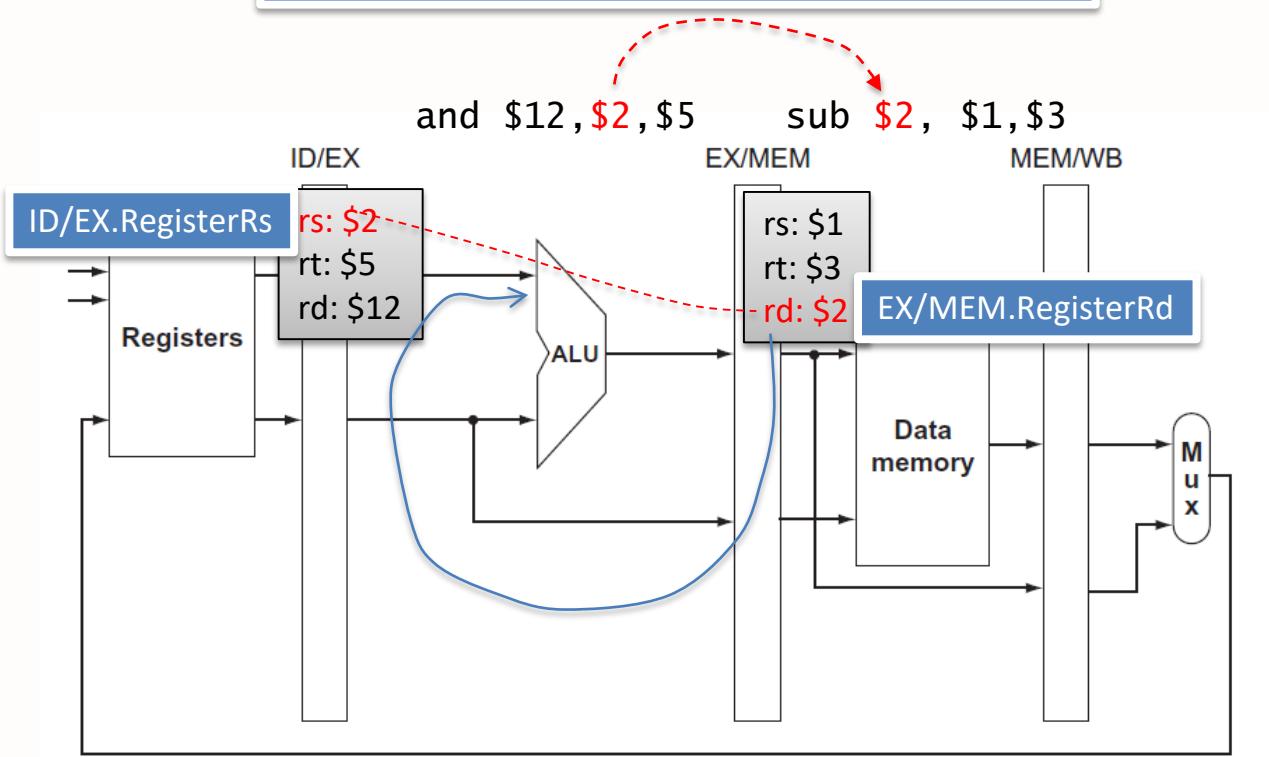
Fwd from
MEM/WB
pipeline reg

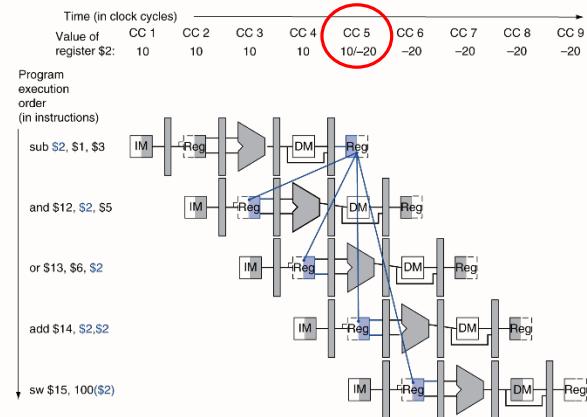




CC4 Forward

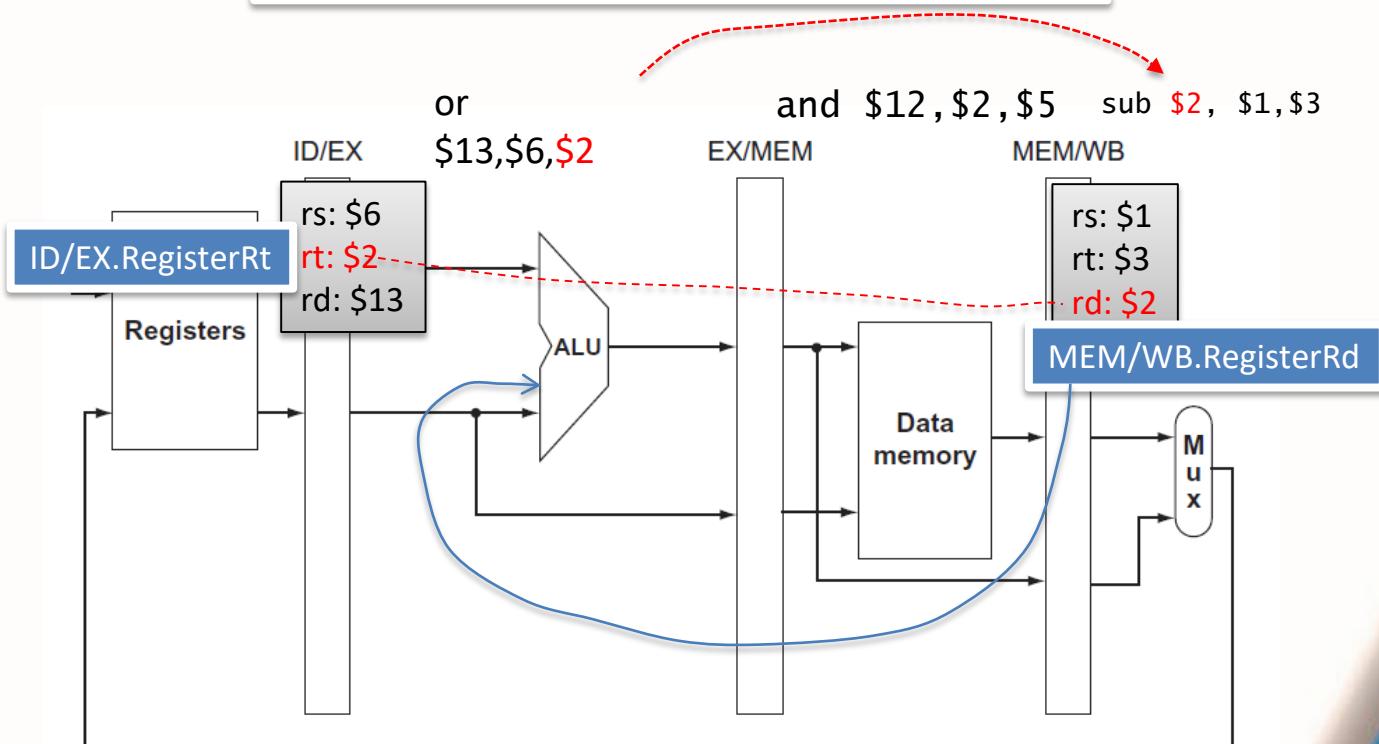
EX/MEM.RegisterRd == ID/EX.RegisterRs



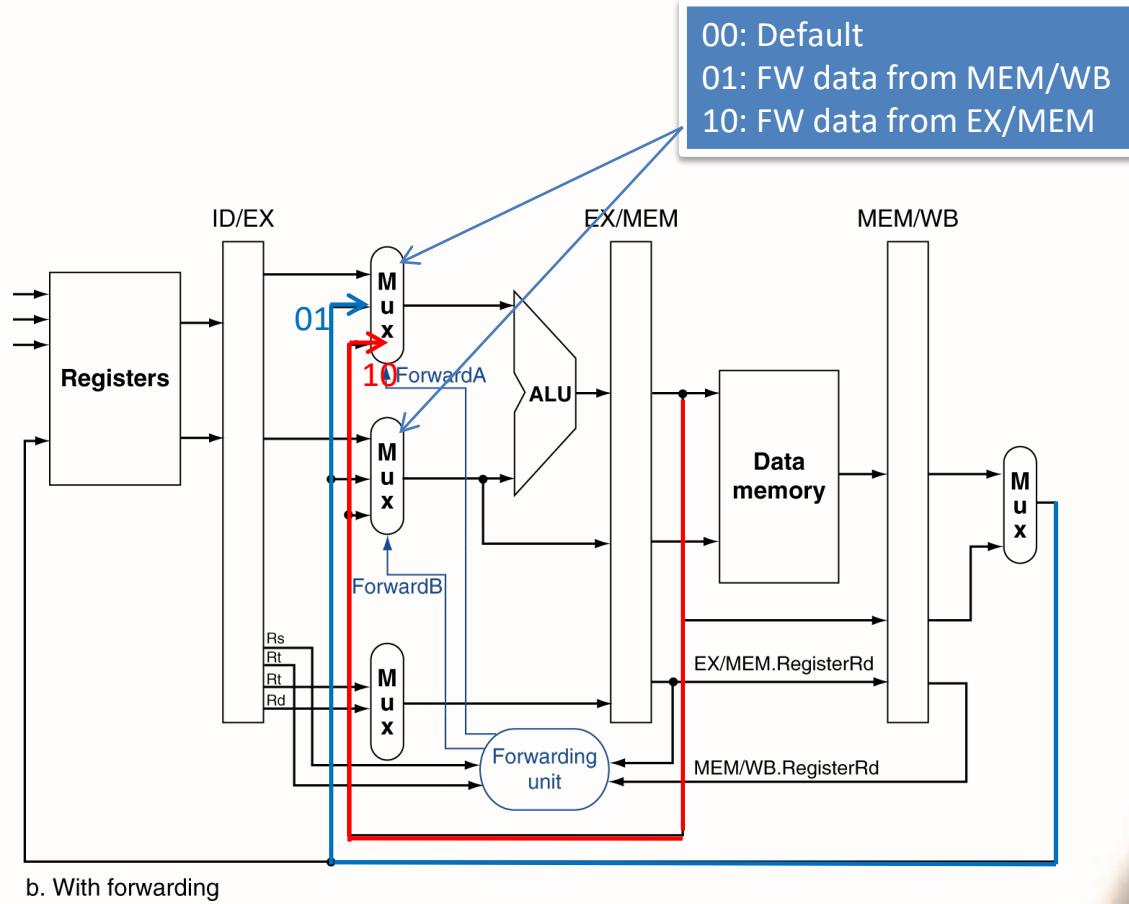


CC5 Forward

MEM/WB.RegisterRd == ID/EX.RegisterRt



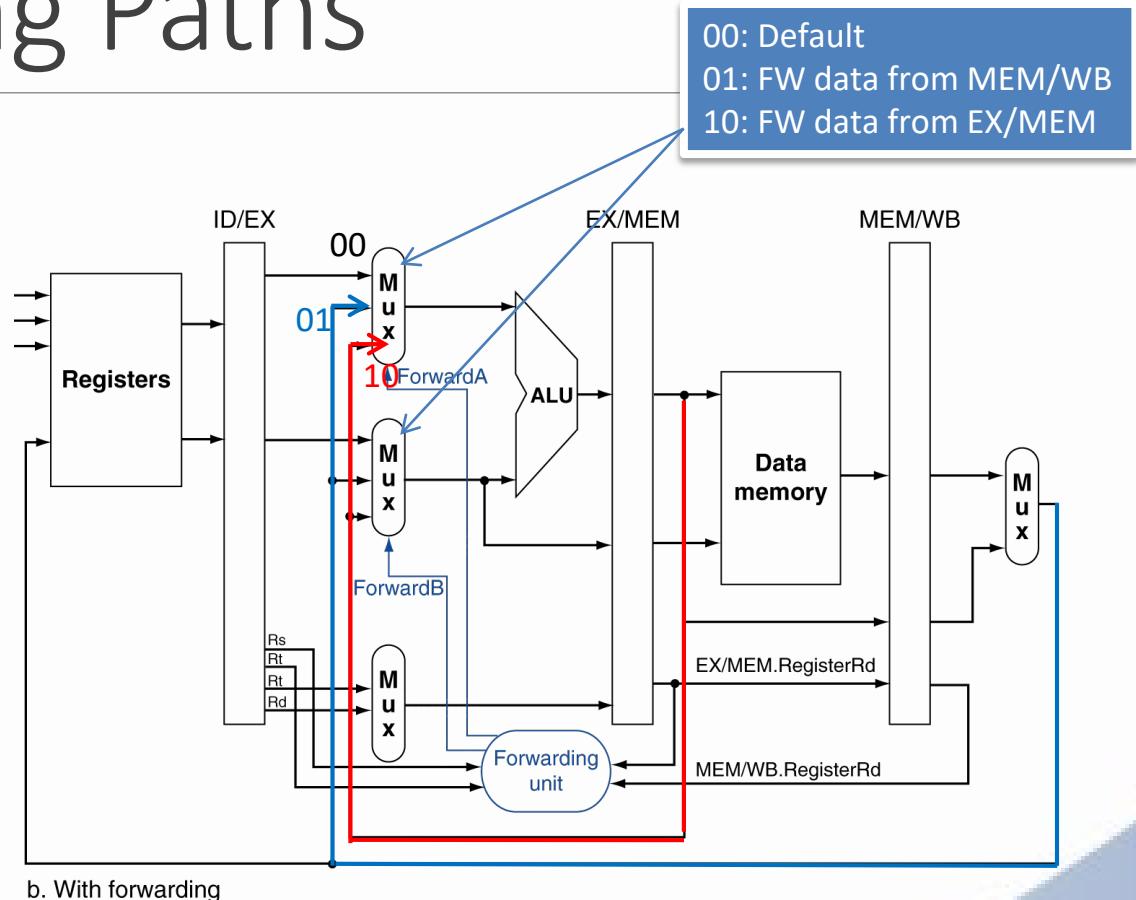
Forwarding Paths



Forwarding Paths

Mux Select Signals

- ForwardA
 - 00: Default
 - 01: FW data from MEM/WB
 - 10: FW data from EX/MEM
- ForwardB
 - 00: Default
 - 01: FW data from MEM/WB
 - 10: FW data from EX/MEM



Two Exceptions in Forwarding

■ No Forwarding

- Do not write data into Register
 - EX/MEM.RegWrite==0 or MEM/WB.RegWrite==0
- If writing data into \$zero (forbidden)
 - EX/MEM.RegisterRd == 0,
MEM/WB.RegisterRd == 0

Forwarding Conditions

❑ EX data hazard

- ❑ if (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)
and (EX/MEM.RegisterRd == ID/EX.RegisterRs))

ForwardA = 10

- ❑ if (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)
and (EX/MEM.RegisterRd == ID/EX.RegisterRt))

ForwardB = 10

❑ MEM data hazard

- ❑ if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)
and (MEM/WB.RegisterRd == ID/EX.RegisterRs))

ForwardA = 01

- ❑ if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)
and (MEM/WB.RegisterRd == ID/EX.RegisterRt))

ForwardB = 01

Mux Select Signals

- ForwardA
 - 00: Default
 - 01: FW data from MEM/WB
 - 10: FW data from EX/MEM
- ForwardB
 - 00: Default
 - 01: FW data from MEM/WB
 - 10: FW data from EX/MEM

Double Data Hazard

- Consider the sequence:

add \$1, \$1, \$2

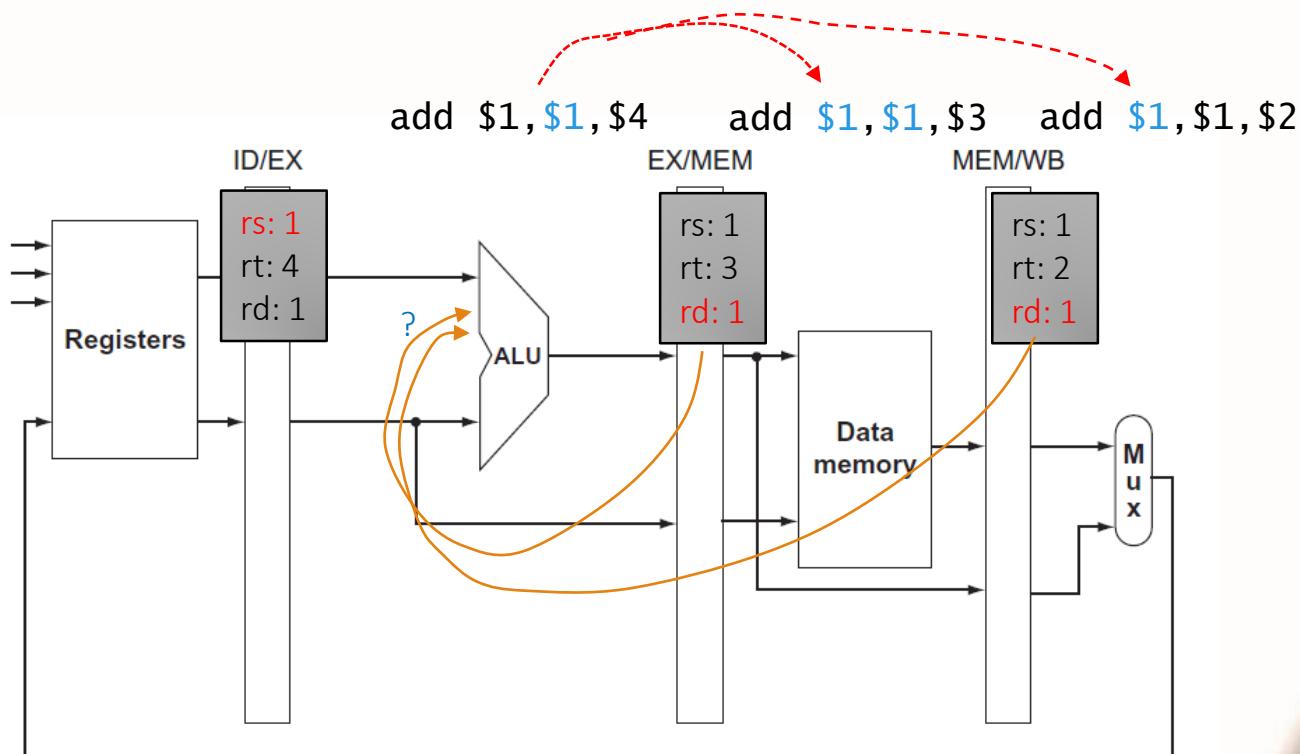
add \$1, \$1, \$3

add \$1, \$1, \$4

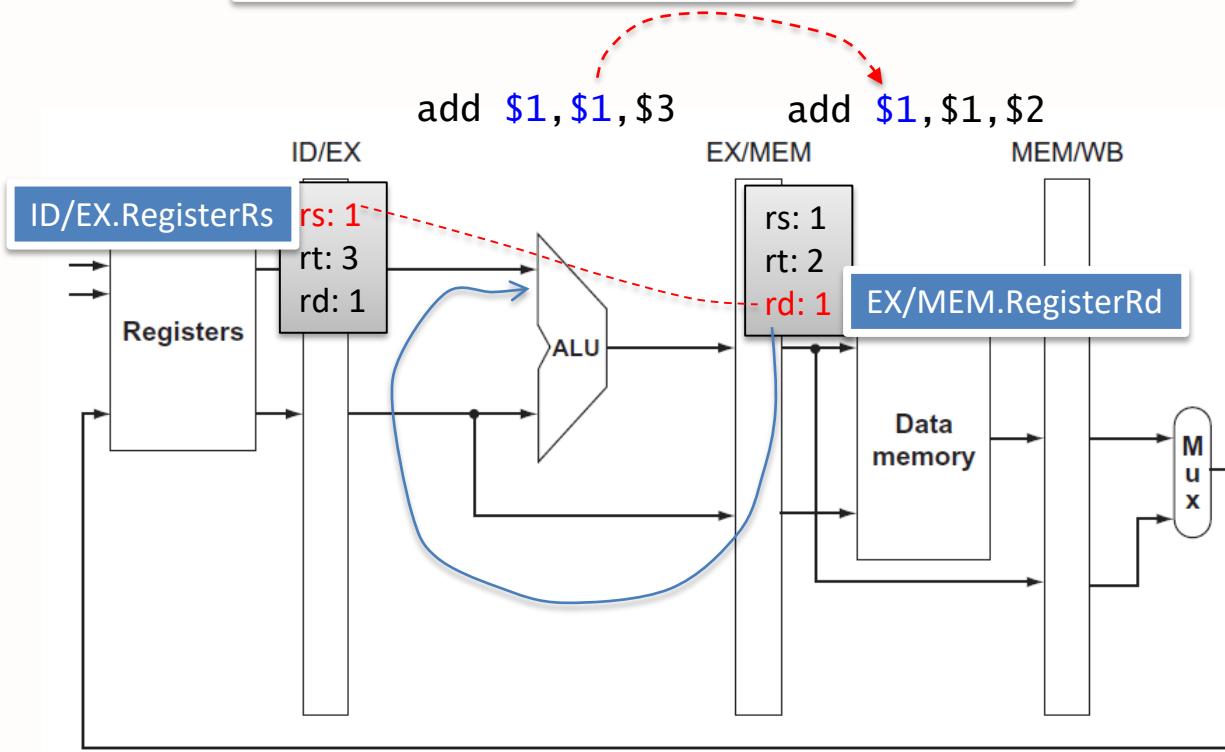
1. EX/MEM.RegisterRd == ID/EX.RegisterRs
2. MEM/WB.RegisterRd == ID/EX.RegisterRs
3. EX/MEM.RegisterRd == ID/EX.RegisterRs

Two data hazard detected:

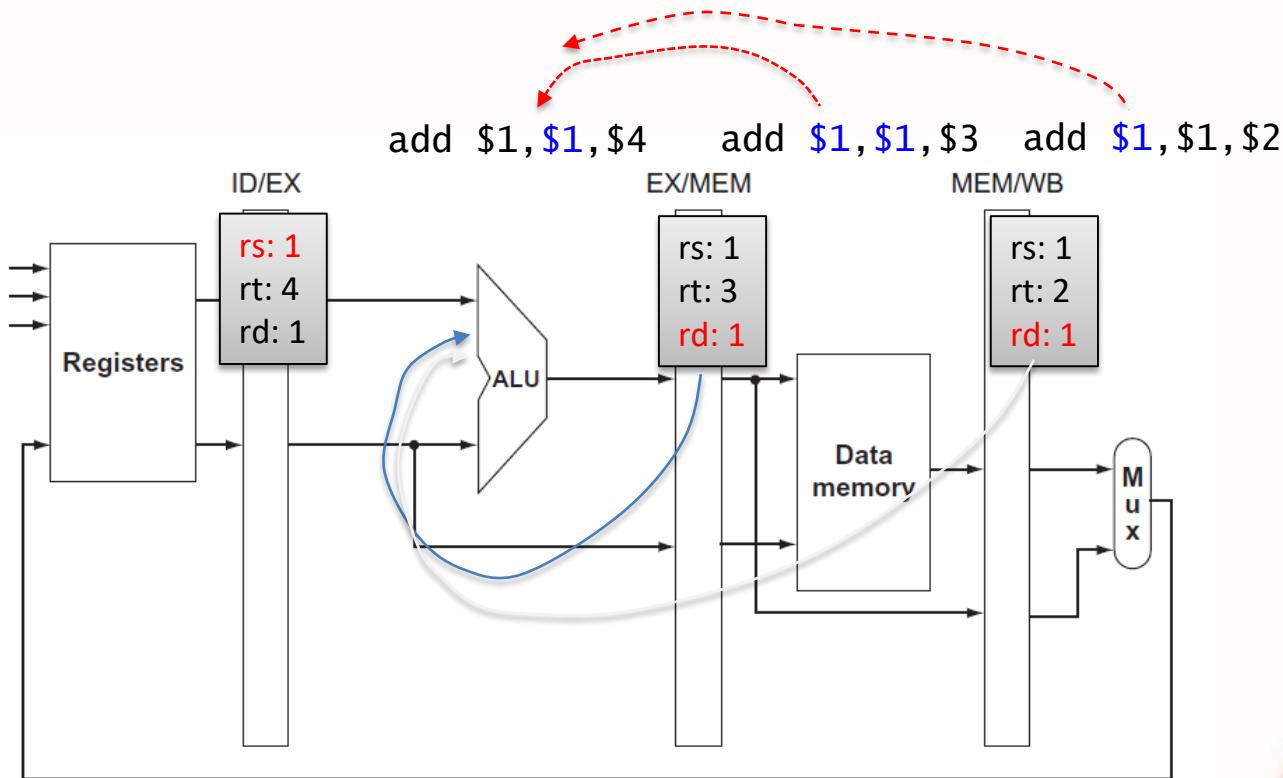
1. MEM/WB: ForwardA = 01
2. EX/MEM: ForwardA = 10



EX/MEM.RegisterRd == ID/EX.RegisterRs



When a double data hazard occurs · Forward data from EX/MEM (most recent result)



Revised Forwarding Condition

MEM hazard

- if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)
and not (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)
and (EX/MEM.RegisterRd == ID/EX.RegisterRs))
and (MEM/WB.RegisterRd == ID/EX.RegisterRs))

ForwardA = 01

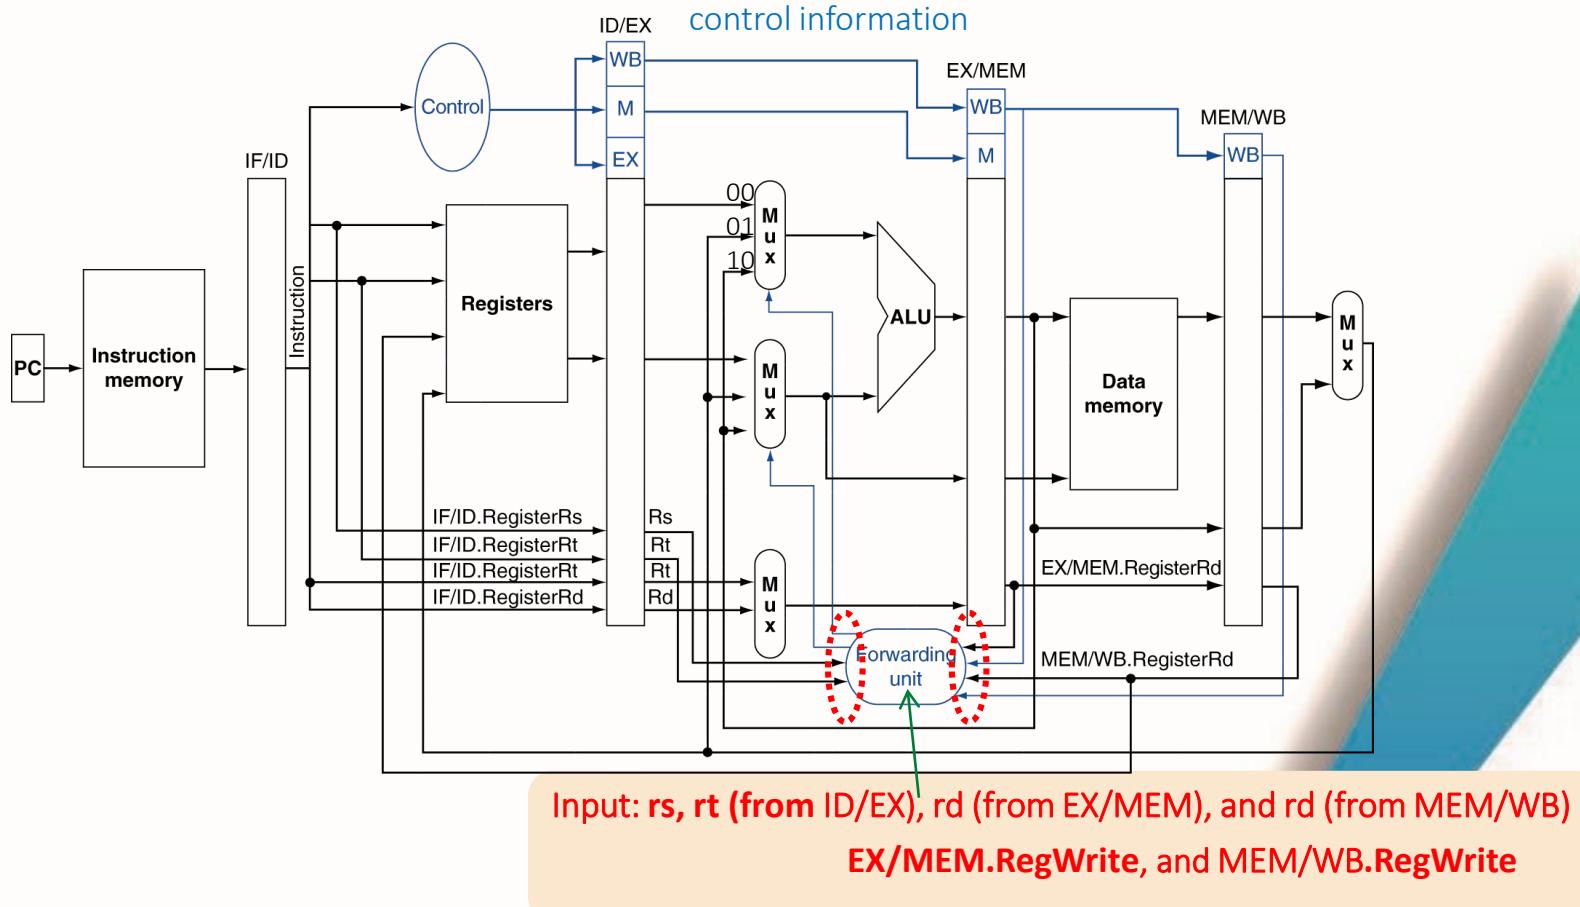
- if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)
and not (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)
and (EX/MEM.RegisterRd == ID/EX.RegisterRt))
and (MEM/WB.RegisterRd == ID/EX.RegisterRt))

ForwardB = 01

Mux Select Signals

- ForwardA
 - 00: Default
 - 01: FW data from MEM/WB
 - 10: FW data from EX/MEM
- ForwardB
 - 00: Default
 - 01: FW data from MEM/WB
 - 10: FW data from EX/MEM

Datapath with Forwarding



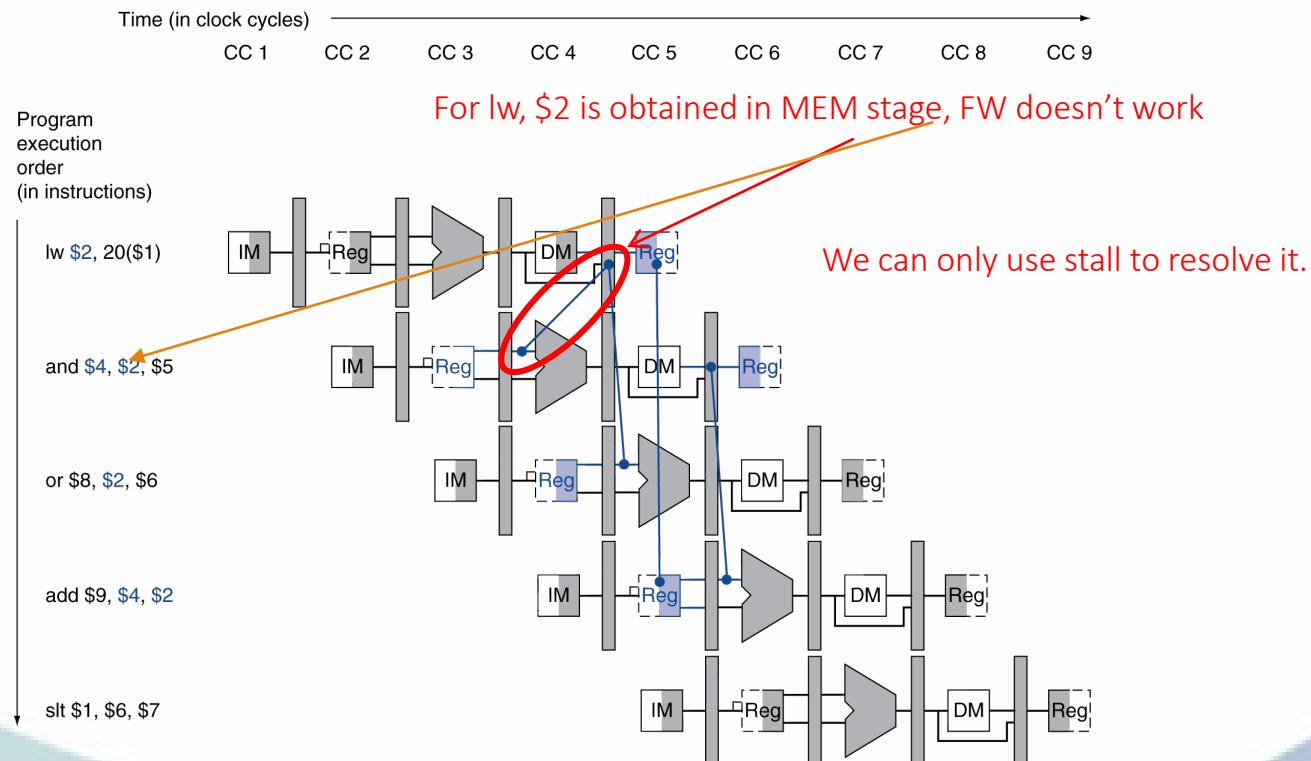
Load/Store

Step name	Action for R-type instructions	Action for memory-reference instructions	Action for branches	Action for jumps
Instruction fetch		IR = Memory[PC] PC = PC + 4		
Instruction decode/register fetch		A = Reg [IR[25-21]] B = Reg [IR[20-16]]		
Execution, address computation, branch/ jump completion	ALUOut = A op B	ALUOut = A + sign-extend (IR[15-0])	if (A == B) then PC = PC + sign-ext(IR[15-0])<<2	PC = PC [31-28] II (IR[25-0]<<2)
Memory access or R-type completion	Reg [IR[15-11]] = ALUOut (WB)	Load: MDR = Memory[ALUOut] or Store: Memory [ALUOut] = B		
Memory read completion		Load: Reg[IR[20-16]] = MDR		

lw:

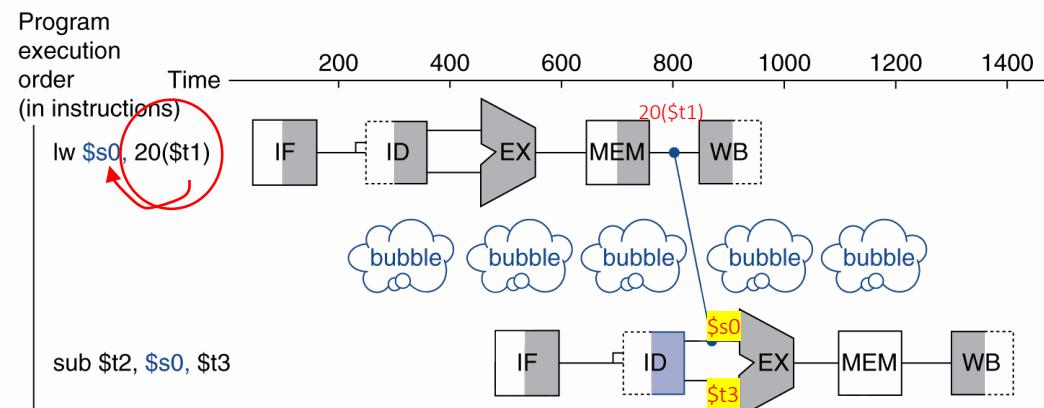
data read from memory loaded into the register in WB stage
(it can access memory in MEM stage)

Load-Use Data Hazard



Load-Use Data Hazards

- It cannot always avoid stalls by forwarding
 - It still incurs a stall even using forwarding
 - Can only access the data as early as in MEM stage

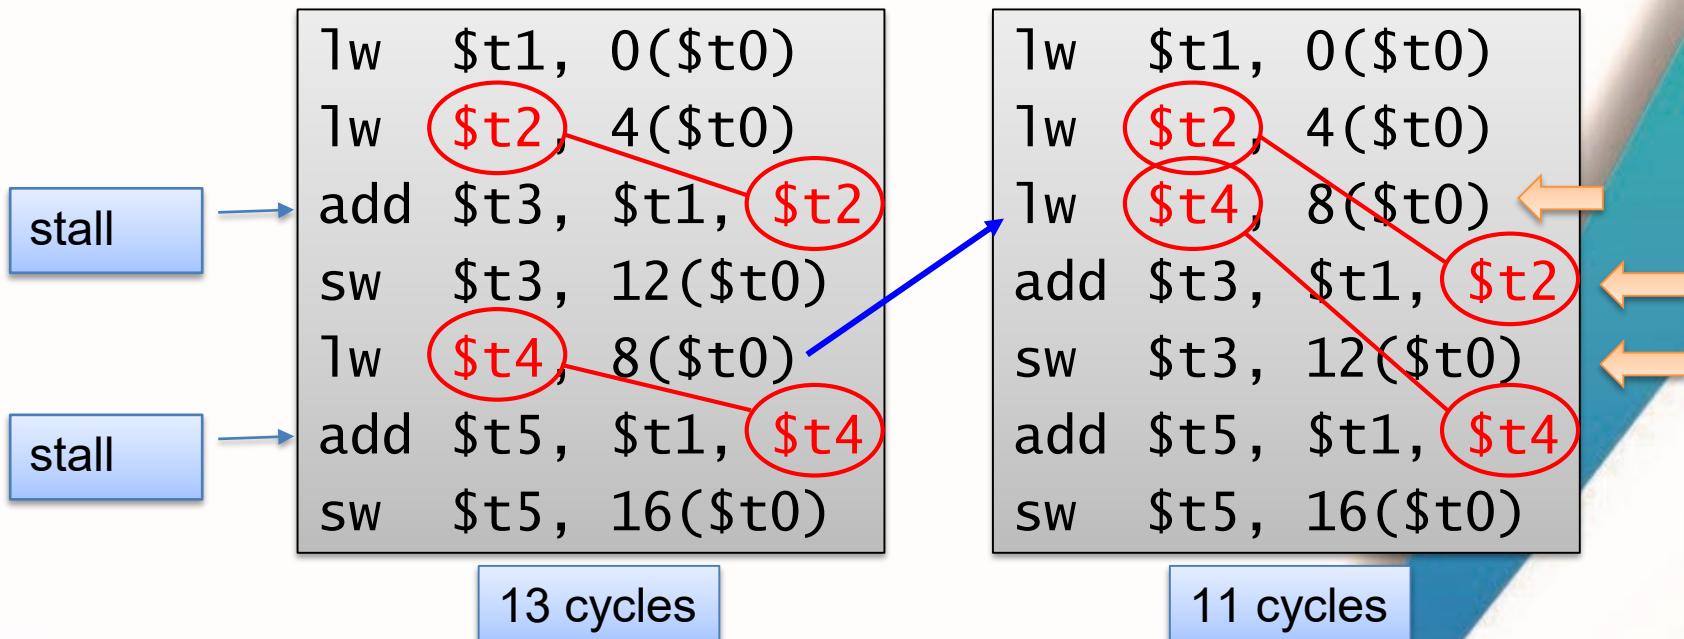


Compared to R-type instructions?

Code Scheduling to Avoid Stalls

- Reorder code to avoid use of the loaded result in the next instruction
- C code : $a = b + e;$
 $c = b + f;$

After reordering the code, the results of two lw's would be used one instruction after the lw instruction



```

lw    $t1, 0($t0)
lw    $t2, 4($t0)
add   $t3, $t1, $t2
sw    $t3, 12($t0)
lw    $t4, 8($t0)
add   $t5, $t1, $t4
sw    $t5, 16($t0)

```

Instr	CC1	CC2	CC3	CC4	CC5	CC6	CC7	CC8	CC9	CC10	CC11	CC12	CC13
lw	IF	ID	EXE	ME M	WB								
lw		IF	ID	EXE	ME M	WB							
add				IF	ID	EXE	ME M	WB					
sw				IF	ID	EXE	ME M	WB					
lw					IF	ID	EXE	ME M	WB				
add						IF	ID	EXE	ME M	WB			
sw							IF	ID	EXE	ME M	WB		

```

lw    $t1, 0($t0)
lw    $t2, 4($t0)
lw    $t4, 8($t0)
add   $t3, $t1, $t2
sw    $t3, 12($t0)
add   $t5, $t1, $t4
sw    $t5, 16($t0)

```

Instr	CC1	CC2	CC3	CC4	CC5	CC6	CC7	CC8	CC9	CC10	CC11
lw	IF	ID	EXE	ME M	WB						
lw		IF	ID	EXE	ME M	WB 4(\$t0)					
lw			IF	ID	EXE	ME M	WB				
add				IF	ID	EXE	ME M	WB			
sw					IF	ID	EXE	ME M	WB		
add						IF	ID	EXE	ME M	WB	
sw							IF	ID	EXE	ME M	WB

Example

For each code sequence below, state whether it must stall, can avoid stalls using only forwarding, or can execute without stalling or forwarding.

Check Yourself

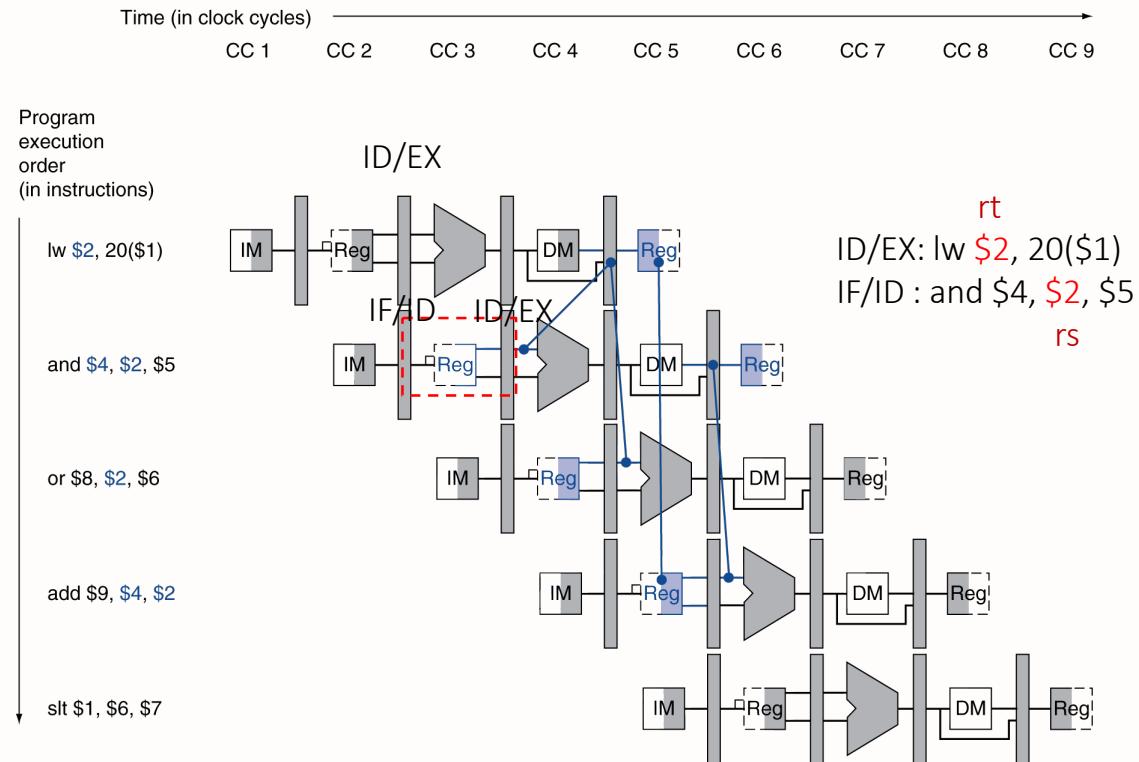
Sequence 1	Sequence 2	Sequence 3
lw \$t0,0(\$t0) add \$t1,\$t0,\$t0	add \$t1,\$t0,\$t0 addi \$t2,\$t0,#5 addi \$t4,\$t1,#5	addi \$t1,\$t0,#1 addi \$t2,\$t0,#2 addi \$t3,\$t0,#2 addi \$t3,\$t0,#4 addi \$t5,\$t0,#5

lw: 1 stall even with forwarding

r-type: no stall with forwarding

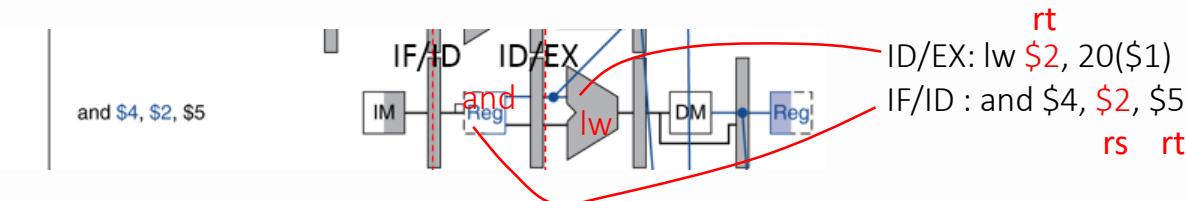
branch: 1 stall

Detect Load-Use Hazard in ID stage (since we need a stall anyway, do it as early as it can)



Load-Use Hazard Detection

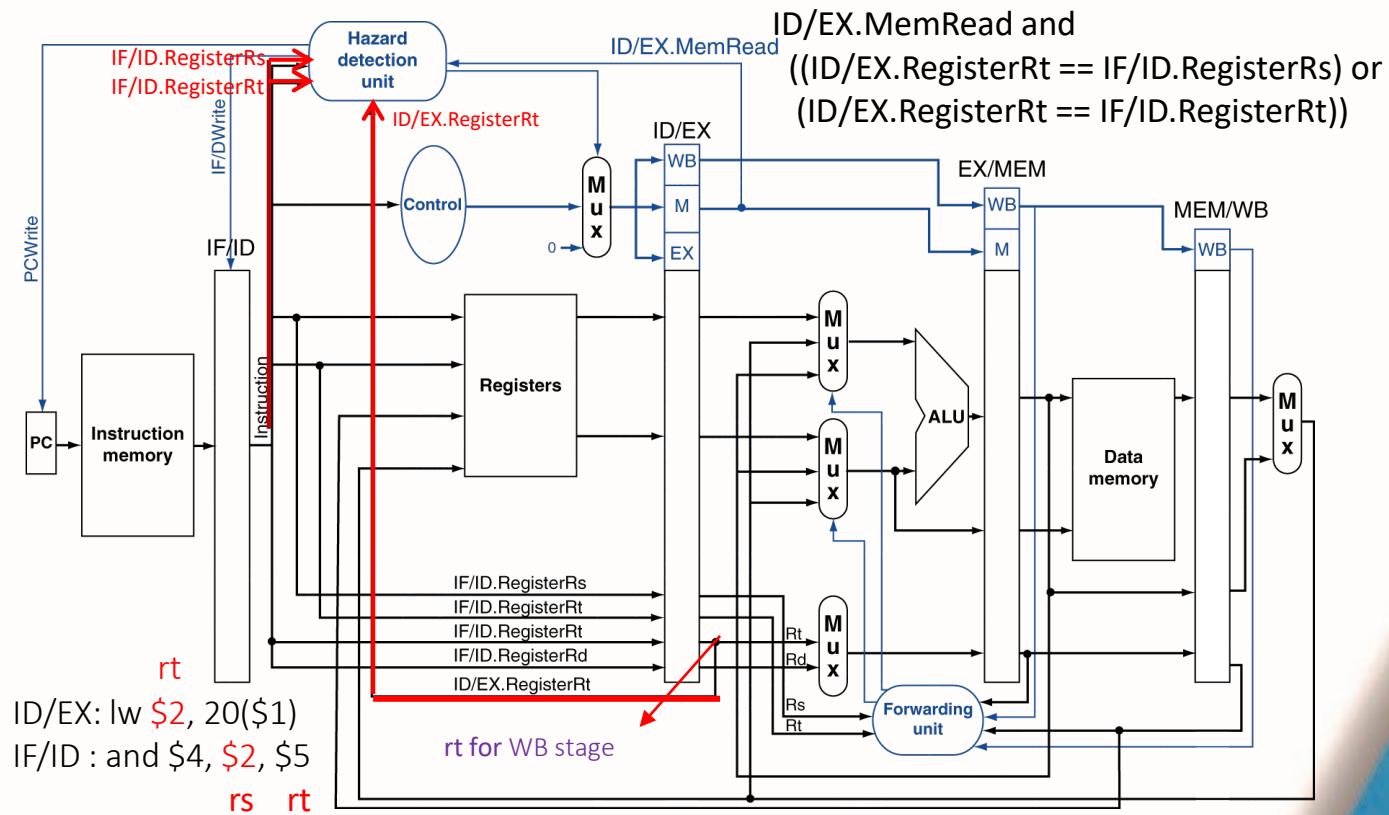
- ❑ Check in ID stage
- ❑ ALU operand register numbers in ID stage are given by
 - ❑ IF/ID.RegisterRs, IF/ID.RegisterRt



- ❑ Load-use hazard when
 - ❑ ID/EX.MemRead and
 $((ID/EX.RegisterRt == IF/ID.RegisterRs) \text{ or } (ID/EX.RegisterRt = IF/ID.RegisterRt))$
- stall the pipeline**

ID/EX.MemRead==1 -> load instruction

Datapath with Hazard Detection

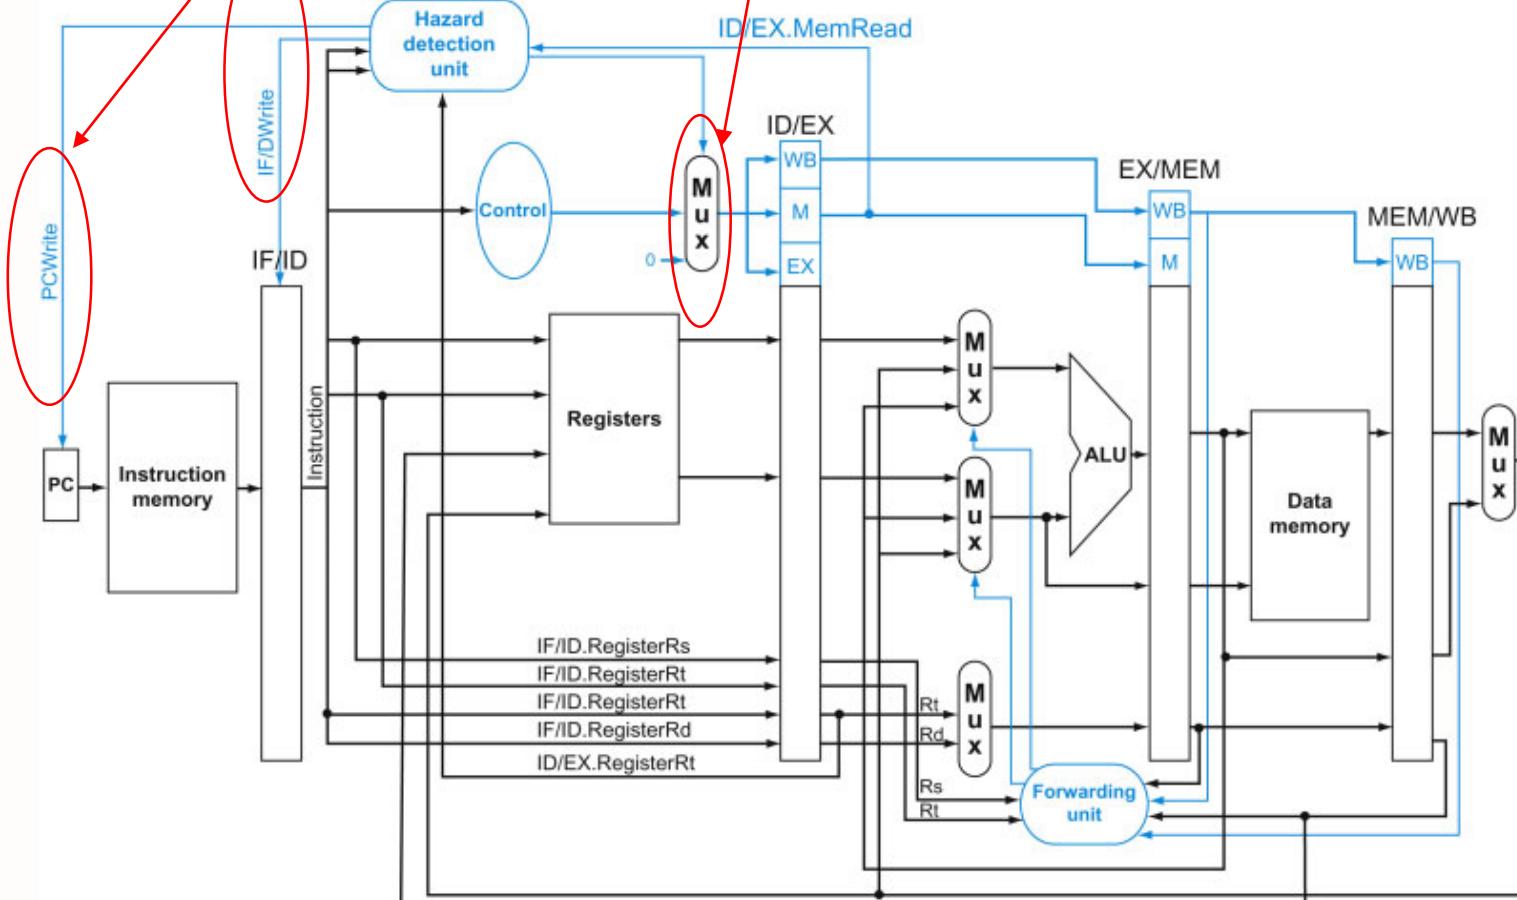


How to Stall the Pipeline

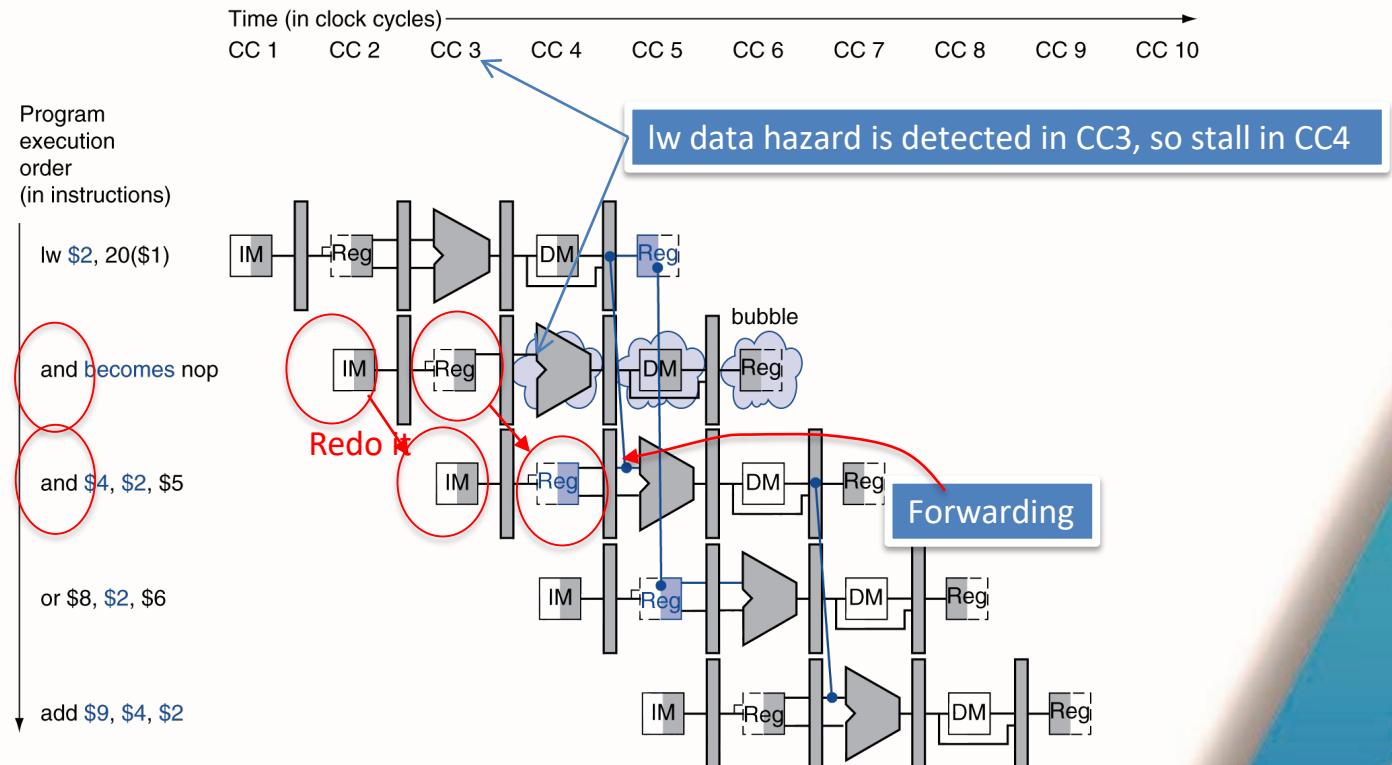
- Force control values in ID/EX register to 0
 - EX, MEM and WB do **nop** (no-operation)
- Prevent update of PC and IF/ID register
 - Current instruction is decoded again
 - Following instruction is fetched again
 - 1-cycle stall allows MEM to read data for 1w
 - Can subsequently forward to EX stage

Prevent update of PC and IF/ID register

Force control values in ID/EX register to 0



Stall/Bubble in the Pipeline (CC3)



Stalls and Performance

- Stalls reduce performance
 - It is required to obtain correct results
- Compiler can re-arrange code to avoid hazards and stalls
 - Requires knowledge of the pipeline structure

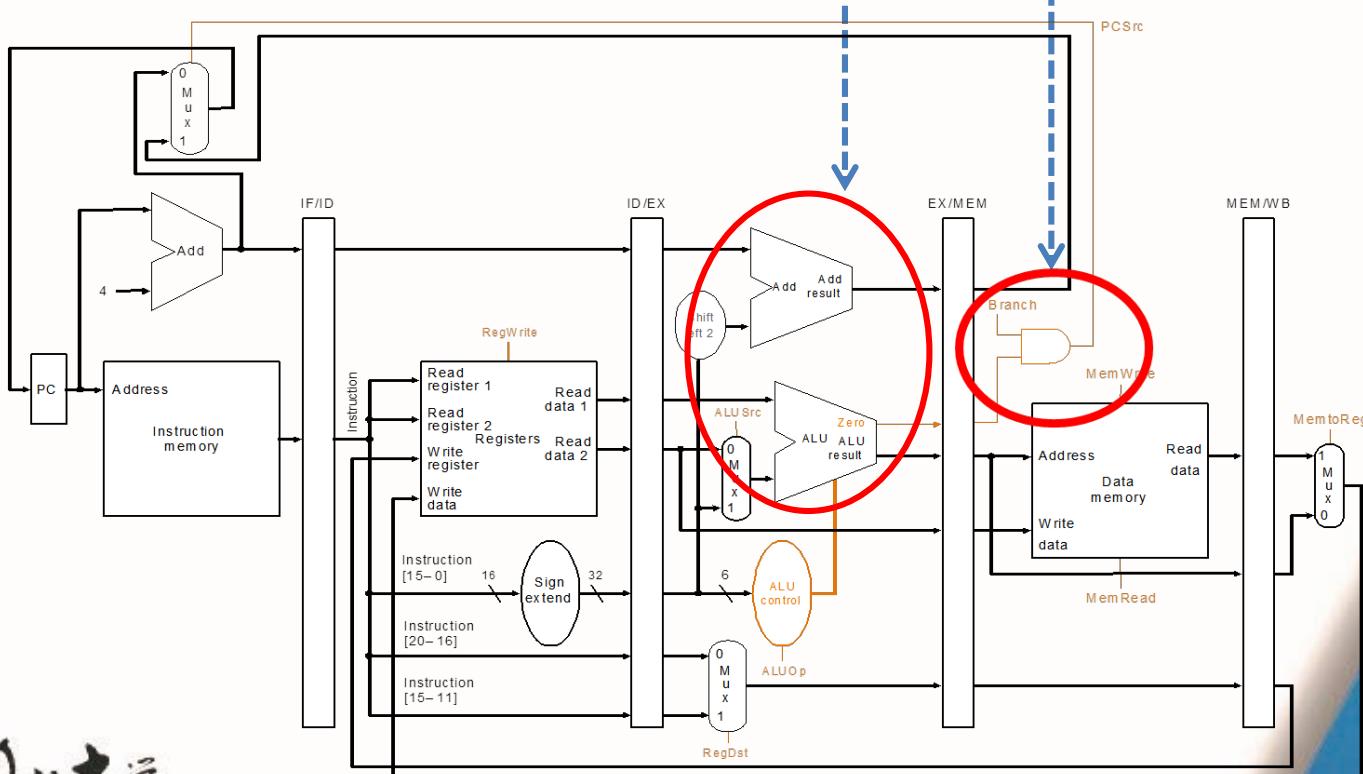
Branch

Step name	Action for R-type instructions	Action for memory-reference instructions	Action for branches	Action for jumps
Instruction fetch			IR = Memory[PC] PC = PC + 4	
Instruction decode/register fetch		A = Reg [IR[25-21]] B = Reg [IR[20-16]]		
Execution, address computation, branch/ jump completion	ALUOut = A op B	ALUOut = A + sign-extend (IR[15-0])	if (A ==B) then PC = PC+ sign-ext(IR[15-0])<<2	PC = PC [31-28] II (IR[25-0]<<2)
Memory access or R-type completion	Reg [IR[15-11]] = ALUOut	Load: MDR = Memory[ALUOut] or Store: Memory [ALUOut] = B		
Memory read completion		Load: Reg[IR[20-16]] = MDR		

Branch Outcome

- Branch in MEM stage

- Know rs-rt == 0 in EXE stage
- Control PCSrc in MEM stage



Control Hazards (Branch Hazards)

- The fetched instruction is not the one that shall be executed because of branch conditions
- Branch determines flow of control
 - Fetching next instruction (IF) depends on branch outcome (MEM)
 - Pipeline can't always fetch correct instruction
- In MIPS pipeline
 - Need to compare registers and compute target early in the pipeline
 - Add hardware to do it in ID stage

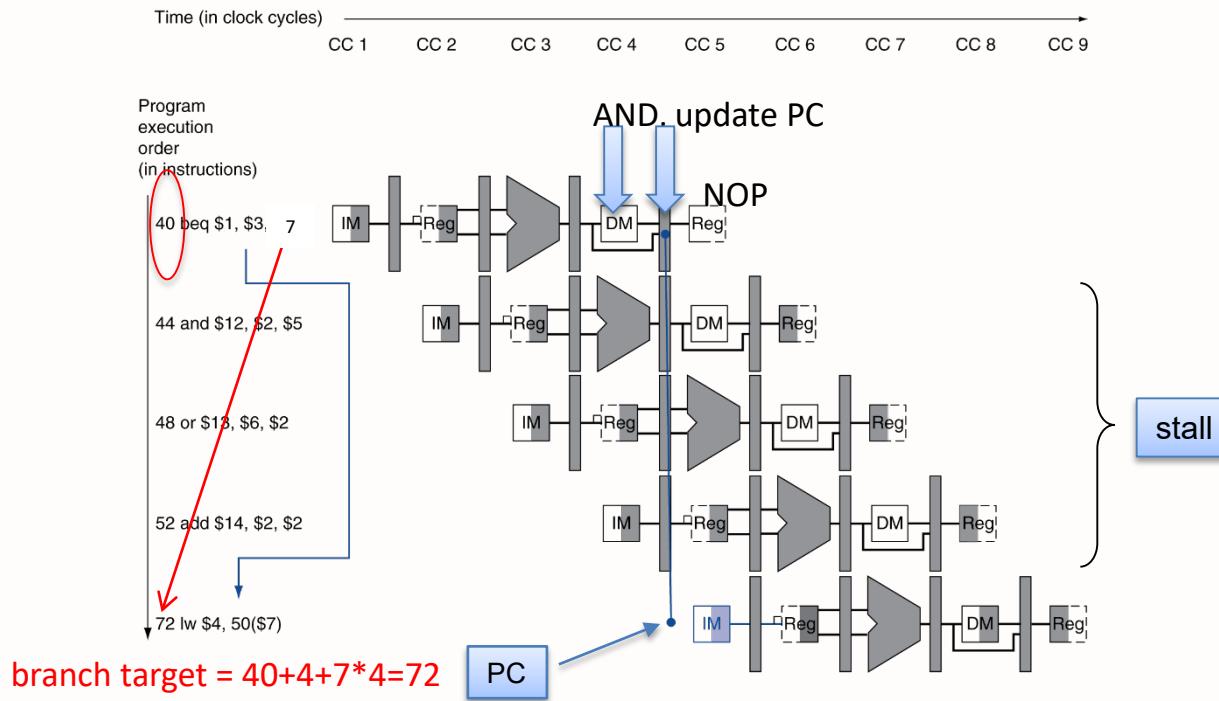
See P.320, Fig.4.62

Add extra hardware to test registers, calculate the branch address, and update the PC during ID stage

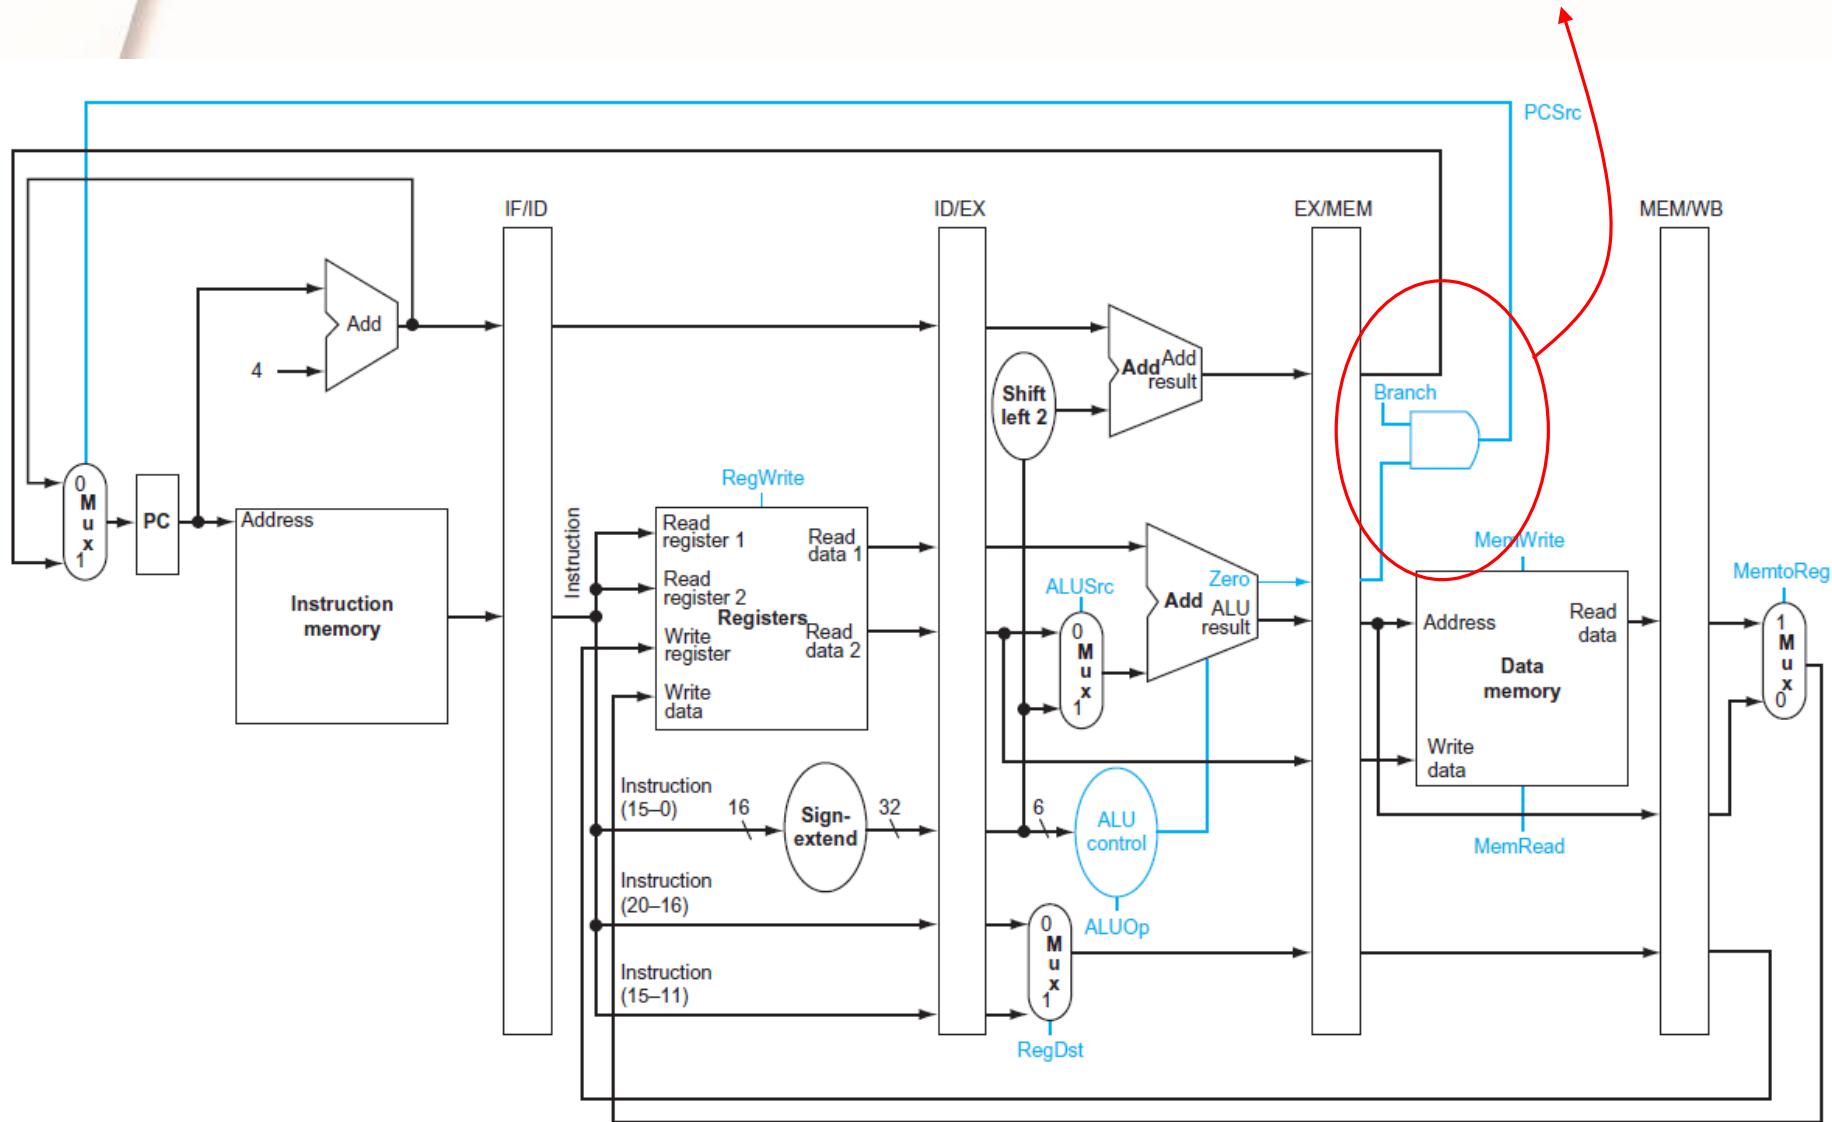
Note that it is often the case for longer pipelines that the branch cannot be resolved in ID stage.
(larger slowdown)

Control Hazards

Branch taken incurs three stalls (3 cycles)



branch finished in MEM



Reducing Branch Delay

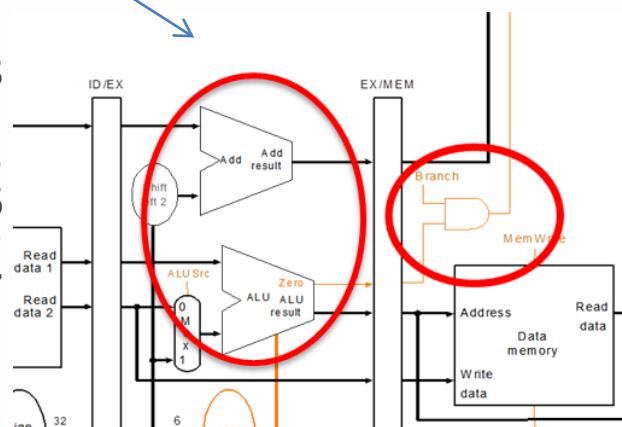
- Determining branch condition in ID stage (using additional hardware)

- Target address adder
- Register comparator

- Example: branch taken

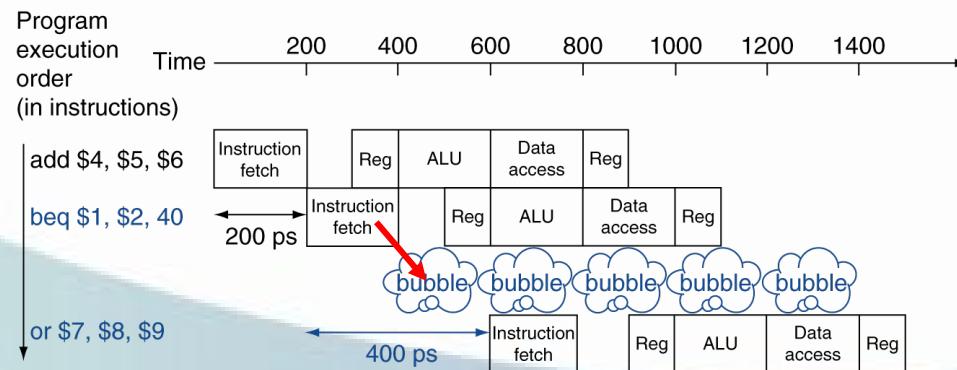
```

36: sub $10, $4, $8
40: beq $1, $3, 7
44: and $12, $2, $5
48: or $13, $2, $6
52: add $14, $4, $2
56: slt $15, $6, $7
...
72: lw $4, 50($7)
    
```

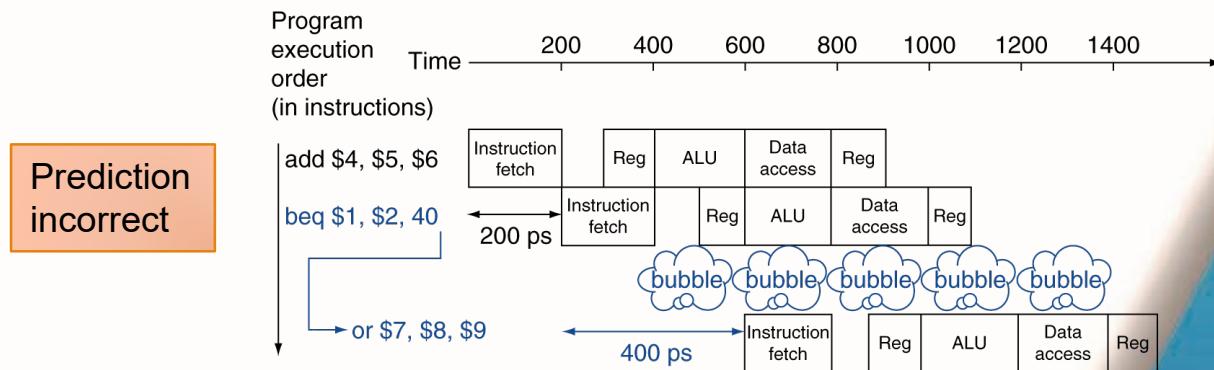
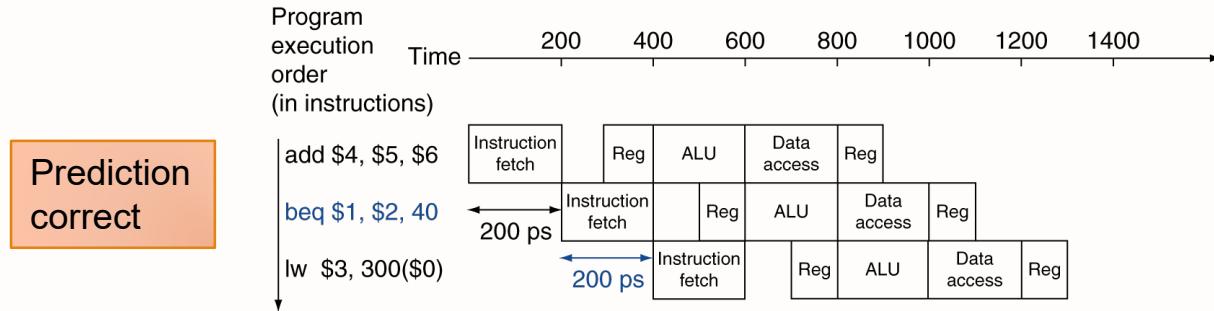


Stall on Branch

- If branch result comes out in MEM stage, we need to flush instructions in IF/ID, ID/EX registers.
- Even if we know it's taken in ID stage, a stall is needed
- Use **prediction** to handle branches
 - Always predict that branches will not be taken.
 - When the prediction is right, the pipeline can proceed at the full speed; when the branches are taken, the pipeline will stall anyway even without prediction.



MIPS with Predict Not Taken



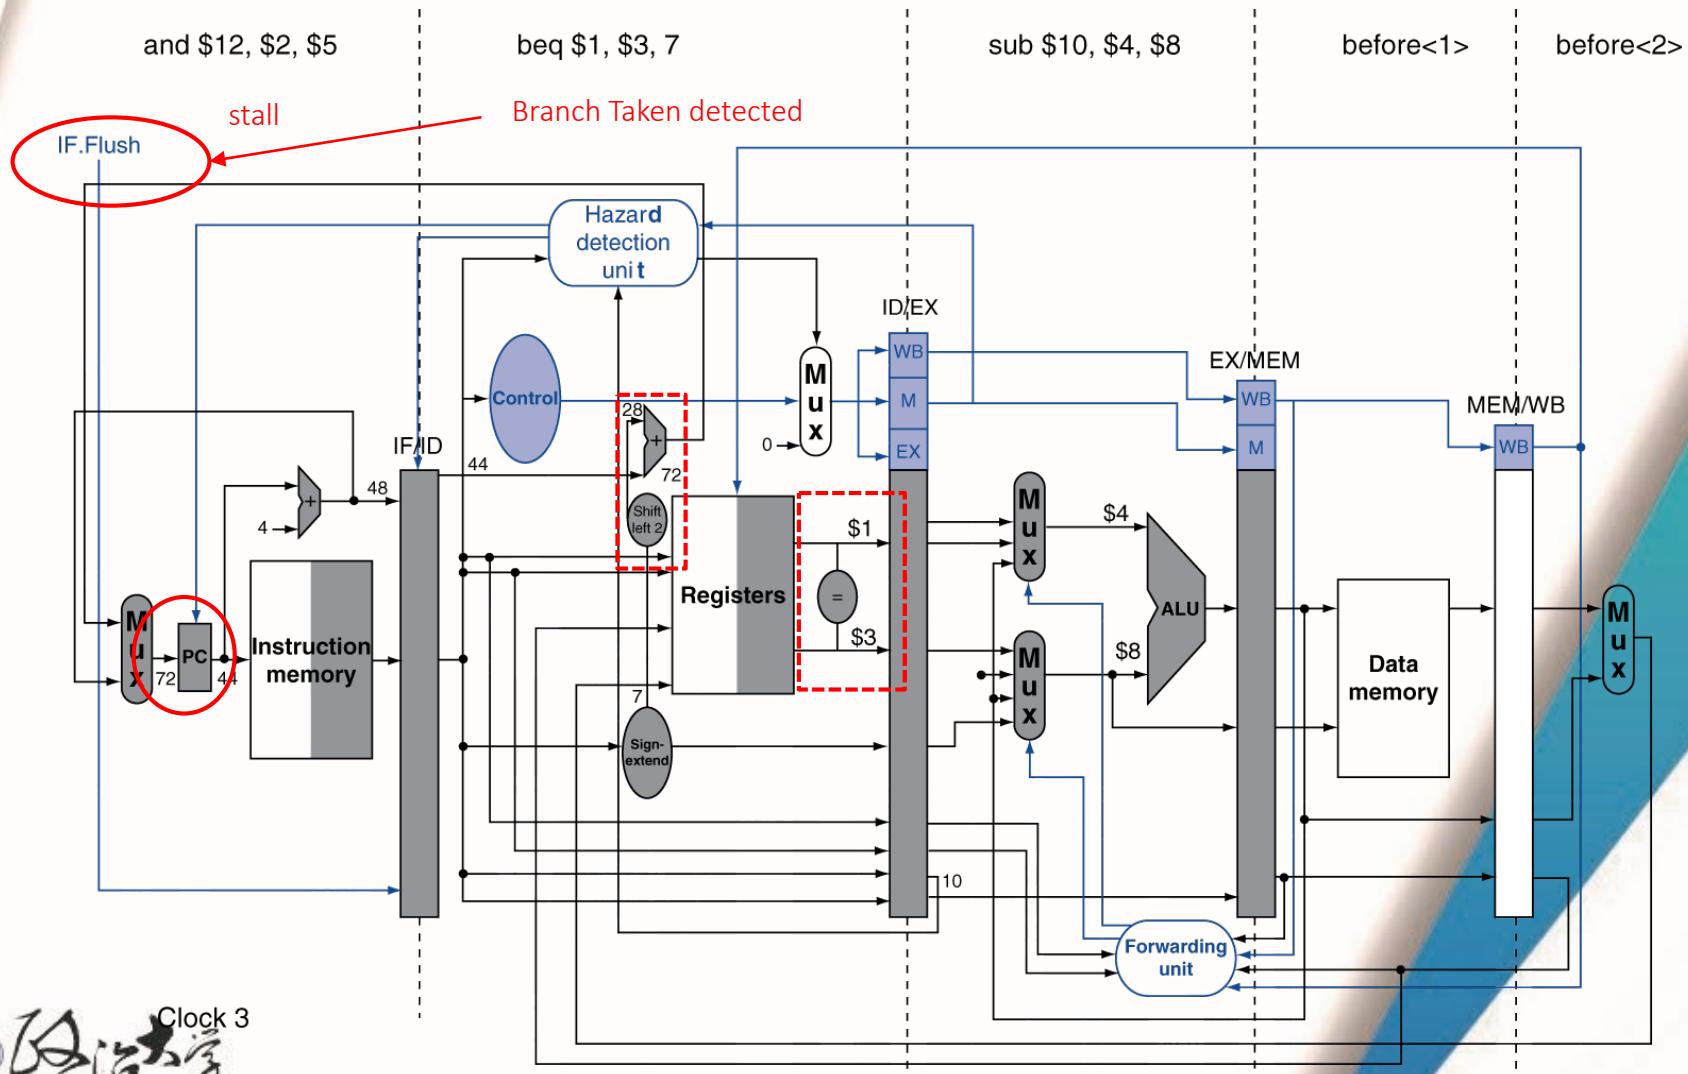
Pipelined Branch

Show what happens when the branch is taken in this instruction sequence, assuming the pipeline is optimized for branches that are not taken and that we moved the branch execution to the ID stage:

```
36 sub $10, $4, $8
40 beq $1, $3, 7 # PC-relative branch to 40 + 4 + 7 * 4 = 72
44 and $12, $2, $5
48 or $13, $2, $6
52 add $14, $4, $2
56 slt $15, $6, $7
...
72 lw $4, 50($7)
```

Example: Branch Taken

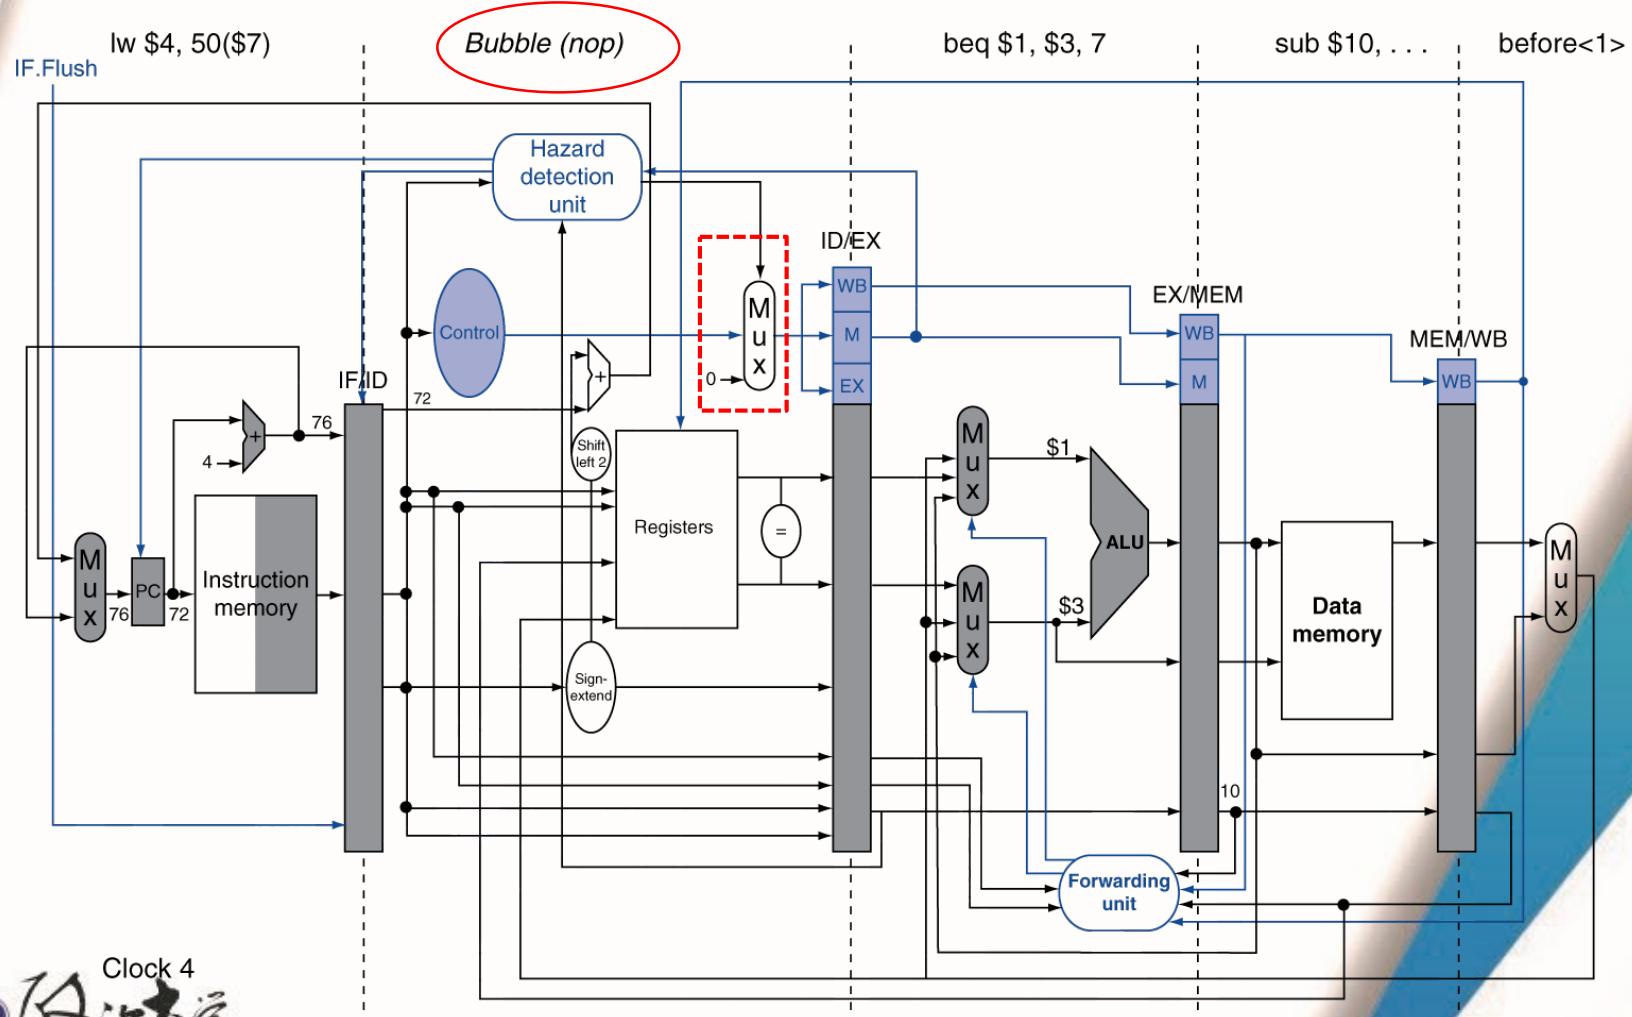
At the ID stage of CC3, branch taken is detected



Example: Branch Taken

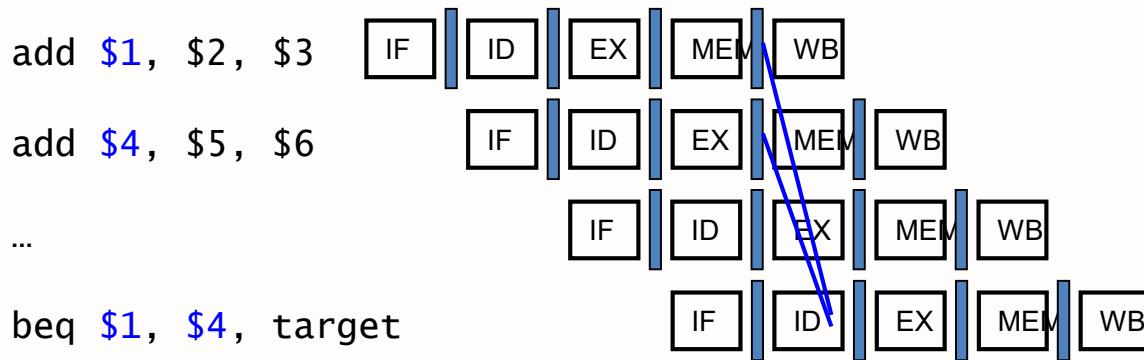
72 is selected as the next PC

Zeros are filled for ID/EX and the instruction (nop: sll \$0, \$0, 0)



Data Hazards for Branches

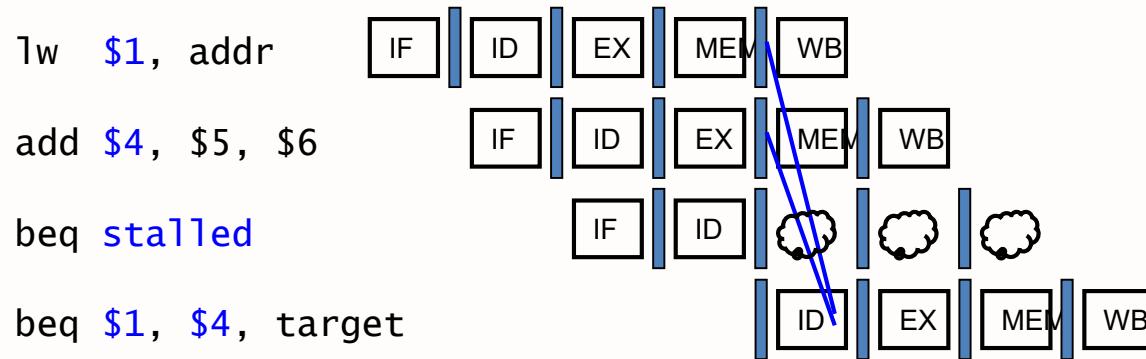
- If a comparison register is a destination of 2nd or 3rd preceding ALU instruction



- Can resolve using forwarding

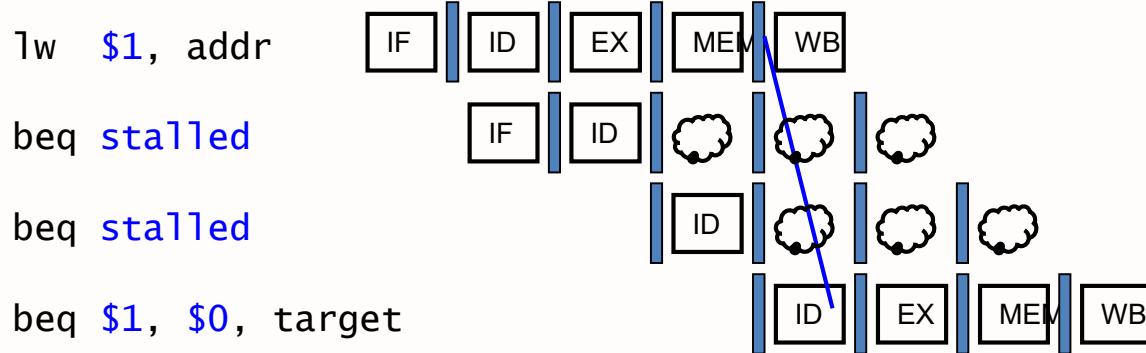
Data Hazards for Branches

- If a comparison register is a destination of preceding R-type instruction or 2nd preceding **load** instruction
 - Need 1 stall cycle



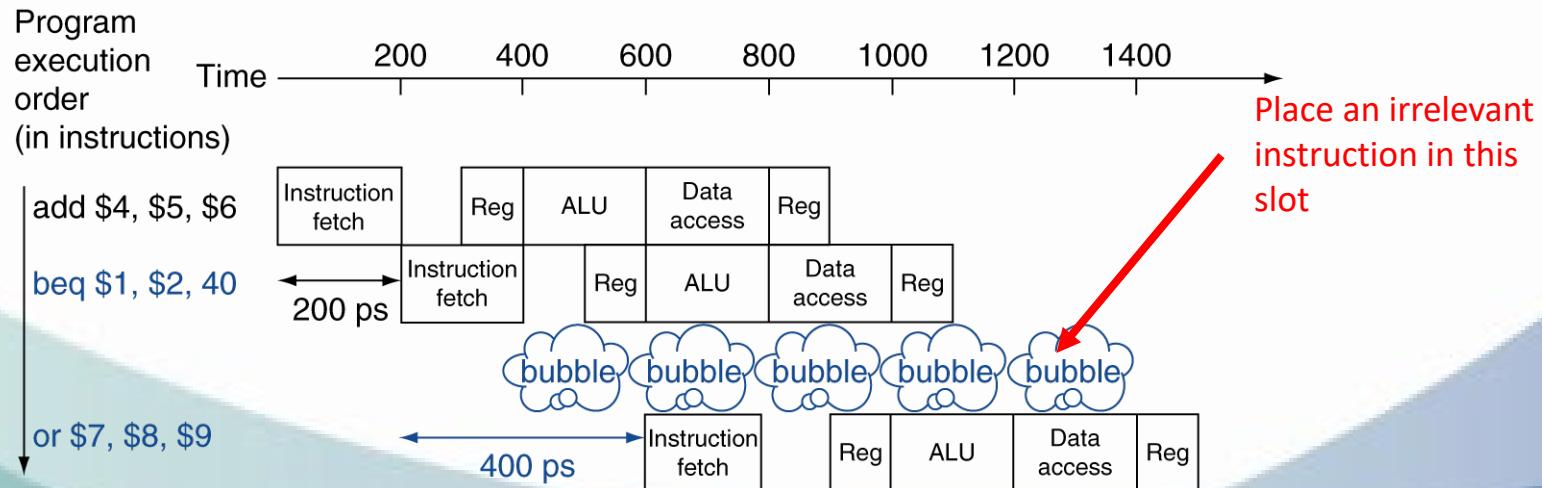
Data Hazards for Branches

- If a comparison register is a destination of immediately preceding load instruction
 - Need 2 stall cycles



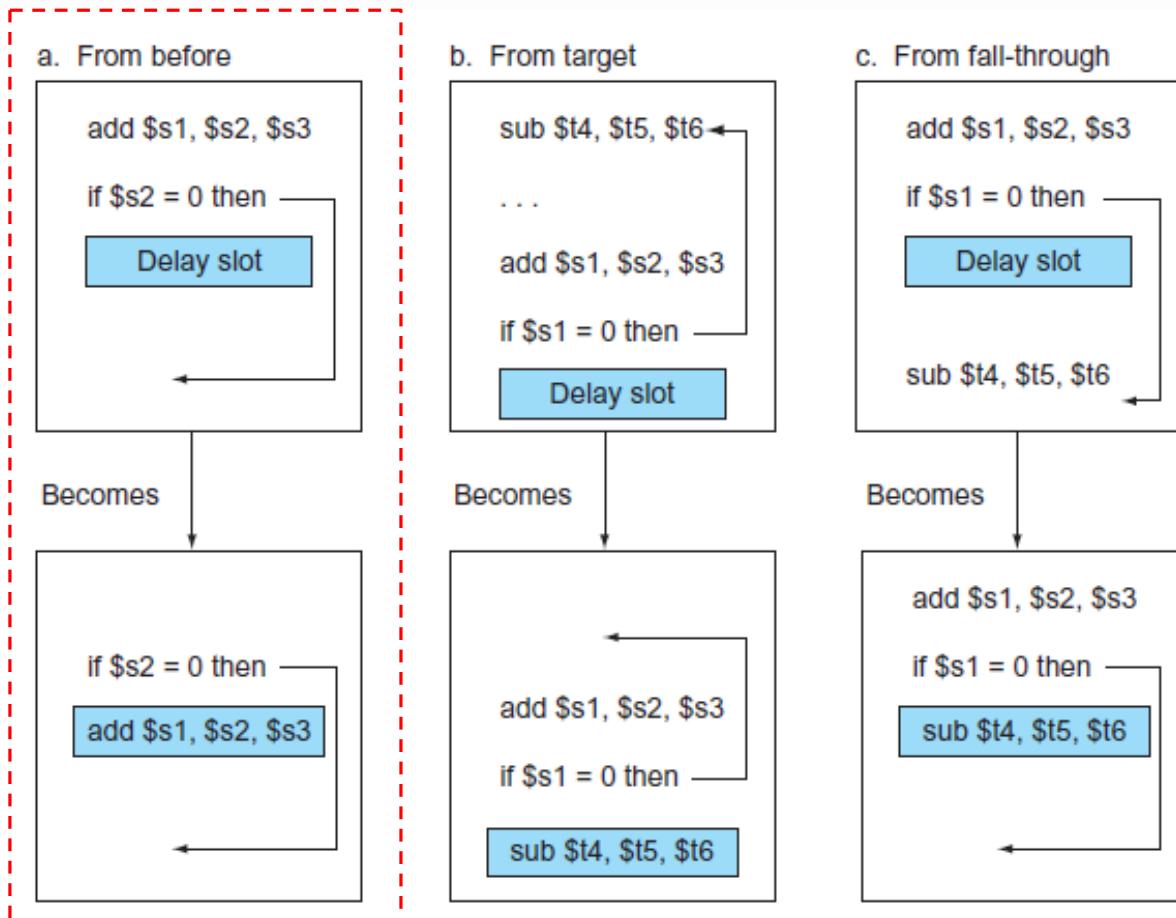
Delayed Branch

- Combine Predict-not-taken and branch decision in ID stage
 - the following instruction is always executed
 - Insert one irrelevant instruction after a branch instruction
 - Compilers and assemblers would try to insert an always-executed instruction immediate after a branch instruction in the branch delay slot.
 - There is no stall needed for a branch instruction if the compiler can find an instruction to put in slot



Scheduling a Branch Delay Slot

Best choice



Sophisticated Branch Prediction

■ Static branch prediction

- Based on typical branch behavior
- Example
 - loop: Predict backward branches taken
 - if: Predict forward branches not taken

■ Dynamic branch prediction

- Prediction of branches at runtime
- Hardware measures actual branch behavior
 - e.g., record recent history of each branch
- Assume future behavior will continue the trend
 - When wrong, stall while re-fetching, and update history

Dynamic Branch Prediction

- Generally, using “**predict branch not taken**” and stall should often resolve a control hazard
- Under some circumstances, it is not good; for example, “for loop,” which is taken almost every time until the stopping criterion is satisfied.
- Dynamic branch prediction
 - branch prediction buffer (also called branch history table)
 - 1-bit Branch prediction buffer
 - a small memory that has a bit to indicate whether the branch was recently taken or not.
 - Using a hash table:<instruction addr, isTaken> to store branch history
 - Do the same action as the previous one did based on the hash table (branch history)

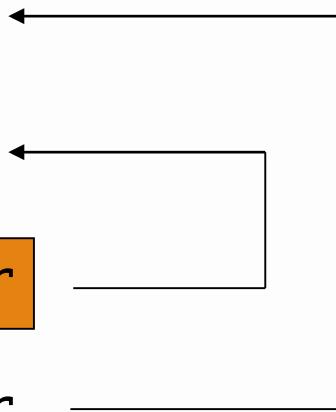
1-Bit Predictor

outer: ...

inner: ...

...
beq ..., ..., inner

...
beq ..., ..., outer



Will lead to twice mispredictions for the Inner loop (in the first and last iterations)

The last iteration: inevitable

1-Bit Predictor

```
for(int i=0;i<1000;i++)  
{  
    for(int j=0;j<1000;j++) {  
        //A  
    } //B  
}
```

0

Mispredict twice for an entire loop
(mispredict on the first and last loop iterations)

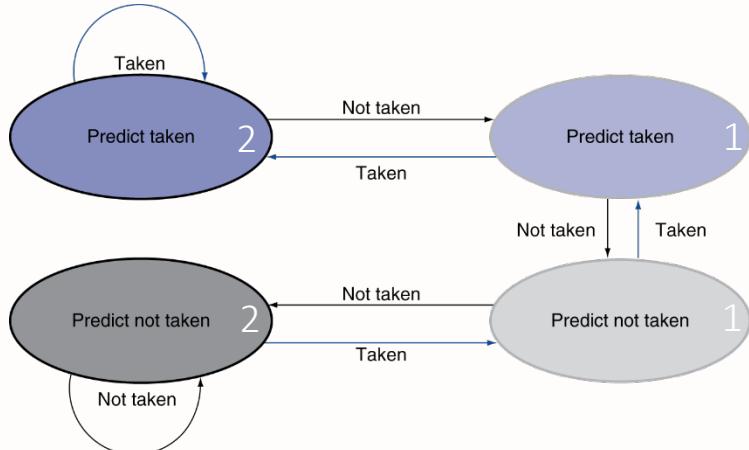
	Predict	Outcome
A	NT	T
A	T	T
A	T	T
...999		
A	T	NT
B	NT	T
A	NT	T
A	T	T
...999		
A	T	NT
B	T	T
A	NT	T
A	T	T
...999	...	
A	T	NT
B	T	NT

inevitable

Can be taken care of

2-Bit Predictor

```
for(int i=0;i<1000;i++)
{
    for(int j=0;j<1000;j++) {
        } //A
    } //B
```



0

	Predict	Outcome
A	NT1	T
A	T1	T
A	T2	T
...999		
A	T2	NT
B	NT1	T
A	T1	T
A	T2	T
...999		
A	T2	NT
B	T1	T
A	T1	T
...999		
A	T2	NT
B	T2	NT

inevitable

← Taken care

Calculating the Branch Target

- Even though having a predictor, we need to calculate the target address (ID stage)
 - 1-cycle penalty for a taken branch
- Branch target buffer
 - Cache of target addresses
 - The target address was stored when the branch was taken last time
 - Indexed by PC when the instruction is fetched
 - If it hits an entry in the table and was predicted taken, the target address can be fetch from the cache.

Conclusion

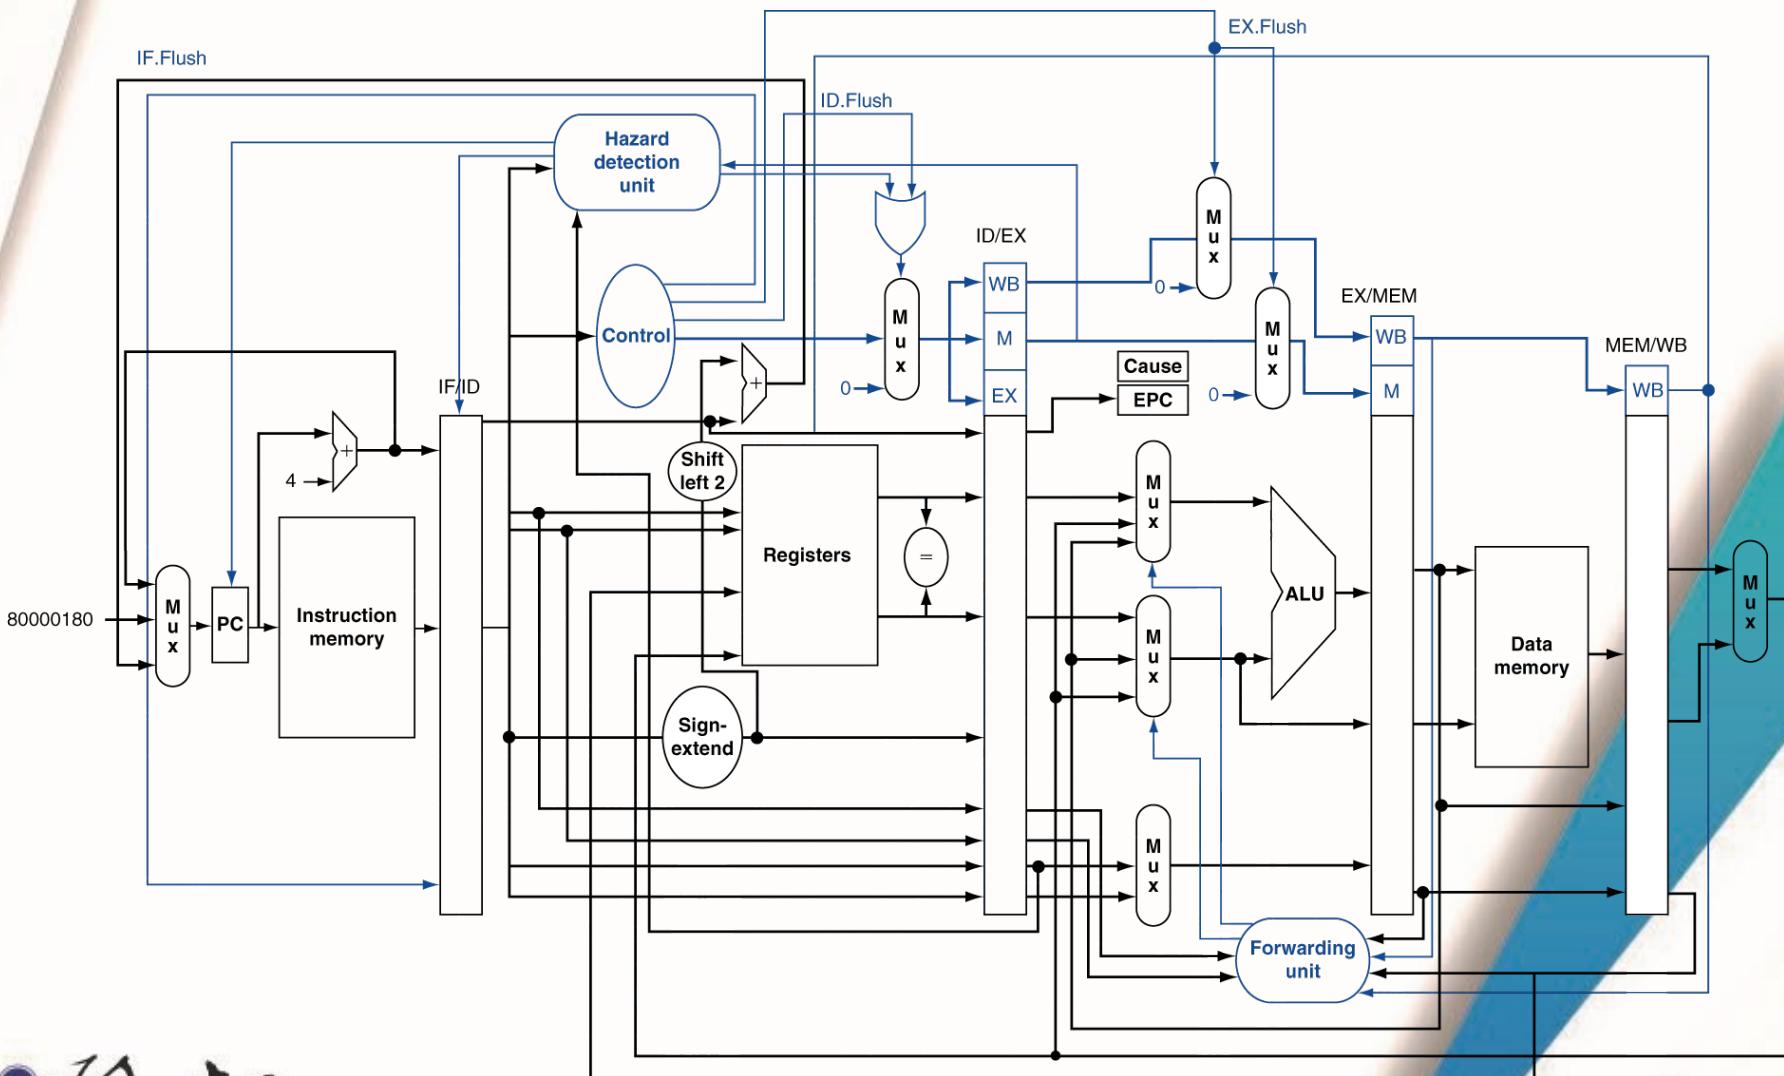
- Pipelining improves instruction throughput instead of latency
 - More instructions executed using parallelism
- Three Hazards
 - structural, data, control
 - Stall, Forwarding

MISC.

Exceptions and Interrupts

- Exceptions raised when unexpected events occur
- Exceptions and interrupts in general
 - Exceptions raised in CPU
 - undefined opcode, overflow, ...
 - Interrupt raised from an external I/O controller
- Exceptions in pipeline is considered a control hazard
- Dealing with them in a pipelined architecture requires
 - Complete previous instructions
 - Flush following instructions
 - Set Cause and EPC register values
 - Transfer control to handler (handler at 8000 00180)

Pipeline with Exceptions



Exceptions

- Require the pipeline flush the instruction
 - Set control signals to 0
- Executing Handler, then returning to the previous instruction
- PC (PC+4) saved in EPC register
 - Identifies causing instruction
 - Exceptions: PC-4 (resume the current instruction)
 - Interrupt: PC+4 (go to the next instruction)

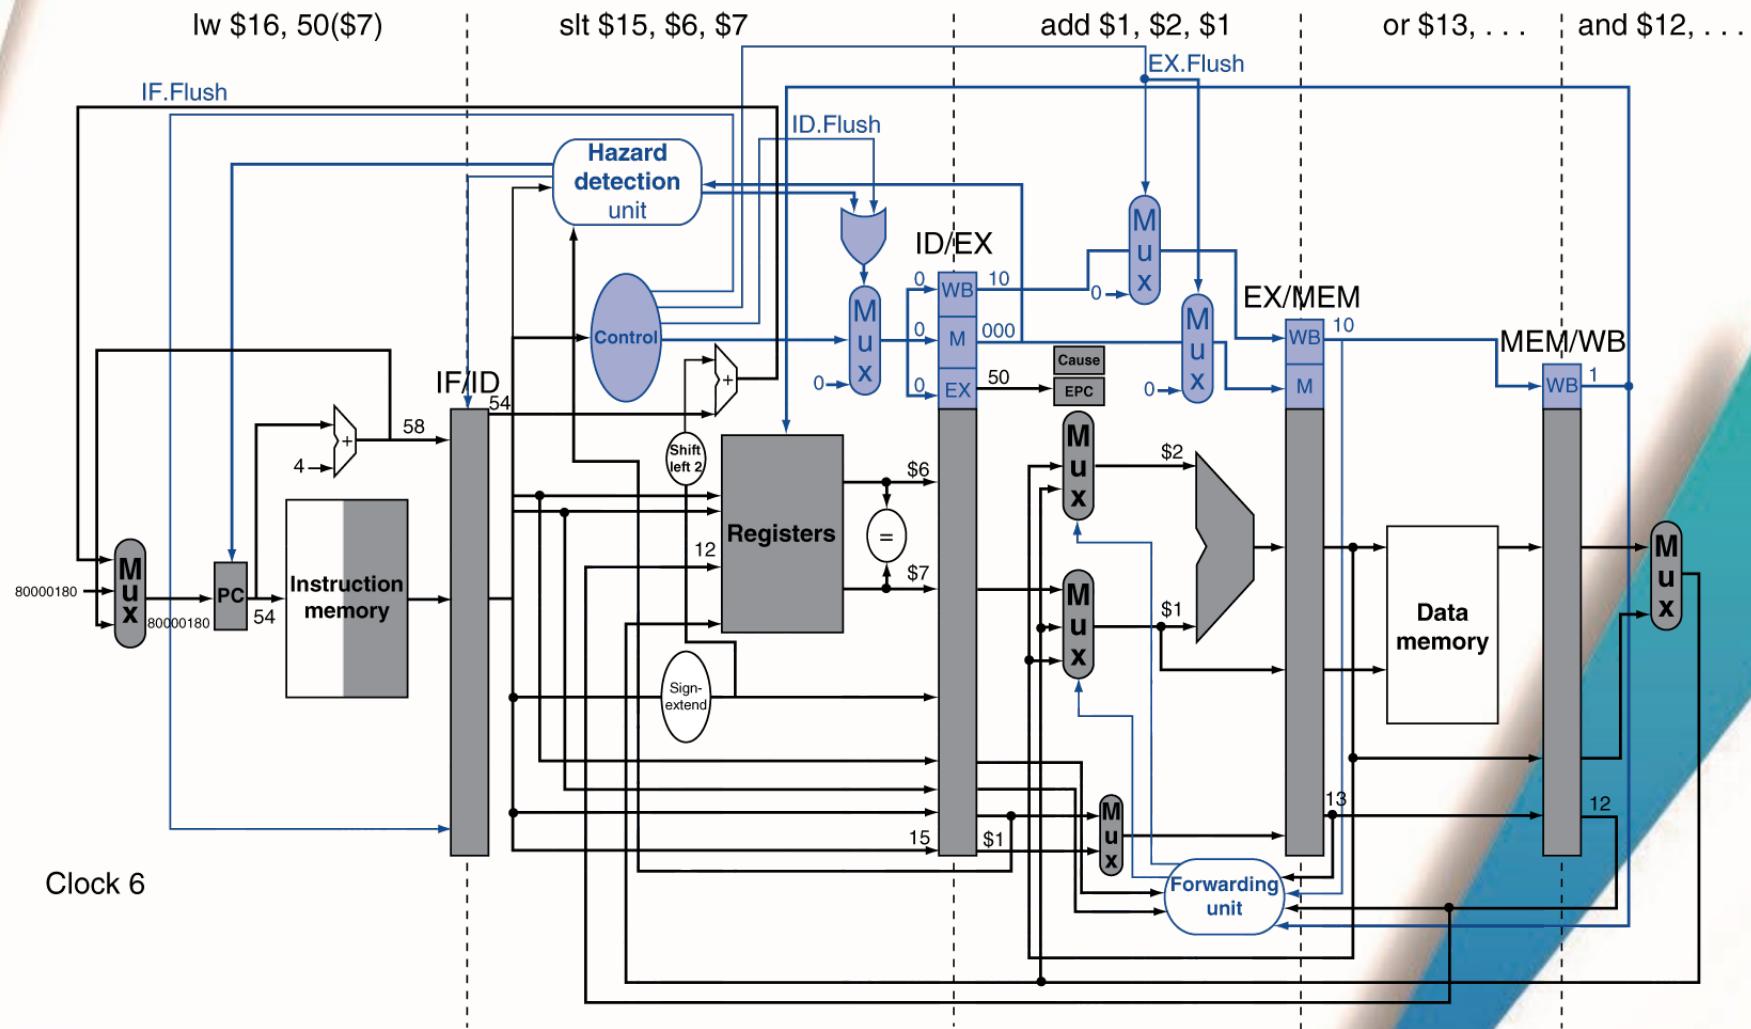
Exception Example

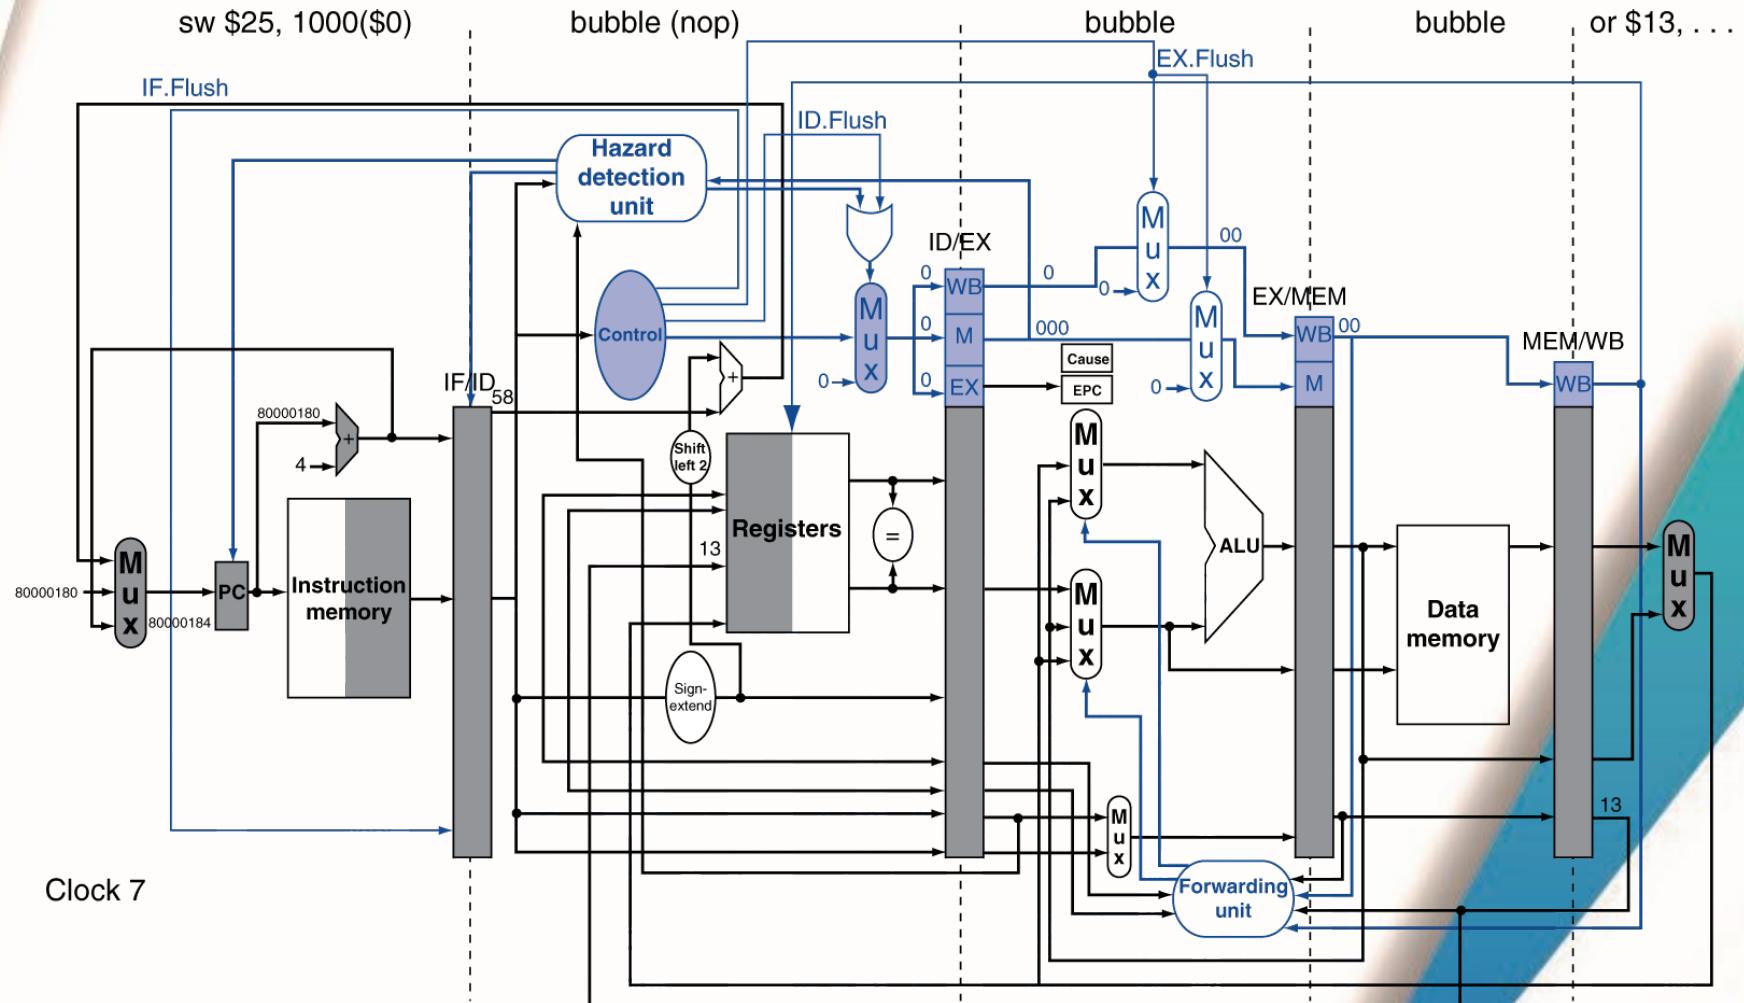
```
40      sub    $11, $2, $4  
44      and    $12, $2, $5  
48      or     $13, $2, $6  
4C      add    $1, $2, $1  
50      slt    $15, $6, $7  
54      lw     $16, 50($7)
```

→ Exception

Handler

80000180	sw	\$25, 1000(\$0)
80000184	sw	\$26, 1004(\$0)
...		





Multiple Exceptions Occur

- Deal with multiple exceptions
 - Handle the exception that comes first
 - Flush the following instructions
 - Cannot deal with out-of-order execution
 - Stop pipeline
 - Save all the related states and exception causes
 - Let the software (OS) to handle

Instruction-Level Parallelism (ILP)

- Pipeline

- Parallel execution for multiple instructions

- To increase ILP

- Increase the number of stages (Depth ↑)
 - Running less tasks per stage, leading to a shorter clock cycle
 - Increase the number of pipeline (Width ↑)
 - Adding hardware to fetch more instructions at once
 - CPI < 1, so use Instructions Per Cycle (IPC)
 - E.g., 2 GHz 4-way multiple instruction parallelism
 - Peak CPI = 0.25, peak IPC = 4, $4 \times 2 \times 10^9 = 8$ billion of instructions per second
 - However, instruction dependencies affect the actual CPI (more than 0.25)

Multiple Instruction Issue

- Static multiple issue

- Instructions are grouped to issue by the compiler
- Compiler deals with hazards

- Dynamic multiple issue

- CPU chooses which instructions to issue each cycle
- Compiler can help by reordering instructions
- CPU deals with hazards

Speculation

- In static and dynamic multiple issues
- Guess on
 - branch target
 - Roll back if path taken is different
 - load
 - Since data not loaded could affect the following instructions that need it, load needs to be executed earlier
 - Roll back if location is updated

More Speculation

- Predict branch and continue issuing
 - Don't commit until branch outcome determined
- Load speculation
 - Avoid load and cache miss delay
 - Predict the effective address
 - Predict loaded value
 - Load before completing outstanding stores
 - Bypass stored values to load unit
 - Don't commit load until speculation cleared

Compiler/Hardware Speculation

- Compiler can reorder instructions
 - e.g., move load before store
 - Can include “fix-up” instructions to recover from incorrect guess

- Hardware can look ahead for instructions to execute
 - Buffer results until it determines they are actually needed
 - Flush buffers on incorrect speculation

Does Multiple Issue Work?

- Yes, but not as much as we'd like
- Programs have real dependencies that limit ILP
- Some dependencies are hard to eliminate
 - e.g., pointer aliasing
- Some parallelism is hard to expose
 - Limited window size during instruction issue
- Memory delays and limited bandwidth
 - Hard to keep pipelines full
- Speculation can help if done well

Power Efficiency

- Complexity of dynamic scheduling and speculations requires power
- Multiple simpler cores may be better

Microprocessor	Year	Clock Rate	Pipeline Stages	Issue width	Out-of-order/ Speculation	Cores	Power
i486	1989	25MHz	5	1	No	1	5W
Pentium	1993	66MHz	5	2	No	1	10W
Pentium Pro	1997	200MHz	10	3	Yes	1	29W
P4 Willamette	2001	2000MHz	22	3	Yes	1	75W
P4 Prescott	2004	3600MHz	31	3	Yes	1	103W
Core	2006	2930MHz	14	4	Yes	2	75W
UltraSparc III	2003	1950MHz	14	4	No	1	90W
UltraSparc T1	2005	1200MHz	6	1	No	8	70W

Conclusions

- ISA influences design of datapath and control
- Datapath and control influence design of ISA
- Pipelining improves instruction throughput using parallelism
 - More instructions completed per second
 - Latency for each instruction not reduced
- Hazards: structural, data, control
- Multiple issue and dynamic scheduling (ILP)
 - Dependencies limit achievable parallelism
 - Complexity leads to the power wall