

Agent

Problem-Solving Agent

觀光景點

大眾運輸

停車場

藥局

自動提款機

抱子腳

萬芳路

坡內坑

文山區

木柵

大柵路三段

歐比咖啡

道南河濱公園

木柵休息站

泉坑

貓纜動物園站

動物園

臺北市立動物園
大貓熊館

象頭塢

魚衡子

台北捷運公司木柵機廠

公館後

國立政治大學

貓纜動物園南站

新興

小坑

指南宮

芋子園坑

石坡坑

指南路三段

炮子崙

公館

猴山岳

岐山

指南路三段

指南路三段

指南路三段

貓空山水客景觀茶館

貓空觀光茶園

樟山寺

老象街45巷

貓空站

found.your.tea

3



Goal Formulation

Problem Formulation

觀光景點

大眾運輸

停車場

藥局

自動提款機

象頭塢

魚衡子

台北捷運公司木柵機廠

抱子腳

萬芳路

動物園站

臺北市立動物園



文山區

木柵

106

秀明路一段

大柵路三段

忠順街二段

歐比咖啡

木新路二段

政治大學後門

國立政治大學

動物園南站

臺北市文山區指南
實驗國民小學

貓空纜車指南宮站

貓空站

泉坑

樟山寺

老象街45巷

貓空山水客景觀茶館

found.your.tea

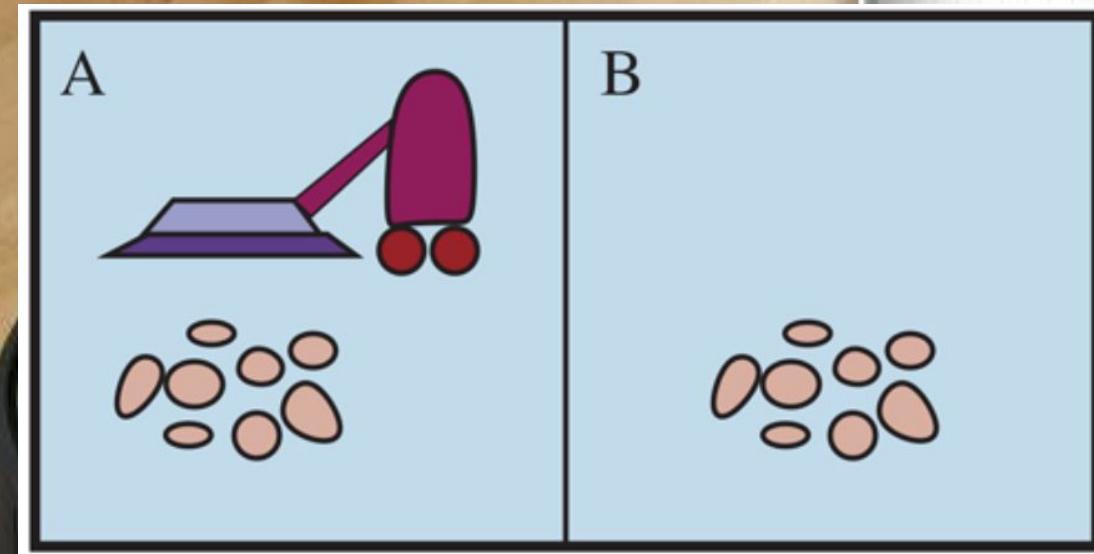
貓空觀光茶園

貓空

7

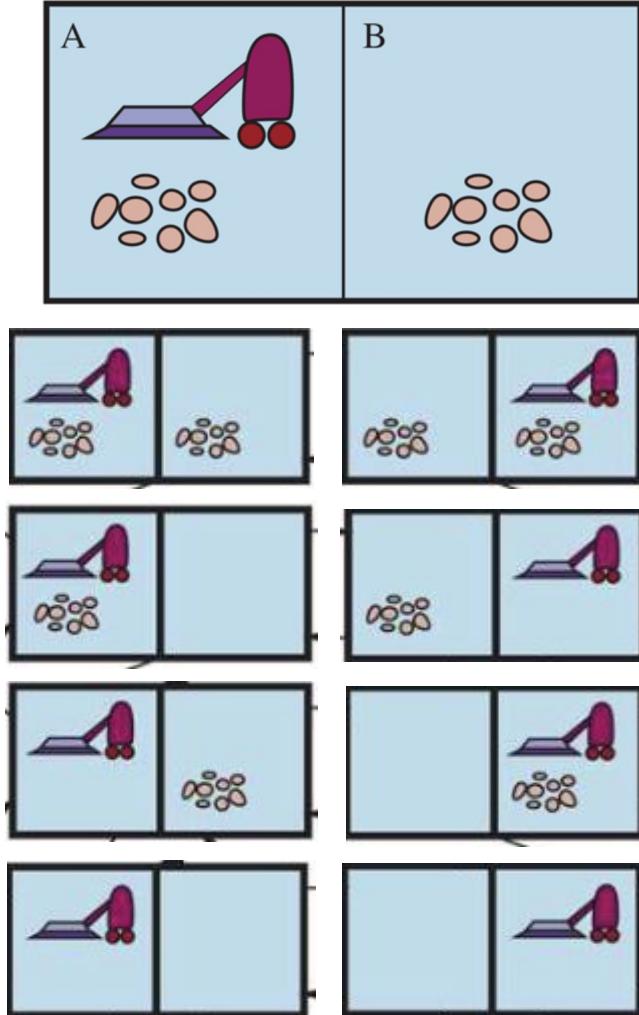
Dining
Room

Entryway



Example: Vacuum World

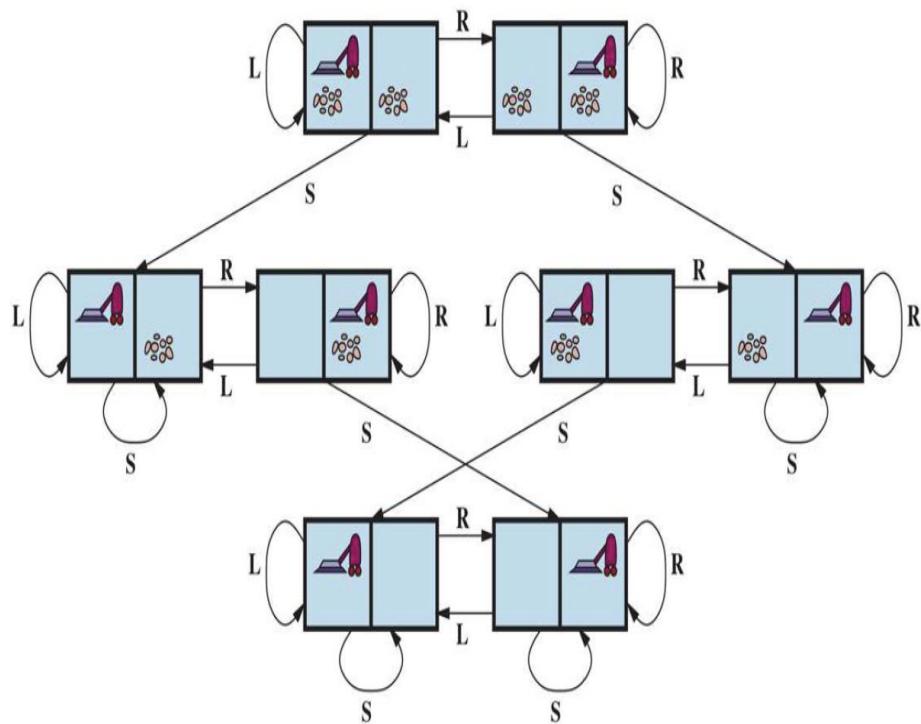
- Percepts
 - Location: A or B
 - Status: clean or dirty
- States
 - $2 \cdot 2^2 = 8$ states
- Actions
 - Move left
 - Move right
 - Suck



Example: Vacuum World

- States
 - 8 states
- Actions
 - Move left, move right, suck
- Transition model
- Initial state
- Goal test
 - Every cell is clean

State Space Graph



**Agent plans a sequence of actions
that form a path to a goal state**

Solving Problems by Searching

Course Topics

- Intelligent agents (AIMA Ch. 2)
- **Search (AIMA Ch. 3, 4, 6, 5)**
- Reasoning (AIMA Ch. 7-9, 12-15)
- Machine learning (AIMA Ch. 19-20)
- Deep learning (AIMA Ch. 22)
- Natural language processing (AIMA Ch. 24)
- Generative AI (Optional)

Search

L.-Y. Wei

Spring 2024

Search Topics

- Search problems (Ch. 3)
- Uninformed search (Ch. 3)
- Informed search (Ch. 3)
- Local search (Ch. 4)
- Adversarial search (Ch. 6)
- Constraint Satisfaction Problems (CSPs) (Ch. 5)

Search Topics

- **Search problems** (Ch. 3)
- Uninformed search (Ch. 3)
- Informed search (Ch. 3)
- Local search (Ch. 4)
- Adversarial search (Ch. 6)
- Constraint Satisfaction Problems (CSPs) (Ch. 5)

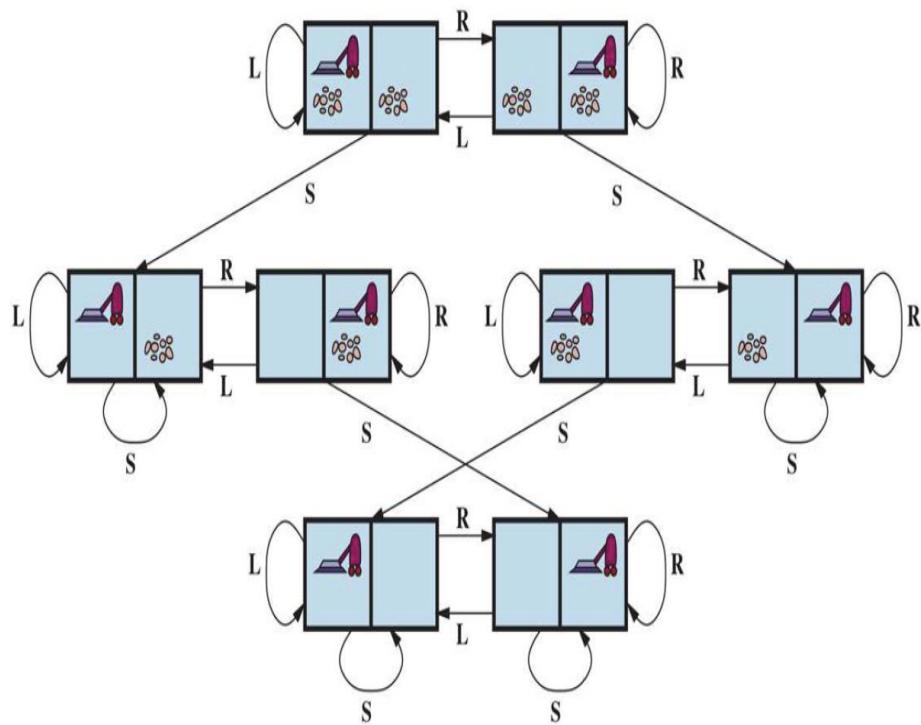
Search Problems

- A **search problem** consists of
 - State space
 - Initial state
 - Goal test (goal states)
 - Actions
 - Transition model
 - Action cost function
- A **solution** is a sequence of actions (path) which transforms the initial state to a goal state
- An **optimal solution** has the lowest path cost among all solutions

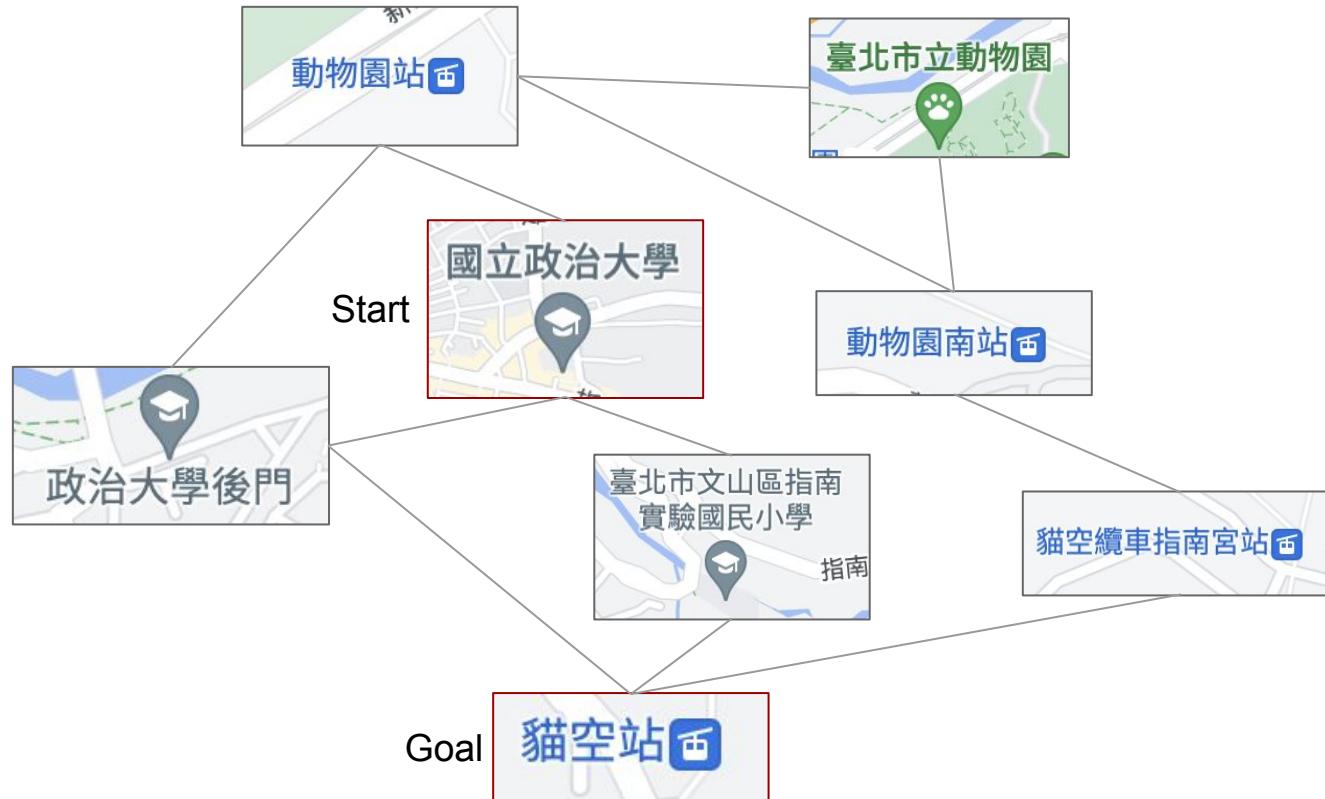
Example: Vacuum World

- States
 - 8 states
- Actions
 - Move left, move right, suck
- Transition model
- Initial state
- Goal test
 - Every cell is clean
- Action cost function
 - Each action costs 1

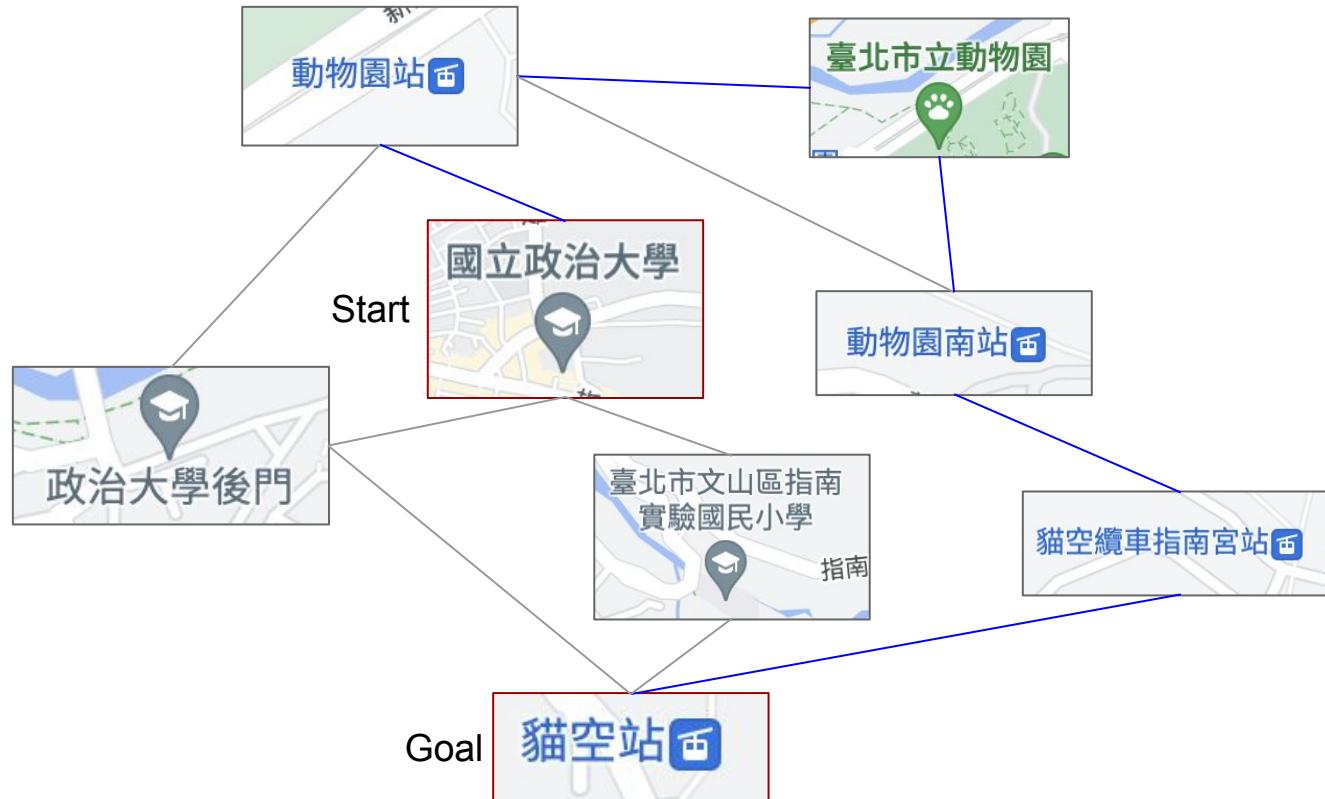
State Space Graph



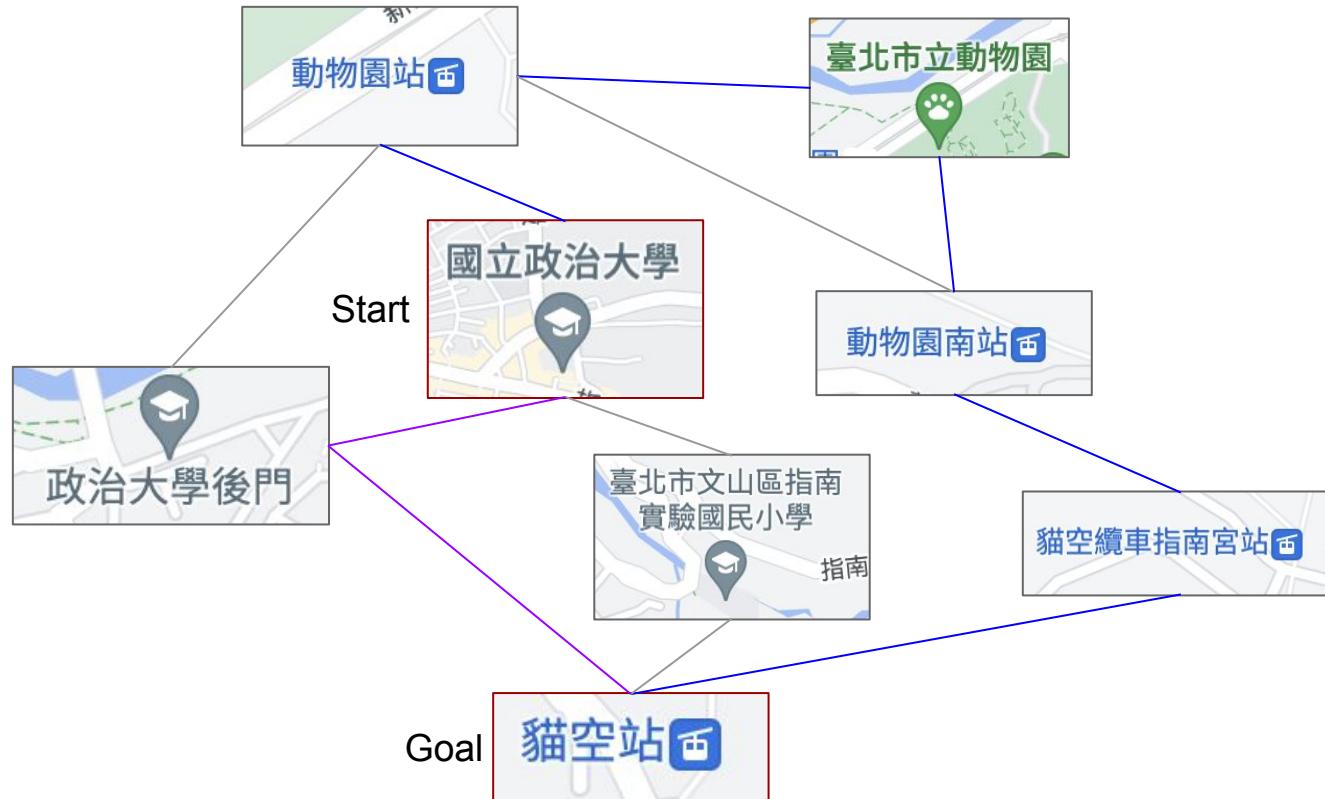
Example: Route-Finding Problem



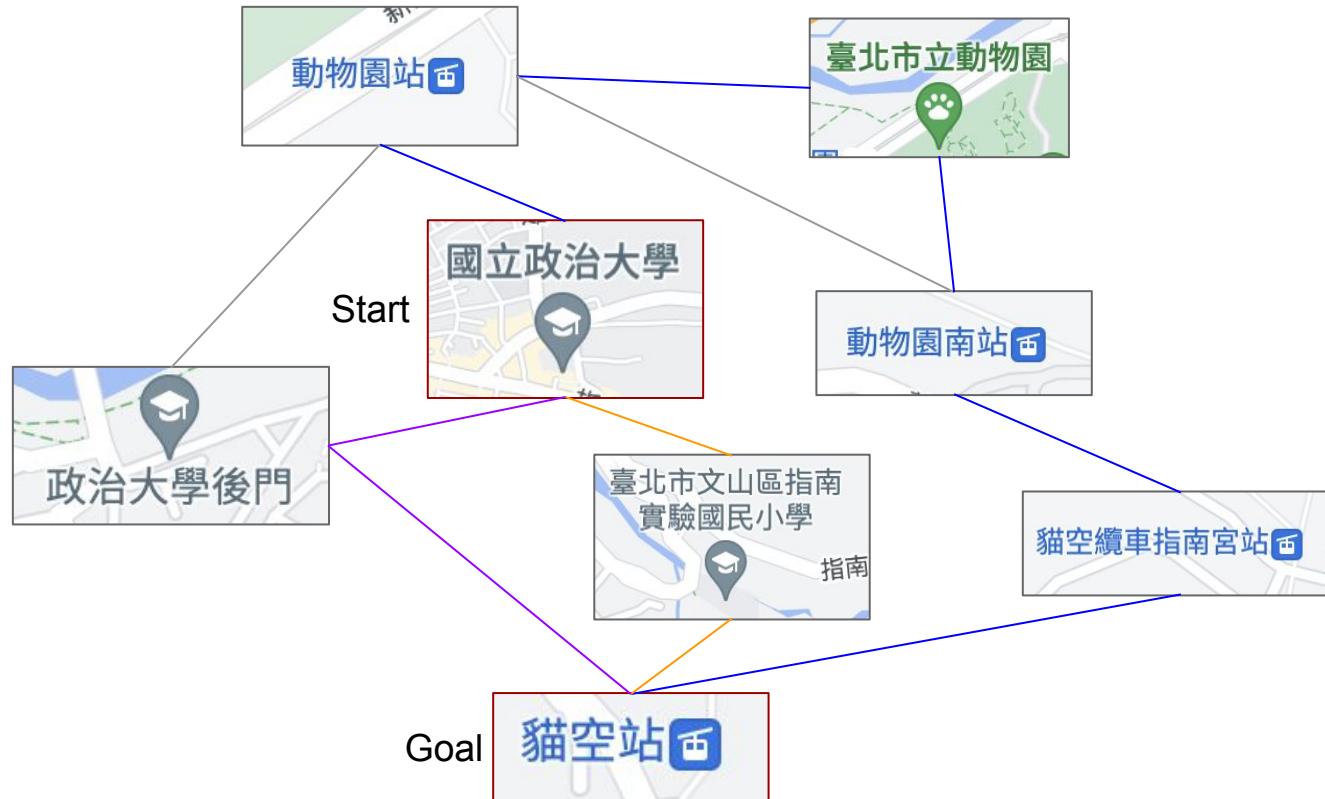
Example: Route-Finding Problem



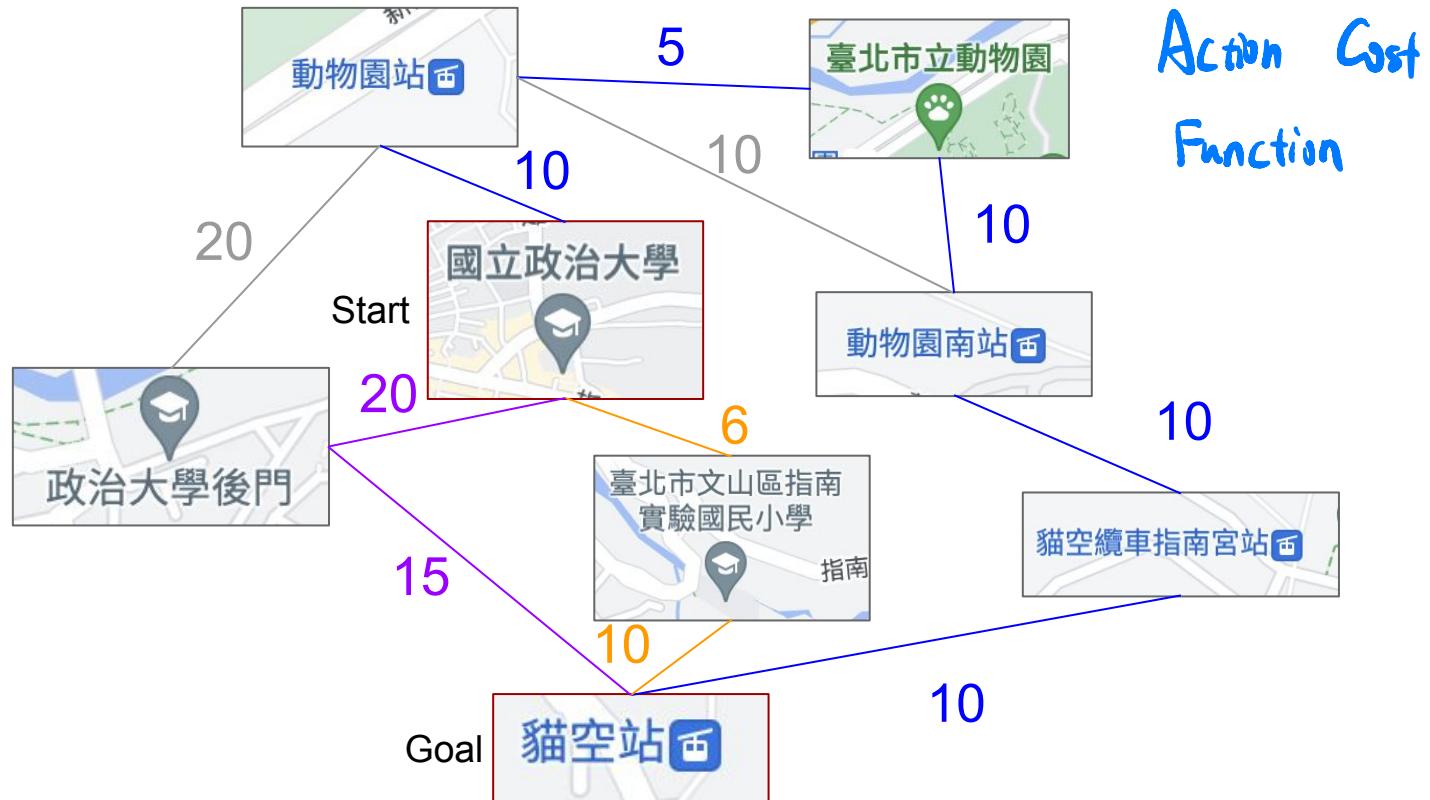
Example: Route-Finding Problem



Example: Route-Finding Problem



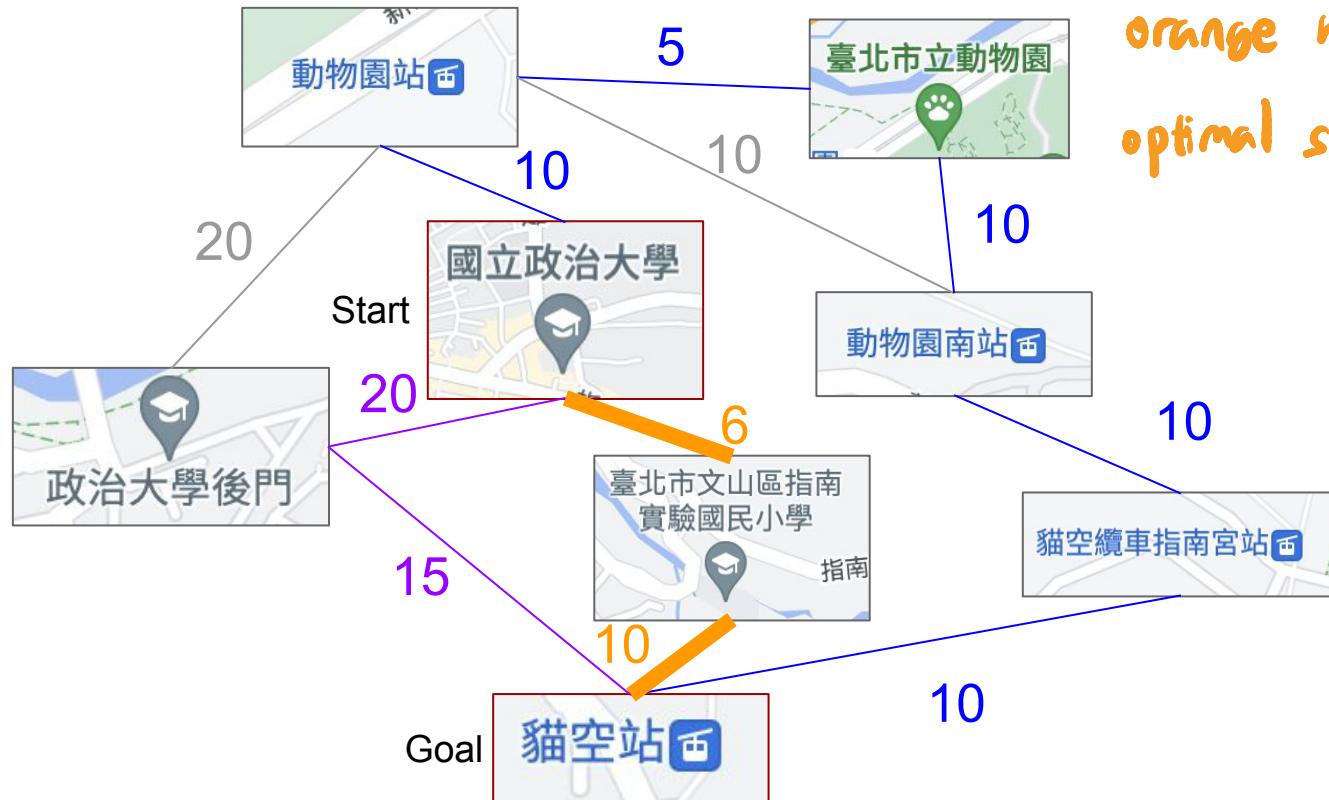
Example: Route-Finding Problem



Example: Route-Finding Problem

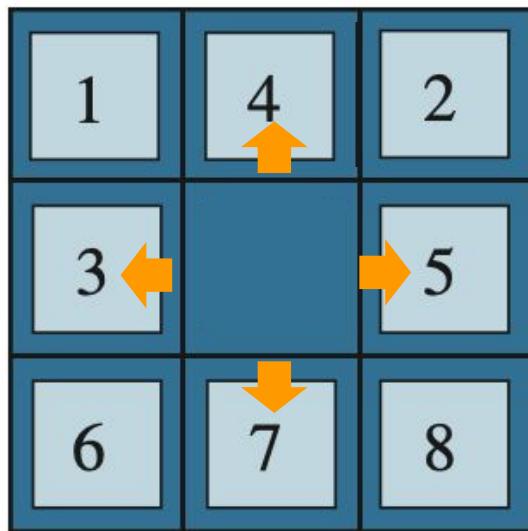
each route : solution

orange route :
optimal solution

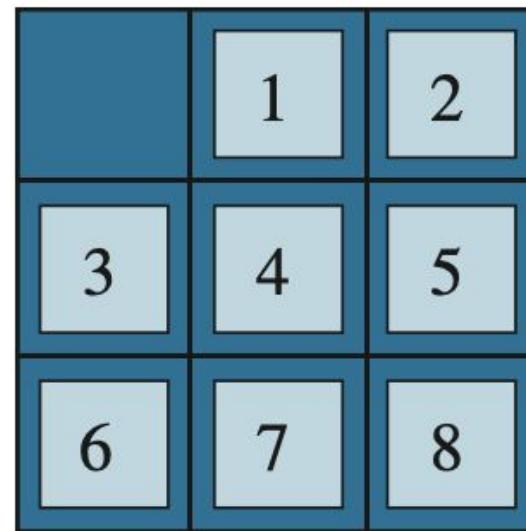


Example: The 8-Puzzle Game

Action

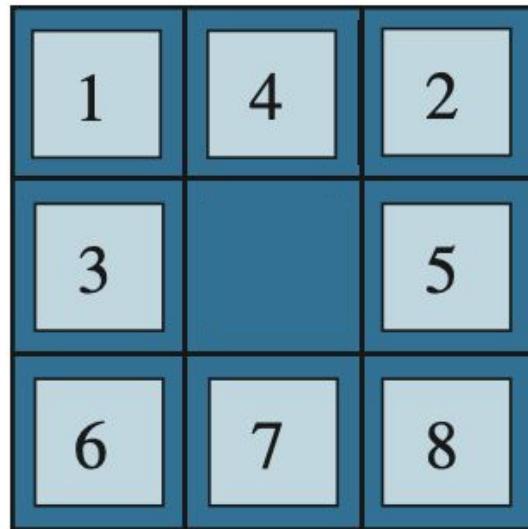


Start State



Goal State

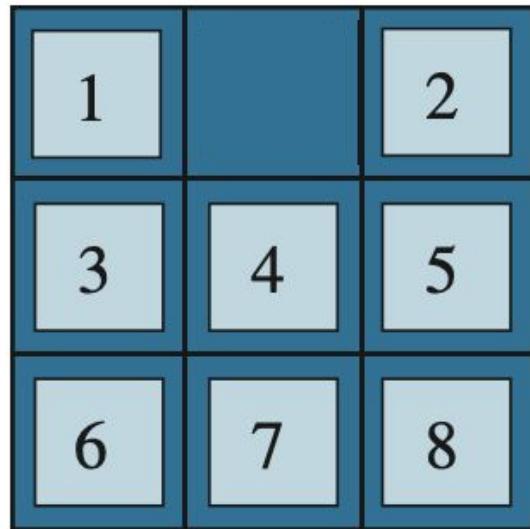
Example: The 8-Puzzle Game



+ Action: ↑

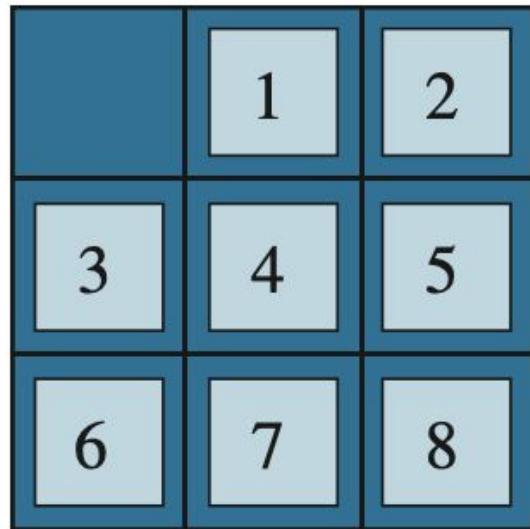
Start State

Example: The 8-Puzzle Game



+ Action: ←

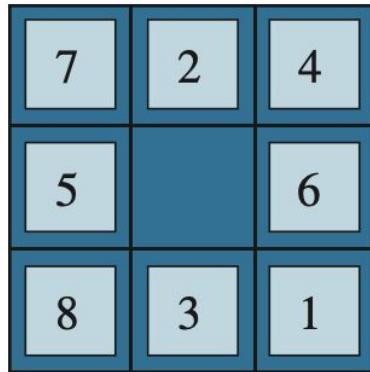
Example: The 8-Puzzle Game



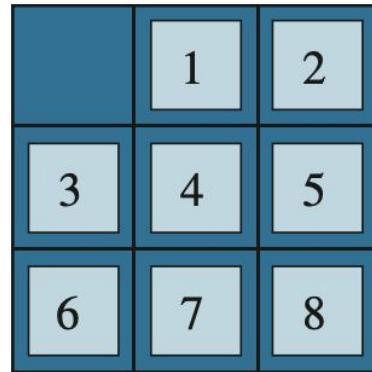
Goal State

Example: The 8-Puzzle Game

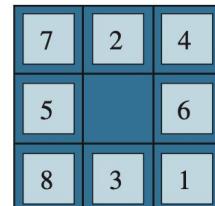
- States
 - A state description specifies the location of each of the tiles
- Initial state
- Goal state
- Actions
 - The blank space moves Left, Right, Up, or Down
(If the blank is at an edge or corner then not all actions will be applicable)
- Transition model
 - Maps a state and action to a resulting state
- Action cost
 - Each action costs 1



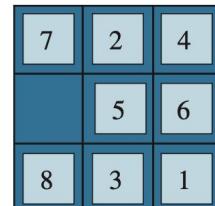
Start State



Goal State



, Left →



Search Problems

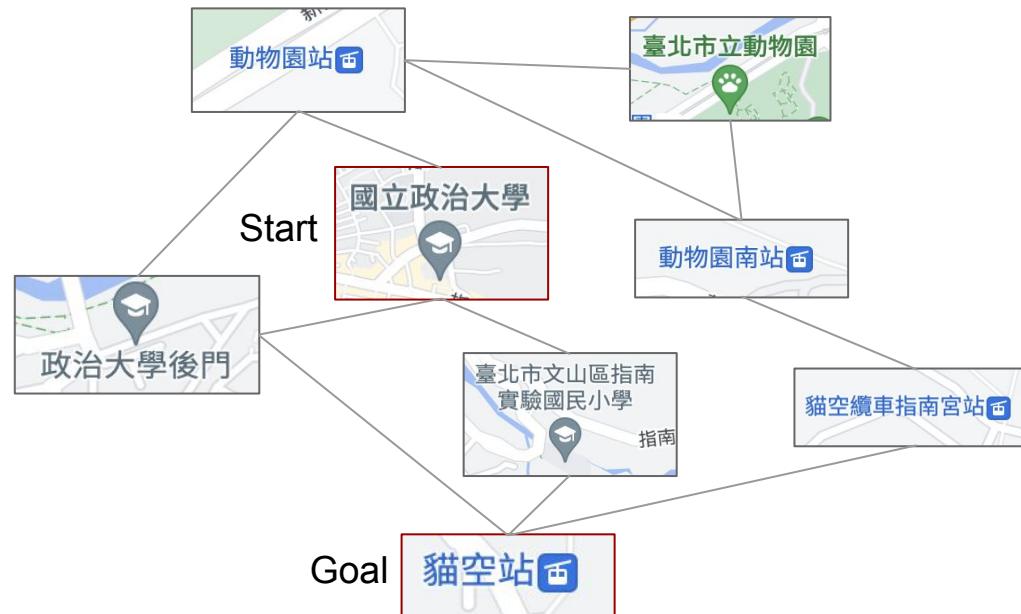
How to solve Search Problems?

Search Algorithm

- A search algorithm takes a search problem as input and returns a solution, or an indication of failure

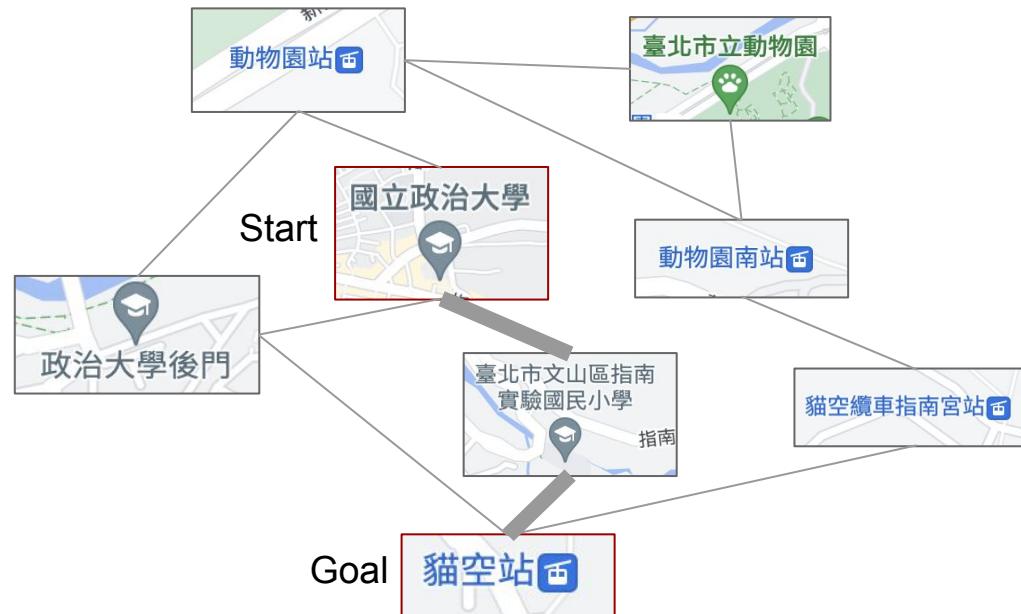
Search Algorithm

- A search algorithm takes a search problem as input and returns a solution, or an indication of failure

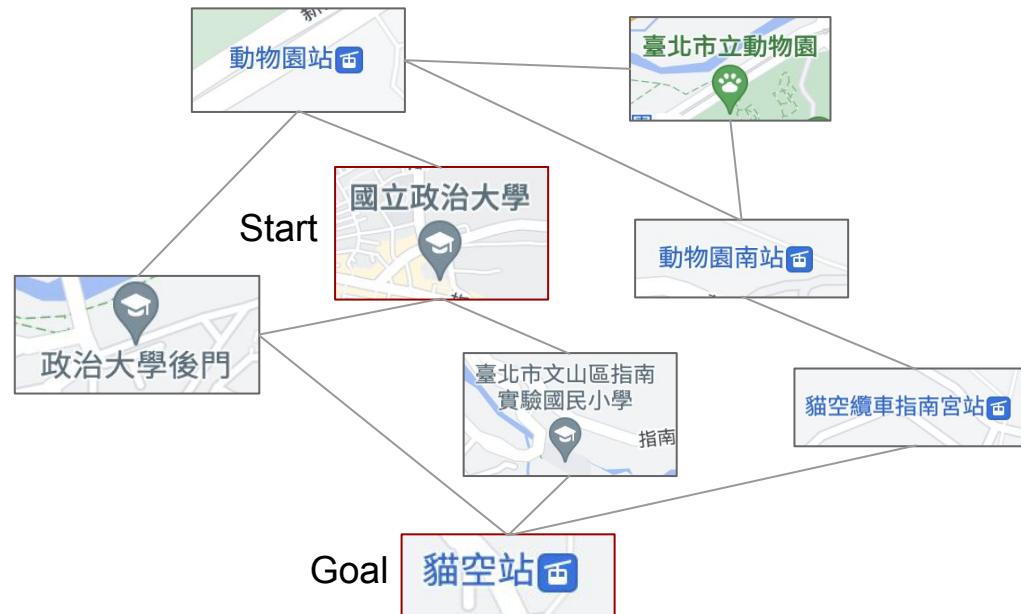


Search Algorithm

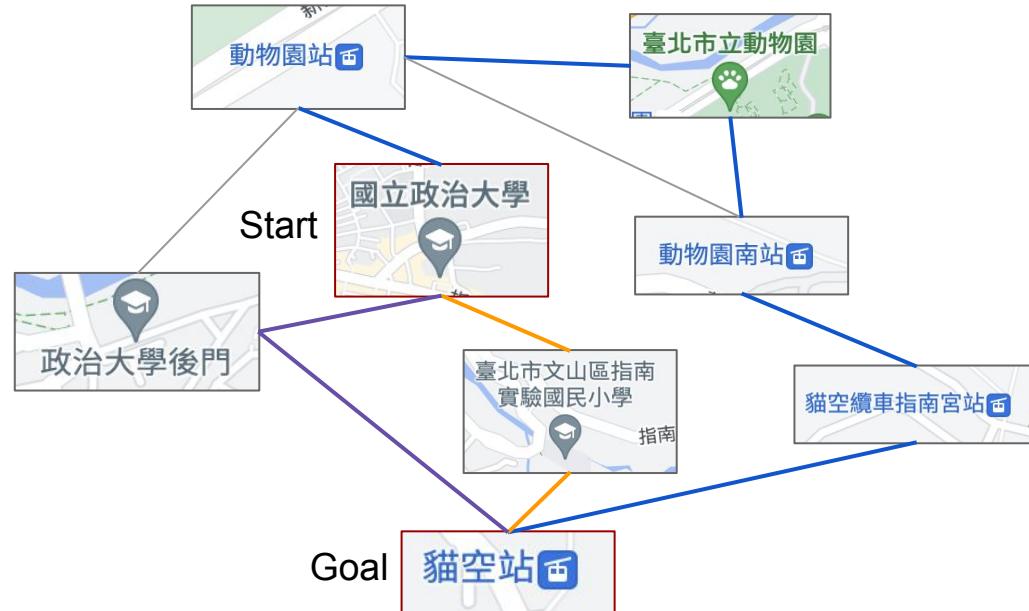
- A search algorithm takes a search problem as input and returns a solution, or an indication of failure



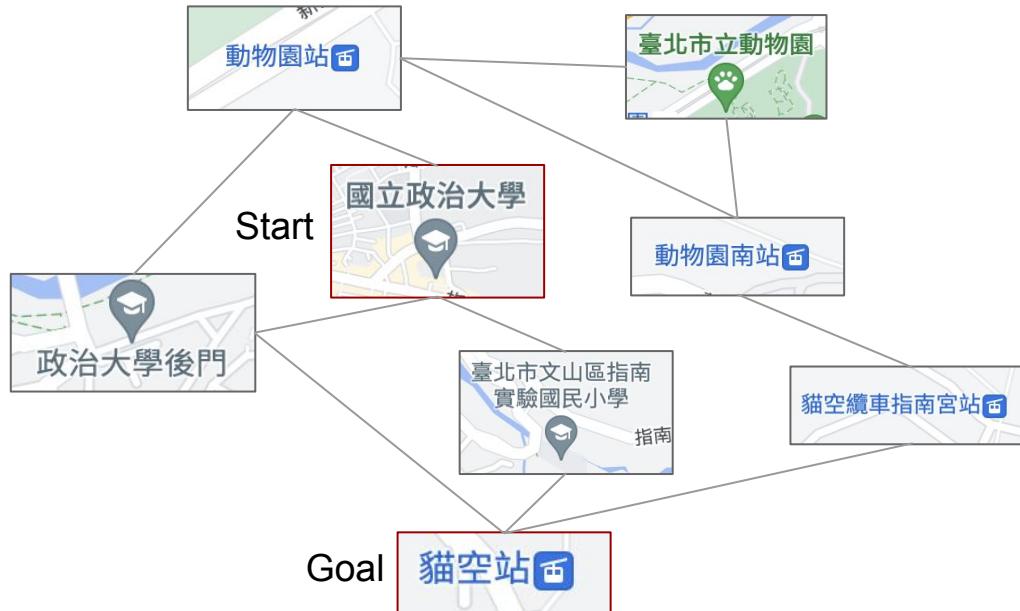
Redundant Path Issue



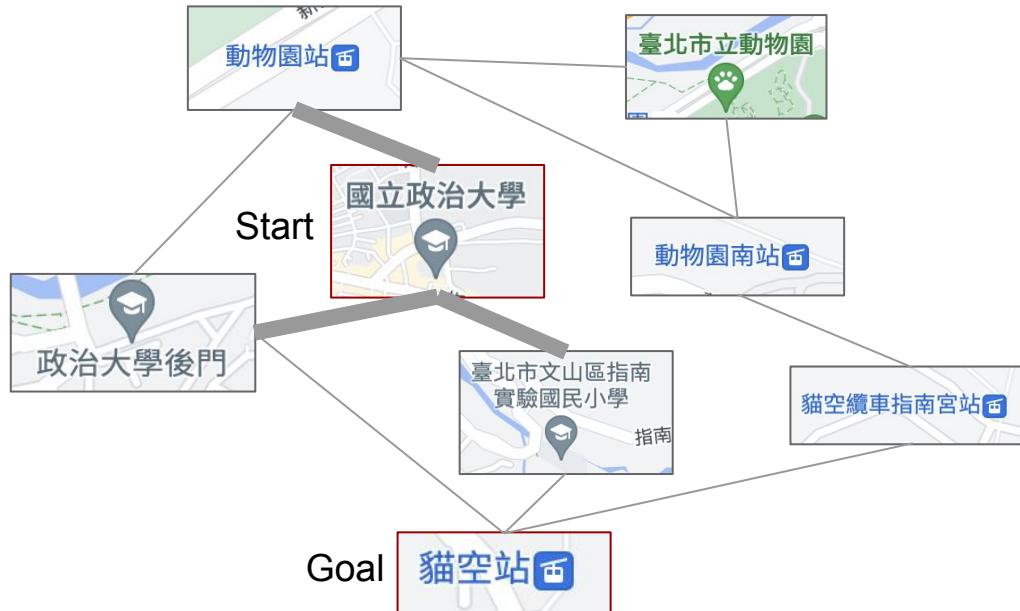
Redundant Path Issue I



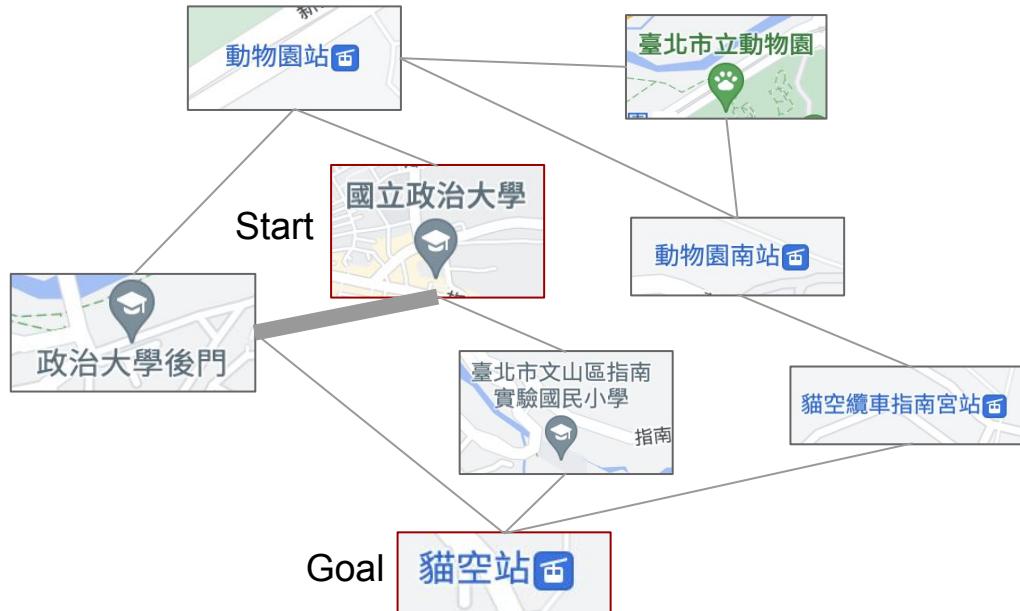
Redundant Path Issue II



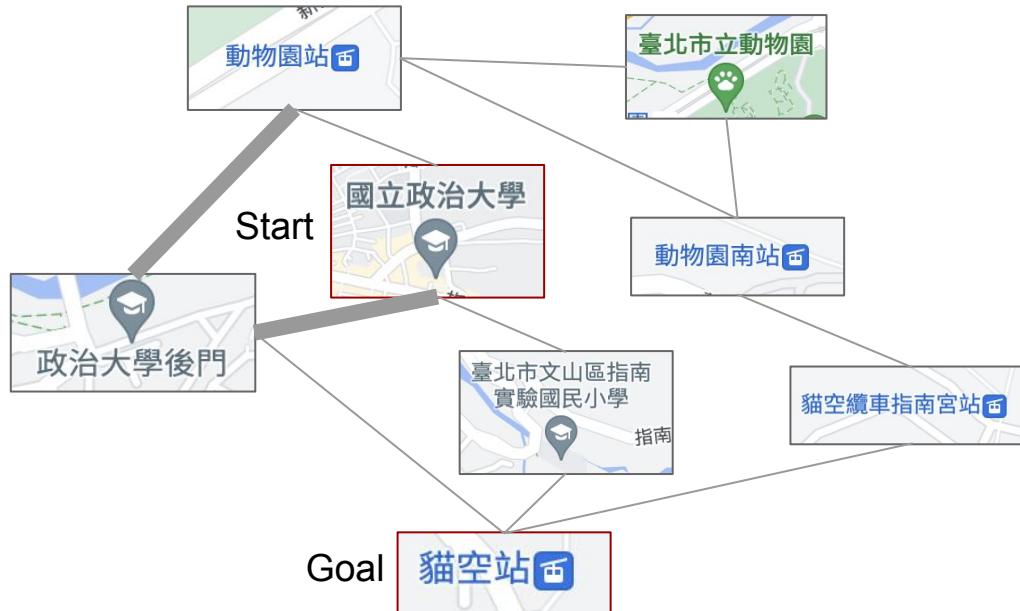
Redundant Path Issue II



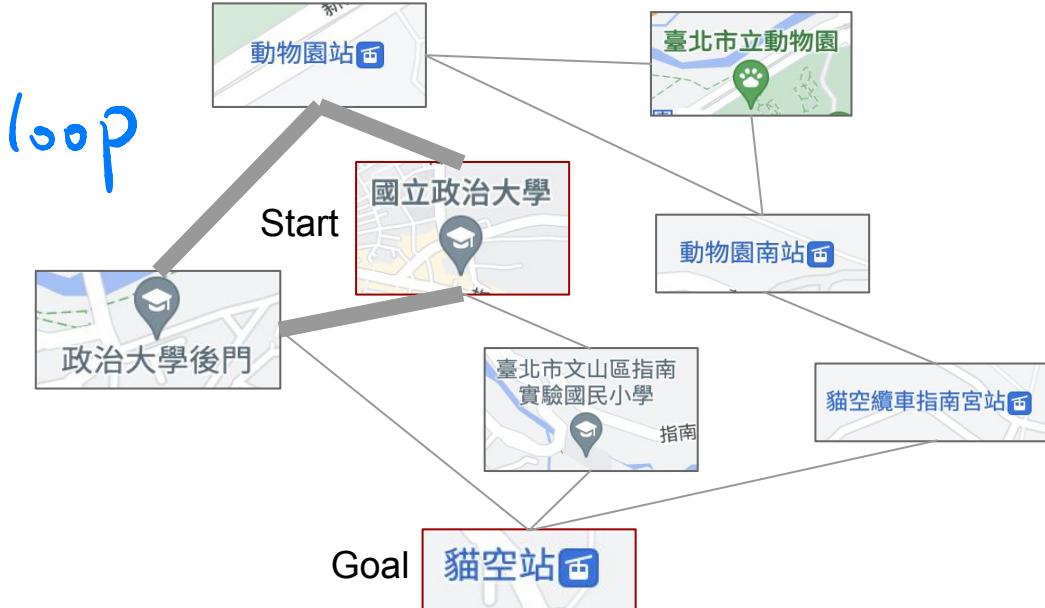
Redundant Path Issue II



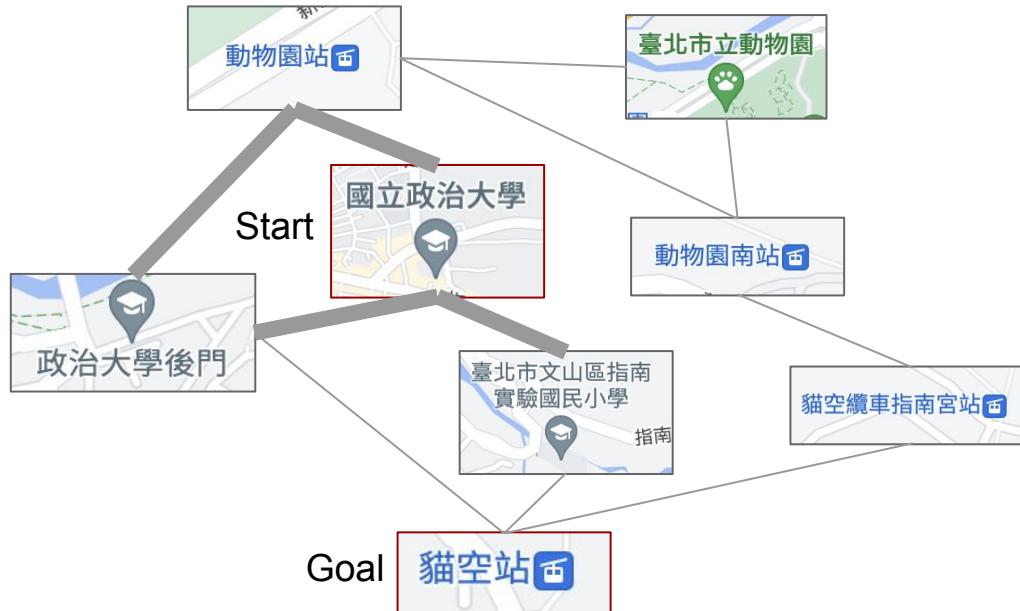
Redundant Path Issue II



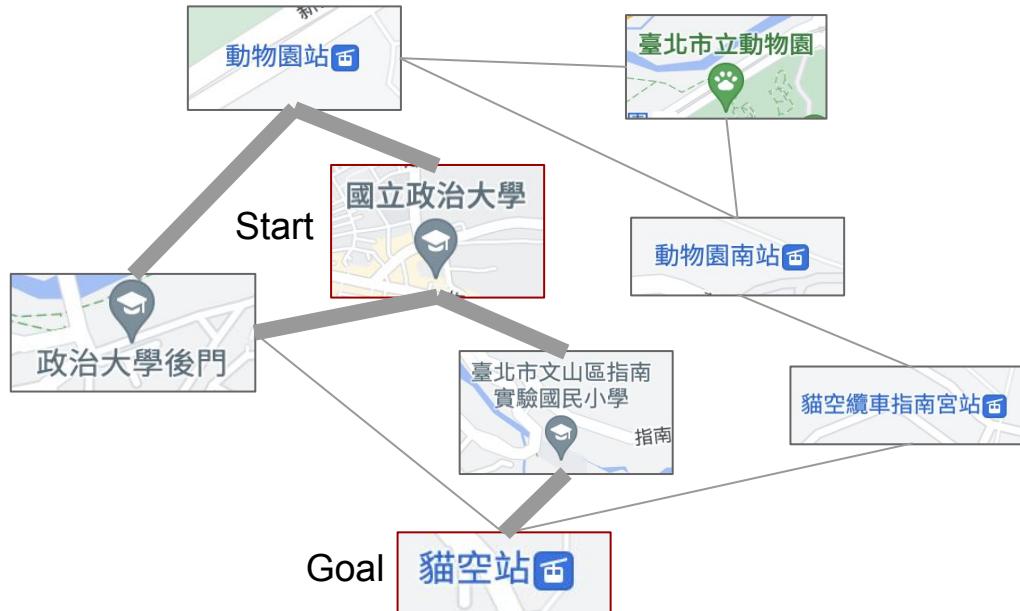
Redundant Path Issue II



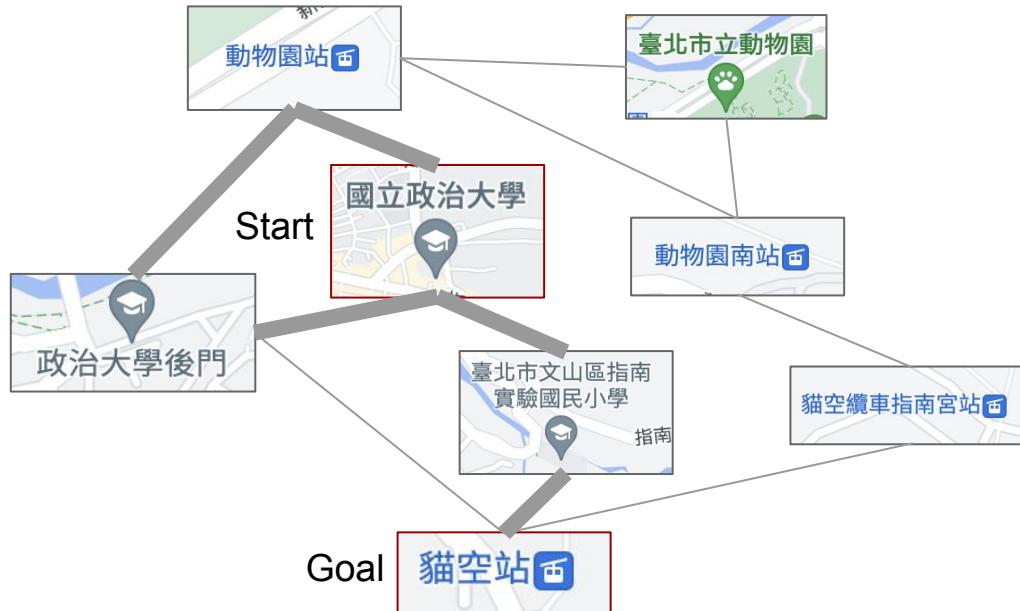
Redundant Path Issue II



Redundant Path Issue II



Redundant Path Issue II: Cycle (Loopy Path)



Three Approaches to the Redundant Path Issue

- Remember all previously **reached** states and keep only the best path to each state
- Check redundant paths
 - Graph search
 - A search algorithm checks for redundant paths
- Check for cycles but not for redundant paths
 - Since each node has a chain of parent pointers, we can check for cycles by following up the chain of parents to see if the state at the end of the path has appeared earlier in the path

Good Search?

Performance of a Search Algorithm

- Completeness
 - Guaranteed to find a solution when there is one
- Cost optimality
 - Guaranteed to find a solution with the lowest path cost
- Time complexity
 - How long does it take to find a solution?
- Space complexity
 - How much memory is needed to perform the search?

Search Topics

- Search problems (Ch. 3)
- **Uninformed search (Ch. 3)**
- Informed search (Ch. 3)
- Local search (Ch. 4)
- Adversarial search (Ch. 6)
- Constraint Satisfaction Problems (CSPs) (Ch. 5)

Uninformed Search

L.-Y. Wei

Spring 2024

Environment

Fully Observable
Deterministic
Static

Uninformed Search (Blind Search)

- An uninformed search algorithm is given **no clue** about how close a state is to the goal
(i.e., only the information available in the problem definition)

Uninformed Search Strategies

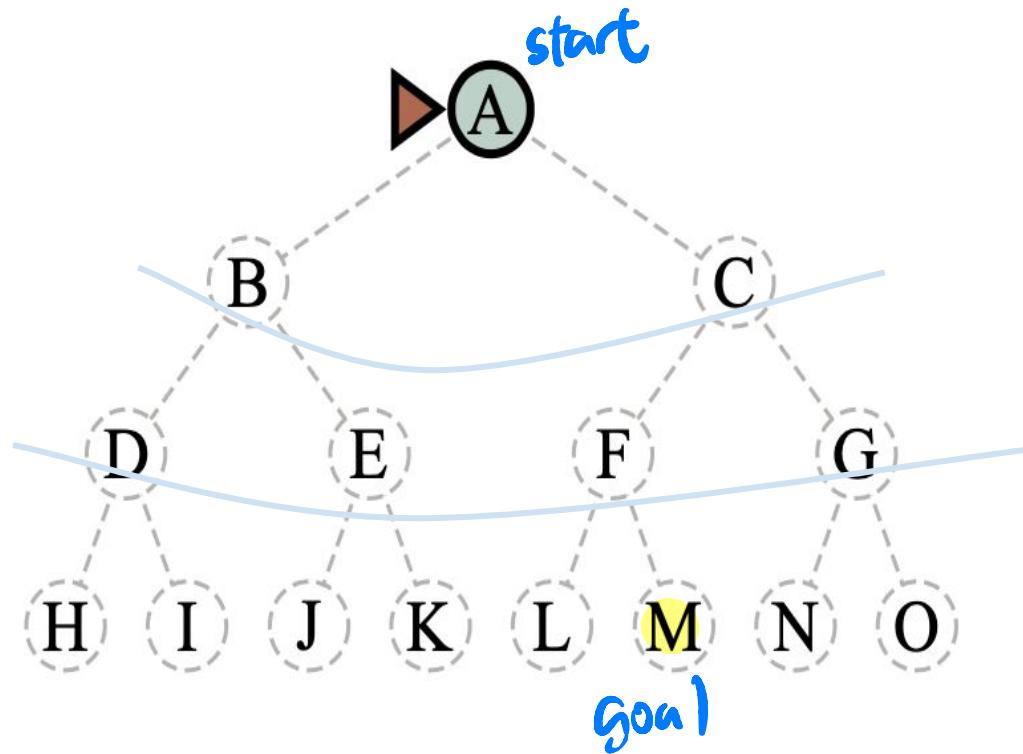
- Breadth-first search
- Depth-first search (DFS)
- Depth-limited search (DLS)
- Iterative deepening search (IDS)
- Uniform-cost search (UCS)
- Bidirectional search

Uninformed Search Strategies

- **Breadth-first search**
- Depth-first search (DFS)
- Depth-limited search (DLS)
- Iterative deepening search (IDS)
- Uniform-cost search (UCS)
- Bidirectional search

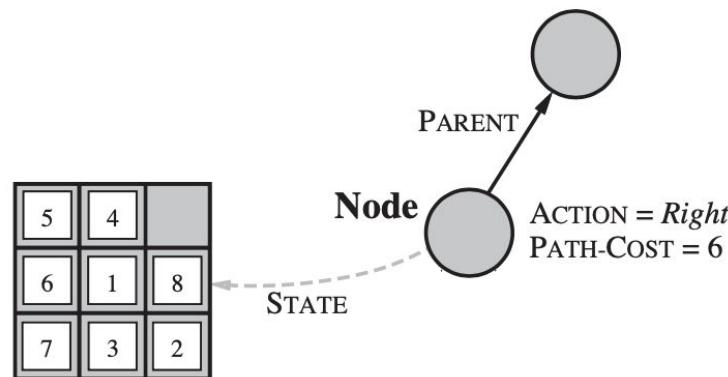
Example: Breadth-First Search on a Simple Binary Tree

- Start: A
- Goal: M



Implementation Issue

- A node in the tree is represented by a data structure with four components
 - `node.STATE`: the state to which the node corresponds
 - `node.PARENT`: the node in the tree that generated this node
 - `node.ACTION`: the action that was applied to the parent's state to generate this node
 - `node.PATH-COST`: the total cost of the path from the initial state to this node



Implementation Issue (cont.)

```
function EXPAND(problem, node) yields nodes
  s  $\leftarrow$  node.STATE
  for each action in problem.ACTIONS(s) do // apply all actions
    s'  $\leftarrow$  problem.RESULT(s, action)
    cost  $\leftarrow$  node.PATH-COST + problem.ACTION-COST(s, action, s') // now cost
    yield NODE(STATE=s', PARENT=node, ACTION=action, PATH-COST=cost)
```

A function used to generate all successor nodes by applying each possible actions to the current state

Breadth-First Search Algorithm

function BREADTH-FIRST-SEARCH(*problem*) **returns** a solution node or *failure*
node \leftarrow NODE(*problem.INITIAL*)

if *problem.IS-GOAL(node.STATE)* **then return** *node*: edge case

frontier \leftarrow a FIFO queue, with *node* as an element : queue

reached $\leftarrow \{\text{problem.INITIAL}\}$: keeps track of the reached states

while not Is-EMPTY(*frontier*) **do**

node \leftarrow POP(*frontier*)

for each *child* in EXPAND(*problem, node*) **do**

s \leftarrow *child.STATE*

if *problem.IS-GOAL(s)* **then return** *child* :

if *s* is not in *reached* **then**

 add *s* to *reached*

 add *child* to *frontier*

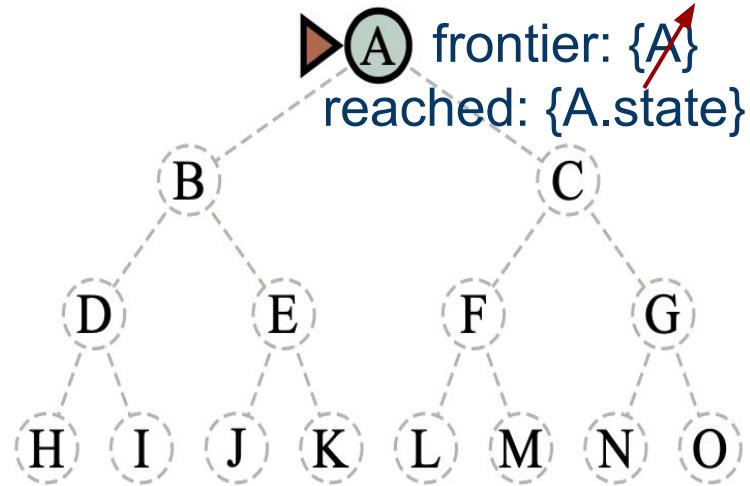
return *failure* : fails to reach the goal

→ See the previous slide
(try each action)
goal reached

Example: Breadth-First Search on a Simple Binary Tree

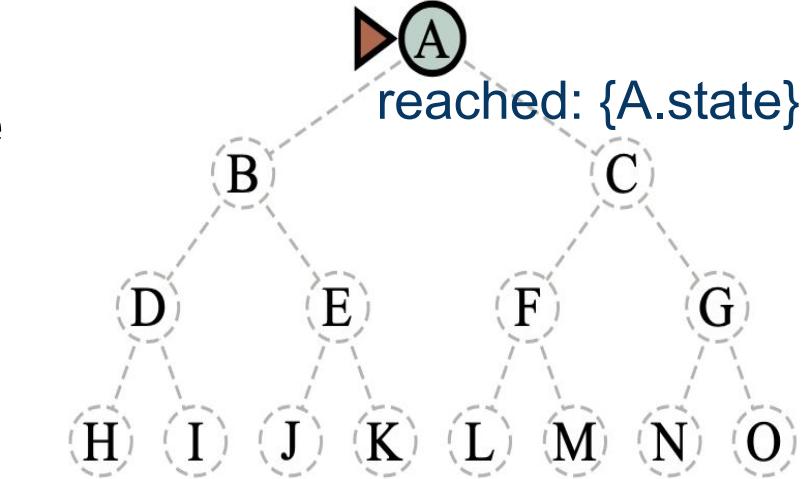
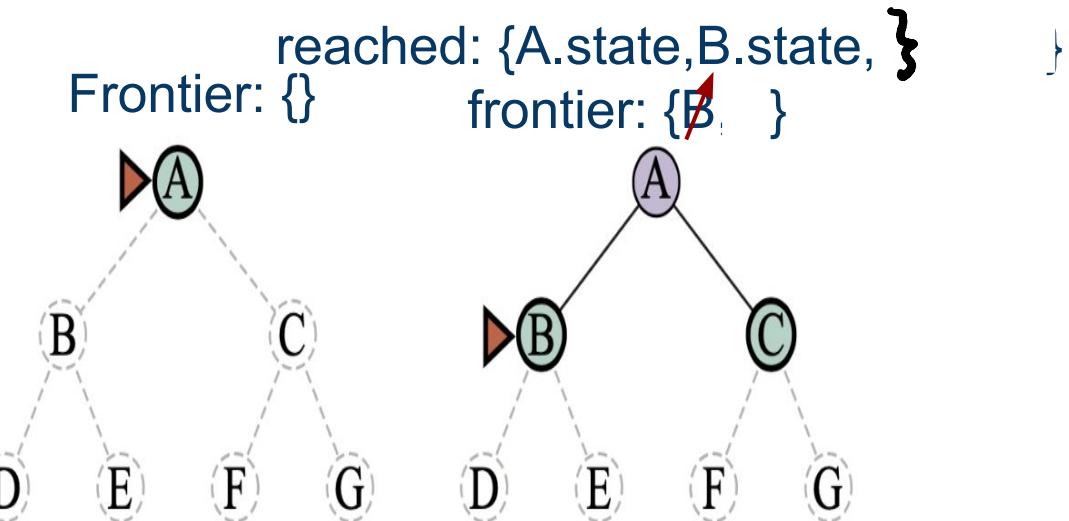
- Start: A
- Goal: M

Frontier: {}



Example: Breadth-First Search on a Simple Binary Tree

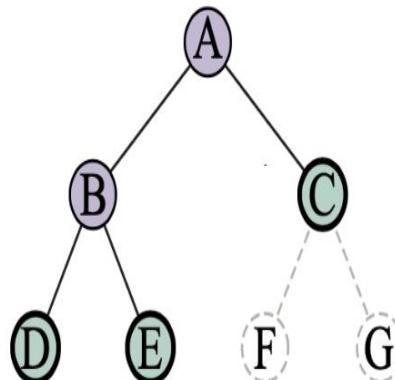
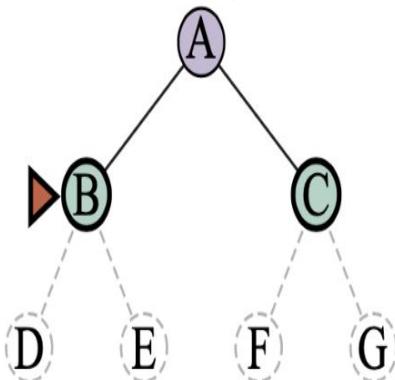
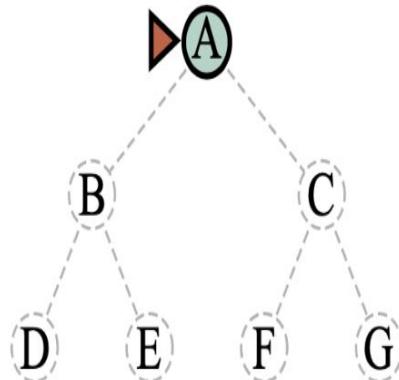
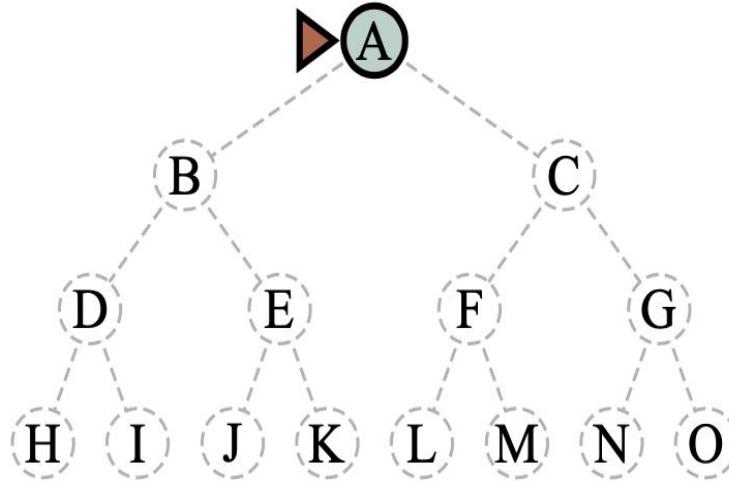
- Start: A
- Goal: M



Example: Breadth-First Search on a Simple Binary Tree

- Start: A
- Goal: M

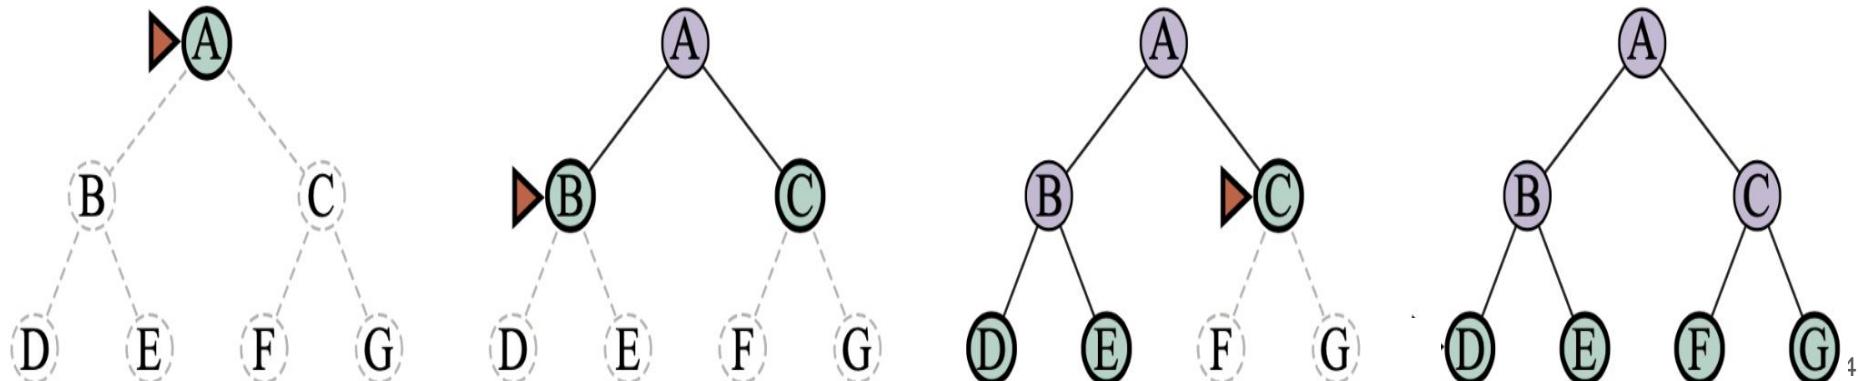
reached: {A.state,B.state,C.state, D.state, E.state}
frontier: {B,C} Frontier: {C,D,E}



Example: Breadth-First Search on a Simple Binary Tree

- Start: A
- Goal: M

reached: {A.state,B.state,C.state, D.state, E.state,F.state,G.state}
Frontier: {C,D,E} Frontier: {D,E,F,G}



Properties of Breadth-First Search

Note. b is the branching factor and d is the depth of the shallowest solution

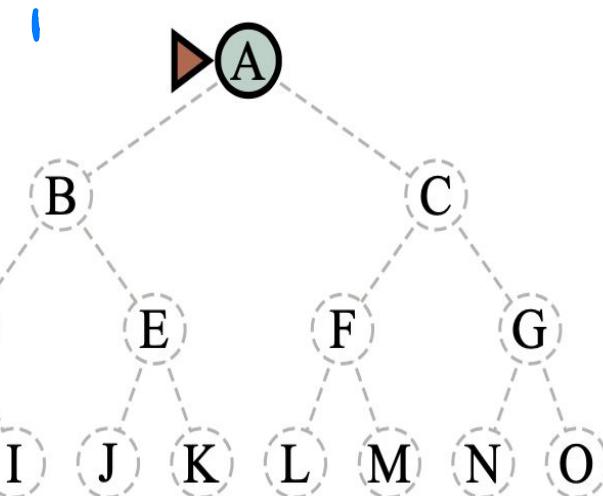
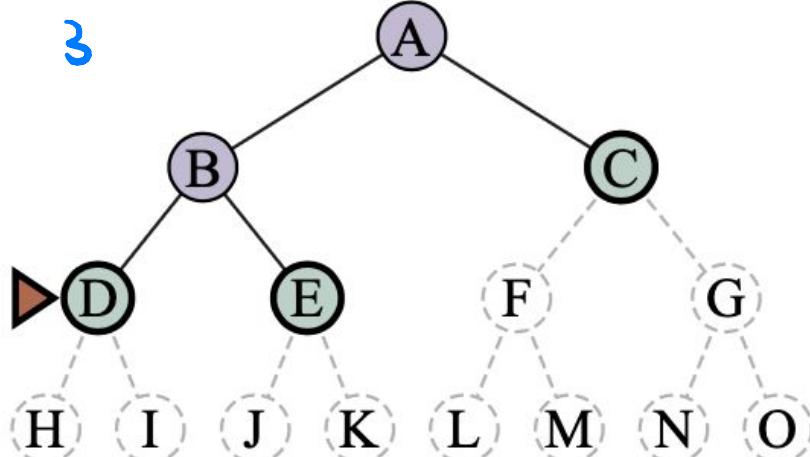
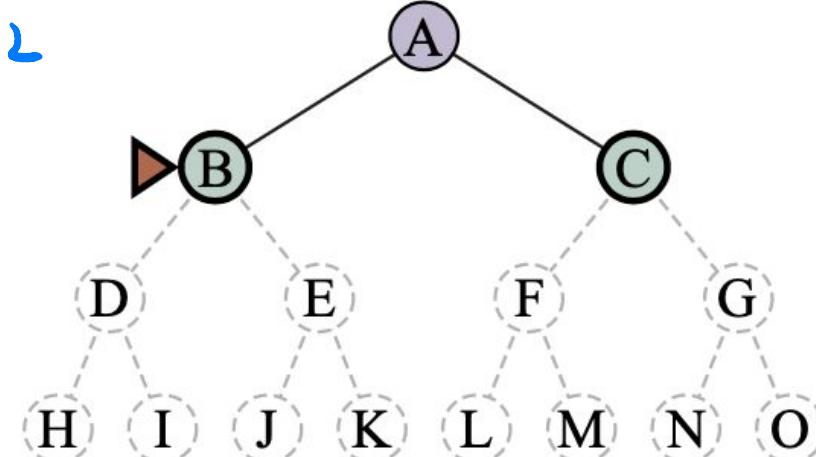
- Complete?
 - Yes (if b is finite and the state space either has a solution or is finite)
∴ BFS performs search level by level
- Optimal cost?
 - Yes (if action costs are all identical)
- Time complexity?
 - $O(b^d)$
In the worst case, BFS travels all nodes at each level until it finds the shallowest solution
 - $b^0 + b^1 + b^2 + \dots + b^d = O(b^d)$
- Space complexity?
 - $O(b^d)$
 $b^0 + b^1 + \dots + \underline{b^d}$ # of nodes in the last level

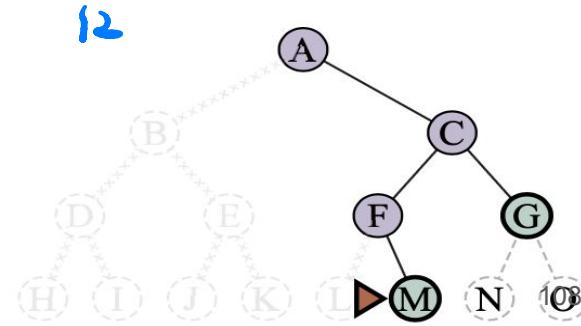
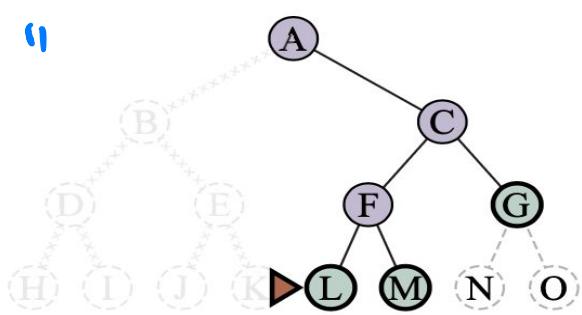
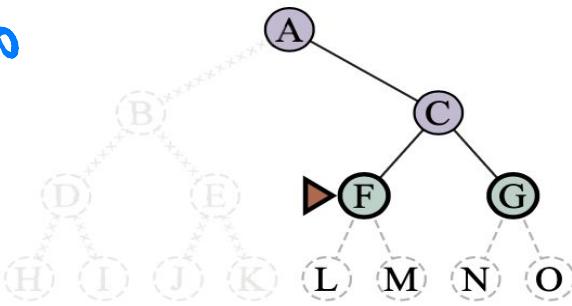
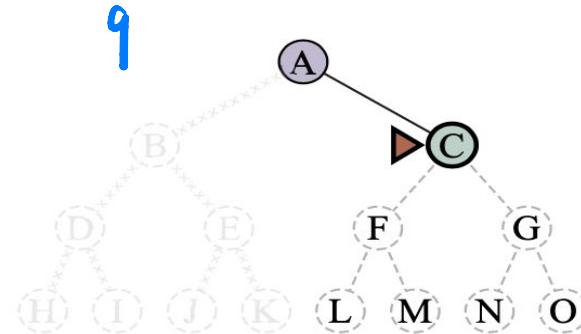
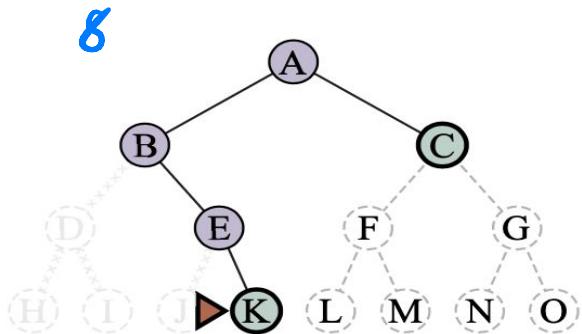
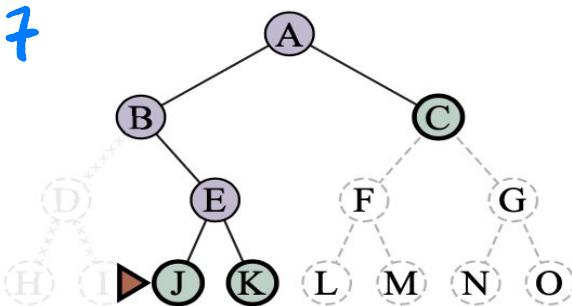
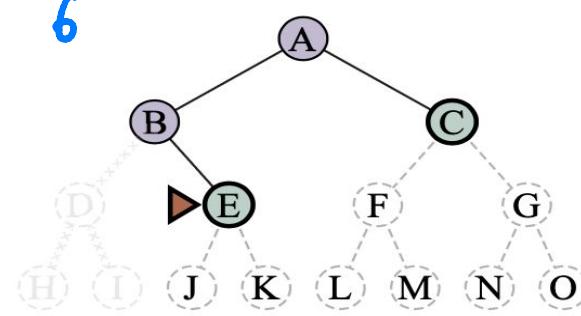
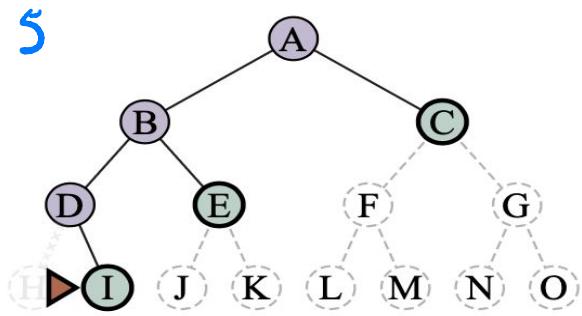
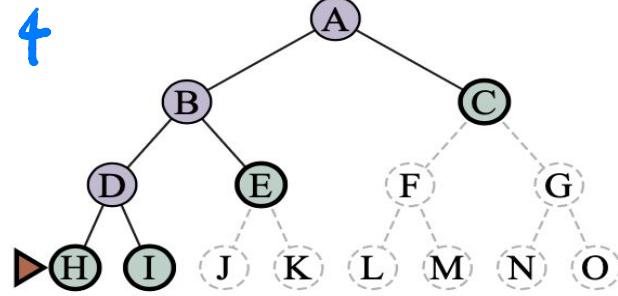
Uninformed Search Strategies

- Breadth-first search
- **Depth-first search (DFS)**
- Depth-limited search (DLS)
- Iterative deepening search (IDS)
- Uniform-cost search (UCS)
- Bidirectional search

Example: DFS on a Simple Binary Tree

- Start: A
- Goal: M





Pseudocode for DFS in Graph

DFS(G, u)

 u.visited = true

 for each $v \in G.\text{Adj}[u]$

 if $v.\text{visited} == \text{false}$

 DFS(G, v)

init() {

 For each $u \in G$

 u.visited = false

 For each $u \in G$

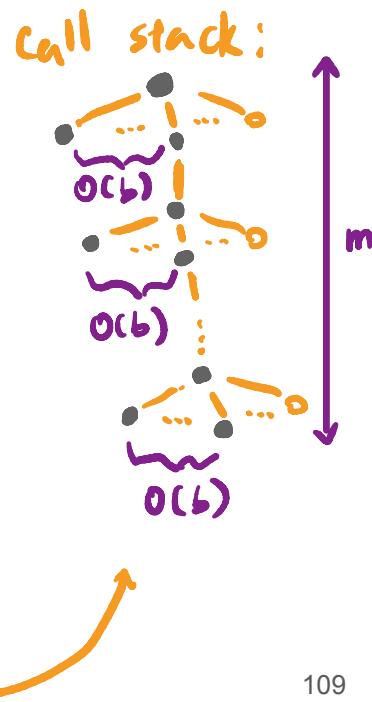
 DFS(G, u)

}

Properties of DFS

Note. where b is the branching factor and m is the maximum depth of the search tree

- Complete?
 - No (in infinite-depth spaces, spaces with loops)
Modify to avoid repeated states along path
⇒ complete in finite spaces
- Optimal cost? *In the worst-case, DFS may traverse all the way down to the maximum*
- Time complexity?
 - $O(b^m)$ → *depth of the tree before backtracking*
- Space complexity?
 - $O(bm)$



Uninformed Search Strategies

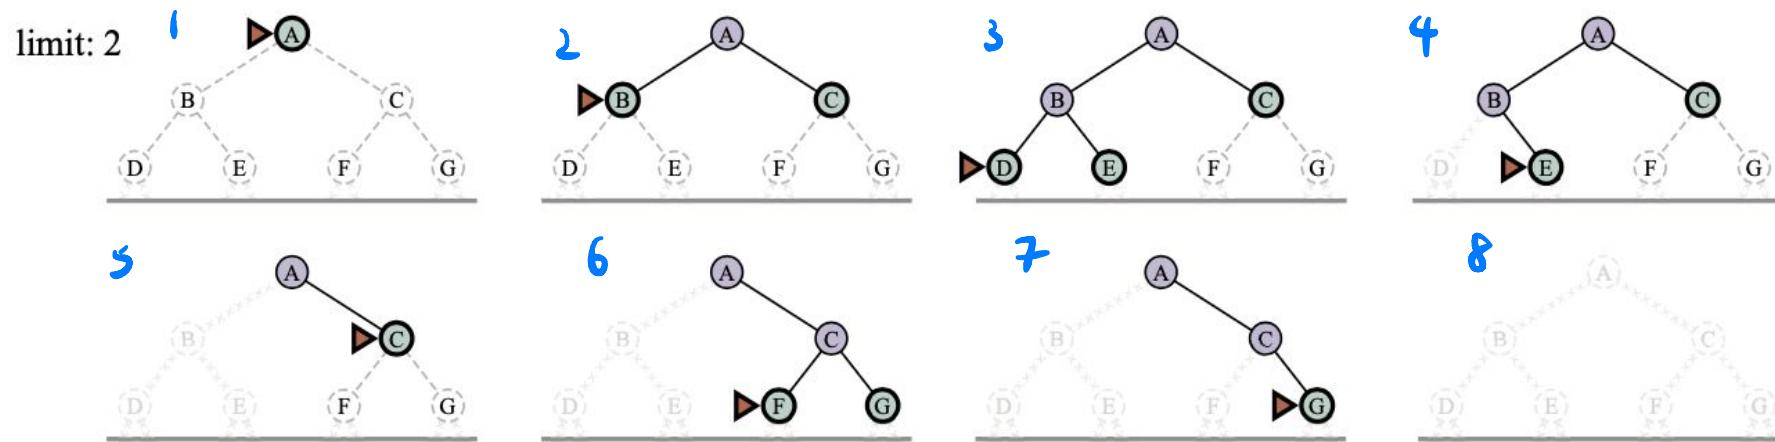
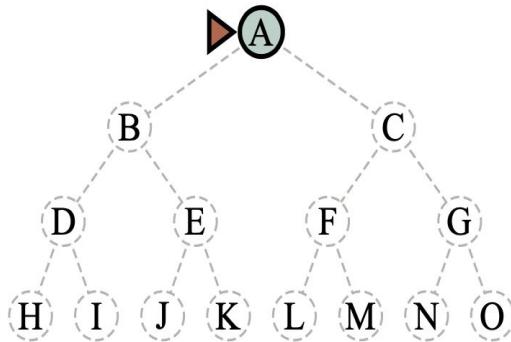
- Breadth-first search
- Depth-first search (DFS)
- **Depth-limited search (DLS)**
- Iterative deepening search (IDS)
- Uniform-cost search (UCS)
- Bidirectional search

Depth-Limited Search (DLS)

- A version of **depth-first search** in which we supply a **depth limit ℓ** and treat all nodes at the depth ℓ as if they had no successors

Example: DLS on a Simple Binary Tree

- Start: A
- Goal: M
- $\ell: 2$



Depth-Limited Search (DLS) Algorithm

```
function DEPTH-LIMITED-SEARCH(problem,  $\ell$ ) returns a node or failure or cutoff
  frontier  $\leftarrow$  a LIFO queue (stack) with NODE(problem.INITIAL) as an element
  result  $\leftarrow$  failure
  while not Is-EMPTY(frontier) do
    node  $\leftarrow$  POP(frontier)
    if problem.IS-GOAL(node.STATE) then return node
    if DEPTH(node)  $>$   $\ell$  then result  $\leftarrow$  cutoff ) limit search depth
    else if not IS-CYCLE(node) do
      for each child in EXPAND(problem, node) do
        add child to frontier
  return result
```

: solution found

Properties of DLS

Note. b is the branching factor and ℓ is the depth limit

- Complete?
 - No
- Optimal cost?
- Time complexity?
 - $O(b^\ell)$
- Space complexity?
 - $O(b\ell)$

\Rightarrow :: Depth Limit

$$O(1 + b + b^2 + \dots + b^\ell)$$

Refer to DFS

Uninformed Search Strategies

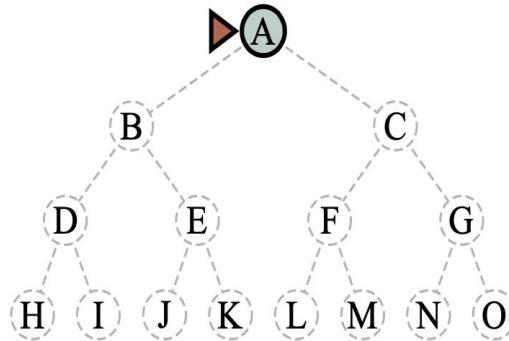
- Breadth-first search
- Depth-first search (DFS)
- Depth-limited search (DLS)
- **Iterative deepening search (IDS)**
- Uniform-cost search (UCS)
- Bidirectional search

Iterative-Deepening Search (IDS)

- Iterative deepening search solves the problem of picking a good value for ℓ by trying all values: first 0, then 1, then 2, and so on — until either a **solution is found**, or the depth-limited search returns the **failure** value

Example: IDS on a Simple Binary Tree

- Start: A
- Goal: M

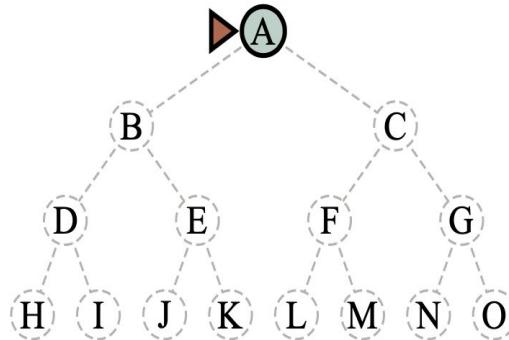


limit: 0

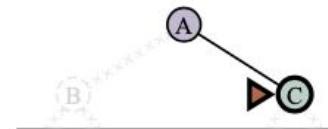
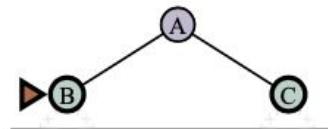
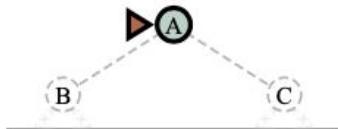


Example: IDS on a Simple Binary Tree

- Start: A
- Goal: M

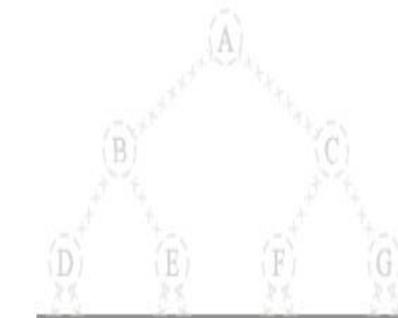
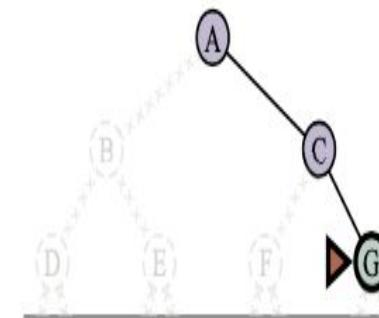
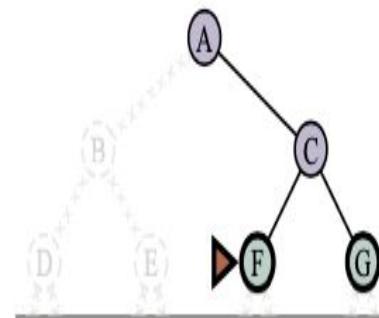
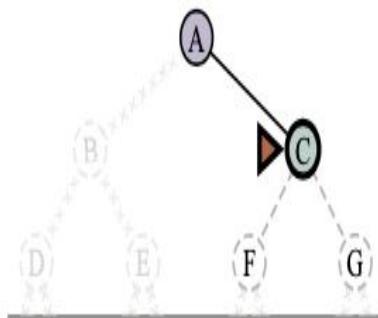
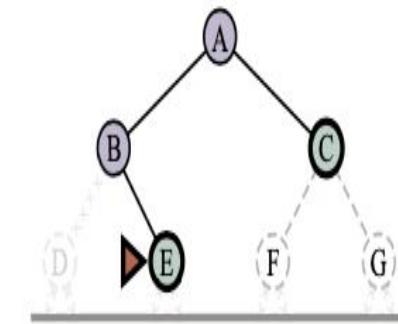
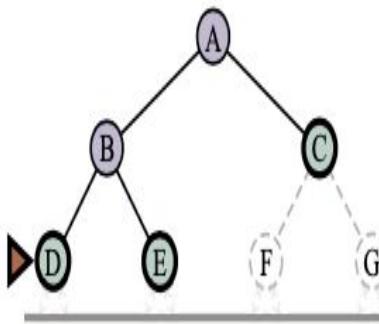
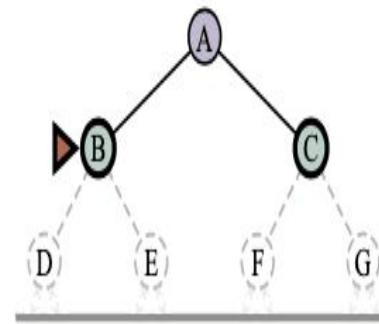
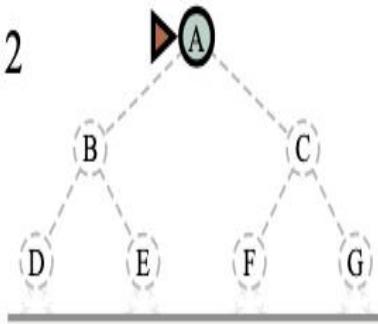


limit: 1

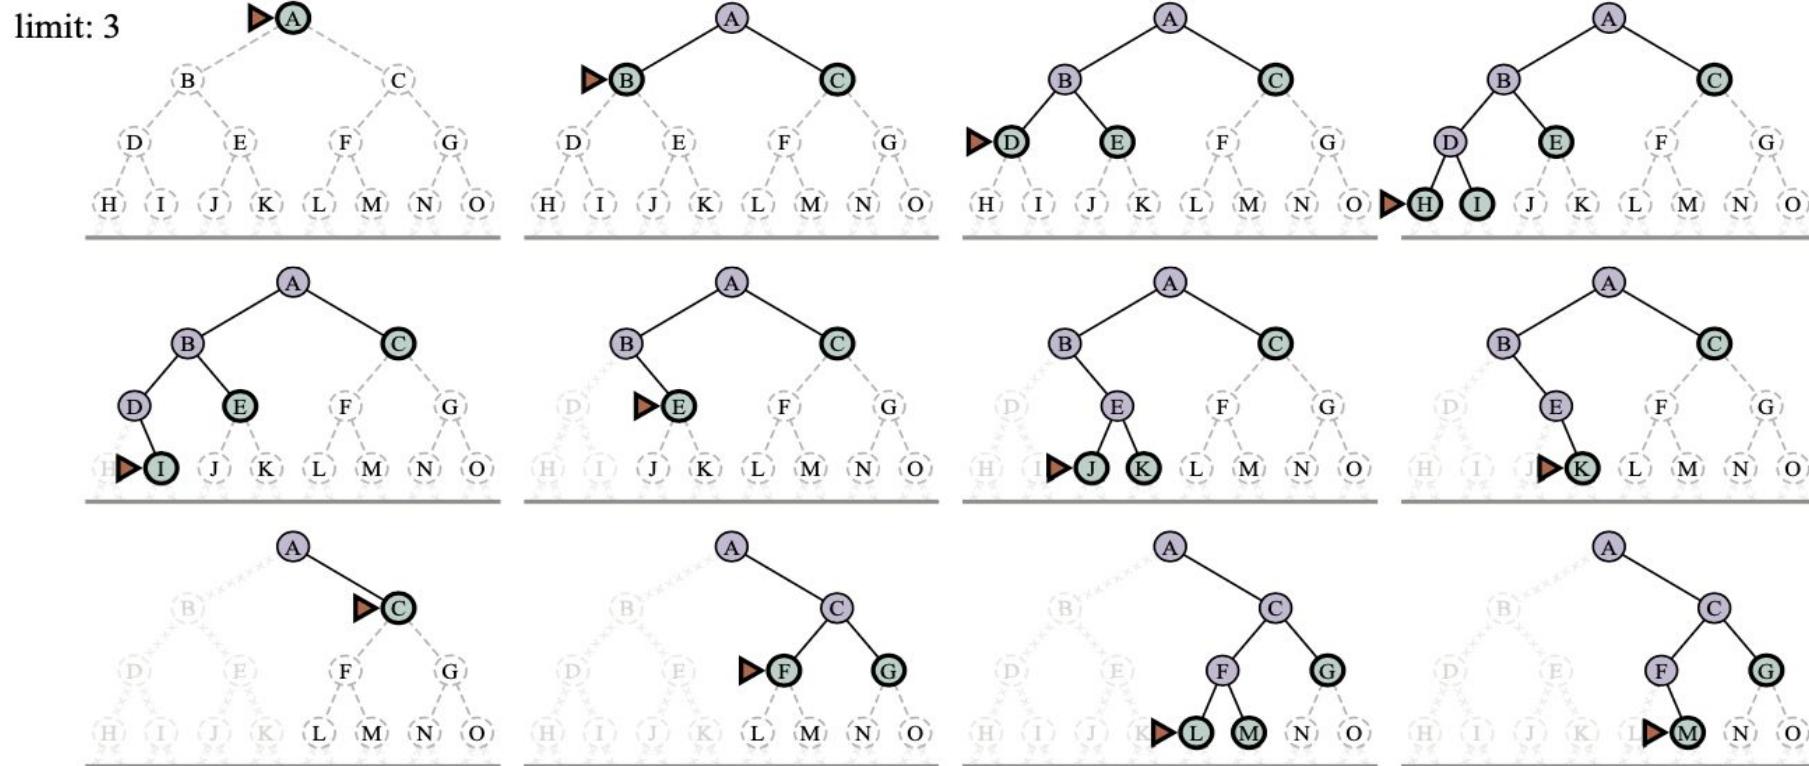


Example: IDS on a Simple Binary Tree

limit: 2



Example: IDS on a Simple Binary Tree



Iterative-Deepening Search (IDS) Algorithm

```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution node or failure
  for depth = 0 to  $\infty$  do
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)
    if result  $\neq$  cutoff then return result
```

```
function DEPTH-LIMITED-SEARCH(problem,  $\ell$ ) returns a node or failure or cutoff
  frontier  $\leftarrow$  a LIFO queue (stack) with NODE(problem.INITIAL) as an element
  result  $\leftarrow$  failure
  while not IS-EMPTY(frontier) do
    node  $\leftarrow$  POP(frontier)
    if problem.IS-GOAL(node.STATE) then return node
    if DEPTH(node)  $>$   $\ell$  then
      result  $\leftarrow$  cutoff
    else if not IS-CYCLE(node) do
      for each child in EXPAND(problem, node) do
        add child to frontier
  return result
```

Properties of IDS

Note. b is the branching factor and d is the depth of the shallowest solution

- Complete?
 - Yes if the branching factor b is finite
- Optimal cost?
 - Yes if action costs are all identical
- Time complexity?
 - $O(b^d)$
 - $(d + 1)b^0 + (d)b^1 + (d - 1)b^2 + (d - 2)b^3 + \dots + b^d = O(b^d)$
- Space complexity? *(depth = 0) is searched up to (d+1) times*
 - $O(bd)$

Refer to DFS

Pros of IDS

- (Like DFS) Its memory↓
 $\text{DFS: } O(bm)$ $(d \leq m)$
 $\text{IDS: } O(bd)$
- (Like BFS) It is **optimal** for problems where all actions have the same cost and is **complete** on finite acyclic state spaces (or on any finite state space when we check nodes for cycles all the way up the path)

- IDS is the preferred uninformed search method when
 - 1) the search state space is ^{larger} than can fit in memory and
 - 2) the depth of the solution is not known