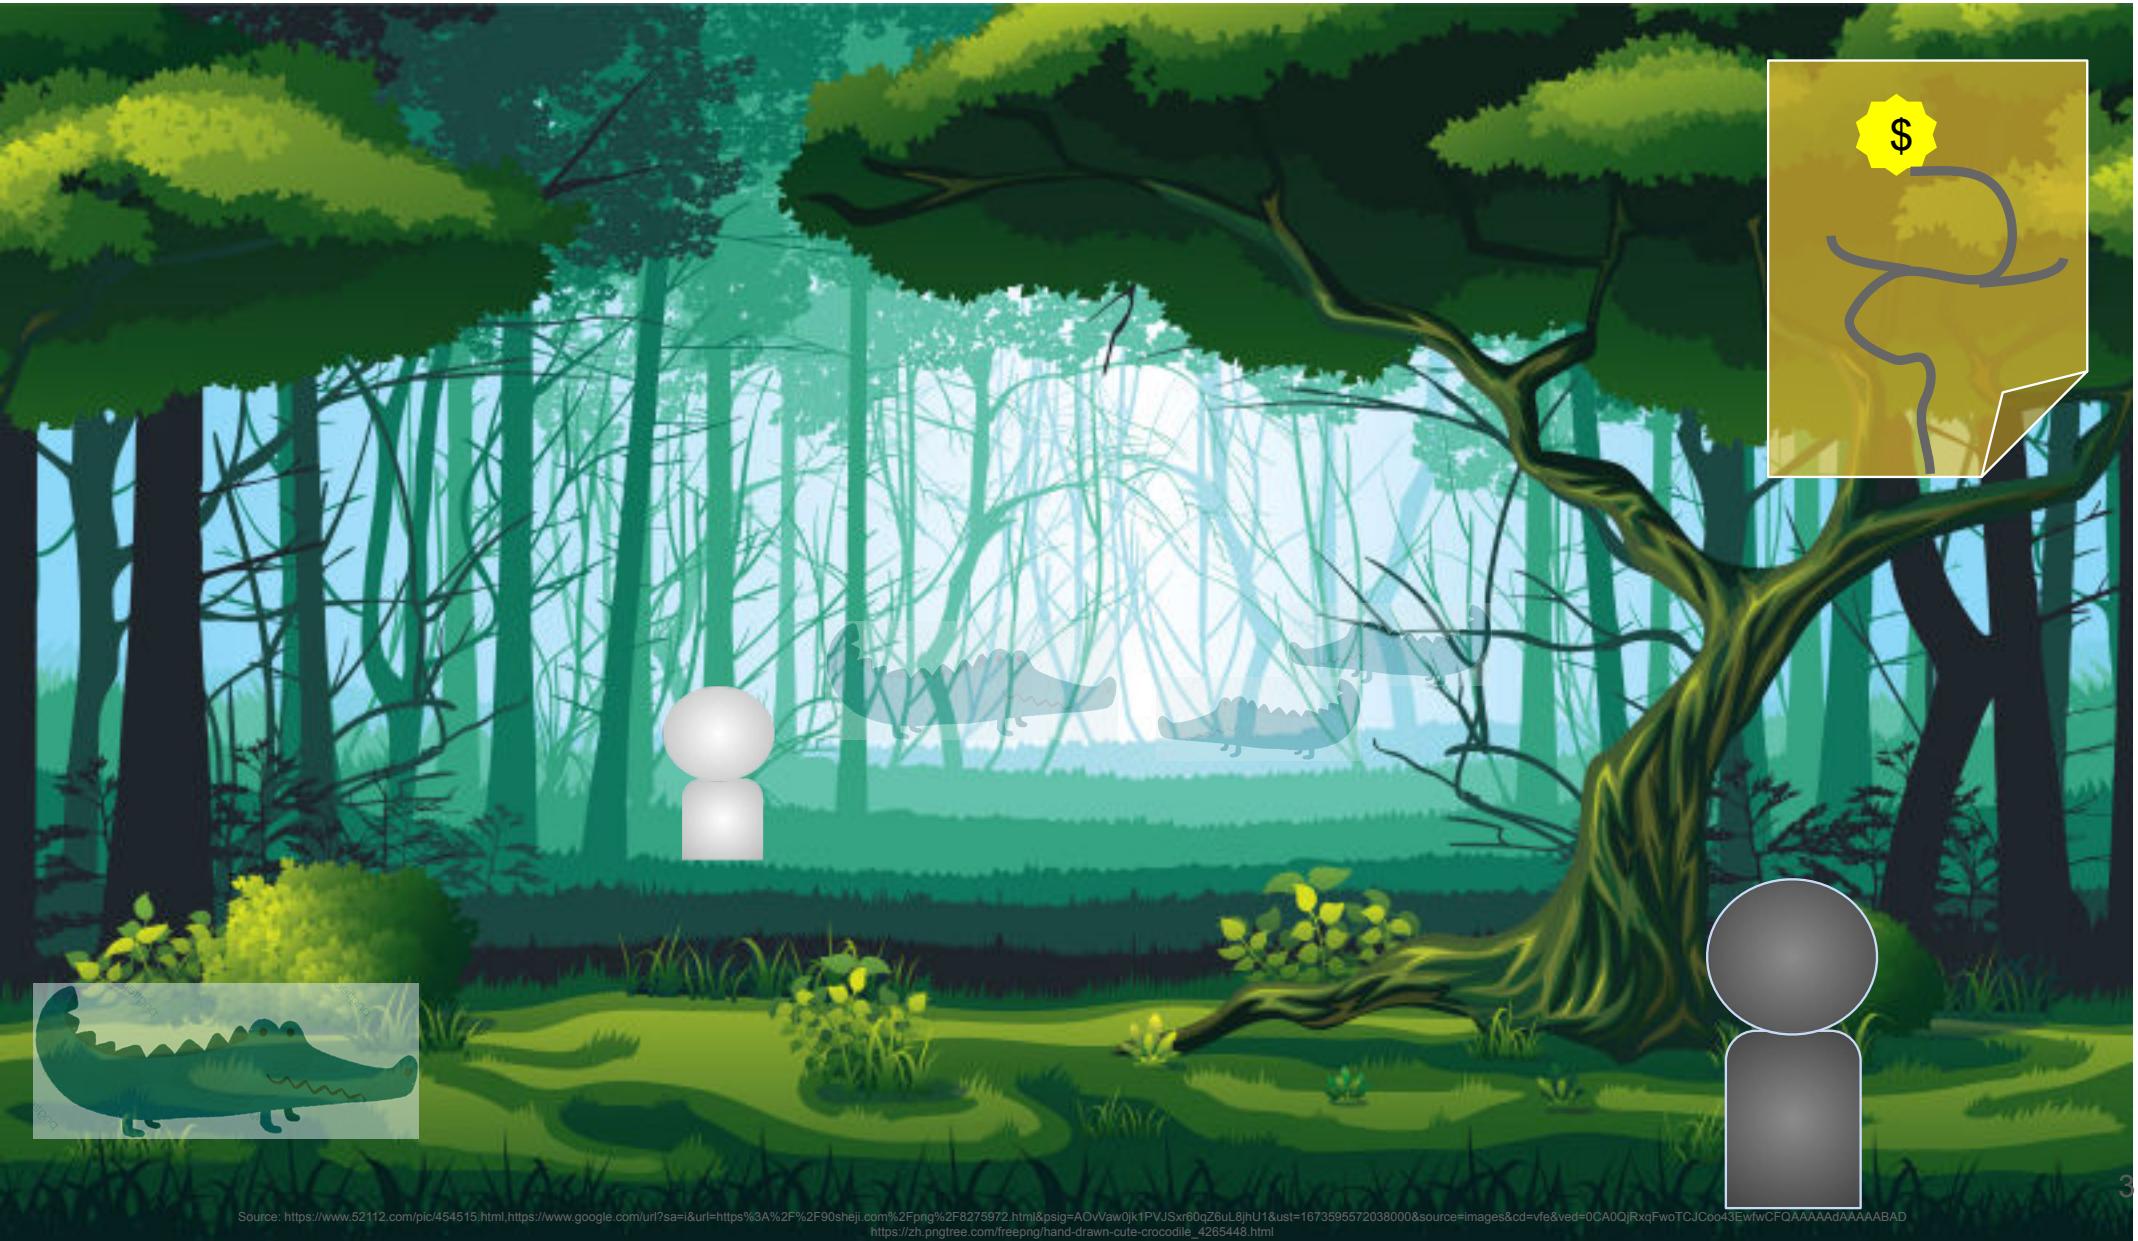


# Grading Policy

- Final report [10%]
  - # of team members
    - If there is a graduate student in the team, you could have 2-3 members
    - Otherwise, you could have 3-4 members
  - Submit your team member list **by April 7** (or assigned randomly)
  - Announce the paper list on **April 16**
    - International conference papers published in 2023-2024
  - Submit your top 3 papers **by April 23** (or assigned randomly)
  - Announce the assigned paper before **April 30**
  - Presentation on **June 4/June 11**



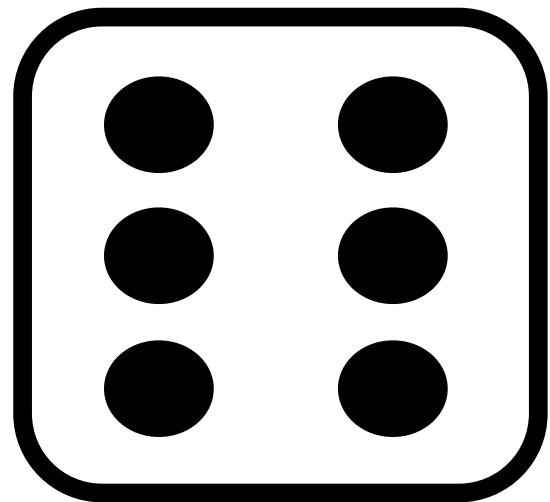
# Uncertainty

# Real World

Partial observability

Nondeterminism

Adversaries





22

°C | °F

Precipitation: 0%  
Humidity: 70%  
Wind: 13 km/h

Weather  
Monday 8:00 AM  
Cloudy

Temperature

Precipitation

Wind

0%

0%

10%

5%

5%

0%

0%

0%

9AM

12PM

3PM

6PM

9PM

12AM

3AM

6AM

Mon



28° 20°

Tue



32° 23°

Wed



33° 23°

Thu



27° 21°

Fri



24° 21°

Sat



26° 21°

Sun



27° 21°

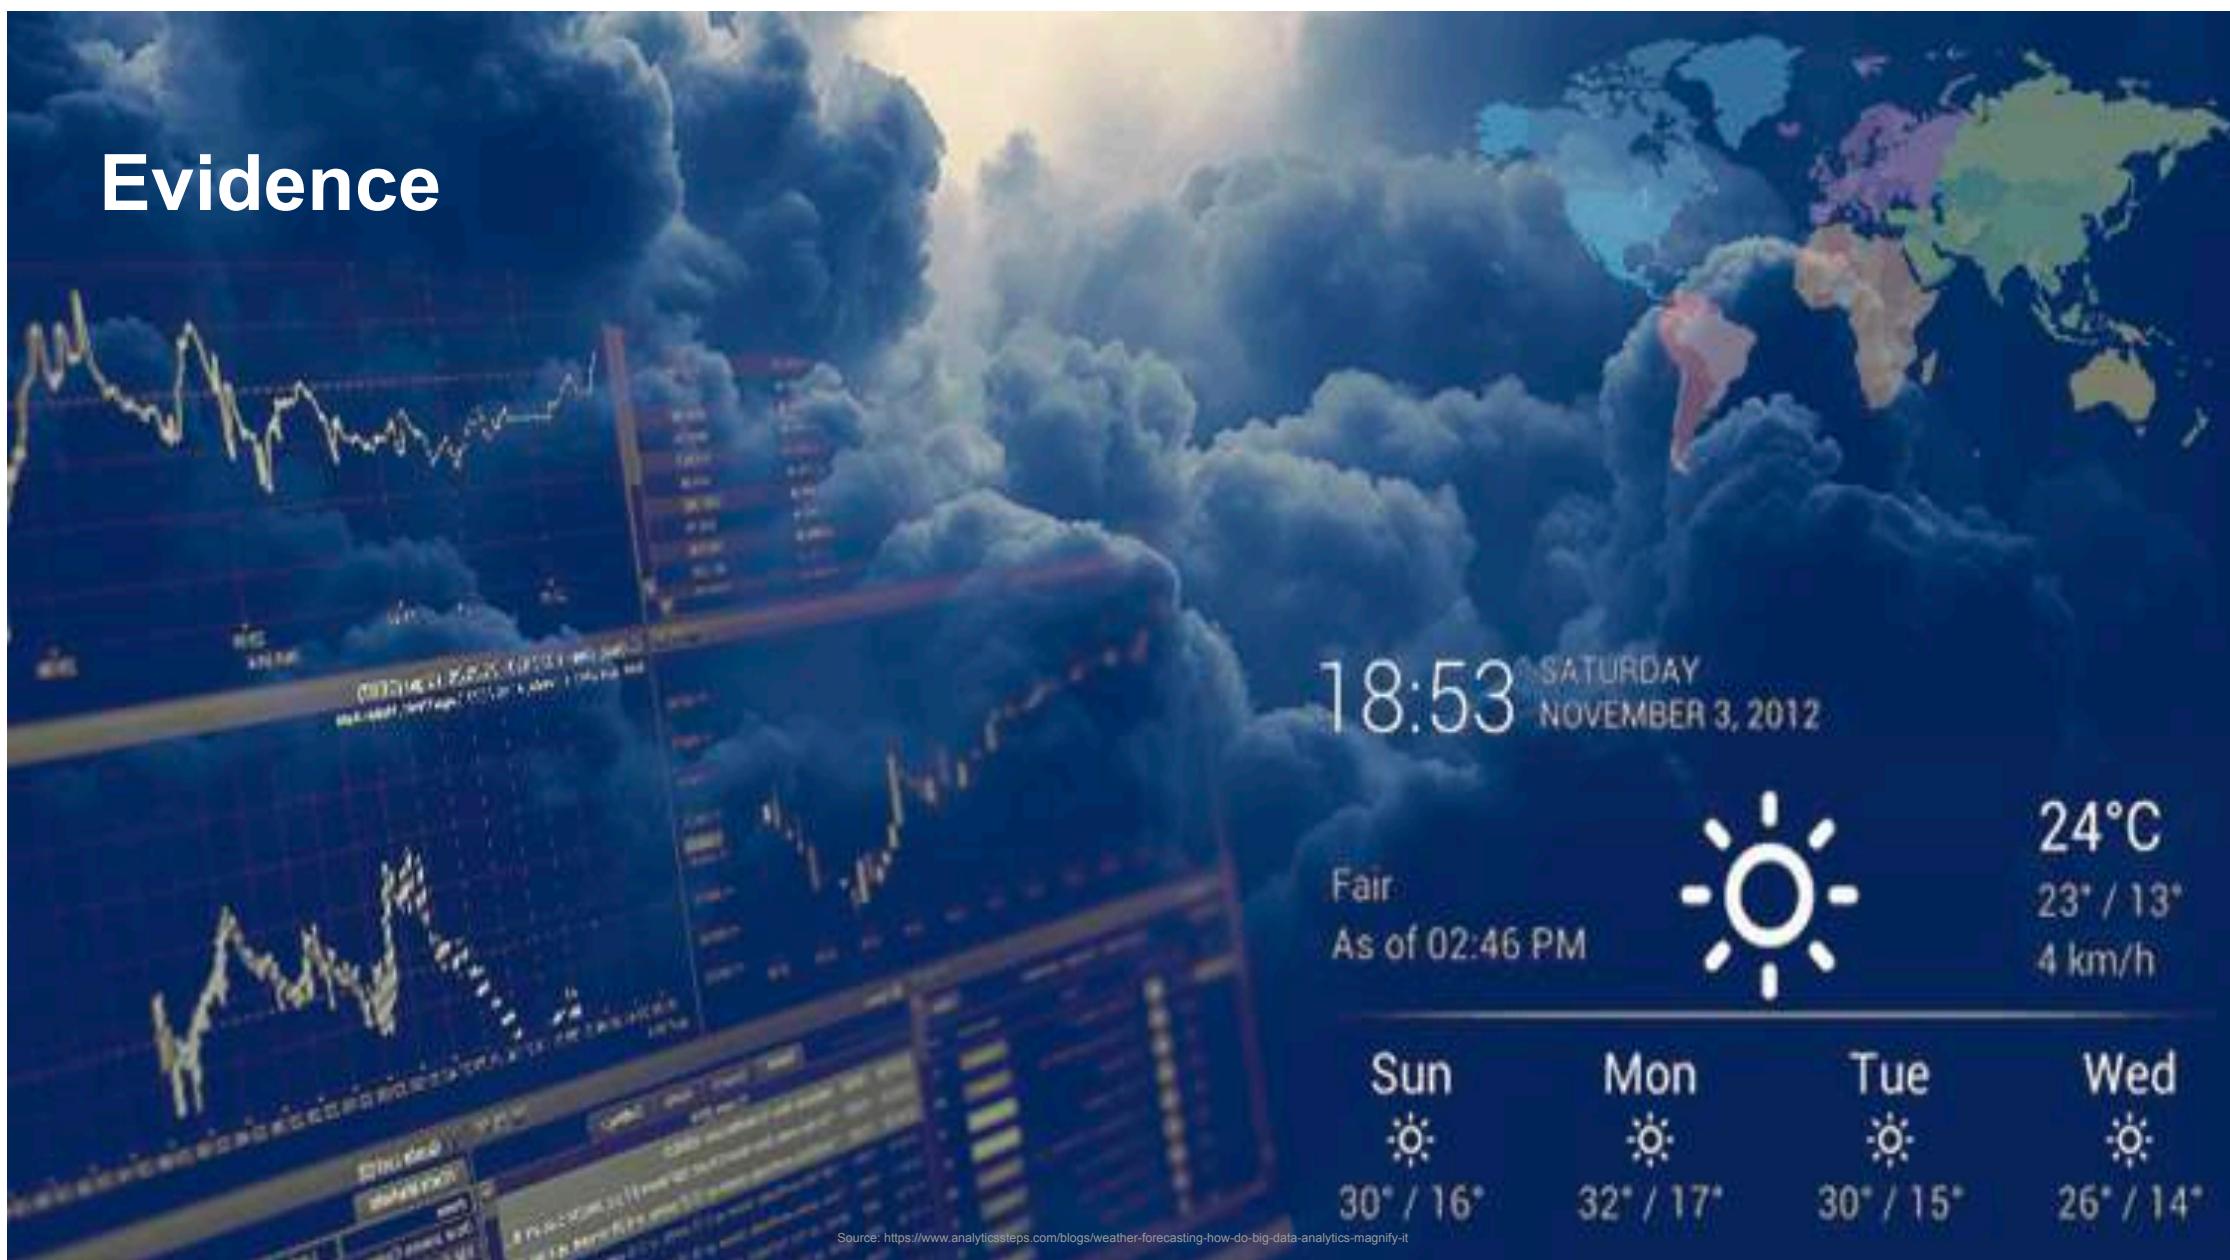
Mon



26° 18°

Credit: Google

# Evidence



# Reasoning

# Probabilistic Reasoning

# Possible World

$\omega$

eg. 1

# All Possible Worlds

{ $\omega$ }

{1, 2, 3, 4, 5, 6}

# All Possible Worlds

$\Omega = \{\omega\}$

(Sample Space)

# Probability Model

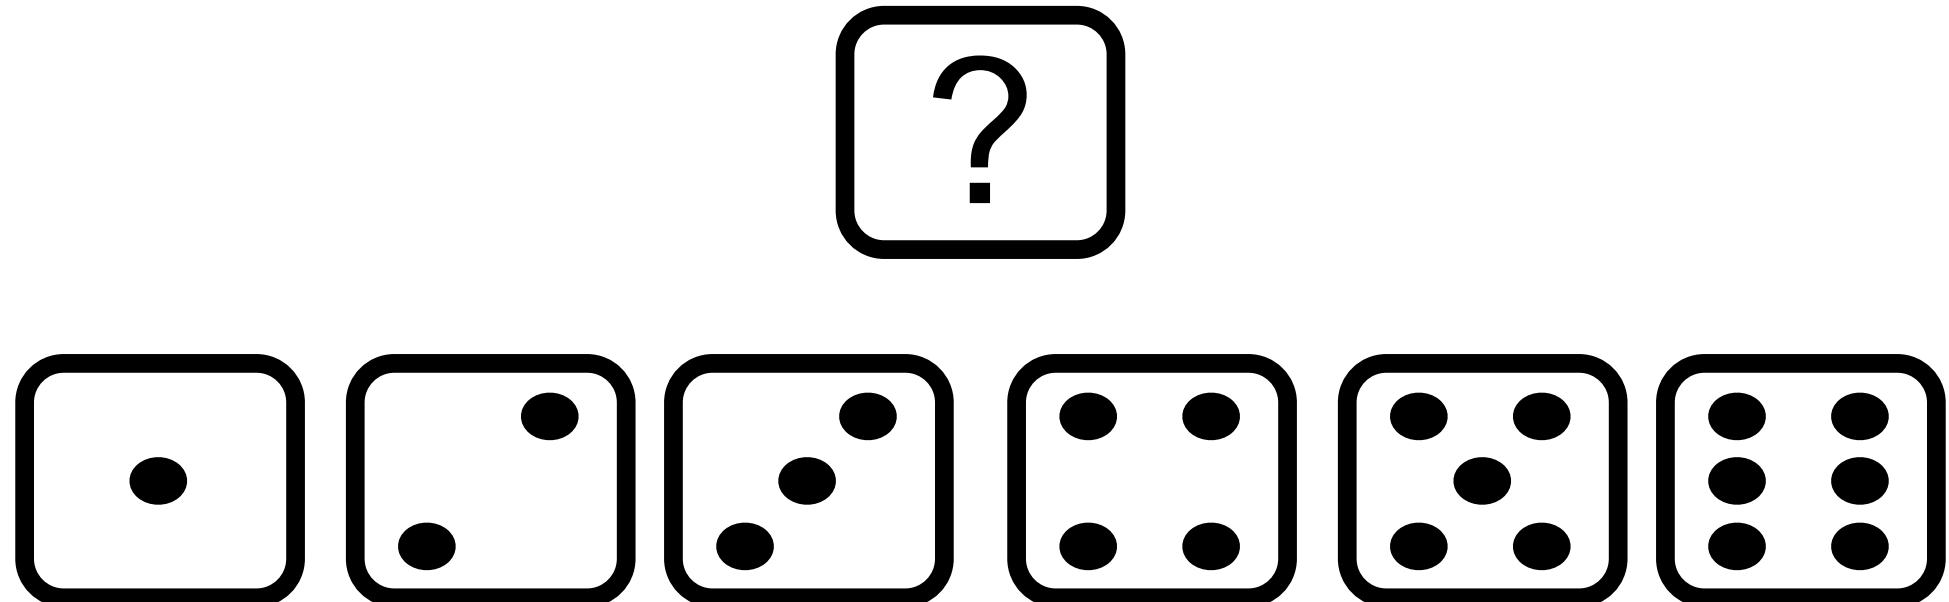
$P(\omega)$

# Probability Axioms

$$0 \leq P(\omega) \leq 1$$

$$\sum_{\omega \in \Omega} P(\omega) = 1$$

## Example: Rolling Fair Dice



$0 \leq \frac{1}{6} \leq \frac{1}{6}$     $\frac{1}{6}$     $\frac{1}{6}$     $\frac{1}{6}$     $\frac{1}{6}$     $\frac{1}{6}$



Sum = 1

# Random Variable

A variable in probability theory with a domain of possible values it can take on.

# Random Variable

Dice

1, 2, 3, 4, 5, 6

# Random Variable

Weather

sun, rain, cloud, snow

# Random Variables

- A **random variable** is some aspect of the world about which we have uncertainty, and their names begin with **an uppercase letter**
  - Dice
  - **R** = Is it raining?
  - **T** = Is it hot or cold? *temperature*
  - **W** = How's the weather?
- Random variables have domains and **names for values** are always **lowercase**
  - Dice in {1, ..., 6}
  - R in {true, false}
  - T in {hot, cold}
  - W in {sun, rain, cloud, snow}

Shorthand notation:

- **R = true**, abbreviated as **r** *real*
- **R = false**, abbreviated as **¬r**
- **T = cold**, abbreviated as **cold**

# Probability Distributions

- A probability distribution for the random variable
  - Associate a probability with each value
- Example
  - Weather
    - $P(W = \text{sun}) = 0.6$
    - $P(W = \text{rain}) = 0.1$
    - $P(W = \text{cloud}) = 0.29$
    - $P(W = \text{snow}) = 0.01$
  - Shorthand notation:

$P(\text{sun}) = 0.6$	
$P(\text{rain}) = 0.1$	
$P(\text{cloud}) = 0.29$	
$P(\text{snow}) = 0.01$	

an abbreviation

$$\mathbf{P}(W) = \langle 0.6, 0.1, 0.29, 0.01 \rangle \text{ (vector)}$$

where we assume a predefined ordering  $\langle \text{sun}, \text{rain}, \text{cloud}, \text{snow} \rangle$

Must have:

- 1)  $\forall x P(X = x) \geq 0$
- 2)  $\sum_x P(X = x) = 1$

# Joint Distributions

- A joint distribution over a set of random variables  $X_1, X_2, \dots, X_n$  specifies a real number for each assignment (or outcome)

$$P(X_1 = x_1, X_2 = x_2, \dots, X_n = x_n) \text{ or } P(x_1, x_2, \dots, x_n)$$

- Must have
  - $P(x_1, x_2, \dots, x_n) \geq 0$  and
  - $\sum_{(x_1, x_2, \dots, x_n)} P(x_1, x_2, \dots, x_n) = 1$

# Example

Given two variables

- Weather W in {sun, rain, cloud, snow}
- Cavity C (do I have a cavity?) in {true, false}

The joint probability distribution of Weather and Cavity is

$$P(W = \text{sun} \wedge C = \text{true})$$

$$P(W = \text{sun} \wedge C = \text{false})$$

$$P(W = \text{rain} \wedge C = \text{true})$$

$$P(W = \text{rain} \wedge C = \text{false})$$

$$P(W = \text{cloud} \wedge C = \text{true})$$

$$P(W = \text{cloud} \wedge C = \text{false})$$

$$P(W = \text{snow} \wedge C = \text{true})$$

$$P(W = \text{snow} \wedge C = \text{false})$$

$$4 \times 2 = 8$$

# Probabilistic Models

- A probabilistic model is a joint distribution over a set of random variables
- Probabilistic models
  - Random variables with domains
  - Assignments are called outcomes
  - Joint distributions say whether assignments (or outcomes) are likely
  - Normalized: sum to 1

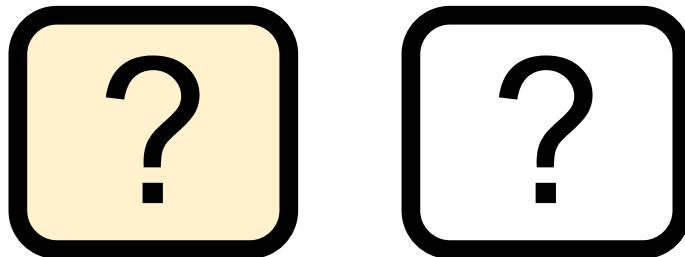
# 命題 -

## Proposition (Event)

- Proposition (or Event) is a set of worlds (or outcomes)
- For each proposition, the corresponding set contains just those possible worlds in which the proposition holds
- The probability associated with a proposition  $\Phi$ , is defined to be

$$P(\Phi) = \sum_{\omega \in \Phi} P(\omega)$$

## Example: Rolling Two Fair Dice



1	1	2	1	3	1	4	1	5	1	6	1
1	2	2	2	3	2	4	2	5	2	6	2
1	3	2	3	3	3	4	3	5	3	6	3
1	4	2	4	3	4	4	4	5	4	6	4
1	5	2	5	3	5	4	5	5	5	6	5
1	6	2	6	3	6	4	6	5	6	6	6

## Example: Rolling Two Fair Dice

- $\Phi$ : The two dice add up to 11 (Total=11)
- $P(\text{Total}=11) = ?$

1	1	2	1	3	1	4	1	5	1	6	1
1	2	2	2	3	2	4	2	5	2	6	2
1	3	2	3	3	3	4	3	5	3	6	3
1	4	2	4	3	4	4	4	5	4	6	4
1	5	2	5	3	5	4	5	5	5	6	5
1	6	2	6	3	6	4	6	5	6	6	6

## Example: Rolling Two Fair Dice

- $\Phi$ : The two dice add up to 11 (Total=11)
- $P(\text{Total}=11) = 1/36 + 1/36 = 1/18$

1	1	2	1	3	1	4	1	5	1	6	1
1	2	2	2	3	2	4	2	5	2	6	2
1	3	2	3	3	3	4	3	5	3	6	3
1	4	2	4	3	4	4	4	5	4	6	4
1	5	2	5	3	5	4	5	5	5	6	5
1	6	2	6	3	6	4	6	5	116	6	6

# Full Joint Probability Distribution

- A probability model is completely determined by the joint distribution for all of the random variables
- Example
  - Boolean variables: Cavity, Toothache, Catch\*
  - Full joint probability distribution,  $\mathbf{P}(\text{Cavity}, \text{Toothache}, \text{Catch})$  where the joint distribution is represented by

{ bold  $P$ : distribution  
non-bold  $P$ : probability

		toothache		$\neg\text{toothache}$	
		catch	$\neg\text{catch}$	catch	$\neg\text{catch}$
cavity	catch	0.108	0.012	0.072	0.008
	$\neg\text{catch}$	0.016	0.064	0.144	0.576

| Knowledge Base | =  $2^3 = 8$

\*the dentist's nasty steel probe catches in my tooth

# Probabilistic Inference

- Compute a desired probability from other known probabilities
- Example
  - Variables: Toothache, Cavity, Catch\*

		<i>toothache</i>	$\neg$ <i>toothache</i>	
		<i>catch</i>	$\neg$ <i>catch</i>	<i>catch</i>
<i>cavity</i>		0.108	0.012	0.072
$\neg$ <i>cavity</i>		0.016	0.064	0.144
				0.576

$$P(\text{cavity} \vee \text{toothache}) = 0.108 + 0.012 + 0.072 + 0.008 + 0.016 + 0.064 = 0.28$$

$$P(\text{cavity}) = 0.108 + 0.012 + 0.072 + 0.008 = 0.2$$

- called the **marginal probability** of cavity

# Marginal Distributions

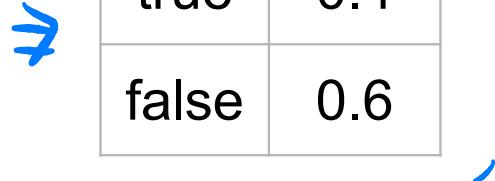
- Marginal distributions are sub-tables which eliminate variables
- **Marginalization (summing out)**
  - Combine collapsed rows by adding
    - Extract the distribution over some subset of variables or a single variable
    - Sum up the probabilities for each possible value of the other variables
- General **marginalization rule** for any sets of variables Y and Z:

$$P(Y) = \sum_z P(Y, z = z)$$

# Example 1

- Variables
  - R: is it raining?
  - T: is it hot or cold?
- $P(R) = ?$   $P(T) = ?$

$P(R, T)$		
R	T	P
true	hot	0.1
true	cold	0.3
false	hot	0.4
false	cold	0.2



$P(R)$	
R	P
true	0.4
false	0.6

$P(T)$	
T	P
hot	0.5
cold	0.5

## Example 2

*bold  $\Rightarrow$  distribution*



*uppercase  $\Rightarrow$  random variable*

- The distribution  $\mathbf{P}(\text{Cavity}) = \langle ?, ? \rangle$

		<i>toothache</i>	<i><math>\neg</math>toothache</i>	
		<i>catch</i>	<i><math>\neg</math>catch</i>	<i>catch</i>
<i>cavity</i>	0.108	0.012	0.072	0.008
<i><math>\neg</math>cavity</i>	0.016	0.064	0.144	0.576

$$\begin{aligned}\mathbf{P}(\text{Cavity}) &= \mathbf{P}(\text{Cavity}, \text{toothache}, \text{catch}) + \mathbf{P}(\text{Cavity}, \text{toothache}, \neg\text{catch}) \\ &+ \mathbf{P}(\text{Cavity}, \neg\text{toothache}, \text{catch}) + \mathbf{P}(\text{Cavity}, \neg\text{toothache}, \neg\text{catch}) \\ &= \langle 0.108, 0.016 \rangle + \langle 0.012, 0.064 \rangle + \langle 0.072, 0.144 \rangle + \langle 0.008, 0.576 \rangle \\ &= \langle 0.2, 0.8 \rangle.\end{aligned}$$

# Unconditional Probability

(Prior Probability, Priors)

Degree of belief in a proposition  
in the absence of any other  
evidence

# Conditional Probability

Degree of belief in a proposition given some evidence that has already been revealed

Bayes' Theorem :  $P(A|B) = \frac{P(B|A) P(A)}{P(B)}$

*posterior*      *likelihood*      *prior*  
*evidence*

# Conditional Probability

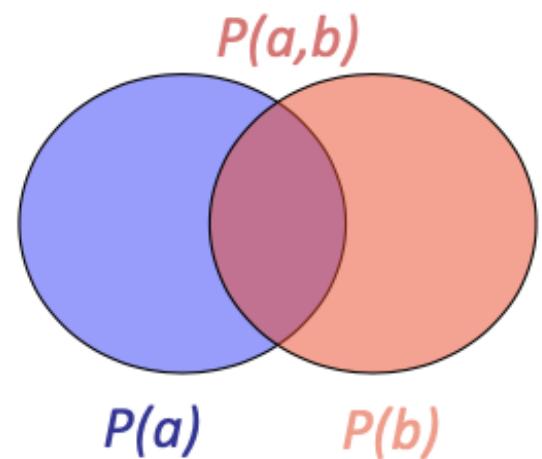
- For any propositions a and b, we have

$$P(a | b) = \frac{P(a \wedge b)}{P(b)}$$

which holds whenever  $P(b) > 0$

- Product rule

$$P(a \wedge b) = P(a | b)P(b)$$



# Example 1

- Variables
  - R: is it raining?
  - T: is it hot or cold?
- $P(R=\text{false} \mid T=\text{cold}) = ?$

$P(R, T)$

R	T	P
true	hot	0.1
true	cold	0.3
false	hot	0.4
false	cold	0.2

$$\begin{aligned}P(R = \text{false} \mid T = \text{cold}) &= P(R = \text{false}, T = \text{cold}) / P(T = \text{cold}) \\&= 0.2 / 0.5 \\&= 0.4\end{aligned}$$

## Example 2

↙ non bold  $\Rightarrow$  probability

- $P(\text{cavity} \mid \text{toothache})=?$

		toothache		$\neg\text{toothache}$	
		<i>catch</i>	$\neg\text{catch}$	<i>catch</i>	$\neg\text{catch}$
<i>cavity</i>	<i>catch</i>	0.108	0.012	<i>catch</i>	0.072
	$\neg\text{catch}$	0.016	0.064	<i>catch</i>	0.144

$$P(\text{cavity} \mid \text{toothache}) = \frac{P(\text{cavity} \wedge \text{toothache})}{P(\text{toothache})}$$

## Example 2

- $P(\text{cavity} \mid \text{toothache})=?$

		<i>toothache</i>		$\neg\text{toothache}$	
		<i>catch</i>	$\neg\text{catch}$	<i>catch</i>	$\neg\text{catch}$
<i>cavity</i>	<i>catch</i>	0.108	0.012	<i>catch</i>	0.072
	$\neg\text{catch}$	0.016	0.064	<i>catch</i>	0.144

$$\begin{aligned} P(\text{cavity} \mid \text{toothache}) &= \frac{P(\text{cavity} \wedge \text{toothache})}{P(\text{toothache})} \\ &= \frac{0.108 + 0.012}{0.108 + 0.012 + 0.016 + 0.064} = 0.6. \end{aligned}$$

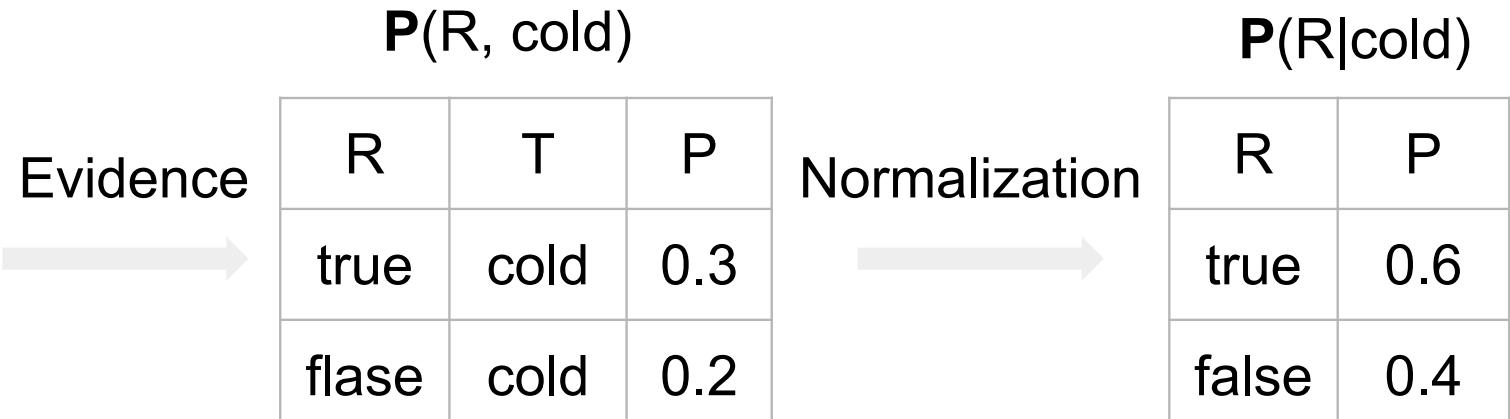
# Conditional Distributions

- Conditional distributions are probability distributions over some variables **given fixed values of others**

# Example 1

- Variables
  - R: is it raining?
  - T: is it hot or cold?
- The distribution  $P(R|\text{cold}) = \langle ?, ? \rangle$      $\langle 0.6, 0.4 \rangle$   
 $P(R, T)$

R	T	P
true	hot	0.1
true	cold	0.3
false	hot	0.4
false	cold	0.2



## Example 2

- The distribution  $\mathbf{P}(\text{Cavity} \mid \text{toothache}) = \langle ?, ? \rangle$

		<i>toothache</i>	$\neg\text{toothache}$		
		<i>catch</i>	$\neg\text{catch}$	<i>catch</i>	$\neg\text{catch}$
<i>cavity</i>		0.108	0.012	0.072	0.008
$\neg\text{cavity}$		0.016	0.064	0.144	0.576

Select evidence:

$$\begin{aligned} & [\mathbf{P}(\text{Cavity}, \text{toothache}, \text{catch}) + \mathbf{P}(\text{Cavity}, \text{toothache}, \neg\text{catch})] \\ &= [\langle 0.108, 0.016 \rangle + \langle 0.012, 0.064 \rangle] = \langle 0.12, 0.08 \rangle \end{aligned}$$

Normalization: ( $\alpha$ : a constant)

$$\alpha \langle 0.12, 0.08 \rangle = \langle 0.6, 0.4 \rangle$$

# Independence

- The knowledge that one proposition occurs does not affect the probability of the other proposition
- Independence between propositions a and b can be written as

$$P(a | b) = P(a)$$

or  $P(b | a) = P(b)$

or  $P(a \wedge b) = P(a)P(b)$

$$P(a|b) = P(a) \xrightarrow[\text{rule}]{\text{Product}} \frac{P(ab)}{P(b)} = P(a) \Rightarrow P(ab) = P(a)P(b)$$

# Bayes' Rule (Bayes' Law or Bayes' Theorem)

For any two propositions a and b,

$$P(b|a) = \frac{P(a|b)P(b)}{P(a)} = \frac{P(a|b)}{\frac{P(a)}{P(b)}} P(b)$$

posterior      likelihood      prior  
evidence

Update Factor

Prof. By product rule, we have

$$P(a \wedge b) = P(a|b)P(b) \quad \text{and} \quad P(a \wedge b) = P(b|a)P(a)$$

$$\rightarrow P(a|b)P(b) = P(b|a)P(a) \quad (= P(ab))$$

$$\rightarrow P(b|a) = \frac{P(a|b)P(b)}{P(a)}$$

$$\text{Update Factor} = \frac{\text{likelihood}}{\text{evidence}} = \frac{P(a|b)}{P(a)} : \text{If indicates how the occurrence of } b \text{ changes the occurrence of } a.$$

(a)  $0 = 1$  :  $b$  does not change the likelihood of  $a \Rightarrow$  independent

②  $> 1$  : positive association between  $a$  &  $b$ , suggesting  $a$  is more likely when  $b$  occurs than otherwise.

③  $< 1$  : negative association, suggesting  $a$  is less likely when  $b$  occurs than otherwise

(b)  $0 \gg 1$  ( $P(a|b) \gg P(a)$ ) :  $b$  tells us significant amount of  $a \Rightarrow$  high information

②  $\ll 1$  ( $P(a|b) \ll P(a)$ ) :  $a$  is very unlikely given  $b$   
 $\Rightarrow$  high information

③  $\approx 1$  ( $P(a|b) \approx P(a)$ ) : low information

## Bayes' Rule (Bayes' Law or Bayes' Theorem)

- General case for multivalued variables can be written as follows:

$$\mathbf{P}(Y | X) = \frac{\mathbf{P}(X | Y)\mathbf{P}(Y)}{\mathbf{P}(X)}$$

- General version conditionalized on some background evidence  $e$ :

$$\mathbf{P}(Y | X, \mathbf{e}) = \frac{\mathbf{P}(X | Y, \mathbf{e})\mathbf{P}(Y | \mathbf{e})}{\mathbf{P}(X | \mathbf{e})}$$

## Bayes' Rule (Bayes' Law or Bayes' Theorem)

- General form of Bayes' rule with normalization is

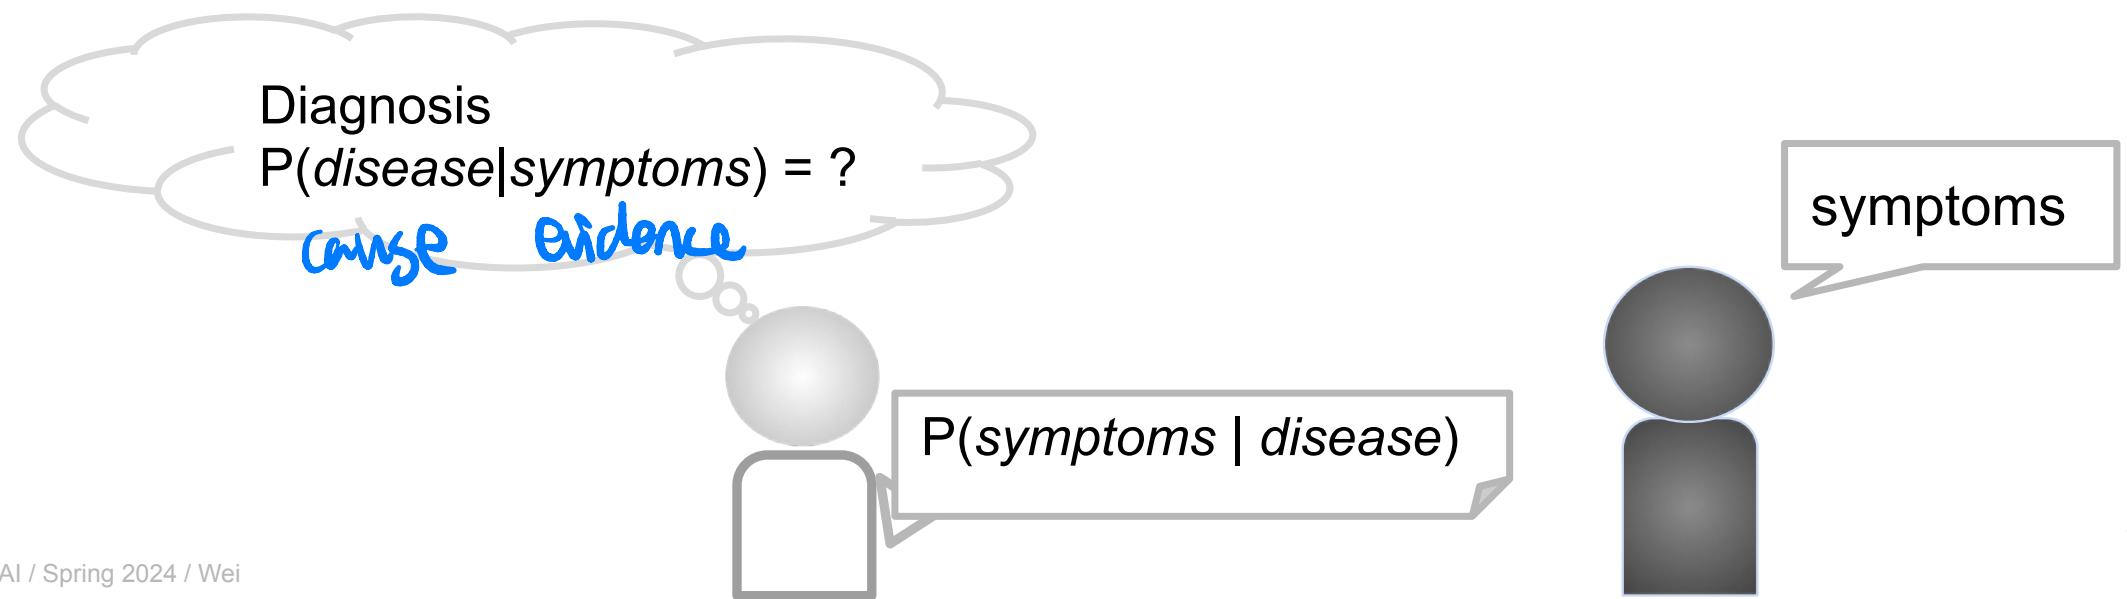
$$P(Y | X) = \alpha P(X | Y)P(Y) \quad \stackrel{\text{normalization constant}}{=} \frac{1}{P(X)}$$

where  $\alpha$  is the normalization constant needed to make the entries in  $P(Y | X)$  sum to 1

# Inference: Applying Bayes' Rule

Often, we perceive as evidence the **effect** of some unknown **cause**, and we would like to determine that *cause*

- Doctor



## Inference: Applying Bayes' Rule (Cont.)

Often, we perceive as evidence the **effect** of some unknown **cause**, and we would like to determine that *cause*

$$P(\text{cause} \mid \text{effect}) = \frac{P(\text{effect} \mid \text{cause})P(\text{cause})}{P(\text{effect})}$$

**Diagnostic**

# Example: Applying Bayes' Rule

- Doctor
  - The disease covid-19 causes a patient to have a running nose, say, 70% of the time
  - $c$ : Any patient has covid-19 is  $1/50000$
  - $R$ : Any patient has a running nose is 1%
- Question:  $P(\text{covid-19} \mid \text{running nose}) = P(c|r)$ ?

$$P(r|c) = 0.7 \quad \text{likelihood}$$

$$P(c) = 1/50000 \quad \text{prior}$$

$$P(r) = 0.01 \quad \text{evidence}$$

$$P(c|r) = \frac{P(r|c)P(c)}{P(r)} = \frac{0.7 \times 1/50000}{0.01} = 0.0014.$$

*posterior*

# Naive Bayes' Model (Bayesian Classifier)

- A single cause directly influences a number of **effects**, all of which are **conditionally independent**, given the cause
- The full joint distribution can be written as

$$\mathbf{P}(Cause, Effect_1, \dots, Effect_n) = \mathbf{P}(Cause) \prod_i \mathbf{P}(Effect_i | Cause).$$

# Conditional Independence

- Definition
  - **Conditional independence** of two variables  $X$  and  $Y$ , given a third variable  $Z$ , i.e.,  $X \perp\!\!\!\perp Y | Z$ , is

$$P(X, Y | Z) = P(X | Z)P(Y | Z)$$

## Example

$$P(\text{Cavity} \mid \text{toothache} \wedge \text{catch})$$

$$= \alpha P(\text{toothache} \wedge \text{catch} \mid \text{Cavity}) P(\text{Cavity}) \quad (\text{By Bayes' Rule})$$

$$= \alpha P(\text{toothache} \mid \text{Cavity}) P(\text{catch} \mid \text{Cavity}) P(\text{Cavity})$$

(If Toothache and Catch are conditional independence given Cavity)

## Inference: Applying Naive Bases Models

- Given some **observed** effects  $\mathbf{E} = \mathbf{e}$  while the remaining effect variables  $\mathbf{Y}$  are **unobserved**, the probability of the cause is

$$\mathbf{P}(Cause | \mathbf{e}) = \alpha \mathbf{P}(Cause) \prod_j \mathbf{P}(e_j | Cause)$$

$$\mathbf{P}(Cause, Effect_1, \dots, Effect_n) = \mathbf{P}(Cause) \prod_i \mathbf{P}(Effect_i | Cause)$$

Prof.

$$\mathbf{P}(Cause | \mathbf{e})$$

$$= \alpha \mathbf{P}(\text{Cause}, \mathbf{e})$$

$$= \alpha \sum_{\mathbf{y}} \mathbf{P}(Cause, \mathbf{e}, \mathbf{y}).$$

$$= \alpha \sum_{\mathbf{y}} \mathbf{P}(\text{Cause}) \mathbf{P}(\mathbf{y} | Cause) \left( \prod_j \mathbf{P}(e_j | Cause) \right) \text{(by Naive Bayes)}$$

$$= \alpha \mathbf{P}(\text{Cause}) \left( \prod_j \mathbf{P}(e_j | Cause) \right) \sum_{\mathbf{y}} \mathbf{P}(\mathbf{y} | Cause)$$

$$= \alpha \mathbf{P}(\text{Cause}) \prod_j \mathbf{P}(e_j | Cause)$$

# Example: Text Classification with Naive Bayes

Newspaper articles:

- Stocks rallied on Monday, with major indexes gaining 1% as optimism persisted over the first quarter earnings season.
- Heavy rain continued to pound much of the east coast on Monday, with flood warnings issued in New York City and other locations.

Task:

Classify each sentence into a Category:

*news, sports, business, weather, or entertainment*

## Example: Text Classification with Naive Bayes (Cont.)

To categorize a new document, we check which key words appear in the document and then apply

$$\mathbf{P}(Cause | e) = \alpha \mathbf{P}(Cause) \prod_j \mathbf{P}(e_j | Cause)$$

(Cause) Category

(Effect) Keywords - HasWord<sub>i</sub>

If 9% of articles are about weather,

$$P(\text{Category}=\text{weather}) = 0.09$$

If 37% of articles about weather contain word 6, "rain",

$$P(\text{HasWord}_6=\text{true} | \text{Category}=\text{weather}) = 0.37$$

Then take the one with the highest probability

# Probabilistic Models

- Reason about unknown variables, given evidence
- Applications
  - Explanation (diagnostic reasoning)
  - Prediction (causal reasoning)